

Practical Game Theory

See the Game Behind the Conversation

PAUL MASCETTA

Practical Game Theory

See the Game Behind the Conversation

Copyright © 2026 Paul Mascetta. All rights reserved.

No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means—electronic, mechanical, photocopying, recording, or otherwise—without the prior written permission of the author, except for brief quotations in reviews and articles.

This book is intended for educational purposes. The examples are illustrations, not guarantees of outcome, and nothing in it constitutes legal, financial, or psychological advice.

First edition

Contents

Introduction.....	5
PART I — SEE THE STRUCTURE BENEATH THE CONVERSATION	
Chapter 1 — See the Game	17
Chapter 2 — See the Payoffs	43
Chapter 3 — Think in Responses.....	71
Chapter 4 — Manage Predictability	93
Chapter 5 — See the Equilibrium.....	117
PART II — SEE WHAT ISN'T OBVIOUS	
Chapter 6 — Map Information & Beliefs.....	145
Chapter 7 — Separate Words From Credibility.....	171
Chapter 8 — Commitment Changes the Game.....	207
PART III — THINK BEYOND THE MOMENT	
Chapter 9 — The Shadow of the Future.....	241
Chapter 10 — Diagnose the Coordination Problem	275
Chapter 11 — Change the Game	311
Conclusion — Strategic Judgment.....	359

Introduction

I've had a pretty interesting life. Maybe not as interesting as some people's, but probably more interesting than average.

I grew up as an only child with two loving parents on a dead-end block in the 1980s. It was a great childhood, and I wouldn't trade it for anything.

Our block was always alive. We had huge block parties, and at the time it seemed like there were hundreds of kids outside playing. My mother is one of sixteen children, and at one point two of her brothers lived on our block with their eight children. My grandparents lived just down the street. There were always relatives coming and going, kids running around, people visiting, and something happening.

But there was something different about being an only child.

At the end of the day, everybody else seemed to go home to brothers and sisters. I went home alone.

And because of that, I think I became very aware, very early, of my need to connect with other people. I wanted

friends. And somewhere along the way, I realized that if I wanted people to want me around, I had to figure out how to get them to like me.

The problem was that I didn't really know how.

When you grow up in a house with two adults and don't have siblings your own age constantly teaching you the unwritten rules of childhood interaction, you improvise.

For me, that often meant trying to be funny.

I acted out. I got in trouble in school. I tried to make people laugh. Sometimes it worked. Other times I took it much too far and became incredibly disruptive.

But underneath all of that, something else was happening.

I was learning to **read people**.

I started noticing when a joke would work and when it wouldn't. I learned that the exact same thing could make one group of people laugh and annoy another. I learned about timing. Delivery. Tone. I learned to feel what was happening in a room before deciding what to say.

I didn't have language for any of this at the time.

I was simply learning the unwritten rules of human interaction.

Then life gave me a much more serious test.

After high school, I had absolutely no idea what I wanted to do with my life. I eventually took a job on Wall Street. I didn't initially understand that what was happening at the firm was illegal. Once I realized it, I quit and found another job.

But about two years later, the FBI showed up at my door.

They wanted me to cooperate with them in their investigation of the people who owned the firm. I refused.

And at twenty years old, I found myself in federal prison.

Another guy and I were the two youngest inmates on the entire unit.

Suddenly, reading people wasn't about getting a laugh in a classroom.

I was surrounded by people from completely different cultures, races, backgrounds, ages, and life experiences. There were rules—some written, many unwritten—and I had to figure them out quickly.

Who could I trust?

Who should I stay away from?

When should I talk?

When should I keep my mouth shut?

How would someone interpret what I was about to say?

If I did one thing, how would another person respond?

What was really happening underneath the obvious interaction?

Again, I didn't think of any of this formally.

I simply navigated it.

And somehow, I navigated it well enough that I made some very good friends and built genuine relationships with people whose lives were completely different from mine.

When I came home, I faced another problem.

I was still only twenty years old. I had a federal conviction, no college degree, and almost no meaningful professional experience.

My options were limited.

The jobs I could get were mostly sales jobs—positions where nobody particularly cared about my résumé as long as I could produce results. If I wanted to make money, I had to convince people to buy.

Once again, my livelihood depended on understanding human beings.

I had to listen.

I had to read people.

I had to understand what mattered to them.

I had to recognize resistance.

I had to know when to push, when to stop, when to ask a question, when to explain something differently.

And I had to understand that what I said was only part of the equation. What really mattered was **how the other person responded to it**.

For years, I did all of this mostly by intuition.

Then, in 2004, a friend of mine started an online business and needed a copywriter. He couldn't afford an experienced one, so he came to me.

I was a sales guy. He figured I knew how to sell.

So he asked me if I wanted to write his copy.

I had no formal copywriting training. I simply started writing the way I thought people needed to be spoken to in order to understand an offer and become interested in it.

And it worked.

The copy started making him a lot of money.

Eventually, I realized I was making more from copywriting than I was from my six-figure sales job in Manhattan.

So I went all in.

I studied the craft seriously. Over roughly two years, I wrote more than a hundred sales letters. That gave me an enormous amount of real-world experience in persuasion, copywriting, and marketing.

Eventually, though, I burned out.

I didn't want to spend the rest of my life writing sales letters for other people. I wanted to build something of my own and teach.

The question was:

What did I actually know well enough to teach?

The answer was obvious.

Communication.

By that point, I had spent most of my life trying to understand how people interact, how they influence one another, how they respond, how they make decisions, and how different approaches produce different outcomes.

So in 2011, I started a publishing company focused on influence, persuasion, and human communication.

Since then, I've created and published close to a hundred different trainings.

But for most of that time, I focused on the same thing almost everyone in this field focuses on:

strategies and tactics.

What words should you use?

What questions should you ask?

How should you handle an objection?

How do you build rapport?

How do you persuade?

How do you negotiate?

How do you communicate more effectively?

All of those things matter.

But there was something underneath them that I had never really identified.

Something I had been doing since I was a kid.

I wasn't merely thinking about **what to say**.

I was trying to understand **the situation surrounding what was being said**.

Who are the people involved?

What does each person want?

What choices do they have?

What do they know?

What do they believe?

What are they expecting me to do?

If I make this move, how are they likely to respond?

And if they respond that way, what happens next?

Why does the same approach work with one person and fail completely with another?

Why do people sometimes remain stuck in situations nobody actually wants?

Why can a promise be persuasive in one situation and meaningless in another?

Why does cooperation emerge in some relationships while conflict seems almost inevitable in others?

I had been thinking this way for most of my life.

I just didn't know there was an entire field devoted to studying many of these questions.

Then I discovered **game theory**.

And as I began studying it, something strange happened.

I kept recognizing ideas.

Players.

Payoffs.

Best responses.

Strategic moves.

Information.

Signals.

Credibility.

Commitment.

Repeated interactions.

Cooperation.

Coordination.

Equilibrium.

Again and again, I found formal concepts describing patterns I had spent much of my life noticing intuitively.

And that led me to a question:

Could the concepts of game theory be used to teach people a fundamentally different way of thinking about communication and human interaction?

Not another collection of scripts.

Not another list of persuasion tricks.

Not another book telling you the seven phrases you should use when someone disagrees with you.

Something deeper.

A way of looking at an interaction and seeing the structure underneath it.

Because every important interaction contains more than words.

There are people with different goals.

Different incentives.

Different information.

Different options.

Different expectations.

And every move one person makes can change what the other person wants to do next.

In other words:

There is a game behind the conversation.

I began digging into the research to find out whether the connection I was seeing was actually there.

The deeper I went, the more convinced I became that it was.

Game theory was never created as a theory of better communication. It is a formal study of strategic interaction. And human beings are far messier than the perfectly rational players found in many mathematical models.

But the mental models game theory gives us are extraordinarily useful.

They force us to ask different questions.

Instead of only asking:

What should I say?

we can ask:

What game am I actually in?

Instead of:

How do I convince them?

we can ask:

What do they actually value?

Instead of:

Was that a good move?

we can ask:

Then what? What response is that move likely to create?

Instead of:

Why does this problem keep happening?

we can ask:

What is making this pattern stable?

And sometimes, instead of asking:

How do I play this game better?

we discover that the better question is:

Can I change the game?

That is what this book is about.

I'm going to take you through the game-theoretic concepts that I believe are most useful for understanding everyday human interaction, the research behind them, and the practical ways they change how we think about communication, relationships, leadership, negotiation, conflict, business, and life.

The goal isn't to turn you into a mathematician.

And it certainly isn't to teach you how to manipulate people.

The goal is much simpler:

To teach you to see the game behind the conversation.

Because once you can see the game, you start seeing human interaction differently.

And once you see it differently, you have far more choices about how you play.

PART I

**See the Structure Beneath the
Conversation**

CHAPTER 1

See the Game

Most communication advice begins with a deceptively simple question:

What should I say?

What if that's the wrong place to start?

Imagine you walk into your boss's office and ask for a raise.

Before you say anything, you're probably thinking about the words you're going to use. Maybe you're considering how much to ask for. Maybe you're thinking about how to demonstrate your value. Maybe you're trying to anticipate the objections your boss will raise.

All of that matters.

But there's another way to look at the interaction.

What does your boss actually want?

What do you want?

What happens if your boss says no?

What happens if you say no to a counteroffer?

What alternatives do you have?

What alternatives does your boss have?

What does your boss know that you don't?

What do you know that your boss doesn't?

Who has the authority to make the decision?

Who else might influence it?

What will your boss probably do after you make your request?

And what are you likely to do after your boss responds?

Suddenly, this isn't simply a conversation.

It's a game.

Not a game in the sense that you're trying to beat your boss.

A game in the game-theoretic sense.

A situation in which the outcome depends, at least partly, on the decisions of multiple people.

That distinction is the foundation of everything we're going to explore in this book.

What Does It Mean to "See the Game"?

When most people hear the words game theory, they imagine chess, poker, business competition, or perhaps two countries threatening each other.

Those are certainly games.

But game theory is much broader than that.

At its most fundamental level, game theory studies **strategic interaction**: situations in which the choices of one decision-maker affect the outcomes available to others. The formal framework describes players, their available strategies, their preferences or payoffs, and the resulting outcomes.¹

That means a game doesn't have to involve a winner and a loser.

You can have:

- cooperation,
- competition,
- negotiation,
- coordination,
- conflict,
- compromise,
- friendship,
- business partnerships,
- family decisions.

In fact, Thomas Schelling emphasized that many real-world interactions contain **both conflicting and common interests** rather than being purely competitive.²

That's important because we're not going to use game theory as a cynical philosophy that says:

Everyone is trying to manipulate everyone else.

That's not what I'm proposing.

I'm proposing something much simpler:

1. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/archives/fall2014/entries/game-theory/>

2. Nobel Prize, <https://www.nobelprize.org/uploads/2018/06/advanced-economicsciences2005.pdf>

Other people are decision-makers, and their decisions affect your outcomes.

Once you recognize that, communication starts to look different.

The Conversation Is Only One Part of the Game

Let's go back to the conversation with your boss.

Suppose you say:

"I've been here for three years, I've increased sales by 30%, and I'd like to discuss increasing my compensation."

Most communication books would immediately start evaluating the statement.

Was that the right thing to say?

Should you have phrased it differently?

Should you have established more rapport first?

Should you have asked a question?

Should you have used a different persuasion technique?

Those can be worthwhile questions.

But game theory makes us ask something that comes **before** all of them:

What happens after I say it?

Your sentence doesn't exist in isolation.

Your boss hears it and makes a decision.

Perhaps your boss:

- agrees,

- rejects the request,
- asks for more information,
- makes a counteroffer,
- delays the conversation,
- consults someone else,
- tests how badly you want the raise.

Your move changes their decision problem.

And their response changes yours.

That is the essence of strategic interaction.

In a sequential game, the order in which players act matters because later decisions can depend on what happened earlier. Formal game theory represents these interactions using game trees and histories of previous actions.³

But you don't need a game tree to use the mental model.

You simply need to develop the habit of asking:

Then what?

Who Are the Players?

One of the first mistakes we make in an interaction is assuming that the person we're talking to is the entire game.

Sometimes they aren't.

Imagine you're negotiating with a salesperson.

You think you're negotiating with:

You vs. the salesperson.

3. MIT OpenCourseWare, https://ocw.mit.edu/courses/6-254-game-theory-with-engineering-applications-spring-2010/688240e73afebf4a6293fef0e529398_MIT6_254S10_lec12.pdf

But perhaps the salesperson has to get approval from a manager.

Maybe the company has a pricing policy.

Maybe another customer is interested in the same product.

Maybe you have three alternative suppliers.

Suddenly, the strategic situation looks very different.

The visible conversation is only the surface.

There may be other people whose decisions affect the outcome.

So one of the first questions we should learn to ask is:

Who can actually affect the outcome?

Sometimes the answer is obvious.

Sometimes it isn't.

And sometimes the people who matter most aren't even in the room.

People Are Players, Too

The word player can sound cold when we're talking about human relationships.

It shouldn't.

A player is simply a decision-maker within the interaction.

Your spouse is a player.

Your business partner is a player.

Your customer is a player.

Your employee is a player.

Your boss is a player.

You are a player.

And recognizing someone as a player doesn't mean reducing them to a selfish calculation machine.

Real human beings have complicated motivations.

They care about money, but also about fairness.

They care about winning, but also about relationships.

They care about status, identity, loyalty, security, principles, reputation, and countless other things.

That's precisely why we need to learn to identify what matters to each person rather than simply assuming that everyone wants the same thing.

What Does Everyone Actually Want?

We'll devote the next chapter to this question, but we need to introduce it here.

Imagine two business partners.

They both want the company to succeed.

So are they cooperating?

Yes.

But they might simultaneously disagree over:

- how much each should be paid,
- who has decision-making authority,
- how much money should be reinvested,
- how quickly the company should grow,
- who deserves credit.

Their interests are neither completely aligned nor completely opposed.

They're mixed.

This is actually much closer to the structure of many human interactions.

The Nobel research on Schelling's work emphasizes precisely this mixture of common and conflicting interests in real strategic situations.⁴

So when you're trying to understand an interaction, don't immediately ask:

Who's right?

And don't immediately ask:

Who's trying to win?

First ask:

Where are our interests aligned, and where are they different?

That's a much more useful question.

What Can Each Player Actually Do?

Here's another question we often overlook.

We tend to think about what **we** can do.

Game theory makes us look at what **everyone** can do.

Suppose your employee wants to give you a refund.

But company policy doesn't allow it.

Suppose your boss wants to give you a raise.

4. Nobel Prize, <https://www.nobelprize.org/uploads/2018/06/advanced-economicsciences2005.pdf>

But the budget has already been approved.

Suppose your business partner wants to accept your proposal.

But their other partner has veto power.

The person's preference is only part of the picture.

Their **available actions** matter too.

Formal game theory distinguishes the players' strategies or available choices from their preferences over outcomes.⁵

In ordinary language:

Wanting something and being able to do something are not the same thing.

That's why constraints matter.

Constraints Are Part of the Game

Imagine you're negotiating with someone who says:

"I would love to give you that price, but I simply can't go below \$10,000."

You could immediately assume they're lying.

Or you could ask a different question:

What actually constrains them?

Maybe nothing does.

Maybe they are bluffing.

Maybe the company really does have a minimum price.

Maybe they have authority to go lower but don't want to.

5. Nobel Prize, <https://www.nobelprize.org/uploads/2017/05/nash-lecture.pdf>

Those situations look identical on the surface.

Strategically, they're completely different.

This is one of the reasons later chapters will distinguish between:

what people **say** constrains them,
and

what actually **constrains** them.

What Happens If We Don't Agree?

This is one of the most important questions you can ask in any strategic interaction.

Suppose you're negotiating with a contractor.

You want them to lower their price.

They want you to accept their price.

You might spend all your time trying to persuade them.

But there's another question:

What happens if we walk away?

Maybe you have four other contractors you could call.

Maybe they have ten other customers waiting.

Maybe neither of you has a particularly good alternative.

Maybe one of you desperately needs the deal and the other doesn't.

Those outside options change the interaction.

Your alternatives don't just matter **after** the negotiation fails.

They can change the negotiation itself.

This is why bargaining theory treats outside options as an important part of the strategic environment, although their effects depend on the specific structure of the game rather than following a simplistic rule that “the person with the better alternative always wins.”⁶

So another question belongs in our mental toolkit:

What happens if this interaction doesn’t produce an agreement?

Sometimes that question changes everything.

Dependence Changes the Game

Here’s another way to think about it.

Ask:

Who needs whom more?

You need your employee.

Your employee needs the job.

You need your supplier.

Your supplier needs your business.

You want the relationship.

The other person wants the relationship too.

But perhaps one person has more alternatives than the other.

This doesn’t automatically determine the outcome.

6. MIT OpenCourseWare, <https://ocw.mit.edu/courses/14-126-game-theory-spring-2016/pages/syllabus/>

But it changes the strategic environment.

And this is why “power” isn’t simply something one person possesses.

Often, power comes from the structure of **dependence and alternatives**.

If I have no alternative and you have five, my position is different from what it would be if I had five and you had none.

That’s not a persuasion tactic.

It’s a property of the game.

Timing Is Part of the Game

Here’s another thing that becomes obvious once you start seeing interactions this way:

When you act matters.

Suppose I make you an offer and you have to respond immediately.

That’s one game.

Suppose I have to make an offer, wait for your response, and then decide whether to improve it.

That’s another game.

Suppose you get to see my offer before deciding what yours will be.

Different game.

Suppose neither of us knows what the other is going to do.

Different game again.

Game theorists distinguish simultaneous-move situations from sequential games precisely because the timing and observability of actions change strategic possibilities.⁷

And this gives us a simple practical question:

Who moves when?

Sometimes the most important part of your strategy isn't what you do.

It's when you do it.

The Game Behind the Conversation

Here's where things get even more interesting.

Sometimes the obvious subject of the conversation isn't actually the whole problem.

Someone says:

"I don't care about the money."

But perhaps what really matters to them is recognition.

Someone says:

"I just want to know why you didn't tell me."

But perhaps the deeper issue is trust.

Someone says:

"I don't care who gets credit."

But perhaps they very much care about status.

Now, we need to be careful.

The lesson isn't:

7. MIT OpenCourseWare, https://ocw.mit.edu/courses/6-254-game-theory-with-engineering-applications-spring-2010/688240e73afebf4a6293fef0e529398_MIT6_254S10_lec12.pdf

Assume everyone is lying about their motives.

That's not strategic thinking.

That's paranoia.

The lesson is:

Don't assume the surface description of an interaction tells you everything about the strategic structure.

Sometimes the stated issue really is the issue.

Sometimes it isn't.

The point is to remain curious.

The Game Can Change While You're Playing It

Here's something else that is easy to miss.

A game isn't always static.

Your actions can change the situation itself.

Suppose you ask someone for a favor.

They say yes.

Now you've established something about your relationship.

Suppose you break a promise.

Now they may behave differently the next time.

Suppose you reveal information.

Now the other person knows something they didn't know before.

Suppose you publicly commit yourself to a position.

Now backing away may become more costly.

Your move can therefore change:

- what the other person knows,
- what they believe,
- what they expect,
- what choices remain available,
- and what happens next.

This is why game theory isn't simply about choosing the best move from a static menu.

It is about understanding **interdependent decisions over time**.

A Game Doesn't Have to Be Competitive

This distinction is so important that I want to make it explicit.

You can think about your marriage through game theory without believing your spouse is your opponent.

You can think about a business partnership through game theory without believing your partner is trying to beat you.

You can think about a friendship through game theory without turning it into a calculation.

In fact, understanding the strategic structure of a relationship can sometimes make cooperation **easier**, because you begin to understand why another person's behavior makes sense from their position.

Suppose you and a coworker both want a project completed.

You think they're procrastinating.

They think you're going to take over if they start.

You both want the same outcome.

But your expectations about each other are creating a problem.

That's not necessarily a character flaw.

It may be a game.

And once you see the structure, you may be able to change it.

The Difference Between a Person Problem and a Game Problem

This distinction is going to become increasingly important throughout the book.

Imagine a manager says:

“My employees aren't cooperating.”

You could conclude:

“They're selfish.”

But perhaps the compensation system rewards employees for outperforming one another.

Or perhaps nobody knows who owns the task.

Or perhaps employees believe management will take credit for their ideas.

Or perhaps speaking up has historically been punished.

The behavior may be coming from the structure of the interaction.

That doesn't mean the people have no responsibility.

It means we shouldn't automatically diagnose a **person problem** when there may be a **game problem**.

One of the most powerful things game theory gives us is the ability to step back from the individual behavior and examine the environment producing it.

And Sometimes the Game Is Inside the Game

Let's say two people are arguing about a business decision.

On the surface, they're debating:

Which strategy should we use?

But underneath that might be another game:

Who gets final authority?

Or:

Who gets credit if it succeeds?

Or:

Who gets blamed if it fails?

Or:

Who is trying to establish precedent for future decisions?

The visible disagreement may therefore be only one layer of the interaction.

This doesn't mean we should become obsessed with hidden motives.

It means we should recognize that human interaction can contain **multiple overlapping incentives**.

That's one reason a purely tactical approach to communication can be insufficient.

You can have the perfect words and still be playing the wrong game.

The Game Map

At this point, I want to give you a simple way to begin practicing this way of thinking.

When you enter an important interaction, pause and look at the structure.

WHO?

Who are the relevant players?

Who can actually affect the outcome?

WHAT?

What does each person want?

Where are your interests aligned?

Where are they different?

OPTIONS?

What can each player actually do?

CONSTRAINTS?

What limits those choices?

INFORMATION?

Who knows what?

Who doesn't know what?

TIMING?

Who moves first?

Who responds?

What can each person observe?

ALTERNATIVES?

What happens if there is no agreement?

DEPENDENCIES?

Who needs whom?

Who has alternatives?

THEN WHAT?

If I make this move, what is the other person likely to do next?

This isn't meant to become a rigid nine-step procedure that you mechanically run every time you speak to someone.

That's not the point of this book.

The point is to train your **attention**.

Eventually, you don't consciously run through the list.

You start noticing these things automatically.

That's what a mental model does.

A Simple Test

Let's put this into practice.

Imagine your boss sends you a text:

"We need to talk about your performance."

What do you do?

Most people immediately start preparing what they're going to say.

But stop.

Try seeing the game first.

What do you know?

Not much.

What don't you know?

Almost everything important.

What does your boss want?

You don't know yet.

What options does your boss have?

Several.

What options do you have?

Several.

What happens if the conversation goes badly?

Potentially quite a lot.

What happens if it goes well?

Potentially quite a lot.

What information might your boss have that you don't?

What information might you have that your boss doesn't?

Who else could influence the outcome?

And perhaps most importantly:

Why did your boss choose those particular words?

Now you're entering the conversation differently.

You're not simply preparing a speech.

You're trying to understand the game.

Why This Is Different From Traditional Communication Advice

Most communication instruction starts with the assumption that the problem is primarily linguistic.

If you say the right thing in the right way, you will produce the right result.

Sometimes that's true.

But it isn't always.

Because your words enter a system that already contains:

- incentives,
- expectations,
- information,
- relationships,
- constraints,
- alternatives,
- history,
- timing,
- and other decision-makers.

John Nash's equilibrium concept captures the fundamental idea that each player's behavior must be considered in relation to the behavior of the other players.⁸

That's the shift I want this book to make.

Instead of asking only:

What should I say?

start asking:

What situation am I saying it in?

8. Nobel Prize, <https://www.nobelprize.org/uploads/2017/05/nash-lecture.pdf>

Instead of:

How do I get them to agree?

ask:

What does agreement mean to them?

Instead of:

How do I win this interaction?

ask:

What outcome does each of us actually want?

Instead of:

Why did they respond that way?

ask:

**What about the structure of the interaction made
that response reasonable from their perspective?**

And instead of:

What should I do next?

ask:

What will my move cause them to do next?

Those questions don't give you a script.

They give you something more valuable.

They give you a way of seeing.

You Don't Have to Become a Mathematician

There is a tremendous amount of mathematics behind game theory.

There are formal models of equilibrium, Bayesian games, sequential games, repeated games, bargaining, cooperative

games, mechanism design, and much more. Modern university courses can spend entire semesters working through these models.⁹

We're not doing that here.

We're extracting something different.

We're looking for the **mental models underneath the mathematics** that can help us understand human interaction.

The goal isn't for you to calculate a Nash equilibrium before asking your spouse where she wants to eat.

The goal is for you to start noticing:

We're not actually disagreeing about the restaurant.

We're both worried about choosing something the other person won't like.

That's game theory becoming useful.

The First Mental Shift

Before you discovered this way of thinking, you might have looked at an interaction like this:

Me → What should I say? → Them

Now I want you to start seeing:

Me → My move → Their response → My next move → Their next response → Outcome

And underneath that:

Their incentives.

My incentives.

9. MIT OpenCourseWare, <https://ocw.mit.edu/courses/14-126-game-theory-spring-2016/>

Our information.

Our expectations.

Our alternatives.

Our constraints.

Our history.

The structure of the interaction.

That's the game.

And once you begin seeing it, something interesting happens.

You stop treating every interaction as a problem of finding the perfect words.

Sometimes the words are the problem.

But sometimes the words are merely the most visible part of a much larger system.

The Question to Carry With You

You don't need to memorize everything in this chapter.

If you remember one question, remember this:

What's really going on here?

Who are the players?

What do they want?

What can they do?

What do they know?

What do they expect?

What happens next?

And what happens if nothing gets resolved?

Those questions are the beginning of strategic thinking.

Because before you can make a better move...

you have to see the game you're playing.

What we have established

Game theory gives us a formal language for strategic interaction: players have choices and preferences, and the outcome depends on how those choices interact.¹⁰

But the purpose of this book isn't to turn communication into mathematics.

It's to take the underlying ideas and make them useful.

The first lesson is therefore deceptively simple:

Don't start with the words. Start with the situation.

In the next chapter, we'll take the next step.

Once you've identified the game, the next question is:

What does everyone actually want?

Because if you misunderstand the payoffs, even a brilliant strategy can be aimed at the wrong target.

Yesterday 2:20 PM

proceed

10. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/archives/fall2014/entries/game-theory/>

See the Payoffs

In the last chapter, we started with a simple idea:

Before you decide what to say or do, understand the game you're in.

But identifying the players isn't enough.

You also have to understand what each player is trying to accomplish.

Because here's a mistake we make all the time:

We assume other people want what we want.

And when they don't behave the way we expect them to, we conclude that they're irrational, difficult, selfish, stubborn, or just wrong.

Sometimes they are.

But often, they're simply responding to a different set of incentives.

They're playing the same game...

but they're playing for a different payoff.

The Raise You Didn't Get

Let's go back to the example from the first chapter.

You walk into your boss's office and ask for a raise.

You want more money.

That seems obvious.

But what does your boss want?

You might immediately answer:

“To save money.”

Maybe.

But perhaps your boss also wants:

- to retain good employees,
- to stay within a budget,
- to avoid setting an expensive precedent,
- to keep the rest of the team happy,
- to demonstrate fairness,
- to hit this year's profit target,
- to avoid explaining the raise to a superior.

And perhaps there's something else you haven't considered:

Your boss may actually **want** to give you the raise.

But doesn't believe they can afford to.

Suddenly, the interaction looks very different.

You came into the conversation thinking:

I want more money. My boss doesn't want to give it to me.

But the actual game might be:

We both want me to stay, but we have different constraints and different preferences about how the compensation should be structured.

That's a completely different problem.

And therefore, it requires a completely different way of thinking.

What Is a Payoff?

In game theory, a player's **payoff** is essentially what they value about an outcome.

Economists often represent payoffs numerically because mathematics requires some way of comparing outcomes.

But you don't need numbers to use the concept.

A payoff is simply:

What makes this outcome better or worse for this player?

Money can be a payoff.

But so can:

- time,
- security,
- status,
- freedom,
- convenience,
- fairness,
- reputation,
- control,
- relationships,
- avoiding embarrassment,

- reducing risk,
- keeping a promise,
- protecting one's identity.

And sometimes what someone wants isn't even the outcome itself.

It might be what the outcome **represents**.

A person may fight over \$500 because of what conceding the \$500 means to them.

They may be protecting:

respect,
principle,
status,
or the feeling that they weren't taken advantage of.

If you only look at the dollar amount, you may completely misunderstand the game.

Money Isn't Always the Real Currency

Imagine two business partners arguing over how much money to reinvest in their company.

Partner A wants to reinvest 80%.

Partner B wants to distribute most of the profits.

You might assume the disagreement is about money.

But suppose Partner A is motivated primarily by:

long-term growth.

And Partner B is motivated primarily by:

financial security.

Now their disagreement makes more sense.

Neither person is necessarily irrational.

They are optimizing for different things.

Or consider a parent and teenager arguing about a curfew.

The parent says:

“You need to be home by midnight.”

The teenager says:

“That’s ridiculous.”

What is the parent optimizing for?

Maybe:

safety.

What is the teenager optimizing for?

Maybe:

independence.

If you frame the disagreement as:

parent versus teenager,

you miss the underlying structure.

If you recognize the competing payoffs, you can begin to see why both positions exist.

The Most Important Question

Whenever you’re interacting with another person, ask:

What are they actually trying to get out of this?

Not:

What do I think they should want?

Not:

What would I want if I were them?

And not even necessarily:

What did they say they want?

Ask:

What outcome appears to matter to them?

That distinction is fundamental.

Because people don't make decisions based on what would make you happy.

They make decisions based on what they value.

Don't Assume the Other Person Is Your Opponent

One of the most important lessons in game theory is that interests can have different structures.

Sometimes interests are strongly opposed.

Sometimes they're largely aligned.

And very often they're mixed.

Consider a negotiation.

You want to buy a house for \$500,000.

The seller wants \$550,000.

There's an obvious conflict:

You want a lower price. They want a higher price.

But there's also a powerful area of alignment:

You both want the transaction to happen.

That means this isn't a purely competitive game.

You have:

conflicting interests over price

and

aligned interests over completing the transaction.

This is incredibly common in human interaction.

And it changes the question.

Instead of asking:

How do I beat them?

you can ask:

Where are our interests actually aligned, and where are they opposed?

That question opens up possibilities.

Three Basic Structures of Interest

For our purposes, I want you to recognize three broad patterns.

1. Opposed interests

What benefits me hurts you.

A simple example is a zero-sum game, where one player's gain corresponds directly to another player's loss. In formal game theory, zero-sum games represent a particularly strong form of conflict of interest.

Chess is an obvious example.

If I win, you lose.

There isn't much room for us to create additional value together.

But pure zero-sum situations are not nearly as common in everyday relationships as we sometimes imagine.

2. Aligned interests

We both benefit from the same outcome.

Imagine two employees trying to get a difficult project completed.

If the project succeeds, both benefit.

Their primary problem may not be conflict at all.

It may be coordination.

That's an important distinction because the solution isn't necessarily persuasion or bargaining.

It may simply be figuring out:

How do we get our actions aligned?

We'll explore that much more deeply later in the book.

3. Mixed interests

This is where much of human life actually happens.

We both want the relationship to succeed.

But we disagree about how.

We both want the company to grow.

But we disagree about what to do with the profits.

We both want the deal completed.

But we disagree about price.

We both want the family to be happy.

But we have different ideas about what that requires.

Mixed-motive games are fascinating because they create opportunities for both:

conflict and cooperation.

And that means the strategic thinker isn't constantly looking for a winner.

They're looking for the **structure of the interests.**

The Pie Doesn't Have to Be Fixed

Here's an important consequence.

If you assume every interaction is zero-sum, you start looking at everything as:

Who gets what?

But sometimes the better question is:

Can we create more value before deciding how to divide it?

Imagine you're negotiating a business deal.

You disagree over price.

But you discover that the seller cares enormously about getting paid quickly, while you care more about paying over time.

Now something interesting happens.

You may be able to restructure the deal.

The seller gets:

faster payment.

You get:

better terms.

The total value of the arrangement may increase even though you still have conflicting interests over the final allocation.

This is one of the reasons mixed-motive interactions are so interesting.

Conflict over one dimension doesn't mean conflict over everything.

The Hidden Payoff

Sometimes the payoff you're seeing isn't the payoff that actually drives the decision.

Imagine someone refuses a perfectly reasonable offer.

You think:

"They're being irrational."

But suppose accepting the offer would make them look weak in front of their colleagues.

Now the financial outcome isn't the whole payoff.

There is also:

status.

Or imagine someone refuses to apologize.

You think:

"They don't care about the relationship."

Maybe they do.

But they believe apologizing means admitting they were entirely at fault.

Their payoff isn't simply:

relationship versus no relationship.

It's something more complicated:

relationship + dignity + self-image.

The important lesson isn't:

Everyone has a secret hidden motive.

That would be an unhealthy way to apply this model.

The lesson is:

The visible payoff may not be the complete payoff.

This Is Where Empathy and Game Theory Meet

There's an interesting connection here.

People often think empathy means:

"Imagine how I would feel if I were them."

Strategic thinking asks for something slightly different:

"Try to understand what the situation looks like from their position."

Those aren't necessarily the same thing.

If you were them, but still had **your** preferences, your interpretation might be completely wrong.

Instead, try to reconstruct their game.

What do they value?

What are they afraid of losing?

What alternatives do they have?

What constraints are they facing?

What outcome would make them feel that they won?

What outcome would make them feel that they lost?

You aren't necessarily agreeing with them.

You're trying to understand the incentive structure they're operating within.

The Prisoner's Dilemma

One of the most famous examples in game theory demonstrates why understanding payoffs matters so much.

Two people are given a choice:

cooperate

or

defect.

The important part isn't the labels.

It's the payoff structure.

If both cooperate, both do reasonably well.

If one defects while the other cooperates, the defector does particularly well while the cooperator suffers.

If both defect, both do worse than they could have by cooperating.

That creates a strategic dilemma.

And notice something important:

The players don't have to be bad people.

They don't even have to dislike each other.

The **structure of the payoffs** can create pressure toward a particular behavior.

That's one of the most profound ideas in game theory.

Sometimes the behavior we observe is being produced by the structure of the game rather than the character of the players.

That distinction will become increasingly important later.

The Ultimatum Game

Here's another famous experiment.

One person receives a sum of money and must propose how to divide it with another person.

The second person can either:

accept the proposal,

or:

reject it.

If they reject it, both receive nothing.

From a simplistic economic model, you might predict that the second person should accept almost any positive amount.

If I offer you \$1 out of \$100, you should supposedly prefer \$1 to \$0.

But that's not what people consistently do.

Across many experiments, people frequently reject offers they perceive as unfair, even when rejection leaves them with nothing. The Ultimatum Game became one of the foundational experimental demonstrations that human behavior cannot always be explained by narrow monetary self-interest alone. Güth, Schmittberger, and Schwarze's original 1982 experiment is the classic starting point, and the sub-

sequent literature has replicated and extended the finding across many settings.

Why?

Because the payoff may include something beyond money.

People can care about:

 fairness,
 reciprocity,
 punishment,
 dignity,
 or how another person treats them.

That doesn't mean the mathematical model was useless.

Quite the opposite.

The experiment teaches us something about the **actual payoff structure**.

If people consistently sacrifice money to punish perceived unfairness, then fairness is part of the game.

This Is Why "Rational" Doesn't Mean "Selfish"

This distinction matters enormously.

In everyday language, we often say:

 "That person isn't being rational."

What we frequently mean is:

"That person isn't maximizing the thing I assumed they cared about."

Those aren't the same thing.

Suppose you turn down a lucrative job because it would require you to miss your children's lives.

From a purely financial perspective, the decision may look irrational.

But if your priorities include:

family,

then the decision may make perfect sense.

The problem wasn't necessarily that you made an irrational decision.

The problem may have been that **the observer misunderstood your payoff function.**

This is one of the most useful ways game theory can change how we interpret people.

But Don't Use This as an Excuse for Everything

There's an important boundary.

Understanding someone's incentives does not mean assuming their behavior is justified.

Someone can have a perfectly understandable incentive and still make a bad decision.

Someone can value something unhealthy.

Someone can pursue short-term gratification at enormous long-term cost.

Someone can knowingly harm another person because they value something else more.

Game theory helps us understand:

why a behavior may make sense from someone's perspective.

It doesn't tell us:

whether that behavior is good, moral, or wise.

That distinction will remain important throughout this book.

The Payoff Can Change

Here's where this model becomes even more powerful.

The payoff structure isn't necessarily fixed.

Imagine you have two employees competing for a promotion.

The company rewards individual performance.

Their incentives are competitive.

Now management changes the compensation system so that both employees receive a bonus if the entire team hits its target.

The people didn't necessarily change.

The game did.

Their incentives changed.

This is one of the bridges to our final chapter, where we'll examine how changing rules can change behavior.

But for now, simply notice:

Before trying to change someone's behavior, understand what the current environment rewards.

What Are They Trying to Avoid?

There's another question that can reveal hidden payoffs:

What are they trying to avoid?

People aren't motivated only by what they can gain.

They may also be strongly motivated by what they could lose.

They may want:

- to avoid embarrassment,
- to avoid uncertainty,
- to avoid conflict,
- to avoid losing status,
- to avoid financial risk,
- to avoid looking incompetent,
- to avoid being controlled.

This is particularly important because two people can value the same positive outcome but assign very different weight to the risks involved in achieving it.

For one person:

“Let’s try it.”

For another:

“What if it fails?”

Same potential upside.

Different perceived downside.

Different behavior.

Losses Can Matter Differently From Gains

Behavioral economics adds an important qualification here.

Daniel Kahneman and Amos Tversky’s prospect theory demonstrated that people’s choices often depend not only on objective outcomes but on how gains and losses are framed relative to a reference point. Losses can carry great-

er psychological weight than equivalent gains in many contexts.¹¹

This doesn't mean:

“People always hate losses more than they like gains.”

Human behavior is more complicated than that.

But it gives us another useful lens:

The same objective outcome can produce different decisions depending on what someone perceives themselves as gaining or losing.

That will become especially important when we discuss framing later.

Perception Is Part of the Game

This connects to something we've discussed before.

There is the objective situation.

And then there is:

how each player perceives the situation.

Suppose two people are offered the exact same deal.

One sees:

“I'm getting \$10,000.”

The other sees:

“I'm losing \$5,000 compared with what I expected.”

Same objective outcome.

Different perceived payoff.

11. nobelprize.org, <https://www.nobelprize.org/uploads/2018/06/advanced-economicsciences2002.pdf>

And because decisions are based on perceived outcomes, behavior can differ.

This doesn't mean reality doesn't matter.

It means:

Your perception of the game influences how you play it.

And the other person's perception influences how **they** play it.

The Same Event Can Have Different Payoffs

Imagine a manager gives an employee additional responsibility.

Employee A thinks:

"They trust me."

Employee B thinks:

"They're dumping more work on me."

Same action.

Different interpretation.

Different perceived payoff.

Different likely response.

This is why effective strategic thinking requires us to distinguish:

what happened

from

what each person thinks happened.

That distinction will become much more important when we get to information, beliefs, and framing.

Ask: “What Would Make This a Good Outcome for Them?”

Here’s a practical exercise.

Take a difficult interaction you’ve had recently.

Instead of asking:

“Why are they being so difficult?”

Ask:

“What outcome would make this person feel that they got what they wanted?”

Then ask:

“What outcome would make them feel that they lost?”

You might discover something surprising.

Maybe they don’t actually need the thing you thought they needed.

Maybe they care about something you never considered.

Maybe the conflict isn’t over the stated issue at all.

Or maybe you discover that your interests really are fundamentally opposed.

That’s useful too.

Because now you know what kind of game you’re in.

Don't Confuse Stated Preferences With Revealed Preferences

There's another subtle lesson here.

People can tell you what they want.

But their behavior can provide additional information.

Someone might say:

“Price is my only concern.”

Then reject a cheaper option because it takes longer to deliver.

Now you've learned something.

Perhaps price wasn't actually their only concern.

This doesn't mean people are lying.

People often don't fully understand their own preferences until they face an actual choice.

But strategically, behavior gives us information about what matters.

That's why game theory often analyzes observed choices rather than simply accepting statements at face value.

The Payoff Isn't Always Immediate

One of the easiest mistakes is looking only at what someone gains **right now**.

Imagine someone refuses a small financial benefit because accepting it would damage their reputation.

Or someone takes a short-term loss to preserve a long-term relationship.

Or someone accepts an uncomfortable conversation because they believe it will improve the relationship later.

The immediate payoff isn't the whole payoff.

This will become central in Chapter 9 when we discuss repeated interactions.

For now, remember:

Don't evaluate someone's choice only by what they gain or lose in the current moment.

A Better Way to Look at Disagreement

The next time someone disagrees with you, don't immediately ask:

"How do I convince them?"

First ask:

"What payoff are they protecting?"

Maybe they're protecting:

- money,
- time,
- status,
- control,
- security,
- identity,
- fairness,
- reputation,
- autonomy,
- or a relationship.

Once you understand that, you may discover that the disagreement isn't as large as it initially appeared.

Or you may discover that it is.

Either way, you've learned something important.

Sometimes the Payoffs Really Are Incompatible

We shouldn't turn this into a "there's always a win-win" philosophy.

Sometimes there isn't.

Two people can genuinely want opposite things.

A company may want to pay less.

An employee may want to earn more.

A buyer may want the lowest price.

A seller may want the highest price.

Two political groups may want incompatible policies.

Two competitors may both want the same customer.

There are real conflicts.

The point of understanding payoffs isn't to pretend they don't exist.

It's to **identify them accurately**.

Because if you misdiagnose a conflict as cooperation, you'll make bad decisions.

And if you misdiagnose an opportunity for cooperation as pure conflict, you'll leave value on the table.

The Payoff Map

At this point, you can add a second layer to the game map from Chapter 1.

When entering an important interaction, ask:

What do I want?

What outcomes matter to me?

What do they want?

What outcomes appear to matter to them?

Where do our interests align?

What can we both gain?

Where do they conflict?

What does one person gain at the other's expense?

What are we each trying to avoid?

What losses or risks matter?

What might I be missing?

Could status, identity, fairness, autonomy, or relationships matter here?

And perhaps the most important question:

Am I assuming their payoff is the same as mine?

The Mental Shift

Before this chapter, you might look at someone's behavior and think:

"Why would they do that?"

Now I want you to start asking:

“What payoff would make that behavior make sense?”

That’s a very different question.

It doesn’t mean you’ll agree with them.

It doesn’t mean you’ll excuse them.

It doesn’t mean you’ll always be able to figure them out.

It simply means you’re no longer assuming that your preferences are the universal standard for rational behavior.

You’re trying to see the game from their side.

And This Changes How You Communicate

This is where the connection to communication becomes especially powerful.

If I know what you value, I can understand:

- what you’re likely to respond to,
- what you’re likely to resist,
- what you’re likely to perceive as a threat,
- what you’re likely to see as an opportunity.

But notice what we’re **not** doing.

We’re not saying:

“Find their hot button and push it.”

That’s tactics.

We’re doing something more fundamental:

Understand the incentive structure before choosing your move.

Sometimes that understanding will help you persuade.

Sometimes it will help you negotiate.

Sometimes it will help you cooperate.

Sometimes it will tell you that agreement isn't possible.

And sometimes it will tell you that you're solving the wrong problem entirely.

That's why this is a mental model rather than a technique.

The Question to Carry With You

When you encounter resistance, disagreement, or unexpected behavior, pause before judging.

Ask:

What payoff would make their behavior make sense?

Then go one step further:

What are they trying to gain?

What are they trying to avoid?

Where are our interests aligned?

Where are they genuinely opposed?

What am I assuming they value that they may not?

Because once you understand the payoffs...

you can begin to understand the choices.

And once you understand the choices, you can begin to anticipate the response.

Which brings us to the next mental model:

Don't just think about your move. Think about theirs.

Because in a game, your move is only half the move.

CHAPTER 3

Think in Responses

In the first chapter, we learned to **see the game**.

In the second, we learned to **see the payoffs**.

Now we get to one of the ideas that, for me, is at the heart of this entire book:

Your move is only half of the move.

The other half is what the other person does in response.

And yet this is something we routinely forget.

We think about what we're going to say.

We think about what we're going to do.

We think about what we want to happen.

But if another person is involved, the outcome isn't determined by our decision alone.

They get a move too.

And their move depends, at least partly, on ours.

That's the fundamental idea behind a **best response** in game theory: given what the other player is doing, what ac-

tion gives you the best outcome? Nash equilibrium is built from this idea—an outcome in which each player’s strategy is a best response to the others, so no one benefits by changing their strategy alone.¹²

You don’t need to calculate best-response functions to use the concept.

You simply need to develop the habit of asking:

If I do this, what are they likely to do next?

And then:

If they do that, what will I do?

That’s where strategic thinking really begins.

Stop Thinking in Straight Lines

Most of us naturally think about communication as a straight line:

I say something → they hear it → outcome.

But real interaction looks more like this:

I act → they respond → I respond to their response → they respond to mine.

And sometimes it keeps going.

That means your decision shouldn’t be based solely on:

“What do I want to do?”

It should also be based on:

“What will this cause the other person to do?”

This sounds obvious.

12. MIT OpenCourseWare, https://ocw.mit.edu/courses/14-12-economic-applications-of-game-theory-fall-2025/resources/mit14_12f25_lec05_1080p_mp4/

But it's remarkably easy to forget.

The Raise, Revisited

Let's return to our employee asking for a raise.

Suppose you decide:

"I'm going to tell my boss that I have another job offer."

At first glance, that may seem like a strong move.

It communicates leverage.

But game theory asks us to go one step further.

Then what?

Your boss might:

1. offer you more money,
2. ask to see the offer,
3. call your bluff,
4. tell you they can't compete,
5. become concerned that you're leaving,
6. decide you're no longer committed,
7. start preparing for your departure.

The value of your statement depends partly on which response you expect.

And your expectation should depend on what you know about the other player.

If your boss desperately wants to retain you, the move may work very differently than if they've already been considering replacing you.

Same sentence.

Different game.

The Strategic Question

Before making an important move, ask:

What does this make the other person want to do?

That's subtly different from:

"What do I want them to do?"

You may want them to cooperate.

But does your action actually make cooperation more attractive?

You may want them to trust you.

But does your behavior give them a reason to trust you?

You may want them to compromise.

But does your proposal make compromise better than their alternatives?

You may want them to stop pushing.

But does your response make pushing more or less attractive?

This is the difference between **intention** and **incentive**.

You can intend to produce a particular response without actually creating the conditions that make that response rational.

Best Responses

This is where a formal game-theory concept becomes incredibly useful.

A **best response** is, roughly:

The action that gives a player the best outcome given what the other player is doing.

MIT's game-theory curriculum introduces dominance, rationalizability, best responses, and Nash equilibrium as foundational concepts.¹³

Imagine two people choosing restaurants.

If you know the other person is definitely going to the Italian restaurant, perhaps your best response is to go there too.

If they're definitely going to the steakhouse, perhaps your best response changes.

Your best move isn't determined in isolation.

It depends on **their move**.

That's the key.

Sometimes the Other Person's Best Response Is the Problem

Here's where this becomes especially useful in communication.

Suppose you manage a sales team and tell everyone:

"I want you to collaborate more."

But you've structured compensation so that salespeople get rewarded individually for outperforming their colleagues.

You may sincerely want collaboration.

13. MIT OpenCourseWare, <https://ocw.mit.edu/courses/14-12-economic-applications-of-game-theory-fall-2025/download/>

But you've created a situation in which the individual's best response may be to protect their own leads.

The problem isn't necessarily that your employees didn't hear you.

They may have heard you perfectly.

The problem is:

Their incentives make another response more attractive.

This is one of the most important themes in the book.

People don't respond only to your words.

They respond to the game those words are entering.

Predict the Response Before You Make the Move

This is the habit I want to develop.

Before you act, pause.

Ask:

If I were them, what would I do next?

Not:

"What should they do?"

Not:

"What would a reasonable person do?"

Not:

"What would I do?"

Ask:

Given their goals, information, options, and constraints, what is the most likely response?

That's a much harder question.

And it's much more useful.

But Don't Assume Everyone Is Perfectly Rational

There's an important qualification.

Game theory often begins with models in which players are assumed to choose strategies that best serve their preferences.

Real human beings don't always do that.

We make mistakes.

We misunderstand.

We get emotional.

We have incomplete information.

We miscalculate.

We sometimes act against our own interests.

And sometimes we simply don't understand the game very well.

That doesn't make game theory useless.

It means we need to distinguish between:

**what someone would choose if they were optimizing
based on their information**

and

what they are actually likely to choose.

Those aren't always the same.

In fact, understanding that difference is one of the reasons experimental game theory exists.

Rationalizability: A Useful Middle Ground

This gives us another important concept.

Before trying to predict exactly what someone will do, you can sometimes eliminate actions that simply don't make sense given reasonable assumptions about what they want and what you might do.

Game theory calls this process **iterated elimination of dominated strategies** and, more broadly, **rationalizability**. MIT treats these as foundational concepts alongside Nash equilibrium.¹⁴

You don't need the formal mathematics.

Here's the practical translation:

Don't waste energy predicting moves that make no strategic sense.

Suppose you have two choices.

Choice A produces a worse outcome for you than Choice B **no matter what the other person does.**

Then you don't need to predict the other person.

Choice A can be eliminated.

That's a powerful mental shortcut.

Before asking:

"What are they going to do?"

sometimes ask:

"Is there anything I already know I shouldn't do regardless of what they do?"

14. MIT OpenCourseWare, <https://ocw.mit.edu/courses/14-12-economic-applications-of-game-theory-fall-2025/download/>

If yes, eliminate it.

Then continue.

Think One Move Ahead

Let's say someone sends you an angry email.

Your first impulse is to respond immediately.

But imagine instead that you pause and ask:

If I respond angrily, what will they probably do?

Maybe they'll escalate.

Then:

If they escalate, what happens next?

Maybe now both of you are defending yourselves rather than solving the original problem.

You haven't just sent an angry email.

You've potentially created a **sequence**.

This is the essence of sequential reasoning.

And formal game theory has a powerful tool for analyzing sequential games:

backward induction.

Think Backward From the End

Backward induction sounds complicated.

The underlying idea is simple:

Instead of only asking what happens next, ask where the sequence eventually leads.

In a finite sequential game, you can start with the final decision and reason backward to the beginning. MIT's game-theory courses explicitly teach backward induction as a core method for solving sequential games.¹⁵

Imagine you're negotiating with someone.

You make an offer.

They can accept or reject.

If they reject, you get one final opportunity to respond.

Instead of thinking only:

“What should I offer?”

you can reason:

If they reject, what will I do?

And if you know what you'll do then:

What will they do knowing that?

And if you know that:

What should I do now?

You're reasoning backward.

The Final Conversation

Imagine two people are having a disagreement.

They both know:

This is our last conversation.

There will be no tomorrow.

No future relationship.

15. MIT OpenCourseWare, https://ocw.mit.edu/courses/6-254-game-theory-with-engineering-applications-spring-2010/688240e73afebf4a6293fef0e529398_MIT6_254S10_lec12.pdf

No reputational consequences.

No opportunity to repair the interaction later.

That changes the game.

Now imagine they expect to work together for the next twenty years.

Same disagreement.

Same people.

Same immediate issue.

Different strategic environment.

We'll explore the importance of repeated interaction much more deeply in Chapter 9.

For now, notice the principle:

The best response to a move depends partly on what comes after it.

Threats Reveal Why This Matters

Suppose someone says:

“If you do that, I’m going to walk away.”

Should you believe them?

You can’t determine that merely by how forcefully they say it.

You need to ask:

If I actually do that, will walking away be their best response?

This is one of the foundations of strategic credibility, which we’ll explore later.

But the reasoning starts here.

A threat isn't strategically meaningful simply because someone made it.

You have to examine the incentives at the point where the threat would actually be carried out.

MIT's treatment of dynamic games uses backward reasoning partly to distinguish credible future actions from threats that wouldn't make sense when the decision point actually arrives.¹⁶

This is an extraordinarily useful way of thinking about everyday interactions.

"What Would I Do If I Were Them?"

There's a trap here.

People often say:

"If I were them, I'd never do that."

But that's usually not the right question.

You don't want to put **your preferences** into **their position**.

You want to put **their preferences** into **their position**.

Remember the previous chapter.

Their payoff may be different from yours.

So the better question is:

"Given what they want, what they know, and what they can do, what would make sense for them?"

16. MIT OpenCourseWare, https://live.ocw.mit.edu/courses/14.126-game-theory-spring-2024/mit14_126_s24_lecture_2_refinements.pdf

That is much more powerful.

Perspective-Taking Becomes Strategic

This is where game theory and ordinary empathy begin to overlap.

Suppose your customer refuses your offer.

Instead of thinking:

“They’re being unreasonable.”

Try:

“What response would make sense if I had their objectives and constraints?”

Maybe:

- they don’t trust the delivery date,
- they have another supplier,
- they’re worried about appearing foolish,
- their boss has to approve the purchase,
- they don’t understand the value,
- they are testing how flexible you are.

You don’t know which one is true.

But now you’re looking for explanations in terms of the game rather than immediately assigning a character judgment.

Don’t Just Predict Their Response

There’s another level.

Once you predict their response, ask:

How will I respond to that?

And then:

How will they respond to my response?

This is where interactions become recursive.

You are not simply choosing an action.

You are choosing an action **while anticipating another person's anticipation of your action.**

That sounds complicated.

But you already do it.

Imagine you're negotiating over price.

You think:

"If I offer \$90, they'll probably counter at \$100."

Then:

"They probably expect me to counter again."

So you might start at \$85.

But then you think:

"They know I'm going to start low."

And now you're thinking several moves ahead.

That's strategic reasoning.

But More Thinking Isn't Always Better

There is a danger here.

Once people discover this concept, they can become obsessed with trying to predict every possible move.

That's not the goal.

You don't need to build a 47-step decision tree every time your spouse asks:

"What do you want for dinner?"

The goal is to identify **strategically relevant responses**.

Ask:

What response could materially change the outcome?

Ignore the rest.

Strategic thinking is not about thinking more.

It's about thinking about the **right dependencies**.

Don't Overpredict

This is equally important.

You can become so convinced that you understand another person's strategy that you stop listening.

You think:

"I know exactly what they're going to do."

And then they surprise you.

That's dangerous.

Your model of the other player is a **hypothesis**.

Not a fact.

This is where our larger meta-principle comes in:

Model → act → observe → update.

You make your best prediction.

They respond.

You compare their actual behavior with your prediction.

If you're wrong, don't immediately conclude:

"They're irrational."

Ask:

"What did I misunderstand about the game?"

Maybe their payoff was different.

Maybe they had information you didn't know about.

Maybe you misunderstood their options.

Maybe they were simply unpredictable.

The response is data.

Your Prediction Changes Their Response

Here's an especially interesting wrinkle.

People aren't passive objects waiting for you to predict them.

They can predict **you** too.

Suppose you always respond to pressure by giving in.

Eventually someone may learn:

Pressure works on you.

Your predictable response becomes part of their strategy.

And now your own behavior is affecting the behavior of the other player.

This is where our next chapter—**Manage Predictability**—will take us.

For now, the lesson is:

Don't just predict the other player. Remember that they may be predicting you.

The Beauty of a Simple Example

Imagine two people are walking toward each other in a hallway.

You move slightly to your right.

They move slightly to their right.

You both move again.

Now you're both awkwardly bouncing from side to side.

What happened?

Neither person necessarily wanted the outcome.

Both were responding to the other's movement.

This is a tiny example of strategic interaction.

Your action changed their best response.

Their response changed yours.

And suddenly a simple interaction became a feedback loop.

That's the game.

Feedback Loops

Many human interactions are essentially loops:

I act.

↓

You respond.

↓

I interpret your response.

↓

I act again.

↓

You interpret my action.

↓

You respond again.

The important word here is **interpret**.

Because you're not responding merely to what the other person objectively did.

You're responding to what you believe their action means.

That brings us toward Chapter 6, where we'll examine information and beliefs.

For now, we just need to recognize the loop.

The Escalation Trap

Here's a more serious example.

Someone speaks aggressively to you.

You interpret it as disrespect.

You respond aggressively.

They interpret your response as confirmation that you are hostile.

They escalate.

You escalate.

Soon, both of you are behaving in ways neither of you intended at the beginning.

Each person is responding rationally—or at least understandably—to the previous move.

Yet the overall interaction is getting worse.

This is one of the most important things strategic thinking can reveal:

A bad outcome can emerge from a sequence of locally understandable responses.

Nobody has to wake up and decide:

“Let’s make this interaction terrible.”

The pattern can create itself.

The Equilibrium Is Coming

We’re not going to fully explain equilibrium until Chapter 5.

But you can already see where we’re heading.

Suppose:

My behavior causes your response.

And:

Your response causes my next behavior.

Eventually we may reach a pattern where neither of us has much reason to change what we’re doing given what the other is doing.

That’s an equilibrium.

And sometimes those equilibria are terrible.

Sometimes they’re excellent.

Sometimes they're stable but inefficient.

Sometimes they're difficult to escape.

That's why simply asking:

“What's the best thing for me to do?”

isn't enough.

You also need to ask:

“What pattern will my action create?”

A Practical Exercise

Think of a recent interaction that didn't go the way you wanted.

Don't analyze whether you were right or wrong yet.

Instead, write down the sequence:

What did I do?

↓

What did they do?

↓

What did I do next?

↓

What did they do next?

Now ask:

At which point did the interaction change direction?

Then ask:

What could I have predicted about their response before making my move?

This is one of the simplest ways to begin developing strategic awareness.

You're training yourself to stop viewing the interaction as a series of isolated statements.

You're learning to see it as a **sequence of interdependent decisions**.

The Mental Shift

Before this chapter, you might have thought:

"What should I do?"

Now I want you to start asking:

"What will they do if I do that?"

And then:

"What will I do if they do that?"

And, when necessary:

"Where does that sequence ultimately lead?"

That's a completely different way of thinking.

It's also one of the biggest differences between ordinary communication advice and the framework we're building here.

Communication advice often focuses on the quality of the move.

Game theory forces you to consider the **response to the move**.

And that means the quality of your move can't be judged independently of the game around it.

The Question to Carry With You

The next time you're about to say something important, don't stop at:

"What should I say?"

Ask:

"What is this likely to make them do?"

Then ask:

"If they do that, what will I do?"

And if the stakes are high:

"Where does that sequence lead?"

Because once another person is involved, your move doesn't end when you stop talking.

Your move becomes their situation.

And their response becomes yours.

That's strategic interaction.

And once you start seeing communication that way, you can begin to notice something else:

Sometimes the most dangerous thing you can do is become completely predictable.

CHAPTER 4

Manage Predictability

In the last chapter, we learned to think in responses.

Before you make a move, ask:

What is this likely to make the other person do?

But there's a problem with that way of thinking.

What happens when the other person can predict **you**?

If someone knows exactly how you're going to respond, they can begin planning around you.

Your behavior becomes part of their strategy.

And this leads to another important game-theoretic idea:

Predictability can create strategic vulnerability.

But there's an important distinction I want to make right at the beginning.

This chapter is **not** about being deceptive.

It isn't about lying.

It isn't about manipulating people.

And it certainly isn't about becoming erratic just so nobody knows what you're going to do.

The principle is much more subtle:

Don't allow your patterns to become so predictable that other people can consistently use them against you.

Sometimes that means being unpredictable.

Sometimes it means being extremely predictable.

The skill is knowing **when each is strategically valuable**.

The Habit They Learn

Imagine you have a coworker who knows that every time they ask you for a favor, you say yes.

At first, that's simply your personality.

You're helpful.

You don't mind doing the extra work.

But eventually, your coworker learns something:

If I ask Paul, he'll probably say yes.

Now your predictability has become part of their strategy.

They don't have to wonder:

"Should I ask him?"

They already know the answer.

And that changes the interaction.

The same thing happens in negotiation.

If the other person knows that you always make a concession after their first counteroffer, they don't have to wonder whether they should push.

They know:

Push, and he'll probably move.

You've taught them the game.

The Value of Uncertainty

This is one of the reasons game theory has long been interested in situations where players cannot predict exactly what another player will do.

In formal game theory, this can lead to **mixed strategies**—situations in which a player randomizes among actions rather than choosing the same action every time.

The classic example is rock-paper-scissors.

If you always throw rock, your opponent eventually learns to throw paper.

If you always throw paper, they learn to throw scissors.

If they know your pattern, your strategy becomes exploitable.

The solution isn't to find one action that is always best.

It's to prevent the opponent from knowing exactly which action you're going to choose.

The formal mathematics behind mixed strategies can be complicated.

The mental model isn't:

"Randomize everything."

It's:

If someone can reliably predict your action, they may be able to choose a response that takes advantage of that predictability.

Rock-Paper-Scissors

Consider the simplest possible example.

You and I are playing rock-paper-scissors.

If you know that I choose rock 70 percent of the time, you have useful information.

You don't need to know exactly what I'll choose.

You only need to know enough about my pattern to adjust your own behavior.

My predictable behavior has created an advantage for you.

That's what strategic predictability means.

And this is why optimal play in certain competitive games involves randomization.

If your opponent can observe your pattern and exploit it, the pattern itself becomes a weakness.

But Life Isn't Rock-Paper-Scissors

This is where we need to be careful.

It would be a mistake to take a formal concept like mixed strategies and turn it into:

“Always keep people guessing.”

That's not what game theory says.

Randomization is useful under particular strategic conditions.

And in many human relationships, excessive unpredictability is actually destructive.

Imagine if your spouse never knew whether you were going to keep your promises.

Or your employees never knew whether you'd enforce a rule.

Or your business partners never knew whether you'd honor an agreement.

That's not strategic sophistication.

That's a lack of reliability.

In relationships where trust matters, **predictability can be enormously valuable.**

So the real question isn't:

Should I be predictable or unpredictable?

It's:

Where does predictability help cooperation, and where does predictability create an opportunity for exploitation?

That's a much more useful question.

Predictability Is Not the Enemy

Think about a stoplight.

You want the other drivers to know what you're going to do.

If everyone suddenly became unpredictable at intersections, traffic would become dangerous.

Coordination requires predictability.

You want to know:

red means stop.

green means go.

the other driver will stay in their lane.

Shared expectations make cooperation possible.

This is one reason conventions and norms are so powerful.

They allow people to coordinate because everyone can predict what everyone else is likely to do.

So predictability has **two very different strategic functions**.

It can create:

trust and coordination

or:

exploitation and strategic vulnerability.

The challenge is knowing which environment you're in.

The Predictability Tax

Let's imagine a negotiation.

You always make three concessions.

First:

\$100.

Then:

\$90.

Then:

\$80.

The other person learns your pattern.

They don't need to know your true minimum.

They simply need to know:

You'll move again.

So instead of accepting your first offer, they push.

You move.

They push again.

You move.

By the third round, your pattern has effectively revealed information about your limits.

Your own behavior has told them:

There's more room.

This is one of the hidden costs of predictable behavior.

Your actions can reveal information you didn't intend to reveal.

When Your Pattern Becomes Information

This is a crucial connection to the next major section of the book.

Every action you take communicates something.

If you always answer emails immediately, people may infer that you are highly available.

If you always agree to last-minute requests, people may infer that your schedule is flexible.

If you always threaten to quit but never do, people learn that your threats aren't credible.

If you repeatedly say you'll enforce a rule and then don't, people learn something about the rule.

Your behavior creates information.

Sometimes more information than your words do.

And once people learn your pattern, they can adapt.

The Bluff That Stopped Working

Imagine a poker player who bluffs frequently.

At first, opponents may fold.

The bluff works.

But eventually they notice:

“He's bluffing a lot.”

Now they call more often.

The same bluff that worked earlier stops working.

What changed?

Not the words.

Not the cards.

The **other players' beliefs** changed.

They learned your pattern.

This is one reason strategic interaction is dynamic.

Your behavior doesn't merely affect the immediate outcome.

It teaches other players about you.

And what they learn changes how they respond in the future.

The Danger of a Reputation You Didn't Intend to Create

You may think you're simply behaving consistently.

Other people may experience that consistency as a rule.

Suppose a manager says:

"I really don't like firing people."

That's a personal preference.

But if employees learn:

"He never actually fires anyone,"

the statement has become part of the strategic environment.

Employees may begin taking risks they wouldn't otherwise take.

The manager's predictability has changed their incentives.

Now imagine the opposite manager.

Employees know:

"If you violate this rule, there will be consequences."

That predictability can actually improve cooperation.

So again:

Predictability isn't inherently good or bad.

Its strategic value depends on what it allows other players to do.

Strategic Unpredictability

This brings us to something I initially hesitated to include in our framework.

I don't anymore.

Because there is an important distinction between:

manipulative unpredictability

and

strategic unpredictability.

Suppose you're competing against another business.

Your competitor knows that you always launch new products at the same time every year.

They can prepare for it.

They know your pricing pattern.

They know your promotional schedule.

They know exactly how you respond to their moves.

There may be strategic value in changing those patterns.

That isn't inherently unethical.

Competition is not inherently unethical.

You can compete aggressively while still being honest.

You can protect information without lying.

You can surprise a competitor without deceiving a customer.

You can vary your strategy so that another player cannot simply optimize against a fixed pattern.

That's a legitimate strategic behavior.

The Difference Between Secrecy and Deception

This distinction is worth making explicit.

You don't owe every competitor complete information about your plans.

Keeping your strategy private is not necessarily deception.

If two companies are competing for a contract, one company doesn't have to announce:

"Here is exactly how we're going to respond if you undercut our price."

They can keep their strategy private.

Likewise, in a negotiation, you don't necessarily have to reveal your reservation price.

There is a difference between:

withholding information you are not obligated to disclose

and

deliberately providing false information to induce someone into a decision.

That distinction matters because I don't want the reader to come away thinking:

"Game theory teaches me to manipulate people."

It doesn't.

It teaches you to understand the strategic value of information and uncertainty.

What you do with that knowledge is an ethical question.

Don't Become Predictable in the Wrong Way

Here's a simple example.

Someone repeatedly interrupts you in meetings.

The first time, you let it go.

The second time, you let it go.

The third time, you let it go.

Now they've learned something:

Interrupting Paul has no cost.

You may not have consciously made a decision to establish that rule.

But your repeated behavior established one anyway.

This is a powerful idea:

Your patterns can create rules even when you never intended to create them.

And once other people learn the rule, they can act accordingly.

The Other Side: Don't Become Unpredictable in the Wrong Way

Now imagine you do the opposite.

One day you tolerate an interruption.

The next day you explode.

The next day you ignore it.

The next day you punish the person.

Nobody knows what will happen.

You may have made yourself unpredictable.

But you've also made yourself difficult to trust.

Your employees can't coordinate around you.

Your family can't anticipate your reactions.

Your partners can't make plans.

Unpredictability has now destroyed something valuable:

reliable expectations.

That's why strategic sophistication isn't:

"Never let people know what you're going to do."

It's:

"Understand the strategic consequences of being predictable."

There Are Times When You Want People to Know Exactly What You'll Do

This is where predictability connects directly to commitment.

Suppose you tell your employees:

"If you falsify customer records, you will be terminated."

And then you consistently enforce it.

You've created predictability.

That's good.

Why?

Because employees can now anticipate the consequences of their choices.

Your predictable response becomes part of their decision-making.

They don't have to wonder:

"What will happen if I do this?"

They know.

And that can change behavior before the violation occurs.

This is one of the reasons credible commitments matter so much.

We'll devote an entire chapter to them later.

Sometimes the Best Strategy Is to Become More Predictable

Imagine you're trying to rebuild trust with someone.

They don't know whether you'll keep your word.

What should you do?

Become unpredictable?

Probably not.

You want to become **more predictable**.

If you say you'll call, call.

If you say you'll show up, show up.

If you say you'll do something, do it.

Over time, the other person learns:

His words predict his actions.

That creates trust.

So the strategic lesson isn't:

Hide your intentions.

It's:

Make your predictability work for you.

The Strategic Value of a Pattern

A pattern can be an asset when:

- it creates trust,
- it makes cooperation easier,
- it establishes credible commitments,
- it reduces uncertainty,
- it helps others coordinate with you.

A pattern can be a liability when:

- competitors can exploit it,
- people know exactly how to pressure you,
- your threats are never followed through,
- your concessions are predictable,
- your boundaries are predictable but unenforced,
- your strategy can be easily anticipated.

Same phenomenon.

Different game.

The Mixed-Strategy Lesson

There's a deeper principle hiding underneath all of this.

Sometimes the best strategy is not:

always do X.

It may be:

do X sometimes and Y sometimes, according to a pattern the other player cannot exploit.

In formal game theory, this is what mixed strategies accomplish in games where no single pure strategy is sufficient.

But don't reduce this to random behavior.

A good mixed strategy is not:

"I'll just do something random."

It is often a carefully chosen probability distribution that makes the opponent indifferent among their responses or prevents them from exploiting your pattern.

You don't need to calculate the probabilities to appreciate the underlying idea:

Sometimes uncertainty itself has strategic value.

The Element of Surprise

This applies outside formal competition too.

A manager who always responds to every problem with the same solution can become easy to manipulate.

A negotiator who always uses the same concession sequence can be exploited.

A salesperson who always handles objections the same way can become predictable.

A competitor who always responds to a price cut with another price cut can be trapped into a predictable escalation.

Sometimes changing the pattern is valuable.

Not because randomness is inherently powerful.

Because **predictability gives the other player information they can use.**

But Surprise Isn't Always Better

There's a tendency in strategic thinking to romanticize surprise.

That's a mistake.

Surprise is useful only when the uncertainty it creates is valuable.

If your employee doesn't know whether you'll approve their vacation request, that's not necessarily strategic.

If your spouse doesn't know whether you'll keep your commitment, that's not strategic.

If your customer doesn't know whether your product will actually work, that's not strategic.

In those situations, predictability creates value.

So ask:

Who benefits from knowing what I'm going to do?

If the answer is:

“Both of us,”

predictability may be beneficial.

If the answer is:

“They can use that information to consistently take advantage of me,”

then perhaps some uncertainty is valuable.

A Deeper Example: Negotiation

Imagine you are negotiating with someone who asks:

“What's the lowest price you'll accept?”

You give them the number.

Now they know something extremely valuable.

They know your boundary.

You may think you've simply answered a question.

Strategically, you've changed the information structure.

You've told them:

how far they need to push.

This doesn't mean you should lie.

You might legitimately say:

"I'm not prepared to disclose my bottom line."

That's not deception.

It's protecting strategically valuable information.

And now the other player has to make a decision without knowing exactly where your boundary lies.

What Happens When Everyone Is Predictable?

Here's an interesting thought experiment.

Imagine a negotiation in which both parties know everything about each other:

- their preferences,
- their alternatives,
- their limits,
- their strategies,
- their likely responses.

There is very little uncertainty.

The interaction becomes much easier to analyze.
 But it also becomes much easier to optimize against.
 In many real-world interactions, uncertainty is unavoidable.
 We don't know exactly what others want.
 We don't know exactly what they'll do.
 We don't know what information they have.
 We don't know how they'll interpret our behavior.
 That's not merely a nuisance.
 It is part of the strategic environment.
 And sometimes the goal isn't to eliminate uncertainty.
 It's to manage it.

Your Behavior Is a Signal

This brings us to an important connection with the next chapters.

Suppose you consistently follow through on commitments.

People learn:

“His promises mean something.”

Your past behavior becomes information.

Suppose you repeatedly make threats you don't carry out.

People learn:

“His threats don't mean much.”

Again, your behavior becomes information.

In this sense, **reputation is accumulated evidence about your future behavior.**

We'll explore that much more deeply later.

For now:

Every repeated pattern teaches the other player something about you.

The Feedback Loop

Remember the sequence from the previous chapter?

I act → you respond → I respond → you respond.

Now add learning:

I act → you respond → I observe → you observe → we both update our expectations.

The game is no longer simply:

action → response.

It's:

action → response → learning → adaptation.

And that means your behavior today can change the game tomorrow.

This is one of the reasons repeated interaction is so powerful.

The Prisoner's Dilemma Again

Remember the Prisoner's Dilemma from the previous chapter?

Two players may have an incentive to defect even though mutual cooperation would be better.

Now imagine they play the game repeatedly.

Suddenly, today's decision affects tomorrow's.

If I cooperate today, perhaps you are more likely to cooperate tomorrow.

If I defect, perhaps you punish me next time.

Now predictability becomes complicated.

If you know that I will always retaliate against defection, my response is predictable.

But that predictability may actually **support cooperation**.

This is an important paradox:

The same predictability that makes you exploitable in one game can make you trustworthy in another.

That's why we can't turn any strategic concept into a universal rule.

We have to understand the structure of the game first.

The Goal Isn't to Be Unpredictable

If there is one thing I want you to take away from this chapter, it's this:

Don't confuse unpredictability with strategic intelligence.

The real skill is **managing predictability**.

Sometimes:

Be predictable.

When trust, coordination, credibility, or cooperation are valuable.

Sometimes:

Protect your predictability.

When revealing your pattern gives another player leverage.

Sometimes:

Vary your behavior.

When another player could otherwise exploit a fixed pattern.

And sometimes:

Be deliberately consistent.

When you want another player to understand that a particular consequence is real.

The question is always:

What does my predictability enable the other player to do?

The Mental Shift

Before this chapter, you learned to ask:

“What will they do if I do this?”

Now add another question:

“What are they learning about me from what I repeatedly do?”

That’s the shift.

You’re no longer thinking only about the immediate response.

You’re thinking about the **information your behavior creates**.

And that information can change the future game.

The Question to Carry With You

The next time you notice a repeated pattern in an interaction, ask:

Have I taught this person how to respond to me?

Then ask:

Is that helping me?

Is it helping them?

Is it helping both of us?

Could they exploit this pattern?

Would greater consistency create trust—or greater flexibility create strategic room?

And perhaps the most important question:

What have my past actions taught this person to expect from me?

Because your behavior doesn't disappear after the moment is over.

It becomes information.

It becomes expectation.

And eventually, it can become part of the game itself.

Which brings us to an even deeper question:

If people are making decisions based on what they believe about you, how do you know what they believe—and how can you know whether what they're telling you is actually credible?

That's where we're going next.

Chapter 5 will take us to equilibrium: the patterns of behavior that can become stable even when nobody is particularly happy with them.

CHAPTER 5

See the Equilibrium

So far, we've built three pieces of the framework.

First:

See the Game.

Understand the players, their options, constraints, information, timing, and alternatives.

Then:

See the Payoffs.

Understand what each player actually values.

And then:

Think in Responses.

Understand that your move changes the situation the other person is responding to.

Now we're ready for one of the most important concepts in all of game theory:

Equilibrium.

And this is where game theory becomes especially interesting for understanding human behavior.

Because sometimes people aren't doing what they want.

They're doing what makes sense **given what everyone else is doing**.

When Nobody Wants to Change

Imagine two people who have been arguing for months.

Neither likes the situation.

Both would prefer things to be better.

But neither changes their behavior.

Why?

Because each person's behavior is partly a response to the other person's behavior.

You stop calling because:

"They never call me."

They don't call because:

"I always have to be the one who reaches out."

You become defensive because:

"They're always criticizing me."

They criticize you because:

"They never listen unless I push."

And around and around you go.

Nobody necessarily sat down and decided:

"Let's create a miserable relationship."

Yet you've arrived at a stable pattern.

That's the basic intuition behind equilibrium.

What Does “Equilibrium” Actually Mean?

In game theory, a Nash equilibrium is a combination of strategies in which each player’s strategy is a best response to the strategies of the others.

In simpler terms:

Given what everyone else is doing, no individual player has a reason to change their strategy on their own.¹⁷

Notice something important.

Equilibrium does **not** mean:

everyone is happy.

It doesn’t mean:

everyone got the best possible outcome.

It doesn’t even necessarily mean:

everyone thinks the situation is fair.

It means something much more specific:

The existing pattern is stable because changing unilaterally doesn’t look attractive.

That distinction is enormously important.

Stable Doesn’t Mean Good

This may be the single most important thing to understand about equilibrium.

We tend to use the word equilibrium as though it means:

balance,

17. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/archives/win2021/entries/game-ethics/>

harmony,
optimality.

In everyday language, that's often what we mean.

In game theory, equilibrium is not necessarily any of those things.

A Nash equilibrium can produce an outcome that is terrible for everyone involved.

The classic example is the **Prisoner's Dilemma**.

Each player has an incentive to defect given what the other player is doing, so mutual defection can be the Nash equilibrium even though both players would be better off if they could somehow sustain mutual cooperation.¹⁸

That's a profound idea.

It means:

A bad situation can persist without anyone needing to be irrational.

The Prisoner's Dilemma

Let's revisit it.

Two people are given a choice:

Cooperate

or

Defect.

If both cooperate, they receive a relatively good outcome.

18. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/archives/spr2011/entries/game-evolutionary/>

If one defects while the other cooperates, the defector does especially well.

If both defect, they both do worse than they would have under mutual cooperation.

Now imagine you're one of the players.

You think:

"If the other person cooperates, I'd be better off defecting."

And:

"If the other person defects, I'd better defect too."

So regardless of what they do, defection looks like your best response.

The other player reasons the same way.

The result:

Both defect.

And now you have an equilibrium.

But it's a lousy one.

Both players would prefer:

both cooperate.

The problem isn't that they don't understand that cooperation would be better.

The problem is that **given the structure of the game, each player's individual best response pushes them toward defection.**

That's one of the deepest lessons game theory has to offer.

The Equilibrium Trap

This happens constantly in real life.

Consider two competing companies.

Company A thinks:

“If we don’t lower our prices, they’ll undercut us.”

Company B thinks:

“If we don’t lower our prices, they’ll undercut us.”

So both lower their prices.

Now both make less money.

Neither company likes the situation.

But neither wants to be the one to raise prices first.

That’s an equilibrium-like situation.

Or consider two coworkers.

Each believes:

“If I stop doing this extra work, the other person will take advantage of me.”

So both continue working unsustainable hours.

Neither likes it.

But neither sees a unilateral way out.

Again:

The pattern persists because each person’s response makes sense given the other’s behavior.

The Relationship Equilibrium

Let’s bring this closer to everyday communication.

Imagine a married couple.

One partner feels unappreciated.

So they complain.

The other partner feels criticized.

So they withdraw.

The withdrawal causes more complaining.

The complaining causes more withdrawal.

Eventually they both develop a story:

“They’re the reason we can’t communicate.”

But step back.

What if neither behavior makes sense in isolation?

The complaint is a response to withdrawal.

The withdrawal is a response to complaint.

Each reinforces the other.

You’ve created a loop.

And the loop can become stable.

That doesn’t mean it’s healthy.

It means:

The pattern is self-reinforcing.

That’s the concept I want you to start noticing.

Look for the Pattern, Not Just the Move

This is a major evolution from Chapter 3.

Chapter 3 taught you to ask:

What will they do after I move?

Chapter 5 asks something deeper:

What pattern are our responses creating?

That's a different level of analysis.

One interaction might look like:

I push → you resist.

But if it happens repeatedly, it becomes:

**I push → you resist → I push harder → you resist harder
→ I push harder still.**

Now we're not looking at individual decisions.

We're looking at a **system**.

And systems can develop stable patterns.

Why Equilibrium Matters in Communication

This is where I think the concept becomes especially powerful for this book.

Suppose you keep having the same argument with someone.

Most communication advice might ask:

“What should I say differently?”

That's not necessarily wrong.

But game theory gives you another question:

“What equilibrium have we created?”

Maybe you've developed a pattern where:

- you pursue,
- they withdraw,

- you pursue harder,
- they withdraw further.

Or:

- they criticize,
- you defend,
- they interpret your defense as proof,
- you defend harder.

Or:

- you ask for a concession,
- they make one,
- you ask for another,
- they become resistant,
- you interpret resistance as lack of cooperation,
- you push harder.

The individual moves make sense when viewed locally.

But the **overall pattern** is terrible.

That's the insight.

You Can Be Rational and Still Produce a Bad Outcome

This is one of the most important ideas in the entire book.

We often explain bad outcomes by saying:

“Someone must be acting irrationally.”

Sometimes they are.

But game theory shows that bad collective outcomes can emerge from individually rational responses.

That's exactly what makes the Prisoner's Dilemma so important.

Each player's behavior can be understandable given the incentives and the other player's behavior, while the combined result is inferior to another possible outcome.¹⁹

This gives us a more sophisticated way to think about conflict.

Instead of asking only:

"Who's behaving badly?"

we can ask:

"What interaction is producing this pattern?"

That's a much more interesting question.

Not Every Equilibrium Is Unique

Here's where things get even more interesting.

Some games have:

one equilibrium.

Others have:

multiple equilibria.

That means the same underlying game can support more than one stable pattern.

Consider two people trying to coordinate.

They can both choose Option A.

Or they can both choose Option B.

19. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/archives/spr2011/entries/game-evolutionary/>

If each person's best response is to match the other, then:

A/A

and

B/B

can both be equilibria.

The problem isn't necessarily choosing the wrong strategy.

The problem is:

How do we get to the same one?

This is one reason coordination is its own important category of strategic problem.

And it's one reason our next few chapters will become increasingly interested in information, expectations, and shared beliefs.

The Meeting at Noon

Thomas Schelling famously illustrated coordination with a simple thought experiment.

Imagine you've agreed to meet someone in New York City tomorrow.

You've agreed on the day.

But you haven't agreed on the location or time.

Where do you go?

A famous result of Schelling's experiments was that people frequently converged on **noon at Grand Central Station** even though nothing about those choices was mathematically unique. The location and time functioned as a **focal**

point—something that stands out and helps people coordinate their expectations.²⁰

This is fascinating because it demonstrates something important:

People don't always need to communicate explicitly to coordinate.

They can use context, conventions, salience, and shared expectations.

That will become much more important in Chapter 10.

For now, it gives us another lesson:

Multiple equilibria create a coordination problem.

The Equilibrium You Don't Even Realize You're In

Sometimes the most powerful equilibria are invisible.

Think about workplace meetings.

Everyone complains privately that meetings are too long.

Yet every meeting lasts two hours.

Why?

The manager keeps scheduling them.

Employees keep attending.

Nobody wants to be the person who says:

“I'm leaving after 30 minutes.”

So the pattern continues.

Or imagine a company where everyone waits for senior management before making decisions.

20. TopCat, https://topcat.aeaweb.org/full_issue.php?doi=10.1257%2Fjep.30.4

Management complains:

“Why doesn’t anyone take initiative?”

Employees think:

“Because whenever we do, management overrides us.”

Each side’s behavior reinforces the other.

That’s an equilibrium.

And here’s the interesting question:

What would have to change for the pattern to move?

That’s a much more useful question than:

“Who’s to blame?”

Equilibrium Is About Unilateral Change

This point is worth slowing down for.

Suppose you and I are in a stable pattern.

I could change.

You could change.

But if I change while you don’t, I become worse off.

And if you change while I don’t, you become worse off.

Then neither of us has a unilateral incentive to move.

That’s the heart of Nash equilibrium.

The situation might be:

mutually beneficial,
mutually harmful,
unfair,

inefficient,
or merely acceptable.

What matters is that each person's current strategy is a best response to the others.²¹

This Explains Why “Just Change” Can Be Bad Advice

Imagine telling someone trapped in a dysfunctional relationship:

“Just stop participating in the argument.”

That might sound reasonable.

But what happens next?

If their partner interprets the silence as abandonment, they may escalate.

The person's unilateral change could initially make the situation worse.

Or imagine a company trying to stop a price war.

One company raises its prices.

If the other company doesn't follow, it may lose customers.

So the first company has little incentive to move alone.

The equilibrium isn't necessarily overcome by one person's good intentions.

Sometimes the players need to **change their expectations about what the others will do.**

That's much harder.

And much more interesting.

21. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/entries/logic-power-games/>

The Difference Between an Equilibrium and an Agreement

This is another important distinction.

An agreement is something people explicitly decide.

An equilibrium doesn't require an agreement at all.

Two people can arrive at a stable pattern without ever discussing it.

In fact, many of the strongest social patterns are unwritten.

You know:

- who speaks first,
- who pays,
- who makes the decision,
- who apologizes,
- who gives in,
- who gets interrupted,
- who gets the last word.

Nobody may have formally negotiated these rules.

Yet everyone understands them.

They're part of the game.

And Equilibria Can Be Self-Reinforcing

Once a pattern develops, it can become harder to change because everyone's expectations adjust to it.

Imagine a new employee joins a company where everyone stays late.

At first, they might think:

"Why is everyone still here?"

After a few weeks:

“I guess this is normal.”

After a few months:

“If I leave at five, people will think I’m not committed.”

Now their behavior reinforces the very norm they initially questioned.

The equilibrium has become cultural.

This is one reason social norms can be so powerful.

People don’t merely respond to incentives.

They respond to expectations about what other people will do.

Equilibrium Doesn’t Tell You Which Equilibrium You’ll Get

This is a major limitation of the concept.

Suppose a game has two equilibria.

Knowing that doesn’t necessarily tell us which one people will choose.

That’s called an **equilibrium-selection problem**.

And this is where context becomes important.

Research following Schelling's work has shown that salient labels and contextual cues can influence which equilibrium people coordinate on, although the effectiveness of such focal points depends on the underlying payoff structure. In experiments, even small payoff asymmetries could substantially disrupt coordination around otherwise salient choices.

22

This is important because it keeps us from making an overly simplistic claim:

"Once you know the equilibria, you know what will happen."

Not necessarily.

You may know the possible stable outcomes without knowing which one people will select.

That's where beliefs, expectations, communication, conventions, and focal points enter the picture.

Equilibrium Isn't Prediction

Here's another subtle point.

A Nash equilibrium tells us:

if everyone is behaving this way, no one wants to deviate unilaterally.

It does not necessarily tell us:

this is exactly what real people will do.

Human beings make mistakes.

They have limited information.

22. AEA Web, <https://swlb1.aeaweb.org/articles?id=10.1257%2Faer.104.10.3256>

They may not calculate perfectly.

They may use rules of thumb.

They may learn over time.

Experimental work has found that actual behavior in games can differ systematically from Nash predictions, and researchers have developed alternative behavioral models partly to account for these deviations.²³

This is an important limitation for our framework.

We're using equilibrium as a **mental model**, not a crystal ball.

It helps us identify patterns of strategic stability.

It doesn't give us certainty about human behavior.

What About Learning?

Here's where things get particularly interesting.

Suppose you and I play a game repeatedly.

At first, neither of us knows what the other will do.

We experiment.

We observe.

We adjust.

Eventually, we may settle into a stable pattern.

In some games, experimental evidence shows that learning can produce convergence toward Nash equilibria under certain conditions, although convergence isn't universal and depends on the structure of the game.²⁴

23. AEA Papers and Proceedings, <https://pubs.aeaweb.org/doi/10.1257/aer.98.3.938>

This connects directly to the meta-principle we established earlier:

Model → act → observe → update.

Equilibrium can sometimes be something players **learn into**.

It isn't necessarily something they calculate on day one.

This Changes How We Think About Bad Habits

Consider two people who repeatedly misunderstand one another.

At first, the misunderstandings may be accidental.

But after enough repetitions, each person starts anticipating the other.

You say:

“We need to talk.”

They think:

“Here comes another criticism.”

They become defensive.

You interpret the defensiveness as confirmation that they're unwilling to listen.

So you push harder.

They become even more defensive.

Eventually, the interaction becomes almost automatic.

Neither person needs to consciously choose the pattern anymore.

24. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2F0002828043052349>

They've learned it.

This is what an interpersonal equilibrium can feel like.

The game starts playing you.

That's a phrase I think is worth remembering.

How Do You Break an Equilibrium?

This is where the concept becomes practically useful.

If a pattern is self-reinforcing, simply changing your intention may not be enough.

You have to change something that affects the strategic structure.

Maybe:

- your response changes,
- the sequence changes,
- the incentives change,
- new information appears,
- a commitment is made,
- a third party becomes involved,
- expectations change,
- the available choices change.

Notice how many of those ideas will become later chapters.

That's intentional.

Equilibrium is the bridge between the individual models.

It explains why the pieces we've already learned can lock together into persistent patterns.

A Simple Example

Imagine this interaction:

You ask:

“Can you help me with this?”

They say:

“I’m too busy.”

You respond:

“You’re always too busy.”

They say:

“Because you always ask at the last minute.”

You respond:

“Because you never tell me when you’re available.”

They respond:

“Because you don’t plan ahead.”

And around you go.

Who’s right?

That’s not necessarily the most useful question.

Instead ask:

What pattern have we created?

Maybe the equilibrium is:

You ask late → they refuse → you criticize → they defend → you feel unsupported → you ask late again.

Now we can see something that wasn’t obvious when looking at individual statements.

The interaction itself is producing the outcome.

The Strategic Question Changes

Earlier, you asked:

What should I say?

Then:

What will they do if I say it?

Now:

What pattern are our responses creating?

And eventually:

How could the pattern change?

That progression is the heart of this book.

We're moving from:

words

to

responses

to

patterns

to

systems.

The Equilibrium Audit

When you find yourself stuck in a recurring interaction, ask:

What keeps happening?

Forget the individual incident.

Look for the repeated pattern.

What does my behavior cause them to do?

What does their behavior cause me to do?

Why might neither of us want to change unilaterally?

Is this pattern actually good for either of us?

Are there alternative stable patterns?

And then the most important question:

What would have to change for a different pattern to become attractive?

That final question takes us toward the rest of the book.

The Mental Shift

Before this chapter, you learned to see:

moves.

Now I want you to see:

patterns.

Don't just ask:

"Why did they say that?"

Ask:

"Why do we keep ending up here?"

Don't just ask:

"How do I win this argument?"

Ask:

"What interaction are we collectively creating?"

And don't assume that because a pattern is stable, it is good.

It may simply mean:

Nobody sees a better unilateral move.

That's a very different thing.

The Question to Carry With You

The next time you find yourself stuck in a recurring conflict, ask:

What equilibrium have we created?

Then ask:

Why is it stable?

What does each of us get from maintaining it?

What makes changing it risky?

Are there other equilibria available?

And finally:

What would have to change for us to move toward one of them?

Because sometimes the problem isn't that you haven't found the right words.

Sometimes the problem is that you've both become very good at playing the same bad game.

And if that's true, the solution isn't necessarily to make your next move better.

It may be to change the pattern.

That leads to the next question:

If the game depends on what people believe, what happens when you don't know what the other player knows—or when you don't know what they believe about you?

That's where we go next.

PART II

See What Isn't Obvious

Map Information & Beliefs

In the first chapter, we learned to **see the game**.

In the second, we learned to **see the payoffs**.

In the third, we learned to **think in responses**.

And in the fourth, we saw how our repeated behavior can become information that other people use to predict us.

Now we need to deal with something even more fundamental:

What does everyone actually know?

Because you and I can be standing in exactly the same situation and still be playing very different games if we have different information.

You may know something I don't.

I may know something you don't.

We may both know something—but one of us may not know that the other knows it.

And sometimes, what matters most isn't what is actually true.

It's **what each person believes is true**.

This is one of the major branches of formal game theory. Modern game-theory courses devote substantial attention to Bayesian games, incomplete information, signaling, cheap talk, and common knowledge.²⁵

And this is where communication becomes particularly interesting.

Because communication is, in large part, an attempt to change what another person knows or believes.

The Same Game Can Look Completely Different

Imagine you're negotiating to buy a used car.

The seller knows the car has a serious mechanical problem.

You don't.

You think you're negotiating over:

What is this car worth?

The seller knows you're actually negotiating over:

How much can I get you to pay before you discover the problem?

Same car.

Same conversation.

Different information.

And therefore, a very different game.

Now reverse it.

25. MIT OpenCourseWare, https://ocw.mit.edu/courses/14-12-economic-applications-of-game-theory-fall-2025/video_galleries/video-lectures/

Suppose **you** know you've already been approved for financing elsewhere, but the seller doesn't know that.

Your alternatives look different to you than they look to the seller.

Again:

The information structure changes the game.

Information Is Not the Same as Truth

This distinction is extremely important.

There is:

what is actually true

and there is:

what each player believes is true.

Those aren't necessarily the same thing.

Imagine a company is actually doing very well.

But employees believe layoffs are coming.

What do they do?

They may start looking for other jobs.

They may stop taking risks.

They may save more money.

They may become less committed.

Their behavior is responding to their **belief about the future**, not the objective reality that management knows.

And their behavior may then affect the company.

This gives us a powerful principle:

People make decisions based on the game as they perceive it.

The Classic Example: The Bank Run

Imagine you have money in a bank.

You believe the bank may fail.

You don't know whether that's true.

But you also know something else:

If enough other depositors withdraw their money, the bank **could** actually fail.

So you go to the bank and withdraw your money.

Other people see you doing it.

They become nervous.

They withdraw theirs.

More people see the withdrawals.

They become even more nervous.

Suddenly, a belief about what might happen can help create the very outcome people feared.

This is the strategic power of beliefs.

The outcome isn't determined only by the bank's underlying financial condition.

It's also affected by what depositors believe other depositors are going to do.

This is one reason bank runs have been studied as coordination problems and dynamic games with imperfect information. MIT's game-theory curriculum uses bank runs as an

example of dynamic games in which players have imperfect information.²⁶

The lesson for our purposes is broader:

What people believe other people will do can affect what they themselves do.

Information Has Value

Think back to the negotiation from Chapter 2.

You have an outside option.

But the other person doesn't know what it is.

That's valuable information.

If they believe:

“This is Paul's only option,”
they may negotiate differently.

If they believe:

“Paul can walk away and has three other offers,”
they may behave differently.

So information itself can have strategic value.

That's why people sometimes:

- ask questions,
- investigate,
- gather intelligence,
- remain silent,
- reveal information selectively,

26. MIT OpenCourseWare, <https://ocw.mit.edu/courses/17-810-game-theory-spring-2021/pages/readings/>

- make statements,
- observe behavior,
- test one another.

They're trying to improve their model of the game.

The Information Gap

Imagine you're hiring someone.

The candidate knows:

how hard they actually work.

You don't.

They know:

how good they really are.

You don't.

You have to infer their qualities from:

- their résumé,
- references,
- work samples,
- interviews,
- past results,
- recommendations.

This is an **information asymmetry**.

One player knows something relevant that another player doesn't.

And information asymmetry is everywhere.

A doctor knows more about medicine than the patient.

A mechanic knows more about cars than the customer.

A lawyer knows more about the legal system than the client.

A salesperson may know more about the product than the buyer.

An employee knows more about their own effort than their manager.

A child knows more about what they did five minutes ago than their parent does.

The strategic question is:

How do you make decisions when someone else knows something you don't?

The Problem With "Just Ask"

It sounds easy.

If you don't know something, ask.

But here's the problem:

The other person gets to decide what to tell you.

And they may have an incentive to give you information that benefits them.

This is where communication becomes strategic.

Suppose you're buying a used car.

You ask:

"Has this car ever had any major problems?"

The seller says:

"No."

You've received information.

But should you automatically believe it?

Not necessarily.

The statement itself has to be evaluated in the context of the game.

What does the seller gain if you believe them?

What do they lose if you discover they're lying?

How easy would the lie be to verify?

What reputation do they have?

Are there other ways to verify the claim?

Now we're beginning to move from:

information

to:

credibility.

That will become Chapter 7.

But first, we need to understand the information problem itself.

Private Information Changes Behavior

In formal game theory, players can have private information about their own “type”—characteristics relevant to the game that other players don't observe. MIT's treatment of Bayesian games describes incomplete information as uncertainty about other players' payoffs or types, with types representing private information relevant to decision-making.²⁷

You don't need to worry about the mathematics of Bayesian games.

27. MIT OpenCourseWare, https://ocw.mit.edu/courses/14-16-strategy-and-information-spring-2016/88a1631b5ed6512f32a2ef239830322a_MIT14_16S16_Bayesian.pdf

Think of **type** as:

something important about me that you don't know.

Maybe I'm:

- highly skilled,
- inexperienced,
- wealthy,
- desperate,
- patient,
- impatient,
- trustworthy,
- risk-averse,
- willing to walk away.

You don't observe these things directly.

You have to infer them.

And I know that you're trying to infer them.

Now we have a game.

The Job Interview

Consider a job interview.

The employer doesn't know exactly how capable you are.

You know more about yourself.

So what happens?

You provide information.

You tell them about your experience.

You give examples.

You answer questions.

You provide references.

You show them your work.

But now something interesting happens.

The employer knows that **you have an incentive to present yourself favorably.**

So your claims aren't automatically treated as facts.

The employer has to distinguish:

information

from

self-interested persuasion.

And that creates a strategic communication problem.

Signaling

This is where we encounter one of the most important concepts in game theory:

signaling.

A signal is an action or piece of information that can help another player infer something they don't directly observe.

Education is a classic example in economics.

An employer doesn't directly observe a candidate's underlying ability.

A degree can provide information about the candidate.

Michael Spence's classic job-market signaling model showed how education can function as a signal when workers have different abilities and the costs of acquiring education differ across types. Signaling has since become a major branch of game theory and information economics. MIT

courses treat signaling as a distinct topic alongside incomplete information and Bayesian games.²⁸

But don't make the mistake of thinking:

"A signal is proof."

It isn't.

A signal is evidence.

Its value depends on how informative it actually is.

Why Cost Can Make a Signal More Informative

Suppose I tell you:

"I'm extremely confident."

That costs me almost nothing.

You don't have much reason to update your beliefs.

Now suppose I put \$10,000 of my own money behind my prediction.

That's different.

My action has become costly.

The reason costly signals can be more informative is that different types of players may face different costs or benefits from sending them.

That is one of the central ideas behind signaling theory.

And it gives us a useful everyday question:

What does this person have to lose by making this claim?

28. MIT OpenCourseWare, <https://ocw.mit.edu/courses/14-126-game-theory-spring-2016/pages/syllabus/>

If the answer is:

nothing,

you should be more cautious about treating the statement as strong evidence.

Words Can Be Cheap

Here's a concept that will become very important when we get to communication.

Suppose someone tells you:

"I really value our relationship."

That's a statement.

But what does it cost them to say it?

Almost nothing.

Now imagine they repeatedly demonstrate that value by making sacrifices, keeping commitments, showing up when it's difficult, and behaving consistently over time.

Now you have something much more informative.

This doesn't mean words are worthless.

Words can be extremely valuable.

But their strategic value depends partly on:

what incentives the speaker has to say them.

This is the intuition behind **cheap talk**.

In formal game theory, cheap talk refers to communication that does not directly change the material payoffs of the game. MIT teaches cheap talk as a distinct communication framework.²⁹

And despite the name, cheap talk isn't necessarily useless.

Messages can sometimes convey meaningful information even when sending them is costless—especially when the sender's interests are sufficiently aligned with the receiver's.

The key is:

Why should I believe you?

The Same Statement Can Mean Different Things

Imagine two salespeople tell you:

“This is our best price.”

Salesperson A has a reputation for being completely transparent.

Salesperson B has a reputation for making exaggerated claims.

Same words.

Different information.

Why?

Because you aren't evaluating the words in isolation.

You're evaluating:

the speaker

the incentives

the history

the context

the alternatives

29. MIT OpenCourseWare, https://ocw.mit.edu/courses/14-12-economic-applications-of-game-theory-fall-2025/resources/mit14_12f25_lec24_1080p_mp4/

the consequences

That's why communication cannot be separated from the game surrounding it.

Beliefs About Beliefs

Now we get into something that can become surprisingly deep.

You don't just have beliefs about the other person.

You may have beliefs about **what they believe about you**.

For example:

I believe you think I'm bluffing.

But perhaps:

You believe I think you believe I'm bluffing.

And perhaps:

I believe you believe I think you're bluffing.

This can continue indefinitely.

Game theorists call these **higher-order beliefs** or **belief hierarchies**. Graduate game theory treats these as a formal part of incomplete-information models.³⁰

We don't need to take the mathematics very far.

The practical lesson is enough:

Sometimes you're not responding to what someone knows. You're responding to what you think they know about what you know.

That is a very different kind of interaction.

30. MIT OpenCourseWare, <https://ocw.mit.edu/courses/14-126-game-theory-spring-2016/pages/readings/>

The Poker Table

Poker makes this intuitive.

You don't simply ask:

"What cards does my opponent have?"

You don't know.

You ask:

"What do they think I have?"

Then:

"What do they think I think they have?"

And:

"What do they think I think they're going to do?"

That's strategic reasoning under uncertainty.

And notice how different this is from the first three chapters.

We aren't simply predicting:

what they'll do.

We're predicting:

what they'll do based on what they believe.

The Bluff

This is why bluffing is possible.

If you can convince someone that you have a strong hand when you don't, you've changed their beliefs.

Their behavior changes.

And because their behavior changes, your outcome changes.

But there's a deeper point here.

A bluff doesn't work because the statement is true.

It works because the other player **believes** it—or at least believes it enough to act differently.

That distinction between truth and belief is central to strategic communication.

But Beliefs Can Be Wrong

Obviously.

You can believe I'm bluffing when I'm not.

I can believe you're going to cooperate when you're actually going to defect.

We can both misunderstand each other.

And sometimes those mistaken beliefs can create stable patterns.

Remember the equilibrium from Chapter 5?

Now imagine that the equilibrium is being sustained partly because each person has incorrect beliefs about what the other person wants or intends.

That's where information becomes extremely powerful.

If you change the information, you may change the equilibrium.

A Misunderstanding Can Become Self-Fulfilling

Imagine two departments in a company.

Department A believes:

“Department B will never cooperate with us.”

So A stops sharing information.

Department B sees this and thinks:

“They don’t trust us.”

So B stops sharing information too.

A then says:

“See? We were right.”

Now the belief has become self-reinforcing.

Neither side necessarily started with malicious intent.

The belief itself altered behavior.

And behavior reinforced the belief.

This is a recurring pattern in strategic interaction:

Beliefs can influence actions, and actions can reinforce beliefs.

What If You Could Change Their Belief?

This is where communication enters the game.

Suppose you know the other person incorrectly believes that you are unwilling to cooperate.

You could tell them:

“I’m willing to cooperate.”

But they may not believe you.

Why?

Because the statement is cheap.

So perhaps you need to do something that changes their information.

You might:

- make a commitment,
- provide evidence,
- involve a trusted third party,
- change your behavior,
- offer a verifiable concession,
- take an action that would be costly if you weren't sincere.

Notice where we're heading.

Information leads to credibility.

And credibility leads directly into our next chapter.

Common Knowledge

There's another layer we need to introduce.

Suppose you and I both know that the meeting starts at 9:00.

That's one level of knowledge.

Now suppose:

I know it.

You know it.

But do I know that **you know it**?

And do you know that **I know that you know it**?

This may sound ridiculous in everyday life.

But in strategic situations, it can matter enormously.

Game theory calls this kind of recursively shared awareness **common knowledge**. MIT's game-theory curriculum explicitly treats common knowledge as a major topic, and the concept is particularly important in coordination problems.³

1

You don't need to memorize the formal definition.

Think of it as:

We both know something, and we both know that we both know it.

And so on.

Why Common Knowledge Matters

Imagine you're in a room with ten people.

Someone announces:

"At exactly noon, we're all going to leave the building."

Everyone hears it.

But imagine instead that only half the people hear it.

Those who heard it know what is supposed to happen.

But they don't know whether everyone else heard it.

Now coordination becomes difficult.

The problem isn't necessarily lack of information.

It's lack of **shared knowledge that the information is shared**.

This distinction is subtle but powerful.

31. MIT OpenCourseWare, https://ocw.mit.edu/courses/14-12-economic-applications-of-game-theory-fall-2025/video_galleries/video-lectures/

The Focal Point Revisited

Remember Schelling's meeting example from Chapter 5?

You and I agree:

“Let's meet tomorrow.”

But we don't specify where.

You might choose Grand Central.

I might choose Times Square.

Even if both are reasonable, we're not coordinated.

But if something makes one location particularly salient, it may become a focal point.

This is where beliefs and common knowledge interact.

We may each think:

“They'll probably choose Grand Central.”

And because we both think that, we both choose it.

The shared expectation helps create coordination.

This is one reason rituals, conventions, public announcements, and visible commitments can matter strategically: they can create shared awareness about what people expect others to do.³²

Information Doesn't Have to Be True to Affect Behavior

This deserves careful wording.

A rumor can change behavior.

A misunderstanding can change behavior.

32. JSTOR, <https://www.jstor.org/stable/j.ctt7rx36>

A false belief can change behavior.

A prediction can change behavior.

A credible warning can change behavior.

The strategic effect comes from:

what people believe and how those beliefs affect their choices.

But that doesn't mean truth doesn't matter.

Truth determines whether a belief is accurate.

And inaccurate beliefs can eventually collapse when reality produces contradictory evidence.

So we're not saying:

"Perception is reality."

We're saying something more precise:

Perception is part of the causal structure of strategic interaction.

That distinction matters.

The Information Problem in Relationships

Imagine you text someone:

"Can we talk?"

They don't respond.

What do you know?

Very little.

But your mind immediately starts generating explanations.

Maybe they're angry.

Maybe they're busy.

Maybe they didn't see it.

Maybe they're avoiding you.

Maybe something happened.

You have incomplete information.

And because you don't know, you begin forming beliefs.

Then those beliefs influence your next action.

Maybe you send another text.

Maybe you withdraw.

Maybe you become angry.

And now **your belief about their behavior is influencing their future behavior.**

This is why uncertainty can create conflict even when neither person intends it.

Don't Confuse Uncertainty With Deception

This distinction is important.

Sometimes you don't know what someone is thinking because:

they are deliberately withholding information.

Sometimes:

they don't know themselves.

Sometimes:

the information simply isn't available.

Sometimes:

you haven't asked.

And sometimes:

you've interpreted ambiguous evidence incorrectly.

Strategic thinking doesn't automatically assume deception.

It asks:

What information do I actually have, and how confident should I be in the conclusions I'm drawing from it?

That's a much healthier approach.

Update Your Beliefs

This gives us another important habit.

Don't treat your initial interpretation as fact.

Treat it as a hypothesis.

You think:

"He's angry with me."

Fine.

That's a hypothesis.

Then ask:

What evidence would support that?

What evidence would contradict it?

What else could explain the behavior?

As new information appears, update.

This sounds obvious.

But human beings are remarkably good at forming an interpretation and then looking for evidence that confirms it.

Strategic thinking requires something different:

Your model of the other player should be revisable.

That's especially important because incomplete information is unavoidable.

You will never know everything.

The Information Map

We can now add another layer to our framework.

When you're entering an important interaction, ask:

What do I know?

Separate facts from assumptions.

What don't I know?

Identify the uncertainty.

What do they know that I don't?

Look for information asymmetry.

What do I know that they don't?

Recognize your own informational advantage.

What do I think they believe?

Now you're modeling their beliefs.

What do they think I believe?

If the stakes are high, this can matter too.

What evidence could change my mind?

This protects you from becoming trapped inside your own assumptions.

The Mental Shift

We've now moved another level deeper.

In Chapter 1, we learned to see:

the game.

In Chapter 2:

the payoffs.

In Chapter 3:

the responses.

In Chapter 4:

the patterns others learn about us.

In Chapter 5:

the equilibrium those patterns can create.

Now:

the beliefs underneath the behavior.

Because sometimes the most important thing happening in an interaction isn't visible at all.

It's happening inside the players' models of one another.

You may be responding to what you **think** I believe.

I may be responding to what I **think** you believe.

And both of us may be wrong.

The Question to Carry With You

The next time you find yourself confused by someone's behavior, don't immediately ask:

"Why are they doing this?"

Ask:

“What do they know—or believe—that I don’t?”

Then ask:

What do I actually know?

What am I assuming?

What might they be assuming?

What information would change either of our decisions?

And when someone tells you something important, don’t stop at:

“What did they say?”

Ask:

“Why should I believe it?”

Because information changes the game.

Beliefs change behavior.

And communication is often an attempt to change what another player knows or believes.

But that creates a new problem.

If someone tells you something that benefits them, **why should you believe them?**

That’s where our next mental model begins:

Separate Words From Credibility.

Because sometimes the most important question isn’t:

“What did they say?”

It’s:

“What makes that statement believable?”

Separate Words From Credibility

In the last chapter, we dealt with information.

We learned that people don't make decisions based only on what is objectively true. They make decisions based on:

what they know,
what they don't know,
what they believe,
and what they think other people believe.

That immediately creates another problem.

Suppose I know something you don't.

I tell you what I know.

Now what?

Should you believe me?

That depends.

Because I am not simply a source of information.

I'm also a **player in the game**.

I have interests.

I have incentives.

And what I want you to believe may affect what you decide to do.

So one of the most important strategic habits we can develop is learning to separate:

what someone says

from

what makes what they're saying credible.

Those are not the same thing.

"Trust Me"

Imagine someone says:

"Trust me."

What information did you actually receive?

Very little.

The words tell you that the person wants you to trust them.

But they don't necessarily tell you whether you **should**.

Now imagine two different situations.

Your longtime business partner, who has kept every major commitment they've made to you for fifteen years, says:

"Trust me. I'll take care of this."

And a stranger you're about to wire \$50,000 to says:

"Trust me. I'll take care of this."

Same words.

Completely different credibility.

Why?

Because credibility isn't contained entirely inside the sentence.

It comes from the **game surrounding the sentence**.

History.

Incentives.

Evidence.

Knowledge.

Reputation.

Consequences.

Available verification.

What the person stands to gain.

What they stand to lose.

That's why persuasion and credibility are not the same thing.

A message can be extremely persuasive and still not be very credible.

And a message can be highly credible without being persuasive at all.

The Strategic Communication Problem

Game theorists have studied this problem formally.

Imagine two players.

One knows something the other doesn't.

We'll call the informed player the **sender**.

The other is the **receiver**.

The sender communicates a message.

The receiver then makes a decision.

But here's the important part:

The sender cares about what decision the receiver makes.

Now communication becomes strategic.

Vincent Crawford and Joel Sobel's classic model of strategic information transmission formalized exactly this situation: a better-informed sender communicates with a receiver whose later action affects both players. One of the model's central insights is that the amount of information that can be credibly transmitted depends partly on how closely the two players' interests are aligned.³³

That's an enormously useful idea.

Because it tells us:

You shouldn't evaluate a message without considering the sender's incentives.

Alignment Changes the Meaning of Words

Suppose you're visiting a city you've never been to.

You ask a stranger:

"Which direction is the train station?"

They say:

33. Haverford College, <https://www.haverford.edu/sites/default/files/CrawfordSobel1982.pdf>

“Two blocks east.”

Why would they lie?

Maybe they could.

But under normal circumstances, their payoff doesn't depend much on whether you walk east or west.

Your interests are basically aligned:

you want accurate directions,
and

they have little reason to give you inaccurate ones.
Now imagine you're negotiating to buy their car.

You ask:

“What's the most anyone else has offered you?”

They say:

“\$20,000.”

Should you treat that statement exactly like the directions to the train station?

Probably not.

Why?

Because now the answer affects how much money they may receive.

Their incentives are different.

That does not mean they're lying.

That's very important.

Strategic thinking isn't:

“They benefit from saying it, therefore it must be false.”

That’s paranoia.

The correct conclusion is:

“They benefit from me believing this, so the statement by itself gives me less reason to update my belief than it would if their incentives were neutral or aligned with mine.”

That’s a much more precise way of thinking.

Cheap Talk

In Chapter 6, we briefly introduced a term game theorists use:

cheap talk.

It sounds dismissive, but technically it has a specific meaning.

Cheap talk is communication that is generally costless and nonbinding. The sender can say something without the message itself directly changing the material structure of the game.

The temptation is to assume:

“If talk is cheap, talk is meaningless.”

But that isn’t what the research says.

Joseph Farrell and Matthew Rabin reviewed the game-theory literature on cheap talk and argued that both common extremes are wrong: costless communication does not automatically produce efficient truth-telling, but neither is it ne-

cessarily meaningless. Under the right incentive structure, ordinary talk can transmit substantial information.³⁴

That distinction matters.

Words can be cheap and still useful.

Why Cheap Talk Can Work

Imagine you and I are driving to an unfamiliar location.

You can see the GPS.

I can't.

You say:

"Turn left."

Those words cost you nothing.

Yet they're highly informative because we both want to arrive at the same destination.

Our interests are aligned.

Now change the game.

We're competing to reach the destination first.

You say:

"Turn left."

Same words.

Different strategic meaning.

Now I need to consider:

What do you gain if I believe you?

Communication didn't change.

34. AEA Papers and Proceedings, <https://pubs.aeaweb.org/doi/10.1257/jep.10.3.103>

The game did.

An Experiment on Truthful Communication

This isn't only theory.

Caleb Cox and Brock Stoddard ran an experiment in which players received private information about whether contributing to a public good would be valuable. Before deciding, they could send costless messages to one another.

There was a tension built into the game.

The group benefited from honest information.

But individual players sometimes had incentives to exaggerate.

What happened?

Communication was largely truthful and improved efficiency.

Then the researchers increased the incentive to exaggerate.

Truthfulness declined, and the benefits of communication became smaller.³⁵

That result is almost a perfect illustration of the principle:

Credibility is partly a property of incentives.

People didn't suddenly become different human beings.

The payoff for exaggerating changed.

And communication changed with it.

35. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2Fmic.20180275>

Don't Ask Only, "Are They Honest?"

This is where I think ordinary thinking about credibility is too simplistic.

We often reduce everything to one question:

Is this person honest?

That's certainly important.

But it isn't enough.

There are at least three different questions.

Are they honest?

Are they telling you what they genuinely believe?

Are they informed?

Are they actually in a position to know?

What are their incentives?

What do they gain or lose depending on what you believe?

Someone can be completely honest and completely wrong.

Your friend can sincerely tell you:

"That investment is guaranteed to go up."

Maybe they believe it.

But perhaps they know almost nothing about investing.

Honesty isn't the same thing as accuracy.

Conversely, someone may know far more than you do but have a powerful incentive to influence your decision.

That's a different credibility problem.

Expertise Without Alignment

Imagine you're buying a complicated financial product.

The salesperson knows far more about it than you do.

So:

Knowledge advantage: salesperson.

But the salesperson also receives a large commission if you buy it.

So:

Incentive alignment: uncertain.

That doesn't mean:

"Ignore the salesperson."

They may be extremely useful.

It means you should separate:

How much do they know?

from:

How much does their payoff depend on what I decide?

Those are independent dimensions of credibility.

Alignment Without Expertise

Now reverse it.

Your brother genuinely wants what is best for you.

You trust him completely.

But he knows nothing about the medical decision you're considering.

His interests may be perfectly aligned with yours.

His knowledge isn't.

So you shouldn't ask:

“Do I trust my brother?”

and think you've solved the problem.

You should ask:

“Trust him to do what?”

Trust him to want the best for you?

Probably.

Trust him to diagnose a complex condition?

Maybe not.

Strategic thinking separates these questions.

Credibility Is Not Confidence

This deserves special attention because human beings are easily influenced by confidence.

Someone speaks decisively.

They make eye contact.

They never hesitate.

They sound certain.

We interpret that as:

They know what they're talking about.

Maybe they do.

But confidence is not evidence of accuracy.

And confidence itself can be strategic.

A negotiator may sound completely certain because uncertainty would weaken their position.

A salesperson may speak emphatically because hesitation would reduce the likelihood of purchase.

A leader may project certainty because they believe the group needs confidence.

None of this automatically makes the underlying claim true or false.

So another mental habit is:

Don't confuse the delivery of a message with the evidence behind the message.

"Would They Say This Anyway?"

This may be one of the most useful questions in the entire chapter.

Suppose a job candidate tells you:

"I'm extremely hardworking."

Ask:

Would an unmotivated candidate also have an incentive to say that?

Of course.

Almost everyone benefits from saying:

"I'm hardworking."

So the statement itself doesn't distinguish very well between hardworking and non-hardworking candidates.

Now suppose the candidate gives you:

a record of completed projects,

references from previous employers,
 years of observable performance,
 a demanding work sample completed successfully.

Those pieces of evidence may be much more informative.

Why?

Because they're harder for different types of candidates to produce equally.

This takes us from cheap talk to **signaling**.

Signals

Michael Spence's work on signaling, recognized as part of the 2001 Nobel Prize in economics, addressed a fundamental problem created by asymmetric information.

How can someone who possesses private information communicate something meaningful about themselves to someone who cannot directly observe it?

Spence's famous example involved education and employment. Employers cannot perfectly observe an applicant's underlying productivity before hiring them. Education can, under certain conditions, function as a signal that helps employers distinguish among workers.³⁶

The broader idea is what matters for us:

Sometimes behavior carries information that words alone do not.

Words Versus Signals

Imagine two contractors.

36. Nobel Prize, <https://www.nobelprize.org/prizes/economic-sciences/2001/spence/facts/>

Both say:

“We stand behind our work.”

Contractor A requires full payment upfront, provides no references, offers no warranty, and disappears after the job.

Contractor B has hundreds of verifiable reviews, provides a written warranty, has operated in the same community for twenty years, and fixes problems at their own expense.

The verbal claim is identical.

The surrounding signals are very different.

And you probably don’t need any game theory training to intuitively understand why.

What game theory does is help us understand **why**.

The second contractor has created evidence that would be harder to maintain if their claim were repeatedly false.

That makes the information more diagnostic.

Cost Alone Doesn’t Make a Signal Credible

Here’s an important mistake to avoid.

People hear the idea of costly signaling and conclude:

“If someone pays a large cost, they must be telling the truth.”

No.

Imagine a terrible company spends \$20 million on advertising.

The advertising was costly.

That doesn’t prove the product is good.

A person can spend a fortune creating an image.

A dishonest company can offer an impressive office.

Someone can make a dramatic gesture for strategic reasons.

Cost matters only if the structure of the signal helps distinguish between different underlying types.

The better question is:

Why would this signal be easier, cheaper, or more worthwhile for someone whose claim is true than for someone whose claim is false?

That's the deeper logic of signaling.

Experiments Show Signaling Is Imperfect

This is another place where the clean theory and actual human behavior need to be distinguished.

Dorothea Kübler, Wieland Müller, and Hans-Theo Normann experimentally tested versions of Spence's job-market signaling game.

More productive worker types invested in education more often, and employers offered higher wages to workers who had invested.

So the basic signaling mechanism appeared.

But separation was incomplete, and wages did not converge perfectly to the theoretical equilibrium.³⁷

That's exactly the kind of evidence we want throughout this book.

37. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0899825608000067>

The theory gives us a lens.

Experiments show us that actual people approximate the logic imperfectly.

So:

Signals can contain information.

But:

Signals do not perfectly reveal reality.

Actions Speak—But Not Always Truthfully

We've all heard:

"Actions speak louder than words."

Sometimes.

But that phrase is still too simplistic.

Actions can also be strategic.

Suppose a company makes a large public investment in a new product.

That could signal:

"We have tremendous confidence in this."

Or perhaps management is trying to convince investors that the product has a future.

Or perhaps they're already too committed to the project and don't want to admit they were wrong.

Same action.

Several possible interpretations.

The correct question isn't:

Words or actions?

It's:

**What does this particular action actually tell me,
given the incentives surrounding it?**

Behavior Is Evidence

One of the most important shifts in this chapter is learning to treat behavior as **evidence** rather than automatically treating it as proof.

Suppose someone tells you:

"I value honesty."

Then you watch how they behave when someone tells them an uncomfortable truth.

Do they punish the person?

Do they become defensive?

Do they retaliate?

Now you've received more information.

Not because:

"Actions always beat words."

But because the behavior occurred in a situation where the person's stated value was tested.

That makes the behavior relevant evidence.

The Manager Who Wants Honest Feedback

Imagine a manager tells the team:

“I want everyone to challenge me when you think I’m wrong.”

Great.

An employee disagrees with the manager during a meeting.

The manager becomes angry.

Later, that employee gets excluded from an important project.

The following month, the manager says again:

“I really want more honest feedback.”

What do employees believe?

Probably not the words.

Because another piece of information has entered the game:

The observed consequence of disagreement.

The manager may sincerely want honest feedback.

But their behavior has created a different incentive structure.

Their actions are communicating:

Disagreement may be costly.

This is where credibility connects back to payoffs.

Credibility Lives in the Whole Game

This is the deeper lesson.

A statement doesn’t exist by itself.

Its credibility depends on things we’ve already learned to notice:

Payoffs

What does the speaker gain if I believe this?

Information

Are they actually in a position to know?

History

What has happened when they've made similar claims before?

Constraints

What prevents them from behaving differently later?

Evidence

Can the claim be independently verified?

Signals

Does their behavior contain information related to the claim?

Reputation

What does their previous behavior suggest?

That is a much richer way of evaluating communication than:

“Did they sound convincing?”

Persuasiveness Is Not Credibility

This distinction could save people from a tremendous number of bad decisions.

Someone may be:

- charismatic,
- confident,

- articulate,
- emotionally compelling,
- persuasive.

None of those qualities automatically tell you whether what they're saying is accurate.

And the reverse can also happen.

Someone can be:

- awkward,
- nervous,
- poor at explaining themselves,

and still be completely correct.

If we're interested in truth, we have to separate:

How convincing was the message?

from:

How credible is the underlying claim?

Those are two different questions.

A Persuasion Failure May Actually Be a Credibility Failure

This matters from the speaker's side too.

Imagine you're trying to convince a client:

"We're going to deliver this project on time."

You explain it beautifully.

You tell a story.

You use powerful language.

You handle every objection.

The client still doesn't believe you.

You conclude:

"I need a more persuasive argument."

Maybe not.

What if you've missed the real problem?

Perhaps you've been late on the last three projects.

The client understands your words perfectly.

They simply don't believe them.

More rhetoric may accomplish nothing.

The strategic problem isn't:

How do I say this better?

It's:

How do I make this claim more credible?

Those are very different problems.

Sometimes the Solution Is Verification

Suppose I tell you:

"The roof is in excellent condition."

You're buying my house.

I benefit from you believing me.

You could trust my statement.

Or you could hire an independent inspector.

Now you've changed the information structure.

The inspector doesn't eliminate all uncertainty.

But you've introduced verification from a player whose incentives may be better aligned with accurate assessment.

This is why:

- audits,
- certifications,
- inspections,
- references,
- independent reviews,
- third-party measurements,

can be strategically powerful.

They aren't merely additional words.

They alter the credibility structure.

Third Parties Can Change Credibility

Imagine I tell you:

"I'm one of the best surgeons in the country."

Maybe I am.

But I have an obvious incentive to make the claim.

Now imagine:

independent medical boards,
peer-reviewed outcomes,
hospital credentials,
other surgeons,

all provide evidence supporting it.

The claim has changed.

Not because my sentence changed.

Because **the information environment changed**.

This is why reputation systems, professional certifications, guarantees, and institutional verification exist.

They help reduce uncertainty when parties cannot simply rely on self-report.

Reputation Is Evidence Accumulated Over Time

This creates a connection to a later chapter.

Suppose you've known someone for fifteen years.

They've consistently:

kept their word,
shown up,
honored agreements,
told you uncomfortable truths,
admitted mistakes.

Now they make a claim that you cannot immediately verify.

Their history matters.

Why?

Because history gives you evidence about how this person behaves across situations.

Reputation is, in one sense:

a summary of information from previous rounds of the game.

We'll explore that more deeply when we reach **The Shadow of the Future**.

For now:

Credibility can accumulate.

And it can collapse.

Credibility Can Collapse Faster Than It Builds

Imagine someone has told you the truth 100 times.

Then you catch them fabricating something important.

What happens?

You don't merely update your belief about that one statement.

You may update your model of:

the person.

Now earlier claims may be reconsidered.

Future claims may require verification.

The informational structure of the relationship has changed.

This is why trust is strategically valuable.

When trust is high, people don't need to verify everything.

Transactions become easier.

Communication becomes faster.

Cooperation requires less monitoring.

Destroy credibility and all of those costs can return.

Influence Without Credibility

Here's another important distinction.

A statement can change your behavior even if it isn't credible evidence.

Someone shouts:

“If you don’t agree, you’ll regret it.”

You may be intimidated.

You may comply.

But compliance doesn’t mean you believed the factual content of the statement.

Likewise, repetition can influence familiarity.

Authority can influence behavior.

Emotion can influence behavior.

Social pressure can influence behavior.

So:

A message can influence without informing.

That’s important because people often confuse:

“It changed my mind”

with:

“It gave me good reason to change my mind.”

Those aren’t necessarily the same.

Truth and Reliability Are Different

Suppose someone tells you something false.

Does that automatically mean they lied?

No.

A lie requires some form of intentional deception.

They may simply be wrong.

So credibility analysis needs another distinction:

Truthfulness

Do they report what they actually believe?

Accuracy

Is what they believe actually correct?

Someone can be truthful but inaccurate.

Someone can occasionally be accurate while being dishonest.

That's why we shouldn't reduce credibility to a personality judgment.

It's a problem of:

knowledge,
incentives,
evidence,
and reliability.

The Speaker Has a Credibility Problem Too

So far, we've mostly looked at credibility from the listener's side.

But suppose **you** are the one trying to communicate something true.

You know you're telling the truth.

That doesn't mean the other person should automatically believe you.

This is one of the hardest things for people to accept.

You say:

"But I'm telling the truth!"

Fine.

How would **they** know?

From their position:

- you may have an incentive to say it,
- they may lack access to your information,
- they may have been misled before,
- they may have competing evidence.

So instead of becoming frustrated, ask:

What reason do they have to believe me?

That's a much more useful question.

"I Know I'm Honest" Isn't Evidence for Them

Imagine you're selling your car.

You know:

the transmission is perfectly fine.

Buyer asks:

"Any transmission problems?"

You say:

"None."

And you become offended when they want a mechanic to inspect it.

From your perspective:

"Why don't they trust me?"

From theirs:

"The seller knows more about this car than I do and benefits from the sale."

Both perspectives make sense.

The inspection isn't necessarily an accusation.

It's a response to an information asymmetry.

That's strategic thinking.

Make Truth Easier to Verify

This gives us a constructive principle:

When you are communicating truthfully, make the truth easier for the other person to evaluate.

That may mean:

- evidence,
- demonstrations,
- references,
- guarantees,
- documentation,
- independent verification,
- consistency between words and behavior.

Not because every truthful claim must be proven.

But because sometimes the listener's skepticism is strategically reasonable.

This is especially important when:

you benefit from being believed.

Credibility Versus Commitment

There's another distinction we need to make before the next chapter.

Suppose I tell you:

“I will not accept less than \$100,000.”

What makes that statement credible?

Maybe my history.

Maybe my alternatives.

Maybe my incentives.

Maybe nothing.

Now suppose I take an action that makes it impossible—or extremely costly—for me to accept less than \$100,000.

Something has changed.

I haven’t simply **signaled** what I’m likely to do.

I’ve changed what I can rationally do later.

That’s **commitment**.

The difference is subtle but important.

Signaling

What does my action tell you about me?

Commitment

How does my action change my future options?

Sometimes one action does both.

But they’re not the same idea.

And Chapter 8 is entirely about the second.

The Credibility Test

When someone makes an important claim, ask:

Do they know?

Are they actually in a position to possess the information?

What do they want me to believe?

Identify the message.

What happens if I believe them?

How does your response affect their payoff?

Would they say this even if it were false?

If yes, the words alone may not be highly informative.

What evidence supports the claim?

Separate assertion from verification.

What does their behavior suggest?

Treat behavior as evidence, not proof.

What does their history suggest?

Reputation can matter.

What would make this claim costly to fake?

Look for genuine signals.

You don't have to run this checklist every time someone tells you where the restroom is.

The point is to develop instinct for **high-stakes credibility problems**.

Don't Become Cynical

This is an important warning.

Once you learn to analyze incentives, it can become easy to distrust everything.

Someone compliments you.

You think:

“What do they want?”

Someone helps you.

You think:

“What’s their angle?”

Someone makes a promise.

You think:

“They’re strategically signaling.”

That’s not the goal.

Most human communication is not some elaborate deception game.

People cooperate.

People tell the truth.

People help each other.

People speak honestly even when lying might benefit them.

The mental model is not:

Trust nobody.

It’s:

Don’t confuse trusting someone with failing to think.

You can trust and still understand the game.

The Ethical Side of Credibility

This chapter could easily become:

“How to make yourself look credible.”

That would miss the entire point.

We’re not trying to manufacture the appearance of truth.

Our ethical standard remains:

Truth. Autonomy. Dignity. Informed choice.

If you’re the speaker, the goal is:

Make truthful claims easier to evaluate.

If you’re the listener:

Evaluate important claims intelligently.

Using fake reviews, fabricated evidence, staged signals, or deliberately exploiting someone’s trust is not strategic excellence.

It’s deception.

Game theory can explain why deception sometimes works.

That doesn’t make it ethical.

What the Research Actually Supports

We should be precise here.

Strategic-communication theory does **not** say:

“People always lie when interests conflict.”

Crawford and Sobel's model shows that informative communication can still occur under partially conflicting preferences, although the degree of informative communication depends on how far the players' interests diverge.³⁸

The cheap-talk literature also does **not** say:

"Costless words are meaningless."

Farrell and Rabin explicitly argued against that extreme.³⁹

And experiments show something very human:

People often communicate truthfully even when they possess some incentive to exaggerate—but stronger incentives to distort can reduce truthfulness.⁴⁰

Signaling theory also doesn't say:

"Observable actions perfectly reveal hidden qualities."

Experimental job-market signaling studies find behavior broadly consistent with signaling logic while still showing incomplete separation and departures from theoretical equilibrium.⁴¹

So our practical conclusion is more modest:

When evaluating a message, consider the strategic environment that produced it.

That's the mental model.

38. Haverford College, <https://www.haverford.edu/sites/default/files/CrawfordSobel1982.pdf>

39. AEA Papers and Proceedings, <https://pubs.aeaweb.org/doi/10.1257/jep.10.3.103>

40. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2Fmic.20180275>

41. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0899825608000067>

The Mental Shift

Before this chapter, you might hear an important statement and ask:

“Do I believe them?”

Now I want you to ask something better:

“What makes this believable?”

That’s a completely different question.

It moves credibility away from gut feeling alone.

You’re now asking:

Do they know?

What are their incentives?

What evidence exists?

Would they say this regardless?

Does their behavior support it?

Is there anything costly about faking the signal?

And from the speaker’s side:

What reason does the other person have to believe me?

That may be even more important.

The Question to Carry With You

The next time someone makes a claim that matters, don’t focus only on the sentence.

Zoom out.

Look at the game around it.

Ask:

What makes this message credible?

Not:

“Did they sound sincere?”

Not:

“Do I like them?”

Not:

“Was the argument persuasive?”

But:

What information did I actually receive?

Because words can inform.

Words can coordinate.

Words can persuade.

Words can deceive.

And sometimes words can do absolutely nothing.

Their strategic value depends on the game they’re spoken inside.

But there is one way to make a statement far more powerful than merely saying it.

You can change your own future options so that the other person has reason to believe you.

And that leads to one of the strangest and most powerful ideas in game theory:

Sometimes having fewer options makes you stronger.

CHAPTER 8

Commitment Changes the Game

In the last chapter, we asked:

What makes a message credible?

Sometimes the answer is evidence.

Sometimes it's reputation.

Sometimes it's aligned incentives.

Sometimes it's a signal that's difficult to fake.

But there is another way to make a statement credible.

You can change the game so that **doing something different later becomes difficult, costly, or impossible.**

That is commitment.

And it leads to one of the most counterintuitive ideas in game theory:

Sometimes having fewer options can make you stronger.

That sounds wrong at first.

More options should be better.

More flexibility should give you more power.

Usually, it does.

But when another person is making a decision based partly on what they expect **you** to do later, your freedom to change your mind can sometimes become a weakness.

And deliberately limiting that freedom can change what they rationally choose to do now.

The Problem With “Final”

Imagine we’re negotiating.

I tell you:

“\$100,000 is my absolute minimum. I will not accept a dollar less.”

What do you do?

That depends on whether you believe me.

Maybe \$100,000 really is my minimum.

Or maybe I would take \$90,000 if you held firm long enough.

Or \$80,000.

You don’t know.

My statement is only a statement.

So you might test me.

You offer:

\$85,000.

I say no.

You wait.

I eventually come back at:

\$95,000.

What did you learn?

My “absolute minimum” wasn’t absolute.

My words weren’t binding me.

Now imagine a different situation.

For some reason, I genuinely **cannot** accept less than \$100,000.

Maybe I’ve entered into an agreement that makes a lower price impossible.

Maybe another party has an enforceable claim.

Maybe I’ve delegated the decision to someone who has no authority to go lower.

Whatever the mechanism, one thing has changed:

My future choice set is different.

You no longer need to ask:

“Will he cave?”

If the constraint is real and visible, continuing to push below \$100,000 may no longer make sense.

I have fewer options.

And yet, strategically, my position may have become stronger.

Thomas Schelling made this paradox central to the study of strategic commitment. The Nobel committee summarized one of his key insights this way: a bargaining player can

sometimes strengthen its position by openly making its own alternatives worse.⁴²

That's a strange idea.

But once you see it, you start seeing it everywhere.

Intention Is Not Commitment

We use the word commitment loosely in everyday life.

We say:

“I'm committed to getting in shape.”

“I'm committed to this business.”

“I'm committed to making this relationship work.”

Those statements can be sincere.

But that's not quite the game-theoretic meaning we're interested in.

We need to separate three things.

Intention

I plan to do something.

Promise

I tell you that I plan to do something.

Strategic commitment

I take an action that changes the incentives or options surrounding what I can do later.

That's different.

If I say:

42. Nobel Prize, <https://www.nobelprize.org/prizes/uncategorized/the-sveriges-riksbank-prize-in-economic-sciences-in-memory-of-alfred-nobel-1969-2007/>

“I’m definitely waking up at 5:00 tomorrow,”
that’s intention.

If I tell my training partner:

“I promise I’ll meet you at the gym at 5:30,”
that’s a promise.

If I arrange for my training partner to pick me up at 5:15
and I’ve already paid for a session I lose if I don’t attend,
I’ve altered the future decision.

Getting out of bed now has a different payoff.

I didn’t just become more motivated.

I changed the game my future self will face.

Why Would Less Freedom Ever Help?

Imagine two players facing each other across a bargaining
table.

Player A can:

accept,
reject,
delay,
compromise,
or walk away.

Player B knows this.

Suppose A says:

“If you reject my offer, I’ll walk.”

But when the time comes, walking away would be extremely
costly for A.

Player B may reasonably think:

“You say you’ll leave, but when you actually have to make that decision, you probably won’t.”

A’s flexibility undermines the threat.

Now imagine A genuinely eliminates the ability to continue negotiating after a deadline.

Perhaps the offer expires because another commitment becomes binding.

Player B’s decision changes.

The important thing isn’t that A sounds more serious.

It’s that:

the future game has changed.

This is the logic behind Schelling’s insight.

A commitment can influence behavior today by changing what another player expects will happen tomorrow.

Burning the Bridge

The classic image is an army eliminating its own route of retreat.

Imagine two opposing forces.

One army still has a bridge behind it.

If battle becomes costly, retreat remains available.

The opposing side knows this.

Now imagine the first army destroys its own bridge.

That seems insane.

It has eliminated one of its own options.

But it has also changed what the opponent believes.

Retreat is no longer part of the future game.

The army has made its determination more credible precisely because it has reduced its own freedom.

This kind of logic is often used to illustrate Schelling's strategic commitment insight: deliberately restricting your own alternatives can alter the other player's incentives.⁴³

Now, that does **not** mean:

"Burn every bridge."

That would be terrible advice.

The important insight is only:

Flexibility has value, but flexibility can also undermine credibility.

That's the tradeoff.

A Threat Is Only as Good as the Future Decision

Suppose a parent tells a child:

"If you don't clean your room, you're grounded for six months."

The child may immediately know:

"That's never happening."

The punishment is severe.

But severity is not the same thing as credibility.

Why?

Because when the moment arrives, carrying out six months of punishment may be miserable for the parent too.

43. Nobel Prize, <https://www.nobelprize.org/prizes/uncategorized/the-sveriges-riksbank-prize-in-economic-sciences-in-memory-of-alfred-nobel-1969-2007/>

So the child anticipates:

They'll back down.

Now compare that with a smaller, clearly defined consequence the parent consistently enforces.

Perhaps:

"If the room isn't cleaned by six, there's no gaming to-night."

That consequence may be less dramatic.

But if the child believes it will actually happen, it can have more strategic effect.

The important question is not:

How severe is the threat?

It's:

Will carrying it out still make sense when the time comes?

Credible Threats and Credible Promises

Commitment works in both directions.

A threat says:

"If you do X, I will do Y."

A promise says:

"If you do X, I will reward you with Y."

Both create the same strategic problem:

Will the other player believe you'll follow through later?

Imagine a company tells an employee:

“If you hit this target, we’ll promote you.”

The employee hits it.

Now management realizes promotion would be inconvenient.

If the promise was nonbinding and management has a history of changing its mind, future employees may stop responding to such promises.

The promise loses strategic value.

A commitment mechanism can make the promise more credible.

Perhaps the bonus is written into a contract.

Perhaps the promotion criteria are published and externally reviewed.

Perhaps the money has already been placed in an incentive pool.

The exact mechanism differs.

The principle is the same:

The future action becomes harder to reverse.

That changes today’s behavior.

Commitment Can Increase Bargaining Power

This isn’t merely a theoretical curiosity.

Laboratory bargaining experiments have tested situations in which players can credibly commit themselves to a bargaining position.

In one set of experiments, Kurtis Swope, John Cadigan, and Pamela Schmitt found that credible commitment could in-

crease the payoff to the person making the commitment, particularly in bargaining environments where other parties otherwise had incentives to hold out for better terms.

But there was a cost.

Commitment could also increase the probability of bargaining failure and reduce overall efficiency.⁴⁴

That is exactly the tradeoff the theory would lead us to expect.

Commitment may strengthen:

your position

while making:

agreement harder.

Those are not contradictory results.

They're two sides of the same mechanism.

You gained power by eliminating flexibility.

But flexibility was also one of the things that might have allowed a deal to happen.

The Power and Danger of "I Can't"

There's a major strategic difference between:

"I don't want to."

and:

"I can't."

Suppose you're negotiating on behalf of a company.

44. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S2214804314000160>

The other side asks you to approve another 10 percent discount.

You say:

“I’d rather not.”

That invites pressure.

Why?

Because they know:

your preference can change.

Now imagine you truthfully say:

“My authority stops at this number. Anything beyond it requires approval from the board.”

That’s different.

The constraint matters.

Not because the board is magical.

Because your available actions have changed.

This is one reason delegation can create commitment.

A decision-maker can sometimes strengthen a bargaining position by transferring authority to someone whose choices are more restricted.

Schelling’s commitment logic has been applied to exactly this kind of delegation and institutional constraint.⁴⁵

But “My Hands Are Tied” Can Also Be Fake

Of course, people know this.

45. Nobel Prize, <https://www.nobelprize.org/prizes/uncategorized/the-sveriges-riksbank-prize-in-economic-sciences-in-memory-of-alfred-nobel-1969-2007/>

Which means they sometimes pretend to be constrained.

You've probably heard:

"I'd love to help you, but corporate won't let me."

Maybe that's completely true.

Maybe "corporate" is sitting ten feet away and would approve the request instantly.

This takes us back to Chapter 7.

A claimed commitment is only strategically useful if it's **credible**.

Ask:

What actually binds them?

Is there:

- a contract?
- a rule?
- a third-party authority?
- a legal restriction?
- a public commitment?
- a financial consequence?
- a reputational cost?
- an irreversible investment?

Or merely:

words?

That distinction matters.

Commitment Must Usually Be Observable to Affect the Other Player

Suppose I secretly make it impossible for myself to accept less than \$100,000.

You don't know.

Will your behavior change?

Probably not.

My private commitment may affect **my future behavior**.

But it can't affect **your present decision** unless you know—or can credibly infer—that it exists.

That's a crucial distinction:

A hidden commitment constrains you.

An observable commitment can constrain you and change them.

That's what makes public commitments, contracts, delegated authority, and visible irreversible investments strategically powerful.

They alter other people's expectations.

A Real-World Field Experiment

A useful example comes from a field experiment by Emily Breza, Martin Kanz, and Leora Klapper.

Factory workers were offered a commitment savings account. Opening the account could signal that a worker expected to remain with the company.

Workers were more likely to open the accounts under conditions where the employer would observe the decision. And

in a separate experiment, managers were more willing to invest in workers who held the accounts.⁴⁶

This example sits right at the intersection of the last chapter and this one.

The account acted partly as a **signal**:

“I expect to stay.”

But it also involved a **financial commitment**.

That’s important because real actions often serve more than one strategic purpose at the same time.

They can:

 reveal information

and

 change future incentives.

Signal or Commitment?

Let’s make the distinction clear.

Imagine I put \$50,000 into a new business.

That action may tell you:

I believe in this business.

That’s signaling.

But suppose the investment is irreversible and now gives me a strong incentive to continue supporting the business.

That’s commitment.

One action can do both.

46. AEA, <https://swlib2.aeaweb.org/articles?id=10.1257%2Fpandp.20181011>

But conceptually they answer different questions.

Signaling asks:

What does this action tell the other player about me?

Commitment asks:

How does this action change what I can or will do later?

Keeping those separate makes strategic analysis much cleaner.

Commitment Can Reduce Uncertainty

Suppose you're negotiating with someone who keeps saying:

"Maybe I'll accept. Maybe I won't."

You don't know how long they'll continue bargaining.

You don't know whether the next offer will be final.

You don't know whether waiting will improve your outcome.

Now suppose they credibly commit:

"This offer expires Friday at noon, and after that the asset goes to another buyer."

If true, uncertainty drops.

You know more about the future game.

That can actually make agreement **easier**.

This is another reason commitment is not simply an aggressive bargaining weapon.

It can also provide clarity.

The same bargaining research that finds commitment can increase impasse also notes that credible ultimatums can reduce uncertainty and eliminate incentives for strategic delay in some settings.⁴⁷

Again:

The structure matters.

A New Experiment Shows Human Commitment Is Messier Than the Theory

A 2026 experiment examined commitment before sequential ultimatum bargaining.

The clean theory predicts particular ways players should use precommitment to gain strategic advantage.

Actual subjects did use commitment strategically—but not exactly as the theory predicted. Commitments were often less aggressive, and concerns such as reciprocity and equal opportunity appeared to affect behavior. In later rounds of some conditions, the availability of commitment produced very high efficiency.⁴⁸

That’s useful for us because it reinforces something we’ve seen repeatedly throughout this book:

Game theory gives us the strategic mechanism. Human beings add psychology, norms, fairness, mistakes, and learning.

So we should never interpret:

“Commitment can create bargaining power”

as:

47. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S2214804314000160>

48. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0167268125004779>

“Humans will use commitment exactly as the mathematical model predicts.”

They won’t always.

Commitment Can Create Disaster

Now let’s look at the downside.

Suppose two countries are in a confrontation.

Country A publicly commits:

“If you cross this line, we will retaliate.”

Country B publicly commits:

“We will cross the line.”

Both commitments may be highly credible.

That doesn’t mean the outcome is good.

It may mean that both sides have removed the paths that would have allowed them to de-escalate.

This is the dangerous side of commitment.

A credible commitment can lock you into a terrible outcome.

Game theory does not teach:

“Commit more.”

It teaches:

Understand what commitment does to the game.

The Commitment Trap

Imagine a business owner publicly announces:

“We will never raise our prices.”

It sounds admirable.

Customers love it.

Years later, costs double.

Now the company has a problem.

The original commitment created credibility.

But it also destroyed flexibility.

Or imagine a manager declares:

“Anyone who makes this mistake is automatically fired.”

Then the best employee in the company makes an honest mistake under unusual circumstances.

Now the manager must choose between:

breaking the commitment

or

carrying out a punishment that may no longer make sense.

This is why strategic commitment should never be confused with stubbornness.

A good commitment considers:

What happens if the world changes?

Uncertainty Makes Commitment Riskier

The less you know about the future, the more valuable flexibility can become.

Suppose you're entering a game where:

- the payoffs may change,
- new information may emerge,

- the other players may change,
- the rules may change,
- your own preferences may change.

Committing strongly today may make tomorrow's decision worse.

So another strategic question becomes:

How much uncertainty am I locking myself into?

If the future is reasonably predictable, commitment may be valuable.

If the future is extremely uncertain, flexibility may be worth more.

There's no universal answer.

That's why this is a mental model rather than a rule.

Don't Confuse Commitment With the Sunk-Cost Fallacy

This distinction is essential.

Suppose I invest \$100,000 in a project.

That investment may function strategically.

It may signal seriousness.

It may create incentives to continue.

It may alter what competitors expect me to do.

But imagine six months later we discover the project is terrible.

The \$100,000 is gone.

Should we continue merely because:

“We’ve already invested too much to quit”?

No.

That’s the sunk-cost fallacy.

A strategic commitment is:

using today’s action to deliberately change tomorrow’s incentives or expectations.

The sunk-cost fallacy is:

allowing yesterday’s irreversible cost to distort today’s decision when that cost should no longer matter to the choice.

Those are very different things.

A smart strategist can make commitments.

A smart strategist must also know when a past commitment has become a liability.

Boundaries Are Commitments

Now let’s bring this into everyday human interaction.

Suppose someone repeatedly speaks disrespectfully to you.

You tell them:

“If this continues, I’m ending the conversation.”

Then it happens again.

You stay.

It happens again.

You stay.

What has the other person learned?

Your statement doesn’t constrain your behavior.

It's cheap talk.

Now imagine instead that when the behavior occurs, you calmly say:

"We'll continue this when we can talk respectfully,"
and leave.

If you do that consistently, the other player begins to understand:

The boundary changes the game.

This is one of the clearest applications of commitment.

A boundary is not:

"I demand that you behave differently."

A boundary is:

"If this occurs, here is what I will do."

And it becomes strategically meaningful only when your future behavior makes it credible.

Don't Threaten What You Won't Do

This is one of the simplest practical lessons in the chapter.

People often make threats in moments of emotion.

"I'll never speak to you again."

"I'm quitting."

"I'm cancelling everything."

"If this happens again, you're done."

Then they don't follow through.

Each time this happens, something changes.

The other person updates.

The next threat carries less information.

This isn't merely about personal integrity.

It's game theory.

You are teaching the other player:

Your stated consequence does not predict your actual behavior.

If you're not prepared to follow through, don't create the threat.

Use a consequence you can actually enforce.

Promises Work the Same Way

A manager repeatedly promises:

"If you do this, I'll reward you."

Then doesn't.

What happens?

Eventually employees stop responding to the promise.

The problem may not be motivation.

It may not be communication.

It may be credibility destroyed by failed commitment.

The same principle applies to:

- parenting,
- partnerships,
- sales,
- leadership,
- friendships,
- relationships.

Every promise that isn't kept alters the future game.

Public Commitment

One reason public commitments can be powerful is that they increase the cost of reversal.

Suppose I privately say:

“I'm going to finish this project by October 1.”

If I miss it, maybe nobody knows.

Now suppose I publicly announce the launch date to my customers.

Reversal may now carry:

- reputational costs,
- financial costs,
- embarrassment,
- customer disappointment.

I've made failure more expensive.

That can increase my incentive to follow through.

But notice the danger.

If new evidence appears showing the launch would be unsafe, the public commitment may pressure me to do something foolish just to save face.

So public commitment can create:

discipline

and

rigidity.

Both need to be considered.

Commitment Against Yourself

There's another form of commitment worth mentioning.

Sometimes the other strategic player is, in a sense:

your future self.

You know that tonight you'll want one thing.

Tomorrow morning, you'll want another.

Tonight:

"I'm getting up early."

Morning:

"Five more minutes."

So tonight you change tomorrow's game.

You put the alarm across the room.

You arrange to meet someone.

You prepay for the session.

You remove tempting food from the house.

These are commitment devices.

You're using your current preferences to constrain the options available to your future self.

A field experiment with university students found that offering a nonbinding "soft commitment" to stay on track toward graduation increased exam registration, participation, and passing; simple reminders did not have the same effect, with procrastinating students benefiting especially.⁴⁹

That's not exactly the same strategic problem as bargaining with another person.

49. AEA, <https://swlb2.aeaweb.org/articles?id=10.1257%2Fapp.20170288>

But the underlying insight is related:

Sometimes constraining future freedom improves the outcome.

Commitment Is Not About Being Tough

This is another misconception worth clearing up.

A person who never changes their position isn't necessarily strategically committed.

They may simply be stubborn.

A strategic commitment has a purpose.

It changes:

incentives,
expectations,
options,
or credibility

in a way that affects the game.

If refusing to adapt provides no strategic benefit and creates enormous cost, that's not sophistication.

That's rigidity.

Commitment is valuable only when the constraint improves the strategic environment enough to justify the loss of flexibility.

The Strongest Commitment Isn't Always the Best Commitment

Suppose we're negotiating.

I commit:

“I will accept nothing under \$100.”

You commit:

“I will pay nothing above \$90.”

Congratulations.

We’ve created:

no deal.

Both commitments may be perfectly credible.

The problem is that they’re incompatible.

So another crucial lesson is:

Never evaluate a commitment solely by how strongly it binds you.

Evaluate it by:

the game it creates.

This is exactly why experiments find that commitment can increase both bargaining power and bargaining failure.⁵⁰

The Other Player May Respond Before You Expect

Remember Chapter 3:

Think in responses.

A commitment is itself a move.

So before committing, ask:

How will they respond to the fact that I’m limiting my options?

Maybe they concede.

Maybe they walk away.

50. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S2214804314000160>

Maybe they make their own commitment.

Maybe they search for an alternative.

Maybe they retaliate.

Maybe they stop negotiating.

The whole point is that the commitment changes **their** decision problem.

If you don't think about their response, you're missing half the game.

The Commitment Test

Before making an important commitment, ask:

What future option am I giving up?

Be precise.

Why am I giving it up?

What strategic benefit does the constraint create?

Can the other player observe or verify the commitment?

If not, it may not influence them.

How will their best response change?

That's the point of the commitment.

What happens if new information appears?

How costly is the lost flexibility?

Can this commitment produce an impasse?

What if the other side also refuses to move?

Is the commitment ethical?

Does it preserve truth, autonomy, dignity, and informed choice?

Those questions protect us from turning commitment into bravado.

The Ethical Problem

Commitment can be used badly.

Imagine someone deliberately creates a crisis and then says:

“I’ve made it impossible for either of us to turn back, so now you have to do what I want.”

That can become coercive.

This is where the ethical foundation of our framework matters.

Understanding commitment does not give us permission to trap people.

Our standard remains:

Truth. Autonomy. Dignity. Informed choice.

A contract can be a commitment.

A boundary can be a commitment.

A transparent deadline can be a commitment.

A public promise can be a commitment.

Those can help people coordinate and trust one another.

But intentionally destroying another person’s meaningful ability to choose is something different.

Strategy and ethics remain separate questions.

We need both.

A Powerful Distinction

By this point, you should be able to see three different statements:

"I'm going to do it."

That's an intention.

"I promise I'll do it."

That's communication.

"I've arranged things so that doing anything else would be extremely costly."

That's commitment.

And the third statement can affect behavior even before the future arrives.

Because other players know something about what you're likely to do when that moment comes.

That's why commitment changes the game.

What Does the Evidence Really Say?

The theory is strong.

Schelling's work established the strategic importance of commitment and the paradox that worsening your own options can sometimes improve your bargaining position.⁵¹

51. Nobel Prize, <https://www.nobelprize.org/prizes/uncategorized/the-sveriges-riksbank-prize-in-economic-sciences-in-memory-of-alfred-nobel-1969-2007/>

Experiments show that credible commitments can indeed increase the committer's bargaining payoff in some settings—but they can also increase impasse and reduce efficiency.⁵

2

More recent experimental work shows that real people don't use commitment exactly as clean theory predicts; fairness, reciprocity, and other behavioral considerations matter.⁵³

And field evidence shows that observable commitments can affect how other people interpret and respond to us.⁵⁴

So the evidence does **not** support:

"Commit harder and you'll win."

It supports something much more interesting:

Changing your future options can change another player's current incentives—but commitment creates tradeoffs between credibility, bargaining power, flexibility, and the possibility of agreement.

That's the model.

The Mental Shift

Before this chapter, you may have assumed:

More choices give me more power.

Now you know to ask:

Could reducing my choices make my position more credible?

And when another person tells you:

52. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S2214804314000160>

53. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0167268125004779>

54. AEA, <https://swlb2.aeaweb.org/articles?id=10.1257%2Fpandp.20181011>

“I have no choice,”
you can ask:

What actually binds them?

Those two questions can reveal an enormous amount about a strategic interaction.

The Question to Carry With You

Before making a threat, promise, deadline, boundary, or “final” position, ask:

What makes this commitment real?

Then ask:

What option am I giving up?

How does giving it up change their decision?

What flexibility will I wish I still had if circumstances change?

What happens if they refuse to move?

Because sometimes strength comes from having more options.

And sometimes strength comes from credibly removing one.

But there is another force that can change today’s behavior without removing any options at all.

Simply knowing that there will be a **next time** can change how people behave now.

Today’s actions become tomorrow’s reputation.

Today’s betrayal can become tomorrow’s retaliation.

Today’s cooperation can become tomorrow’s trust.

And that brings us to the next mental model:

The Shadow of the Future.

What changes today because we expect to deal with each other again tomorrow?

PART III

Think Beyond the Moment

The Shadow of the Future

In the last chapter, we saw that a commitment can change today's behavior by changing what becomes possible tomorrow.

Now we're going to look at the other side of time.

Sometimes you don't have to eliminate any future option.

The mere fact that **there is a future** can change the game.

Imagine you're buying something from a street vendor in a city you'll never visit again.

Now imagine you're buying the same thing from the store around the corner from your house—a store you expect to visit every week for the next twenty years.

Same transaction.

Same product.

Same price.

But not necessarily the same game.

In the first interaction, both of you may reasonably think:

This is probably the only time we'll ever deal with each other.

In the second:

Whatever happens today may affect what happens next week.

Your behavior today can influence how they treat you tomorrow.

Their behavior today can influence whether you ever return.

Trust matters.

Reputation matters.

Retaliation matters.

Repair matters.

The future has entered the payoff.

Game theorists sometimes describe this as **the shadow of the future**.

And once that shadow becomes long enough, behavior that makes little sense in a one-time interaction can make a great deal of sense in a relationship that continues.

The One-Time Game

Let's return to the Prisoner's Dilemma.

In the basic one-shot version, two players each choose whether to cooperate or defect.

The payoff structure makes defection individually attractive.

If you cooperate, I can do better by defecting.

If you defect, I can protect myself by defecting.

So we can end up with:

defect / defect

even though:

cooperate / cooperate

would leave both of us better off.

We've already used this game to understand payoffs and equilibrium.

But now change one thing.

Instead of playing once, suppose we are going to play again tomorrow.

And the day after that.

And perhaps again after that.

Now my choice today doesn't only determine today's payoff.

It can change:

what you expect from me,
how you treat me next time,
whether you trust me,
whether you punish me,
whether you cooperate with me again.

The immediate payoff hasn't necessarily changed.

The **continuation of the relationship** has.

And that can completely change the strategy.

Tomorrow Changes Today

Imagine we're business partners.

There is \$10,000 on the table.

I discover a loophole that would allow me to take an extra \$1,000 at your expense.

If this is the last interaction we'll ever have, I might evaluate the choice mainly as:

\$1,000 versus \$0.

But suppose we expect to do millions of dollars of business together over the next ten years.

Now taking the \$1,000 may also produce:

less trust,
more monitoring,
tighter contracts,
fewer opportunities,
resentment,
retaliation,
or the end of the partnership.

The extra \$1,000 is no longer just worth \$1,000.

It comes bundled with possible future consequences.

This is one of the central ideas of repeated games:

Today's move can change the value of tomorrow's game.

Experimental Evidence: The Future Really Does Matter

Pedro Dal Bó tested this idea experimentally.

Participants played repeated Prisoner's Dilemma games in the laboratory. The probability that another round would occur was varied, which allowed him to change how long the "shadow of the future" was expected to be.

The result was clear: a longer expected future significantly reduced opportunistic behavior, and the effect broadly followed the direction predicted by repeated-game theory.⁵⁵

That's important because it allows us to make a more precise claim than:

"People are nicer when they know they'll see you again."

The strategic claim is:

When future interaction becomes more likely, the future consequences of current behavior become more important, and that can make cooperation more attractive.

But there's an important word in that sentence:

can.

Repeated interaction does not automatically produce cooperation.

We'll come back to that.

Cooperation Becomes Possible for a Different Reason

When people talk about cooperation, they often explain it in terms of personality.

"She's cooperative."

"He's selfish."

"They're trustworthy."

Those things may matter.

But repeated-game theory shows us another possibility.

55. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2F000282805775014434>

Sometimes cooperation isn't just a personality trait.

It is a **strategic response to the structure of the relationship**.

Suppose I could make \$100 by betraying you today.

But doing so causes me to lose \$1,000 of future business with you.

Even if I cared only about my own payoff, cooperation could make sense.

That doesn't make the cooperation fake.

It simply means there is more than one path to cooperative behavior.

Someone might cooperate because of:

generosity,
morality,
loyalty,
friendship,

or because:

future consequences make cooperation strategically valuable.

Real human behavior may contain several of those motivations at once.

Axelrod and the Evolution of Cooperation

Robert Axelrod became famous for exploring this problem through computer tournaments built around the **iterated Prisoner's Dilemma**.

He invited researchers to submit strategies that would repeatedly play against one another.

Some of the strategies were complicated.

They tried to model opponents.

They used statistical inference.

They attempted to exploit patterns.

Yet the strategy that performed best in Axelrod's first tournament was remarkably simple:

Tit for Tat.

It began by cooperating.

After that, it did whatever the other player had done on the previous round.

If you cooperated, it cooperated.

If you defected, it defected.

When Axelrod ran a larger second tournament with 62 entries, Tit for Tat won again. Axelrod and William Hamilton later analyzed how reciprocal cooperation could become stable when players had a sufficiently high probability of interacting again.⁵⁶

There's something beautiful about the result.

Tit for Tat wasn't successful because it dominated everyone.

It couldn't exploit a cooperative player and then keep getting away with it.

It succeeded partly because it created an environment where cooperation could continue while still responding to exploitation.

But we need to be careful here.

56. University of Michigan Public Websites, <https://public.websites.umich.edu/~axe/research/Axelrod%20and%20Hamilton%20EC%201981.pdf>

The lesson is **not**:

“Tit for Tat is the universally best strategy for life.”

It isn’t.

Its success depends on the environment, the other strategies present, the payoff structure, how mistakes are handled, and whether interactions continue.

Axelrod himself and later researchers explored limitations of simple Tit for Tat, especially when misunderstandings or errors enter the game.⁵⁷

The broader lesson is more useful:

Once people expect a future together, reciprocity becomes strategically possible.

You Can Respond Tomorrow to What Happened Today

That’s really what repetition gives us.

In a one-shot game, today’s move is today’s move.

In a repeated game:

today becomes history.

And history can affect tomorrow.

If you cooperate with me:

I can reward that cooperation later.

If you exploit me:

I can protect myself later.

If you keep your promise:

57. University of Michigan Public Websites, <https://public.websites.umich.edu/~axe/research/Axelrod%20Dion%20Further%20EC%20Science%201988.pdf>

I can trust your next promise more.
If you repeatedly break promises:

I can stop relying on them.
Repeated interaction creates memory.
And memory creates consequences.

Reputation Is a Prediction

We tend to think of reputation as something vague:

“He has a good reputation.”
But strategically, reputation serves a practical function.
It helps other people answer:

What is this person likely to do next?

If you’ve kept your promises for twenty years, that history becomes evidence.

If you’ve broken every agreement you’ve entered, that becomes evidence too.

Reputation is not perfect.

People can change.

Reputations can be wrong.

A good reputation can be manufactured.

A bad reputation can be unfair.

But when people have incomplete information about how you will behave, past behavior can influence their expectations.

In that sense:

Reputation is accumulated information about future behavior.

That's why reputation can become an economic and relational asset.

Trust Changes the Cost of Interaction

Imagine doing business with someone you completely distrust.

You may want:

- detailed contracts,
- deposits,
- verification,
- monitoring,
- documentation,
- third parties,
- constant follow-up.

Now imagine doing business with someone who has demonstrated reliability across hundreds of interactions.

Perhaps you still use contracts.

But many of the transaction costs fall.

You don't need to verify every small statement.

You don't need to protect yourself against every possible betrayal.

Trust makes interaction easier.

That's not merely sentimental.

It can have strategic value.

The more reliable my behavior has been, the less uncertainty you face in dealing with me.

And that can make cooperation cheaper for both of us.

Today's Response Teaches Tomorrow's Expectation

Remember our manager from earlier chapters?

The manager tells employees:

“Bring me bad news early.”

An employee discovers a serious problem and immediately reports it.

The manager becomes furious.

They humiliate the employee.

They blame them for the problem.

What happens next month when another employee discovers bad news?

The manager may be surprised that nobody says anything.

But the employees are not simply responding to the manager's words.

They're responding to the history of the game.

The manager said:

Report problems.

The manager's behavior said:

Reporting problems is dangerous.

Which message will shape future behavior?

Probably the second.

That's the shadow of the future.

Today's response is teaching tomorrow's players what to expect.

You Are Always Training the Repeated Game

This applies to far more than leadership.

Imagine a child asks for something.

You say no.

They complain.

You still say no.

They escalate.

Eventually you give in.

What happened?

You solved the immediate problem.

The tantrum stopped.

But you may also have changed the next interaction.

The child has learned:

Escalation can work.

You optimized for this round.

You may have made the next round harder.

Or imagine the opposite.

Someone makes an honest mistake.

You punish them mercilessly.

They learn:

Never admit mistakes.

Again, you may feel like you dealt with today's problem.

But you also shaped tomorrow's incentives.

This gives us a powerful question:

What am I teaching this person to expect from me next time?

Winning the Round Can Lose the Game

This is one of the most important practical lessons of repeated-game thinking.

Suppose you're arguing with your spouse.

They once told you something painful and deeply personal.

During the argument, you realize:

I can use that against them right now.

And it would work.

You could end the argument.

You could embarrass them.

You could make them stop talking.

Maybe you "win."

But what game exists tomorrow?

Perhaps they learn:

Vulnerability is dangerous with you.

So next time they hide something.

They disclose less.

They trust less.

The immediate move succeeded.

The repeated-game strategy failed.

This is why strategic thinking sometimes requires resisting a short-term win.

The Long Game Is Not “Always Be Nice”

This is important.

Repeated-game theory does **not** teach:

“Always cooperate.”

If you cooperate unconditionally with someone who repeatedly exploits you, you may simply make exploitation profitable.

The possibility of future interaction creates room for:

reward,
punishment,
reciprocity,
boundaries,
repair.

Axelrod’s Tit for Tat was cooperative, but it was not passive.

It responded to defection.

That mattered.

A repeated relationship can make cooperation valuable partly because players understand that exploitation may have future consequences.

So:

Thinking long-term does not mean refusing to protect yourself.

Sometimes a firm boundary preserves the long-term relationship better than endless accommodation.

Boundaries Can Protect Cooperation

This may seem paradoxical.

People often treat boundaries as the opposite of cooperation.

But imagine a relationship where one person repeatedly takes advantage of the other.

If the cooperative person never responds, exploitation becomes profitable.

Eventually resentment builds.

The relationship may collapse.

A credible boundary can change the payoff:

“I want to continue cooperating with you, but this behavior changes what I will do next.”

That can make a healthy repeated relationship more sustainable.

Again:

cooperation does not require defenselessness.

Sometimes **conditional cooperation** is more durable than unconditional cooperation.

But Punishment Can Create Its Own Problem

Now we encounter another difficulty.

Suppose you betray me accidentally.

I think it was intentional.

So I punish you.

You believe **my punishment** was unprovoked.

You punish me.

Now I see your punishment as further betrayal.

I punish you again.

What began as a mistake has become a feud.

This is where real human relationships become much harder than clean models.

We don't observe one another perfectly.

We misunderstand.

We forget.

We mishear.

We misread motives.

There is **noise**.

And noise changes which strategies work well.

The Problem of Noise

Masaki Aoyagi, V. Bhaskar, and Guillaume Fréchette experimentally compared repeated Prisoner's Dilemma games under different kinds of monitoring.

In one condition, players could perfectly observe what had happened.

In others, past actions were observed imperfectly.

Subjects sustained cooperation under all three conditions, but their strategies changed when monitoring became noisy. Under imperfect monitoring, players used strategies that

were **more complex and more lenient**, helping reduce the damage caused by observational errors.⁵⁸

That is a deeply useful finding.

It tells us that when misunderstandings are possible, a completely unforgiving strategy can be fragile.

If every apparent violation produces permanent retaliation, one mistake can destroy the relationship.

The Strategic Value of Repair

This gives us a concept that matters enormously in real relationships:

repair.

Suppose I fail you.

If there is no route back to cooperation, then once the relationship breaks, it stays broken.

But if the game contains a way to distinguish:

mistakes
from

exploitation,
and a way to restore cooperation after a mistake, the relationship can become more robust.

That might involve:

apology,
clarification,
restitution,
forgiveness,

58. AEA Papers and Proceedings, <https://pubs.aeaweb.org/doi/abs/10.1257/mic.20160304>

renewed commitment.

Those are not just “soft skills.”

In a repeated interaction with noise, they can perform an important strategic function:

They stop errors from becoming permanent retaliation loops.

Forgiveness Is Not Forgetting

We should be careful with that word.

Strategic forgiveness does not mean:

pretend nothing happened.

Nor does it mean:

allow someone to exploit you indefinitely.

It means:

Do not allow every deviation to make future cooperation impossible.

A good repeated-game strategy may need to do two things at once:

discourage exploitation

and

allow recovery from mistakes.

Too forgiving:

you become exploitable.

Too unforgiving:

you can become trapped in endless conflict.

That tension is very real.

The Feud

Think about two families who have hated one another for years.

Ask either side:

“Why?”

They may start describing something the other side did in response to something their side did in response to something that happened years earlier.

At some point, the original event barely matters.

The retaliation has become the game.

Each side sees its own behavior as:

justified response.

And the other side’s behavior as:

aggression.

This is a repeated-game trap.

The relationship now contains history.

And that history affects the interpretation of every new move.

History Changes Meaning

Suppose someone cancels dinner with you.

If they’ve been reliable for ten years, you might think:

“Something must have come up.”

If they’ve cancelled the last six times, you may think:

“Here we go again.”

Same action.

Different history.

Different belief.

Different likely response.

That's another way repetition changes communication.

Actions don't arrive alone.

They arrive attached to **everything that came before them.**

The End of the Game Changes Behavior

The future only matters if players believe there is a future.

This creates an interesting phenomenon.

Imagine someone has worked at a company for twenty years and plans to remain there another ten.

Reputation inside the company matters.

Now imagine the same person knows:

Friday is my final day.

Some future consequences disappear.

Perhaps they still behave ethically because of values, conscience, relationships, or broader reputation.

But certain incentives have changed.

The game has acquired an endpoint.

Formal game theory makes an important distinction between finitely repeated games and interactions whose ending is uncertain or potentially indefinite, because the expected continuation of the relationship can change what strategies are sustainable.⁵⁹

The practical lesson is straightforward:

**When people believe there will be no next round,
their incentives may change.**

The One-Time Customer

This happens in business all the time.

Imagine a restaurant located beside a highway that serves mostly travelers who will never return.

Compare that with a neighborhood restaurant whose customers live within a mile.

All else equal, the second restaurant has stronger repeated-game incentives.

Mistreat someone today and:

they may never return,
they may tell neighbors,
their family may stop coming.

That doesn't mean the highway restaurant will behave badly.

Morality, law, competition, reviews, and institutional reputation all matter.

59. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2F000282805775014434>

But structurally, the local restaurant has a longer direct shadow of the future.

Again:

The game shapes incentives.

Technology Can Extend the Future

This becomes especially interesting online.

Suppose you buy something from a stranger whom you'll never meet again.

At first glance, that looks like a one-shot game.

But then introduce:

- ratings,
- reviews,
- seller histories,
- public complaints,
- reputation scores.

Now your behavior toward this stranger can affect how **future strangers** treat you.

The repeated game has been extended beyond the same two individuals.

The platform has created memory.

That's what institutions can do.

They can turn behavior in one interaction into information for later interactions.

We will return to this kind of game redesign in the final chapter.

Cooperation Does Not Automatically Emerge

We need to put a firm boundary around this mental model.

It would be very easy to conclude:

“If people interact repeatedly, they’ll cooperate.”

That’s false.

Dal Bó and Guillaume Fréchette ran later experiments examining how cooperation develops as subjects gain experience in indefinitely repeated Prisoner’s Dilemma games.

They found that cooperation can prevail—but the conditions required were more demanding than simply having cooperation available as a theoretical equilibrium.⁶⁰

Their later survey of the experimental literature reached a similar conclusion.

Repeated interaction does make cooperation more likely under the right conditions. But simply showing that cooperation **could** be sustained in equilibrium does not mean most people will cooperate. High cooperation tends to emerge when cooperative strategies are sufficiently robust to what the researchers call **strategic uncertainty**—uncertainty about what the other player is going to do.⁶¹

That’s important.

The future creates **possibilities**.

It doesn’t guarantee outcomes.

60. AEA Papers and Proceedings, <https://pubs.aeaweb.org/doi/10.1257/aer.101.1.411>

61. AEA Papers and Proceedings, <https://pubs.aeaweb.org/doi/10.1257/jel.20160980>

A Long Relationship Can Also Sustain Bad Behavior

There's another reason we shouldn't romanticize repeated games.

The future can sustain cooperation.

But cooperation itself is morally neutral.

A cartel can cooperate.

A corrupt organization can cooperate.

Two people can cooperate to deceive someone else.

Members of a criminal group can maintain loyalty through repeated interaction.

So:

Cooperation is not the same thing as goodness.

Game theory tells us how patterns can become sustainable.

Ethics tells us whether those patterns are worth sustaining.

Our standard remains:

Truth. Autonomy. Dignity. Informed choice.

Reputation Can Trap You Too

A good reputation can create value.

But reputation can also create constraints.

Suppose you've built your identity around:

"I never back down."

Now every conflict creates pressure to maintain that reputation—even when backing down would be wise.

Or:

“I’m always the nice one.”

Now asserting a legitimate boundary may feel inconsistent with the reputation you’ve created.

Reputation changes how others expect you to behave.

But sometimes it also changes how **you feel obligated to behave.**

So the strategic question isn’t simply:

How do I build a strong reputation?

It’s:

What reputation do I want to be bound by?

That’s a much more interesting question.

Reputation and Predictability

This connects back to Chapter 4.

We learned that predictability can either create trust or make you exploitable.

Reputation is essentially a form of strategic predictability.

If your reputation is:

keeps promises,

that’s useful predictability.

If your reputation is:

always gives in eventually,

that’s exploitable predictability.

If your reputation is:

retaliates wildly against any perceived offense,

people may avoid dealing with you altogether.

So don't merely ask:

“What do people think of me?”

Ask:

“What future behavior do they expect from me?”

That is the strategically important part of reputation.

The Value of Being Easy to Deal With

Here's a consequence people often underestimate.

Suppose you are known as someone who:

- keeps agreements,
- communicates clearly,
- handles disagreements fairly,
- does not exploit mistakes,
- owns your own mistakes,
- responds firmly but proportionately to genuine exploitation.

Over time, other people may become more willing to:

- trust you,
- introduce you to others,
- enter deals with you,
- share information,
- cooperate without excessive protection.

You have increased the value of future games.

That's what a good strategic reputation can do.

It doesn't require manipulation.

Quite the opposite.

Often the most valuable reputation is simply:

This person is safe to play repeated games with.

The Cost of Being Hard to Deal With

Now consider the reverse.

You squeeze every penny out of every transaction.

You exploit every mistake.

You change terms whenever you can get away with it.

You use private information opportunistically.

You “win” frequently.

But eventually people learn:

Every interaction with this person requires protection.

Your counterparties begin:

adding contract language,
requiring deposits,
documenting everything,
withholding information,
increasing prices to compensate for risk,
or refusing to interact with you.

You optimized every round.

And made the game itself less valuable.

That’s another version of:

winning the battle and losing the war.

What Are You Teaching Them?

This may be the practical heart of the chapter.

Every important repeated interaction contains a learning process.

People learn:

What happens when I tell you the truth?

What happens when I disagree?

What happens when I say no?

What happens when I make a mistake?

What happens when I keep my word?

What happens when I break it?

What happens when I cooperate?

What happens when I exploit you?

You are answering those questions through your behavior whether you intend to or not.

So ask:

What am I teaching this person to expect from me?

That question applies to:

leaders,

parents,

partners,

customers,

employees,

friends,

spouses.

And yourself.

Don't Judge the Strategy by One Round

Suppose your strategy loses today.

Does that make it bad?

Not necessarily.

You refuse to exploit someone and leave \$500 on the table.

But they later send you \$50,000 of business.

You enforce a boundary and create an uncomfortable conversation.

But the relationship improves afterward.

You admit a mistake and lose face for ten minutes.

But gain credibility that lasts years.

The repeated-game lens teaches us:

A strategy should sometimes be evaluated across the relationship, not the incident.

That's a huge mental shift.

But Don't Use "The Long Term" to Justify Endless Losses

There's an opposite mistake.

Someone treats you badly.

You tell yourself:

"I'm thinking long term."

They do it again.

And again.

And again.

At some point, the long-run data are telling you something.

A repeated game gives you opportunities to learn.

If the other player repeatedly exploits cooperation and never reciprocates, the strategic lesson may not be:

cooperate harder.

It may be:

change your strategy, change the rules, or leave the game.

The shadow of the future matters only if the future game is worth playing.

Sometimes Exit Is the Long-Term Strategy

Imagine a business relationship that repeatedly creates:

broken promises,
distrust,
escalating monitoring,
conflict,
poor outcomes.

You could continue optimizing each new round.

Or you could eventually ask:

Why am I still playing this repeated game?

That's an important strategic question.

The existence of a future does not obligate you to preserve the relationship.

Sometimes the most rational response to a repeated bad game is to end the repetition.

We'll come back to that in **Change the Game**.

The Repeated-Game Audit

When an interaction is likely to continue, don't evaluate it as though it ends today.

Ask yourself:

Is there likely to be a next round?

If there is, how valuable is that future interaction?

Then ask:

What is today's move teaching the other person to expect from me?

What am I learning about them?

What reputation are we each building?

What behavior am I rewarding?

What behavior am I making costly?

If a mistake occurs, is there a path back to cooperation?

And finally:

What happens if we keep playing this way for the next hundred rounds?

That last question can reveal things the immediate interaction hides.

The Mental Shift

Earlier in the book, we learned to ask:

Then what?

That question helped us anticipate the next response.

Now we're expanding the horizon.

Don't only ask:

“What happens after this move?”

Ask:

“What happens next time?”

That's a different level of strategic thinking.

Because a choice can be excellent in a one-shot game and terrible in a repeated one.

And the reverse can also be true.

Think in Relationships, Not Just Transactions

Perhaps the simplest way to carry this model into life is this:

Before making an important move, determine whether you're dealing with:

a transaction

or

a relationship.

A transaction may end when the exchange ends.

A relationship creates another round.

And another.

And another.

Most important parts of life are repeated games:

Your marriage.

Your children.

Your closest friends.

Your employees.

Your employer.

Your business partners.

Your community.

Even your reputation among people you haven't met yet can create future rounds.

That doesn't mean every interaction should be calculated.

It means the future deserves a place in the decision.

The Question to Carry With You

When you're tempted to maximize the immediate outcome, ask:

What changes because there is a next time?

Then ask:

What am I teaching them to expect from me?

What am I learning to expect from them?

If this behavior repeats, where does it lead?

**Am I preserving the possibility of cooperation
without making exploitation profitable?**

**If something goes wrong, is there a way to repair
the game?**

Because sometimes the best move today is not the move that gives you the most today.

It's the move that makes tomorrow's game better.

And once we start thinking about tomorrow, another problem emerges.

Sometimes you and the other player don't actually want opposing things.

You would both be better off if you could act together.

But neither of you knows what the other is going to do.

Nobody wants to go first.

Everyone is waiting for everyone else.

And the bad outcome happens even though nobody wanted it.

That's not necessarily conflict.

It may be something else entirely.

Diagnose the Coordination Problem.

Are we actually fighting—or are we simply failing to get on the same page?

Diagnose the Coordination Problem

In the last chapter, we learned to think about the future.

We saw that when people expect to interact again, today's behavior can affect tomorrow's trust, reputation, retaliation, and cooperation.

Now we need to deal with a different kind of problem.

Sometimes two people end up with a bad outcome even though neither of them wanted it.

Not because one person betrayed the other.

Not because their interests are fundamentally opposed.

Not because one of them failed to communicate persuasively enough.

Sometimes the problem is much simpler:

They failed to coordinate.

And that distinction matters.

Because if you misdiagnose a coordination problem as a conflict problem, you can spend enormous amounts of time trying to persuade someone who already wants essentially the same thing you do.

Conflict or Coordination?

Imagine you and I are trying to meet for dinner.

There are two restaurants.

We both prefer eating together to eating alone.

I slightly prefer the Italian restaurant.

You slightly prefer the steakhouse.

What is the most important thing?

Probably not getting our first choice.

It's:

ending up at the same restaurant.

If I go to the Italian restaurant and you go to the steakhouse, we both lose.

That's a coordination problem.

Now compare that with a situation where there is only one restaurant reservation available and we both want it for separate events.

Now our interests are directly opposed.

Same word:

disagreement.

Very different game.

And that's the first thing I want you to learn from this chapter:

Not every disagreement is a conflict of interests.

Sometimes the people already want compatible outcomes.

They just haven't aligned their actions.

What Is a Coordination Game?

A coordination game is a situation in which players benefit when their actions fit together.

In a pure coordination game, the specific coordinated outcome may matter less than the fact that everyone chooses the same one.

Driving provides an obvious example.

You could build a functioning traffic system where everyone drives on the right.

You could also build one where everyone drives on the left.

The disaster happens when half the people choose one convention and half choose the other.

The Stanford Encyclopedia of Philosophy uses this basic structure to explain coordination games: players do well when they select compatible actions, and pure coordination games can contain multiple equilibria.⁶²

The question isn't necessarily:

Which option is objectively better?

It may be:

62. Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/archives/sum2026/entries/game-theory/>

Which option can we all reliably choose together?

That's a different kind of strategic problem.

The Meeting Problem

Thomas Schelling made this idea famous with a wonderfully simple question.

Imagine you and another person are told:

“Meet somewhere in New York City tomorrow.”

You cannot communicate.

Where do you go?

There are thousands of possible locations.

And yet people often don't choose randomly.

They look for something that seems **obvious**.

A famous landmark.

A central station.

A conventional time.

They ask themselves:

What place will the other person think I will think is obvious?

Schelling called these solutions **focal points**.

His experiments suggested that people can often coordinate at rates far above random chance by using shared culture, context, salience, and common frames of reference. The Nobel committee specifically highlighted this aspect of his work.⁶³

63. Nobel Prize, <https://www.nobelprize.org/uploads/2018/06/popular-economicsciences2005.pdf>

That is a fascinating result because the players don't need to prefer the focal point.

They merely need to believe:

The other person is likely to recognize it too.

Salience Can Solve a Strategic Problem

Imagine I say:

"Pick a positive whole number. Another person will independently pick one too. If you both choose the same number, you win."

What would you pick?

Many people choose:

1.

Why?

There are infinitely many positive integers.

One isn't mathematically superior for the purpose of matching.

But it stands out.

It's the first.

It's simple.

And crucially:

you can imagine the other person thinking the same thing.

That shared sense of obviousness helps solve the coordination problem.

Later experiments by Judith Mehta, Chris Starmer, and Robert Sugden directly investigated focal points and salience in pure coordination games, providing experimental support for Schelling's insight.⁶⁴

This gives us a practical principle:

When several solutions would work, the easiest one for everyone to recognize may be strategically better than the theoretically perfect one.

The Best Solution May Not Be the Most Efficient Solution

This can be hard for highly analytical people.

Imagine a team is choosing a project-management system.

System A is slightly better.

System B is already used by everyone in the company.

Which should they choose?

You could spend weeks proving that A has better features.

But if half the organization stays on B, you've created a coordination failure.

Sometimes:

a slightly inferior standard used by everyone
is more valuable than:

a superior standard used by half the group.
That does not mean quality never matters.

It means coordination itself has value.

64. Research Portal, <https://research-portal.uea.ac.uk/en/publications/focal-points-in-pure-coordination-games-an-experimental-investiga/>

This Is Why Standards Matter

Think about:

language,

electrical outlets,

file formats,

traffic rules,

time zones,

meeting procedures,

measurement systems.

Often the value comes partly from everyone using the same standard.

Imagine inventing a brilliant new alphabet that is 15 percent more efficient than English.

If nobody else can read it, that efficiency isn't especially useful.

A convention solves a repeated coordination problem.

You no longer have to decide each morning:

“Which side of the road should we drive on today?”

The rule has already coordinated expectations.

Conventions Save Strategic Energy

This is something we often overlook.

Every time a group lacks a convention, people have to solve the same coordination problem again.

Who runs the meeting?

Who follows up?

Where do files go?

Who makes the final decision?

When does the meeting start?

How quickly should an email be answered?

If those questions remain ambiguous, everyone spends energy predicting everyone else.

A convention removes the ambiguity.

It says:

This is what we do here.

And once everyone expects everyone else to follow it, coordination becomes much easier.

That's why organizations create processes.

Families create routines.

Communities create customs.

And societies create institutions.

The Project Nobody Owned

Imagine a small company.

Everyone agrees:

“We need to follow up with that customer today.”

The meeting ends.

Tomorrow, someone asks:

“Did anyone call them?”

Silence.

The sales manager thought customer service would do it.

Customer service thought sales owned the relationship.

The owner thought the sales manager had assigned it.

Nobody opposed calling the customer.

Nobody was unwilling to do the work.

Nobody necessarily forgot.

The group simply failed to establish:

who was expected to act.

That's coordination failure.

And notice how easy it would be to misdiagnose.

The owner might say:

"Nobody takes initiative."

The employees might say:

"Management is disorganized."

Those might contain some truth.

But before making a character judgment, ask:

Was there actually a mutually understood assignment?

That's a different question.

Clarity Can Be a Strategic Intervention

Sometimes the solution to a human problem is remarkably boring.

Not motivation.

Not persuasion.

Not incentives.

Just:

Paul owns this by 3 PM.

That one sentence can change the game.

Why?

Because now:

Paul knows he owns it.

Everyone else knows Paul owns it.

Paul knows everyone else knows he owns it.

The expectations have become aligned.

The ambiguity is gone.

A surprising amount of organizational dysfunction comes from people having compatible goals but incompatible expectations.

Communication Can Coordinate Even When It Doesn't Bind

This gives us another reason communication can matter.

Remember **cheap talk** from Chapter 7?

We learned that costless messages aren't necessarily credible evidence.

But communication can serve another function entirely.

It can help players coordinate.

Russell Cooper, Douglas DeJong, Robert Forsythe, and Thomas Ross experimentally studied coordination games where participants could communicate before acting.

The messages were nonbinding.

Yet communication often substantially improved coordination. In some conditions, one-way communication increased selection of a Pareto-superior equilibrium; in another type of coordination game, two-way communication consistently produced the Pareto-dominant equilibrium. Communication was powerful—but it did not eliminate every failure in every environment.⁶⁵

That's important.

Cheap talk may be weak as:

proof.

But powerful as:

coordination.

Those are different jobs.

“Just Talk About It” Works Better for Some Problems Than Others

This explains why communication can feel miraculous in one situation and useless in another.

Imagine two people both want to meet at the airport.

One thinks:

Terminal A.

The other thinks:

Terminal B.

A thirty-second conversation can solve the entire problem.

65. OUP Academic, <https://academic.oup.com/qje/article-abstract/107/2/739/1838346>

Now imagine two people arguing over who receives a \$100,000 payment.

Talking may clarify the disagreement.

It doesn't make the conflict disappear.

So before saying:

“We need better communication,”

ask:

What is communication supposed to accomplish?

Is the problem:

missing information?

credibility?

persuasion?

bargaining?

coordination?

The answer matters.

Communication in the Real World

One of the strongest pieces of evidence for this chapter comes from a study involving agricultural cooperatives in Senegal.

Researchers worked with members of **79 groundnut-producing cooperatives**. They first conducted coordination games with **839 farmers**, then ran a natural field experiment among roughly **2,800 cooperative members**.

The farmers generally understood the advantages of selling collectively through the cooperative.

But there was a strategic problem.

A farmer's willingness to participate depended partly on whether they believed enough **other farmers** would participate too.

The researchers asked members how much they intended to sell through the cooperative and then revealed the group's aggregate intentions.

That relatively simple communication intervention improved collective commercialization. In the field experiment, treated farmers contributed more than twice as much toward collective commercialization on average as the control group, and the effect was larger in bigger groups where coordination was more difficult.⁶⁶

Think about what happened.

The farmers didn't suddenly discover a new financial incentive.

The basic opportunity was already there.

What changed was their information about:

what everyone else intended to do.

That changed the game.

Strategic Uncertainty

There is a useful term for the problem those farmers faced:

strategic uncertainty.

Ordinary uncertainty is:

"Will it rain?"

Strategic uncertainty is:

66. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0304387821001164>

“What are the other players going to do?”

Suppose your company introduces a new software system.

Everyone thinks it could be better **if everyone uses it**.

But you aren't sure whether your coworkers will adopt it.

Switching alone would make your work harder because you'd still need to interact with everyone using the old system.

So you wait.

Your coworker is thinking exactly the same thing.

They wait.

Everyone waits.

Months later management concludes:

“Employees resist change.”

Maybe.

But perhaps every employee was individually willing to change **conditional on everyone else changing too**.

The barrier wasn't motivation.

It was strategic uncertainty.

Sometimes People Aren't Unwilling to Cooperate

They're unwilling to cooperate **alone**.

That distinction is huge.

Imagine five departments all need to switch to a new reporting procedure.

Every department prefers the new procedure if all five use it.

But switching first imposes extra work.

So each one waits for proof that the others are moving.

Now suppose management publicly announces:

“Every department switches Monday.”

And everyone can see the migration plan.

The strategic calculation changes.

Not because management persuaded anyone that the new system is good.

Because the uncertainty about what everyone else will do has been reduced.

That’s coordination.

Shared Goals Are Not Enough

This may seem surprising.

You would think that if everybody wants the same result, cooperation should be easy.

But the Senegal experiment shows why that isn’t necessarily true.

The farmers broadly saw benefits from collective commercialization, yet beliefs about whether others would participate created a barrier.⁶⁷

A shared goal tells you:

where people want to end up.

It does not necessarily tell them:

how everyone is going to get there.

67. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0304387821001164>

That's what coordination solves.

But Be Careful: Shared Goals Don't Always Mean Coordination

This distinction became important when we audited the framework.

Suppose four roommates all prefer:

a clean apartment.

At first glance, that sounds like a coordination problem.

But imagine each roommate thinks:

"The apartment should be clean, but I'd rather somebody else do the cleaning."

Now the structure is different.

If everyone else cleans, your personal best outcome may be:

enjoy the clean apartment without doing any work.

That's a **free-rider problem**.

A public-goods problem.

The group shares a preference for the final outcome, but the individual incentive still pushes each player toward contributing less.

Experimental public-goods research repeatedly finds that cooperation can erode through conditional cooperation and free riding even when participants benefit collectively from contribution.

That's not the same

No problem—you did. I'll pick up exactly where Chapter 10 cut off.

• • •

The group shares a preference for the final outcome, but the individual incentive still pushes each person toward doing less and hoping others do more.

That is **not** a pure coordination problem.

It's a **social dilemma**.

Public-goods experiments are designed around exactly this structure: the group benefits when people contribute, but each individual can often do better personally by keeping more for themselves while benefiting from what everyone else contributes. Research consistently finds that this tension can produce free riding and fragile cooperation.⁶⁸

This distinction matters because the solution changes.

If the problem is coordination, you may need:

- clearer expectations,
- shared information,
- a focal point,
- an agreed standard.

If the problem is free riding, clarity may not be enough.

Everyone may know exactly what they're supposed to do and still have an incentive not to do it.

That may require:

- changing incentives,
- monitoring,
- reputation,
- repeated interaction,

68. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2Faer.100.1.541>

commitments,
or changing the rules.

Which is precisely why diagnosing the game comes before choosing the intervention.

A Simple Test

Here's a useful way to distinguish the two.

Ask:

If I knew for certain that everyone else was going to cooperate, would I want to cooperate too?

If the answer is:

Yes, absolutely,

you may be dealing primarily with coordination.

If the answer is:

I'd still be better off letting everyone else do the work,

you have a social-dilemma or free-rider problem.

That's a very different game.

Consider the roommates again.

Everyone wants the apartment clean.

But if three roommates are definitely going to clean it, the fourth may think:

"Great. Then I don't have to."

That's not failure to coordinate.

They coordinated just fine.

The individual incentive is the problem.

Experimental public-goods research illustrates how this can happen even when many participants are willing to cooperate conditionally. Fischbacher and Gächter found that imperfect conditional cooperation can produce declining contribution over time and eventually widespread free riding even though most participants were not simply behaving as pure selfish maximizers.⁶⁹

Again:

Same desired group outcome. Different strategic problem.

Mixed Coordination and Conflict

Real interactions are often even messier.

They contain coordination **and** conflict simultaneously.

The classic game-theory example is the **Battle of the Sexes**.

The original name is dated, but the structure is useful.

Imagine two people want to spend the evening together.

That's aligned.

But one prefers the concert.

The other prefers the football game.

That's conflict.

Their worst outcome may be:

each going to a different event.

So they need to coordinate.

69. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2Faer.100.1.541>

But they disagree over **which equilibrium** they should coordinate on.

That is a mixed-motive game.

And it looks a lot like ordinary life.

Two founders both want the company to succeed.

But each wants their own strategy adopted.

Two parents both want consistency.

But each prefers a different parenting rule.

Two departments both want the project finished.

But each wants the other department to absorb more of the workload.

This is why saying:

“We all want the same thing”

doesn’t settle the strategic problem.

You need to ask:

Where exactly do our interests align, and where do they diverge?

Common Knowledge

Now we need to go one level deeper.

Sometimes everyone possesses the necessary information and coordination still fails.

Why?

Because there’s a difference between:

everyone knowing something

and

everyone knowing that everyone knows it.

Game theorists call the stronger condition **common knowledge**.

Ariel Rubinstein's famous Electronic Mail Game demonstrated just how sharp this distinction can be. In his theoretical example, players could acquire arbitrarily many levels of:

"I know that you know that I know..."

and yet the strategic situation remained very different from one in which the relevant fact was genuinely common knowledge.⁷⁰

That sounds abstract.

So let's make it ordinary.

The 3:00 Meeting

Suppose a manager tells you privately:

"We're moving the meeting to 3:00."

You know the meeting is at 3:00.

But you don't know whether everyone else was told.

Should you show up at 3:00?

Probably.

But you may still wonder:

"What if everyone else comes at two?"

70. IDEAS/RePEc, <https://ideas.repec.org/a/aea/aecrev/v79y1989i3p385-91.html>

Now imagine the manager announces in a room containing the whole team:

“Everyone, we’re moving the meeting to 3:00.”

You saw everyone hear it.

They saw you hear it.

Everyone can reasonably expect everyone else to act on the same information.

Nothing changed about the meeting time.

What changed was the **mutual awareness of the information**.

That’s strategically important.

Public Information Can Do More Than Private Information

This gives us a useful insight.

Suppose a company tells every employee privately:

“We want everyone to use the new system beginning Monday.”

Technically, everyone may know.

But each employee might wonder:

“Did everyone else get the message?”

“Are they actually going to switch?”

Now imagine a public kickoff where the company announces:

“Every department switches Monday. Here’s the implementation plan. Here are the teams responsible. Here is what everyone else is doing.”

The organization hasn't necessarily changed anyone's preferences.

It's changed what people know about **what everyone else knows and intends to do**.

That can make coordination easier.

This is one reason public announcements, shared dashboards, visible schedules, rituals, and common procedures can have strategic value.

They're not just transmitting information.

They can create shared expectations.

Say It in Front of Everyone

Think about the difference between these two situations.

A manager individually tells five employees:

"You're all responsible for making this change."

Versus gathering all five and saying:

"You five are responsible for this change, and each of you now knows who else is responsible."

The second statement creates something the first doesn't.

Each person can now reason:

"They know I'm responsible."

"I know they're responsible."

"They know that I know they're responsible."

Again, we don't need to turn every office meeting into an exercise in epistemic game theory.

The point is simply:

Sometimes coordination requires making expectations mutually visible.

Focal Points Help When Nobody Can Talk

Common knowledge is one solution.

Focal points are another.

Remember the problem:

several outcomes would work,

but:

nobody knows which one everyone else will choose.

A focal point gives players something mutually recognizable to organize around.

Mehta, Starmer, and Sugden's experiments provided strong evidence that people do use salience in pure coordination games rather than choosing randomly among equivalent options.⁷¹

But focal points have limits.

Once players strongly disagree about which equilibrium they prefer, the thing that seems obvious may no longer be enough to coordinate them. Experimental studies of coordination with asymmetric payoffs find lower coordination rates than in pure coordination settings.⁷²

Again, the lesson is:

Diagnose first.

A focal point can help solve uncertainty.

71. Research Portal, <https://research-portal.uea.ac.uk/en/publications/focal-points-in-pure-coordination-games-an-experimental-investiga/>

72. ScienceDirect, <https://www.sciencedirect.com/science/article/abs/pii/S016726811930349X>

It cannot magically erase incompatible interests.

Fairness Can Become a Coordination Device

Here's a particularly interesting example.

Suppose two people need to divide something.

There are several possible agreements that would benefit both.

Why is:

50/50

so common?

One explanation is fairness.

People may intrinsically prefer equal division.

But there's another possibility.

Fifty-fifty is also extremely easy for both sides to anticipate.

Experiments on tacit bargaining have found that equal divisions can function as powerful focal points—even in situations where an unequal allocation could produce more total value.⁷³

That's fascinating.

It means an outcome can attract people partly because:

they expect the other person to recognize it as obvious.

So fairness can operate both as:

a preference,

and as:

73. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S2214804315001433>

a coordination device.

We don't need to choose only one explanation.

Human behavior can contain both.

"Everyone Knows What to Do" Is Often False

Managers say this constantly.

"They know what they need to do."

Parents say it.

"He knows the rules."

Partners say it.

"She knows what I expect."

Maybe.

But the strategic question is not:

Have I said it?

It's:

Do we actually share the same expectation about what happens next?

Those aren't the same thing.

You may believe:

"Everyone knows the deadline means the work needs to be finished Thursday."

Someone else believes:

"Thursday means that's when we review the draft."

Both heard:

Thursday.

They are not coordinated.

This is why ambiguity is so dangerous in interdependent work.

Coordination Often Fails at the Handoff

Consider a relay race.

The fastest runners in the world can still lose if they can't pass the baton.

Human systems work the same way.

A sales team does its job.

Then operations needs information.

Operations assumes sales sent it.

Sales assumes operations can see it in the system.

The information doesn't transfer.

The customer suffers.

Nobody necessarily failed at their individual task.

The **connection between tasks** failed.

That's a coordination problem.

So when a system breaks, don't just inspect:

each person's performance.

Inspect:

the handoffs.

Where does one person's action depend on another person's?

Where do expectations have to align?

Where can ambiguity create mismatches?

That's where coordination problems often hide.

Roles Are Coordination Technology

A job title can be more than a description.

It can help solve:

Who is supposed to do what?

A deadline can solve:

When are we supposed to act?

A meeting agenda can solve:

What are we coordinating around?

A chain of command can solve:

Who decides when disagreement remains?

Processes are not automatically good.

They can become bureaucratic.

But the underlying reason organizations create them is important:

Repeated coordination is costly.

If the same uncertainty appears every day, eventually someone creates a rule.

The rule becomes a shortcut.

Instead of everyone repeatedly solving the game:

“Who's going to do this?”

the convention answers it.

Coordination and Leadership

Leadership often involves less persuasion than people think.

Sometimes the leader's primary function is to make expectations clear enough that everyone can coordinate around them.

Imagine a team entering a crisis.

Ten people all want to help.

If all ten independently decide what to do, they may interfere with one another.

A leader says:

"You handle this."

"You call them."

"You stay here."

"Everyone meets back here at six."

The leader hasn't necessarily changed anyone's motivation.

They've created:

a shared coordination structure.

That's one reason leadership can become more valuable as group size and interdependence increase.

Communication Isn't Just About Persuasion

This is worth emphasizing because it connects directly to the larger thesis of this book.

We often treat communication as though its purpose is:

getting someone to agree with us.

But communication can serve many strategic functions.

It can:

transmit information,

test beliefs,

signal intent,
establish credibility,
make commitments,
coordinate expectations.

Those are not the same thing.

And treating every communication problem as persuasion is like treating every medical problem with the same medication.

Sometimes nobody needs convincing.

They just need to know:

What are we doing?

Who's doing it?

When are we doing it?

Does everyone else know the same thing?

Why Communication Helps—But Doesn't Always Solve the Problem

The experiments by Cooper and colleagues illustrate this nicely.

Nonbinding communication substantially improved coordination in several games, but it did not eliminate every coordination failure under every condition.⁷⁴

That's important.

Communication is not magic.

People can communicate and still:

misunderstand,
distrust,

74. OUP Academic, <https://academic.oup.com/qje/article-abstract/107/2/739/1838346>

disagree,
 prefer different equilibria,
 free ride,
 or fail to act.

The Senegal field experiment similarly shows how powerful information about others' intentions can be when strategic uncertainty is central to the problem. Revealing aggregate intentions increased collective commercialization, particularly where coordinating larger groups was harder.⁷⁵

But neither result justifies:

"Talk more and every group problem disappears."

The mechanism matters.

Communication helps when it resolves the uncertainty that's actually blocking coordination.

Sometimes the Best Choice Is Simply to Pick Something

This can be surprisingly difficult.

Two people have three equally acceptable options.

They spend twenty minutes debating which is best.

But the cost of failing to coordinate is far larger than the difference among the options.

At some point the strategic problem becomes:

Stop optimizing and coordinate.

This is one reason arbitrary rules sometimes work surprisingly well.

75. ScienceDirect, <https://www.sciencedirect.com/science/article/pii/S0304387821001164>

Alternate.

Flip a coin.

Rotate.

Use alphabetical order.

Let whoever cared most last time defer this time.

The rule doesn't need to reveal some profound truth.

It only needs to create:

a mutually understood method for selecting one equilibrium.

That's useful far beyond game theory.

The Cost of Constant Renegotiation

Imagine a couple has the same argument every week:

"Who's picking up the kids Thursday?"

Every week they renegotiate.

Every week there is confusion.

Every week someone gets irritated.

At some point, almost any workable convention might outperform repeated bargaining.

For example:

"Whoever doesn't do Tuesday handles Thursday."

The particular rule may not matter much.

What matters is:

everyone knows the rule before the conflict arises.

Conventions reduce transaction costs.

They save cognitive energy.

They reduce opportunities for misunderstanding.

And they create predictable expectations.

But Bad Conventions Can Become Stable Too

Of course, once a rule exists, people can become trapped in it.

An organization might say:

“We’ve always done it this way.”

Maybe the convention originally solved a coordination problem.

But the environment changed.

Now the convention creates inefficiency.

This is another example of equilibrium.

The fact that everyone coordinates around a behavior does not mean it’s the best behavior.

It may simply mean:

everyone expects everyone else to keep doing it.

That prepares us for Chapter 11.

Because sometimes the solution isn’t better coordination within the existing system.

It’s redesigning the system.

The Coordination Diagnostic

When a group or relationship is producing a bad outcome, ask:

Do we actually want incompatible things?

If yes, you may have a genuine conflict or bargaining problem.

If no, continue.

Would I want to take the desired action if I knew the others would too?

If yes, you may have a coordination problem.

If no—if you'd rather let everyone else bear the cost—you may have a social dilemma or free-rider problem.

What am I uncertain about?

Their actions?

Their intentions?

Their knowledge?

Their commitment?

What are they uncertain about?

Your intentions may be part of the problem.

What expectation needs to become mutually clear?

Who acts?

When?

According to what standard?

Is there a focal point or convention we can coordinate around?

Sometimes something simple solves the problem.

Does everyone know that everyone knows?

Would making the expectation public change the game?

That's the coordination lens.

The Mental Shift

Before this chapter, an unsuccessful interaction might look like:

“Why won't they cooperate?”

Now I want you to separate at least three possibilities.

Conflict

We genuinely want different outcomes.

Coordination failure

We could both benefit from compatible actions, but we aren't aligned on what each of us will do.

Social dilemma

We all like the collective outcome, but each of us has an individual incentive to let others carry more of the cost.

Those three situations can look remarkably similar from the outside.

But they are not the same game.

And they do not have the same solution.

That's why diagnosis matters.

The Question to Carry With You

When people seem unable to get on the same page, ask:

Are we actually in conflict—or are we failing to coordinate?

If it's coordination, ask:

What expectation needs to become mutually clear?

Do we need:

a shared plan?

a public commitment?

a clear owner?

a standard?

a deadline?

a focal point?

information about what everyone else intends to do?

Sometimes the most strategic thing you can do is surprisingly simple:

Make the next move obvious to everyone.

But what happens when you've done all of that and the same bad behavior keeps returning?

Everyone knows the rules.

Everyone understands the expectations.

Communication is clear.

And yet the system keeps producing outcomes nobody wants.

At that point, perhaps the problem is not the way people are playing.

Perhaps the problem is the **game itself**.

And that brings us to the final mental model:

Change the Game.

Do you need a better move—or a better game?

Change the Game

We've spent ten chapters learning how to see more clearly inside a strategic interaction.

We've learned to ask:

What game am I in?

What does everyone actually want?

How will they respond?

What am I teaching them to predict about me?

What pattern have we created?

What does everyone know and believe?

What makes a message credible?

What commitments are changing future choices?

What changes because there is a next time?

Are we actually in conflict—or failing to coordinate?

Those questions help us make better moves.

But eventually we reach a more powerful question:

Do I need a better move—or a better game?

Because sometimes you can become extraordinarily skilled at playing a game that should never have been designed that way in the first place.

The incentives are wrong.

The rules produce the behavior you're trying to stop.

The information is distributed badly.

Authority sits in the wrong place.

People are rewarded for competing when they need to cooperate.

Everyone has learned to game the metric.

And every time the same bad outcome appears, somebody calls another meeting and says:

“We need people to communicate better.”

Maybe.

Or maybe you've built a game that keeps producing exactly what it was built to produce.

Stop Asking Only How to Play Better

Imagine a sales organization.

Management keeps telling salespeople:

“We want everyone to collaborate.”

Share leads.

Help one another.

Teach each other what's working.

Don't compete internally.

But every month, the company publicly ranks each salesperson against the others.

Bonuses go only to the top performers.

Promotions are heavily influenced by individual ranking.

The lowest performers fear losing their jobs.

Then management is shocked that salespeople guard information.

So they hold a workshop.

They talk about:

teamwork,
company values,
abundance,
collaboration.

Nothing changes.

Why?

Because the game is saying something different from the speech.

Management says:

Cooperate.

The incentive system says:

Your coworkers are your competitors.

At some point, the strategic question is no longer:

“How can we persuade these people to collaborate?”

It’s:

“What behavior does our system actually reward?”

That’s the beginning of game design.

Mechanism Design

Most of the game theory we've explored so far asks something like:

Given the game, what are the players likely to do?

Mechanism design reverses the question.

It asks:

Given the outcome we want, how should we design the game?

This reversal became important enough that Leonid Hurwicz, Eric Maskin, and Roger Myerson received the 2007 Nobel Prize in Economic Sciences for laying the foundations of mechanism design theory. Hurwicz introduced key ideas including **incentive compatibility**; Maskin developed implementation theory; and Myerson extended mechanism design to situations involving private information.⁷⁶

We don't need the mathematics.

The mental shift is what matters.

Instead of:

Here are the rules. What should I do?

we ask:

Here is the outcome we want. What rules would make that outcome more likely?

That's a radically different way to think.

The Cake

Here's a simple example often used to illustrate the idea.

76. Nobel Prize, <https://www.nobelprize.org/prizes/economic-sciences/2007/9267-presentationstal-2007-4/>

Two children need to divide a cake.

Each wants the larger piece.

A parent could try to divide it perfectly.

Then both children may argue:

“Their piece is bigger.”

Or the parent can change the mechanism:

One child cuts. The other chooses first.

Suddenly, the child cutting the cake has an incentive to make the pieces as equal as possible.

Why?

Because if one piece is clearly larger, the other child gets to choose it.

The parent didn’t change the children’s personalities.

Didn’t give a lecture on fairness.

Didn’t attempt to make them less selfish.

The parent changed the rules so that the children’s incentives helped produce the desired outcome. The “I cut, you choose” example is commonly used to explain the mechanism-design approach of starting with a goal and designing rules around it.⁷⁷

That is the essence of this chapter:

Instead of constantly fighting incentives, sometimes redesign them.

77. Leonid Hurwicz, <https://leonidhurwicz.org/mechanism-design/>

Don't Fix the Player Before You Inspect the Game

This connects all the way back to Chapter 1.

We have a strong tendency to explain behavior through character.

Someone isn't cooperating:

"They're selfish."

Employees aren't taking initiative:

"They're lazy."

Customers aren't telling us the truth:

"They're dishonest."

Parents are manipulating a school-choice system:

"They're gaming it."

Maybe.

But strategic thinking asks one more question:

What does the game make attractive?

Because sometimes perfectly ordinary people behave differently when the rules change.

That's one of the great lessons of game theory.

Behavior doesn't come only from personality.

It emerges from the interaction between:

people and structure.

The System May Be Rewarding the Problem

Imagine a customer-service department.

Management says:

“Our number-one priority is solving the customer’s problem.”

But representatives are evaluated primarily on:

average call time.

What happens?

Shorter calls become valuable.

A complicated customer is now strategically expensive.

Employees may:

- rush,
- transfer calls,
- avoid complicated issues,
- end conversations before problems are fully resolved.

Then management complains:

“Why aren’t our people taking ownership?”

Perhaps because you’ve rewarded them for getting customers off the phone.

That’s not a communication problem.

That’s a game-design problem.

Incentives Speak Louder Than Mission Statements

An organization can say:

“Quality matters.”

But reward volume.

It can say:

“Think long-term.”

But measure only this quarter.

It can say:

“Experiment.”

But punish every failure.

It can say:

“Tell us the truth.”

But retaliate against bad news.

It can say:

“Work as a team.”

But reward people only for individual results.

When the stated value and the rewarded behavior conflict, don't be surprised when the reward wins.

This doesn't mean incentives explain everything.

People have values.

People sometimes do what is right even when it costs them.

People resist bad incentives.

But if you repeatedly create a payoff for a behavior, you shouldn't be shocked when more of that behavior appears.

Changing an Incentive Can Change Behavior

There is strong real-world evidence for this.

Researchers Adnan Khan, Asim Khwaja, and Benjamin Olken conducted a field experiment involving property-tax administration in Punjab, Pakistan.

In one study, **525 property-tax inspectors** participated in a system where performance rankings affected the order in which inspectors could choose future job postings.

Instead of simply telling bureaucrats:

“Work harder,”

the system connected performance with something many employees actually valued:

where they would be assigned to work.

The result was substantial: the mechanism increased annual tax-revenue growth by roughly **30 to 41 percent**.⁷⁸

The people hadn’t been replaced.

The work itself hadn’t disappeared.

The structure changed.

And behavior changed with it.

But Incentives Can Produce Exactly What You Measure

Now comes the warning.

Designing incentives is dangerous precisely because incentives work.

Suppose you tell a teacher:

“Your bonus depends on your students’ test scores.”

What do you want?

Probably:

better education.

But what did you actually reward?

higher scores on a particular measurement.

Those are not necessarily the same thing.

78. National Bureau of Economic Research, <https://www.nber.org/papers/w24383>

Paul Glewwe, Nauman Ilias, and Michael Kremer studied a teacher-incentive program in Kenyan primary schools in which rewards were tied to students' test performance.

Scores on the incentivized exams improved.

At first glance:

Success.

But researchers found little evidence that teacher attendance, homework assignment, or classroom pedagogy improved. Teachers increased test-preparation sessions, and score gains were concentrated on the exams linked to the incentives rather than unrelated tests. The improvements also did not persist as broader long-run learning gains.⁷⁹

The incentive worked.

Just not entirely in the way the designers hoped.

That gives us one of the most important questions in this chapter:

What exactly am I rewarding?

Not:

What do I intend to reward?

What behavior actually improves the metric?

People Optimize Against the Scoreboard

Imagine I tell my team:

"I want higher-quality leads."

Then I compensate them based entirely on:

number of leads generated.

79. National Bureau of Economic Research, <https://www.nber.org/papers/w9671>

What will tend to increase?

Probably:

number of leads.

Quality may fall.

Then I accuse the team of missing the point.

But the scoreboard told them what the point was.

This is why metrics are dangerous.

Once a measurement becomes tied to:

- money,
- status,
- promotion,
- punishment,
- reputation,

people begin optimizing against it.

And the measurement may stop being a good representation of the thing you originally cared about.

The Metric Is Not the Goal

This principle applies everywhere.

You want:

student learning.

You measure:

test scores.

You want:

healthy customers.

You measure:

number of appointments.

You want:

useful customer service.

You measure:

call duration.

You want:

good content.

You measure:

clicks.

You want:

productive employees.

You measure:

hours online.

You want:

a healthy relationship.

You measure:

who apologized first.

The measurement may be useful.

But the measurement is not necessarily the goal.

And the moment the metric affects payoffs, people may begin playing **the metric** instead of pursuing **the underlying objective**.

This Is Why Game Design Requires Second-Order Thinking

Remember Chapter 3:

Then what?

Mechanism design requires the same discipline.

You change a rule.

Then what?

People respond.

Then what?

They learn how the rule works.

Then what?

They adapt.

Then what?

They may discover a way to produce the metric without producing the outcome you wanted.

So game design isn't:

Set rule → get outcome.

It's:

Set rule → players respond → players learn → players adapt → system evolves.

Which means:

Designers are players too.

You don't stand outside the game controlling everything.

Once you create rules, everyone else begins responding strategically to them.

Don't Design for the Behavior You Wish People Had

Design for the behavior people are likely to have.

Suppose you create an expense policy that relies entirely on:

“Everyone will behave reasonably.”

Maybe they will.

But if abusing the policy produces large benefits with no meaningful cost, eventually somebody may exploit it.

Now everyone else watches.

If nothing happens, the behavior can spread.

A robust mechanism doesn't assume everyone is selfish.

But it also doesn't require everyone to be a saint.

That leads us to an important mechanism-design concept:

incentive compatibility.

Incentive Compatibility

At a high level, an incentive-compatible mechanism is designed so that following the desired strategy aligns with a participant's own incentives.

Hurwicz's work made incentive compatibility one of the foundational concepts of mechanism design.⁸⁰

The ideal isn't:

Force everyone to fight their incentives.

It's:

80. Nobel Prize, <https://www.nobelprize.org/prizes/economic-sciences/2007/9267-presentationstal-2007-4/>

Where possible, make the desirable action compatible with them.

That's what happened in the cake example.

The child cutting the cake doesn't have to become indifferent to getting the bigger piece.

The mechanism uses that preference.

Because the other child chooses first, the cutter's self-interest creates an incentive for fairness.

That is elegant game design.

Tell the Truth—or Game the System?

School choice provides a much more consequential example.

Boston once used a school-assignment system in which families could benefit from strategically ranking schools rather than simply listing their true preferences.

Sophisticated parents could try to game the mechanism.

Less sophisticated parents might simply report what they genuinely wanted—and could be disadvantaged by doing so.

Researchers Atila Abdulkadiroğlu, Parag Pathak, Alvin Roth, and Tayfun Sönmez documented evidence of both sophisticated and unsophisticated strategic behavior. Some families paid close attention to capacity constraints while others made mistakes that left them disproportionately unsatisfied.

In 2005, Boston replaced the old system with a deferred-acceptance mechanism designed to remove much of the incentive for this kind of strategic ranking.⁸¹

Think about the lesson.

You could tell parents:

“Don’t game the system.”

Or you can ask:

“Why does the system make gaming strategically useful?”

That second question is mechanism design.

Good Rules Can Reduce the Need for Strategy

This is one of the ideas I find most powerful in the entire framework.

The highest form of strategic thinking is not always becoming better at outsmarting everyone.

Sometimes it is:

designing the interaction so that people don’t need to outsmart one another in the first place.

Think about Boston school choice.

A system that rewards strategic misrepresentation forces parents to ask:

“Should I list the school I actually want—or list something else because of how the mechanism works?”

A strategy-proof mechanism tries to reduce that burden.

You want the system to make:

81. National Bureau of Economic Research, <https://www.nber.org/papers/w11965>

truthful preference reporting

a good strategy.

That's better than teaching every parent advanced game theory so they can survive the rules.

The design carries some of the strategic burden for them.

Sometimes the Best Game Is Easier to Play

This is worth carrying far beyond economics.

Imagine a workplace expense system.

Bad design:

forty ambiguous rules,
unclear exceptions,
managers inconsistently interpreting them,
employees constantly figuring out what they can get
away with.

Better design might make the acceptable behavior:

obvious,
easy,
consistent,
verifiable.

The more you can make the desired behavior the easy behavior, the less energy everyone has to spend strategically navigating the system.

This gives us a powerful principle:

Sometimes the highest form of strategy is removing the need for strategy.

Change the Rules

How do you change a game?

Sometimes the most obvious lever is:

the rules.

Imagine two departments constantly arguing over who owns a customer once the sale closes.

You can conduct conflict-resolution training every quarter.

Or you can establish:

“Sales owns the relationship until this event. Then ownership formally transfers to customer success through this handoff.”

You’ve changed the rules.

That doesn’t guarantee harmony.

But you’ve removed one source of strategic ambiguity.

Rules define:

what players can do,
what they cannot do,
and what happens when they do it.

Change those and you may change behavior.

Change the Payoffs

Sometimes the rules are fine.

The incentives aren’t.

Your employees compete destructively because only one can receive the bonus.

Change the compensation structure and you may change the game.

A vendor has no incentive to fix a problem after receiving full payment.

Change the payment schedule.

A customer has no reason to return equipment.

Introduce a refundable deposit.

Again, the point isn't:

Every human behavior should be bought.

It's:

Payoffs matter.

If the system heavily rewards one behavior while leadership asks for another, the contradiction deserves attention.

Change the Information

Sometimes players behave badly because they don't know enough.

Remember the farmers from Chapter 10.

Each person was uncertain about what everyone else intended to do.

Making aggregate intentions visible changed behavior.

So game design may involve changing:

who knows what,

when they know it,

and whether other people know that they know it.

You might create:

- dashboards,
- transparent criteria,
- public schedules,
- independent verification,
- shared reporting,
- visible commitments.

You haven't changed anyone's preferences.

You've changed the information structure.

And that can change the game.

Change the Timing

Timing can also be designed.

Suppose two people make decisions simultaneously.

Neither knows what the other will do.

Now suppose one person must act first and the other responds.

Different game.

Or imagine a purchasing process in which everyone negotiates privately.

Now change it to an auction.

Different game.

Or employees currently receive feedback once per year.

Change it to immediate feedback.

Different game.

Or a customer pays everything before delivery.

Change the payment stages.

Different game.

Timing determines:

- who knows what when,
- who moves first,
- who can respond,
- what commitments become possible.

So sometimes the problem isn't the players.

It's the sequence.

Change the Options

Suppose every conflict has only two available responses:

yes

or

no.

That structure creates a particular game.

Introduce:

- conditional approval,
- trial periods,
- phased implementation,
- partial agreements,
- third-party review,

and suddenly players have new possible moves.

This is one reason creative negotiation can be so powerful.

You aren't merely choosing more intelligently among existing options.

You're **creating new options**.

And new options can alter the equilibrium.

Change Who Decides

Authority is another design variable.

Suppose two business partners disagree constantly.

Every decision requires unanimous agreement.

That creates one game.

Now imagine certain decisions are assigned clearly:

Partner A decides marketing.

Partner B decides operations.

Different game.

Or both defer certain disputes to an independent expert.

Different game.

Or a board decides.

Different game.

Changing who has decision rights can change incentives, timing, information, and bargaining power all at once.

This doesn't mean centralized authority is always better.

It means:

Who gets to decide is part of the game.

Change the Players

Sometimes a third person transforms an interaction.

Two people cannot agree on the value of a property.

Bring in an independent appraiser.

Different game.

Two employees cannot resolve a conflict.

Introduce a mediator.

Different game.

A buyer doesn't trust a seller's claims.

Introduce an inspector.

Different game.

Remember Chapter 7.

A third party can alter credibility because they change the information structure.

Players are part of the architecture.

Change What Is Observable

Imagine employees know that quality matters, but nobody can see who produced poor work.

That's one incentive structure.

Now imagine defects can be traced back to the production stage where they occurred.

Different game.

Or imagine donations are completely private.

Then make contributions public.

Different game.

Or imagine everyone can see one another's progress toward a team goal.

Different game.

Visibility affects:

- reputation,
- accountability,
- signaling,
- reciprocity.

Again, changing information can change incentives without directly changing anyone's reward.

Change the Default

Even something as simple as a default can change the strategic environment.

If every project automatically continues unless someone explicitly stops it, inertia favors continuation.

If every project automatically expires unless someone affirmatively renews it, inertia favors termination.

Neither rule forces a decision.

But they structure what happens when no player acts.

Defaults are easy to underestimate because they don't look like strategy.

But:

what happens when nobody chooses is itself a rule.

And rules influence outcomes.

Change the Feedback

Remember our meta-principle:

Model → act → observe → update → adapt.

A system where people receive feedback six months after an action creates a very different learning environment from one where they see consequences immediately.

If you're trying to teach a skill, change a behavior, improve a process, or coordinate a team, delayed feedback can make the game harder to learn.

Sometimes redesign means shortening the distance between:

action
and

consequence.

You make the game easier for the player to understand.

The Principal-Agent Problem

There is another important family of problems that fits here.

Suppose you own a business.

You hire someone to manage it.

You want:

long-term business value.

The manager may care about that too.

But perhaps their bonus depends mainly on:

this year's revenue.

Now you and the manager are not perfectly aligned.

You—the **principal**—want one thing.

The **agent** who acts on your behalf may have somewhat different incentives.

And crucially:

the agent often knows more about their own actions than you do.

This creates a combination of:

payoff differences

and

information asymmetry.

Principal-agent problems appear everywhere:

- shareholders and executives,
- managers and employees,
- clients and advisors,
- voters and politicians,
- patients and medical providers,
- homeowners and contractors.

Again, the solution is not always:

“Hire better people.”

Sometimes the challenge is to design:

incentives,
monitoring,
authority,
accountability,
information

so that the agent’s behavior better aligns with the principal’s objective.

But Don't Incentivize One Task and Destroy Another

This is where Bengt Holmström and Paul Milgrom's work on **multitask incentives** becomes important.

Most jobs don't contain one output.

A teacher might need to:

- teach knowledge,
- develop reasoning,
- manage behavior,
- encourage curiosity,
- provide feedback,
- mentor students.

If you strongly reward only the easiest measurable task, effort can shift away from the harder-to-measure tasks.

Holmström and Milgrom's multitask principal-agent analysis formalized this problem: strong incentives on one measurable activity can distort effort away from other valuable but poorly measured activities.⁸²

That is exactly what our Kenyan teacher example helps illustrate.

If the reward is tied primarily to test results, you may get:

- more effort toward the test.

That may or may not equal:

- better education.

The designer has to understand the difference.

82. ResearchGate, https://www.researchgate.net/publication/5214308_Multitask_Principal-Agent_Analyses_Incentive_Contracts_Asset_Ownership_and_Job_Design

Every Incentive Creates a Game

The moment you attach a payoff to something, ask:

**How could a smart player maximize this payoff
without giving me what I really want?**

That's not cynical.

That's good design.

If you've created a bonus for:

number of calls,

ask:

Could people make useless calls?

If you've created a reward for:

number of bugs fixed,

ask:

Could people generate easy bugs?

If you reward:

customer reviews,

ask:

Could employees pressure only happy customers to re-view?

If you reward:

quarterly profit,

ask:

Could managers sacrifice long-term investment?

You're not accusing anyone.

You're stress-testing the game.

Design Against Gaming—Without Designing for Monsters

There's a balance here.

One extreme says:

“Everyone will behave perfectly. We don't need safeguards.”

The other says:

“Everyone is trying to cheat us. Build a prison.”

Neither is wise.

Good game design recognizes that human beings contain:

self-interest,
morality,
reciprocity,
fairness,
loyalty,
ambition,
error,
creativity.

You want a system that works reasonably well with ordinary human beings.

Not only with saints.

And not only against villains.

Don't Confuse Control With Design

Mechanism design can sound like:

How do I control people?

That's not what I want you to take from it.

You cannot perfectly control a complex human system.

People adapt.

Contexts change.

Unexpected incentives appear.

New players enter.

Rules interact.

The designer's model can be wrong.

So a better sequence is:

Design → observe → test → update → adapt.

That should sound familiar.

It's the same meta-principle we've used throughout the book.

Your first design is a hypothesis.

Then reality responds.

The System Talks Back

Suppose you redesign compensation.

You expect:

more collaboration.

Instead, people start forming small alliances and excluding outsiders.

That's data.

You redesign a customer policy.

You expect:

fewer complaints.

Instead, complaints fall because customers stop bothering and simply leave.

That's data.

You introduce a productivity metric.

Productivity rises.

Quality collapses.

That's data.

The wrong response is:

“Why aren't people behaving according to my design?”

The right response is:

“What did I misunderstand about the game?”

This is strategic humility.

Sometimes You Should Remove an Incentive

We often assume game design means:

add rewards.

Sometimes the better design is removing the thing creating distortion.

Suppose employees already take pride in doing good work.

Management introduces a competitive bonus for every tiny output.

Now coworkers stop helping one another.

Perhaps the organization didn't need stronger motivation.

It introduced the wrong game.

The solution could be:

simplify.

Remove the competition.

Reduce the measurement.

Restore discretion.

Not every problem requires another rule.

Sometimes the accumulated rules **are** the problem.

Friction Is Part of the Game

Suppose you want customers to cancel a subscription honestly if they no longer want it.

You make cancellation require:

- eight screens,
- a phone call,
- a retention conversation,
- and a written request.

Technically they can cancel.

Strategically you've made one choice costly.

Or imagine the reverse.

Signing up requires one click.

Cancelling requires twenty.

The architecture is influencing behavior.

Again, we need our ethical standard.

Game design is powerful precisely because environments can shape choices.

So:

The fact that you can make something strategically difficult does not mean you should.

Ethical Mechanism Design

Throughout this book we've used four principles:

Truth. Autonomy. Dignity. Informed choice.

They matter especially here.

Because when you design the game, you often have more power than the players participating in it.

You control:

the rules,
the information,
the choices,
the defaults,
sometimes the consequences.

That power creates responsibility.

A good mechanism should not merely ask:

Will this produce the behavior I want?

Ask:

Does the player understand the game?

Do they retain meaningful choice?

Are important facts hidden?

Does the system exploit confusion?

Does it treat people as decision-makers rather than objects to be manipulated?

A dark pattern is still game design.

It's simply game design serving the wrong ethical objective.

Design for Informed Choice

Suppose a customer has three plans.

Ethical design makes:

- price,
- terms,
- tradeoffs,
- cancellation rules

reasonably understandable.

Manipulative design may bury:

- the real price,
- the renewal terms,
- or the exit mechanism.

Both affect behavior.

But only one respects informed choice.

Our goal is not:

Get the outcome by any means available.

It's:

**Create a game in which good outcomes emerge
without violating the people playing it.**

That's a much higher standard.

Sometimes the Rules Should Protect the Weaker Player

Mechanism design is especially important when players have unequal:

- information,
- expertise,

bargaining power.

Remember the used-car problem.

A seller knows far more about the car than the buyer.

A rule requiring:

disclosure,
inspection rights,
warranties,

can change the game.

The solution to information asymmetry doesn't always have to be:

"Buyer, become smarter."

Sometimes institutions can make the game safer.

That's one of the reasons we create rules in the first place.

Change the Game in Relationships

This doesn't apply only to governments and companies.

Imagine a couple repeatedly argues about household responsibilities.

Every Saturday:

"Who's doing what?"

Then conflict.

You could work endlessly on:

tone,
listening,
wording.

Those things matter.

But perhaps the deeper solution is:

Stop renegotiating the same game every Saturday.

Create a system.

You handle these responsibilities.

I handle those.

We rotate this one.

If something unusual occurs, here's how we decide.

Now the conflict doesn't disappear because both people became better negotiators.

The game changed.

Change the Game in Parenting

A child repeatedly asks for more screen time.

Every evening becomes a negotiation.

The parent says:

"No."

The child argues.

The parent argues back.

The child escalates.

Sometimes the parent gives in.

The child learns:

Every night is a bargaining game.

One possible redesign is a clear rule established in advance:

screen time ends at this time,
under these conditions,

with this consequence.

Now the parent isn't renegotiating the payoff structure every evening.

Again, no system eliminates every conflict.

But the repeated strategic environment has changed.

Change the Game in Business

A client repeatedly requests work outside the agreed scope.

You keep debating each request.

Each interaction becomes:

“Will Paul do it this time?”

That's one game.

Now introduce:

a clearly defined scope,
a change-request process,
and a fee for additional work.

Different game.

You haven't become less generous.

You've reduced ambiguity around what happens when the request occurs.

That can actually protect the relationship by preventing resentment from accumulating.

Change the Game With Yourself

Sometimes you're the designer and the player.

You keep telling yourself:

“I need more discipline.”

Maybe.

But inspect the game.

Your phone is beside you.

Notifications are on.

Every distracting app opens instantly.

The difficult work requires six steps to begin.

You can spend your life trying to win a willpower contest.

Or change the environment.

Remove the app.

Turn off notifications.

Make the productive action easier.

Make the distracting action harder.

Schedule another person to work with you.

Use a commitment device.

Now you’re no longer relying entirely on:

making a better decision every five minutes.

You’ve redesigned the choice environment.

Exit Is Also a Strategic Move

There is one more form of game change that people often overlook.

You can sometimes:

stop playing.

Imagine you've spent years trying to optimize your behavior inside a relationship, organization, partnership, or market where:

incentives remain destructive,
commitments aren't credible,
trust cannot be restored,
every equilibrium is bad,
and you lack the authority to redesign the rules.

At some point, ask:

Why am I trying to become better at this game?

Exit changes the game.

If a business relationship is consistently destructive:

leave.

If a negotiation cannot produce an acceptable agreement:

use your outside option.

If a market is structurally unattractive:

compete somewhere else.

If a conversation becomes abusive:

end it.

Not every game deserves another move.

But Exit Has Payoffs Too

Of course, leaving isn't magically correct.

Exit carries costs.

You may lose:

- money,
- relationships,

- status,
- opportunity,
- future upside.

So we return to Chapter 2.

What are the payoffs?

And Chapter 3:

Then what?

And Chapter 9:

What future are you giving up?

This is why the models aren't isolated tricks.

They work together.

The Entire Framework Comes Together Here

Suppose you're managing a dysfunctional team.

You could begin with:

See the Game

Who are the relevant players?

Who actually has authority?

See the Payoffs

What is everyone being rewarded for?

Think in Responses

What behavior does the current system make rational?

Manage Predictability

What patterns have people learned they can exploit?

See the Equilibrium

What recurring behavior has become stable?

Map Information & Beliefs

What do people know?

What are they merely assuming?

Separate Words From Credibility

Which promises and rules do people actually believe?

Commitment Changes the Game

What consequences or promises are genuinely binding?

Shadow of the Future

What reputation and repeated-game incentives exist?

Diagnose the Coordination Problem

Is the problem conflict, coordination, or free riding?

And finally:

Change the Game**Which part of the structure can we redesign?**

Now you can see why I did not want this book to become a collection of communication techniques.

The eleven models form a way of **seeing**.

Don't Change Everything at Once

Once you begin seeing systems, another temptation appears.

You want to redesign everything.

Be careful.

Human systems are interconnected.

Change one incentive and you may affect:

behavior,
trust,
information,
status,
coordination,
reputation.

So start with a diagnosis.

Ask:

What specific mechanism appears to be producing the unwanted behavior?

Then make the smallest change capable of testing that hypothesis.

Observe.

Update.

Adapt.

This is the same discipline we've used throughout the book.

The Game-Change Audit

When the same bad outcome keeps returning, ask:

What behavior is the current game rewarding?

Not what does the policy claim to reward.

What actually pays?

What behavior is it punishing?

Could the desirable action currently be costly?

What metric are people optimizing?

Is it actually the outcome we care about?

What information is missing?

Would greater transparency change decisions?

Who decides?

Is authority located in the right place?

Who moves when?

Would changing the sequence matter?

What options exist?

Could adding or removing choices create a better outcome?

What is currently easy?

What is currently hard?

Friction is part of the design.

How could a smart player game this system?

Assume adaptation.

What unintended behavior might the new rule create?

Think one response further.

And then:

Can I redesign the game so that the behavior I want becomes a more natural response to the structure itself?

That is the central question.

The Mental Shift

When something repeatedly goes wrong, our instinct is often:

“How do I get them to behave differently?”

Now I want you to add:

“What would have to change in the game for different behavior to make sense?”

That’s a profound shift.

From:

fixing people

to:

examining systems.

From:

demanding cooperation

to:

understanding incentives.

From:

telling people what to do

to:

asking what the environment makes attractive.

From:

reacting to bad outcomes

to:

redesigning the structure producing them.

The Designer Is Never Omniscient

But we need to end this chapter with humility.

No system is perfect.

Every mechanism rests on assumptions.

About:

what people value,
what they know,
how they'll respond,
what they'll learn,
what they can observe.

And those assumptions can be wrong.

The Pakistani tax experiments show that changing incentives can generate large behavioral responses.⁸³

The Kenyan teacher experiment shows that those responses can take forms the designer did not fully intend.⁸⁴

Boston's school-choice redesign shows that poorly designed strategic complexity can advantage sophisticated players and that changing the mechanism can reduce the need to game the system.⁸⁵

Taken together, the lesson is not:

Design the perfect game.

There may be no perfect game.

The lesson is:

Pay attention to the game you're creating.

83. National Bureau of Economic Research, <https://www.nber.org/papers/w20627>

84. American Economic Association, <https://www.aeaweb.org/articles?id=10.1257%2Fapp.2.3.205>

85. National Bureau of Economic Research, <https://www.nber.org/papers/w11965>

The Question to Carry With You

When you've tried better arguments...

better wording...

more motivation...

more discipline...

more pressure...

more cooperation...

and the same problem keeps returning...

ask:

Do I need a better move—or a better game?

Then look at:

the rules,
the payoffs,
the information,
the timing,
the authority,
the options,
the feedback,
the metrics,
the players.

Because sometimes the smartest player in the room is not the person who finds the brilliant move inside the existing game.

It's the person who realizes:

We don't have to keep playing it this way.

And that brings us to the end of the eleven mental models.

But it does not bring us to the end of the framework.

Because knowing eleven concepts is not the goal.

You could memorize every chapter in this book and still fail to think strategically.

The real skill is knowing:

which model matters,
when it matters,
how much weight to give it,
and when new evidence should cause you to change your mind.

That is something different from knowledge.

It is **judgment**.

And that's where we'll finish.

CONCLUSION

Strategic Judgment

The goal is not to memorize the game.

The goal is to learn how to see it.

At the beginning of this book, we started with a question that sounds almost too simple:

What if the way we normally think about communication starts too late?

Most communication advice begins after we've already interpreted the situation.

Someone says something.

We decide what it means.

We decide what they want.

We decide how we feel about it.

Then we ask:

What should I say?

But by that point, we've already made half a dozen assumptions about the game.

Maybe they're correct.

Maybe they're not.

Game theory gives us another place to begin.

Before asking what to say, ask:

What's actually going on here?

Who are the players?

What do they want?

What can they do?

What do they know?

What do they believe?

How will they respond?

What happens after that?

What pattern are we creating?

What changes because we'll interact again?

Are we actually in conflict?

And perhaps most importantly:

Is the game itself producing the problem?

Those questions don't give you a script.

They give you something more valuable.

They give you a way to think.

The Models Are Not Eleven Steps

There's an important distinction I want to make before we finish.

The eleven models in this book are not a procedure.

You're not supposed to enter every conversation and think:

"Okay. Step one: identify the game. Step two: identify the payoffs. Step three..."

That's not how good judgment works.

Imagine an experienced driver.

They aren't consciously saying:

"Check mirror. Calculate distance. Estimate speed. Adjust steering angle."

They've developed perception.

They notice things.

The same thing should happen here.

At first, you'll consciously ask the questions.

Eventually, you'll start noticing the structure almost automatically.

Someone makes an aggressive demand and you think:

What payoff are they protecting?

Someone makes a dramatic threat and you think:

Would they actually want to carry that out?

A team keeps failing in the same way and you think:

This is a pattern, not an incident.

Someone tells you something that strongly benefits them and you think:

What makes that credible?

Two people appear to be fighting and you suddenly realize:

They're not actually opposed. They don't know what the other is going to do.

A company keeps trying to motivate people and you notice:

The incentive system rewards the exact opposite behavior.

That's the goal.

Not recitation.

Recognition.

Strategic Judgment Is Choosing the Right Lens

Suppose someone says to you:

"I need your answer by tomorrow. This is absolutely the final offer."

There are many ways you could respond.

You could ask:

"What should I say?"

But now you have other lenses.

See the Game

Who actually has authority here?

Is this person making the decision?

Are there other players?

See the Payoffs

What happens to them if I say no?

Do they genuinely need an answer tomorrow?

Think in Responses

If I decline, what are they likely to do?

If I accept, what happens next?

Manage Predictability

Have I taught them that artificial pressure makes me move?

See the Equilibrium

Is this part of a recurring pattern in our negotiations?

Map Information & Beliefs

What do I actually know about the deadline?

What am I merely assuming?

Separate Words From Credibility

What makes “final offer” believable?

Commitment Changes the Game

Are they actually constrained—or simply saying they are?

Shadow of the Future

Will I negotiate with this person again?

What precedent does my response create?

Diagnose the Coordination Problem

Are we even disagreeing about the underlying outcome?

Change the Game

Do I need to accept their yes-or-no structure at all?

Perhaps I can introduce another option.

You probably don’t need all eleven lenses.

That’s the point.

Judgment means recognizing:

Which two or three matter most right now?

More Analysis Is Not Always Better

There's a danger in a book like this.

Once you learn to see games everywhere, you can begin overanalyzing everything.

Your friend texts:

"Sure."

And suddenly you're thinking:

"What strategic signal does the absence of an exclamation point communicate about their latent payoff structure?"

Don't do that.

Not every interaction requires strategic analysis.

Sometimes:

"Sure"

means:

"Sure."

The purpose of this framework isn't to turn human life into chess.

It's to help when:

- the stakes matter,
- something doesn't make sense,
- a pattern keeps repeating,
- incentives are unclear,
- trust is uncertain,
- multiple people affect the outcome,
- or your normal way of thinking isn't working.

Strategic thinking should reduce confusion.

If your analysis is making ordinary life needlessly suspicious and complicated, you're using the tool badly.

People Are Not Pieces on a Board

This matters enough to say explicitly.

Game theory uses words like:

players,
strategies,
payoffs,
moves.

Those words can make human interactions sound mechanical.

They're not.

People love.

People sacrifice.

People act from principle.

People forgive.

People make mistakes.

People act irrationally.

People change their minds.

People care about fairness.

People choose relationships over money.

People sometimes do the right thing when doing so clearly costs them.

A model that ignores those realities isn't sophisticated.

It's incomplete.

The point of calling someone a "player" is not to reduce them to a piece you move.

It's to remind you:

They are a decision-maker too.

They have agency.

They see the world from a position different from yours.

And your outcome depends partly on what they decide.

That should make us more respectful of other people's agency, not less.

Strategy Without Ethics Is Incomplete

Throughout this book, we've returned to four principles:

Truth. Autonomy. Dignity. Informed choice.

Those are not decorative.

They're constraints on how this framework should be used.

Game theory can help explain:

deception,
coercion,
manipulation,
intimidation,
exploitation.

That does not make those things good strategy in the larger sense.

A strategy can "work" and still be wrong.

You may be able to exploit someone's ignorance.

That doesn't mean you should.

You may discover exactly which insecurity would pressure someone into agreement.

That doesn't mean using it respects their dignity.

You may be able to hide relevant information so someone makes the choice you prefer.

That doesn't mean they made an informed choice.

So we need two different questions:

Will this work?

and

Should I do it?

Never collapse them into one.

Cooperation Is Not Automatically Good Either

The opposite mistake is possible too.

We can romanticize cooperation.

But cooperation is just a strategic structure.

Two people can cooperate to:

build a hospital,

or:

commit a fraud.

A team can coordinate around:

serving customers,

or:

hiding wrongdoing.

A repeated game can sustain:

trust,

or:

corruption.

So the ethical question is not:

Are the players cooperating?

It's:

What are they cooperating toward?

Game theory can help us understand the structure.

It cannot decide our moral ends for us.

Your First Model Will Often Be Wrong

This may be the most important practical point in the conclusion.

Everything we've done depends on your interpretation of the game.

And your interpretation can be wrong.

You think the person cares about money.

They care about status.

You think they're angry.

They're scared.

You think they're bluffing.

They're committed.

You think the problem is free riding.

It's actually confusion about ownership.

You think everyone knows the rule.

Half the group has a different understanding.

You think an incentive will increase quality.

People learn to game the metric instead.

So the framework needs one rule sitting above all eleven models:

Your model of the game is provisional.

Do not fall in love with your diagnosis.

Test it.

Model → Act → Observe → Update → Adapt

This is the loop that holds everything together.

Model

Based on the information available, form your best understanding of the game.

Act

Make a move.

Observe

What actually happens?

Update

Where did reality differ from your prediction?

Adapt

Change your strategy—or change your model.

Then repeat.

Model → act → observe → update → adapt.

That is strategic thinking.

Not perfect prediction.

Not omniscience.

Not always being right.

Learning faster from what happens.

Your Mistakes Are Information

Suppose you predict:

“If I refuse this request, they’ll become angry.”

You refuse.

They say:

“No problem. I completely understand.”

Good.

You were wrong.

That’s valuable.

Perhaps you overestimated their dependency.

Perhaps you misunderstood their payoff.

Perhaps you believed the relationship was more fragile than it was.

Don’t ignore the discrepancy because you want your model to be right.

Update.

Or suppose you think:

“If I explain my reasoning clearly, they’ll agree.”

You explain.

They understand perfectly.

They still refuse.

That tells you something.

Maybe this was never an information problem.

Maybe your interests are genuinely opposed.

Again:

update the model.

Surprise Is a Gift

The moments when another person surprises you are some of the most strategically useful moments you'll get.

You thought they would cooperate.

They didn't.

Ask:

What did I misunderstand?

You thought they'd walk away.

They stayed.

Ask:

What payoff did I underestimate?

You thought the threat was empty.

They followed through.

Ask:

What made the commitment stronger than I realized?

You thought the new rule would improve behavior.

It made things worse.

Ask:

What response did the design create that I failed to anticipate?

Surprise reveals gaps in your model.

A person committed to learning should become curious when surprised.

Don't Label Someone Irrational Too Quickly

One of the easiest ways to stop learning is to say:

“They’re irrational.”

Sometimes people are irrational.

So are you.

So am I.

But before reaching for that explanation, ask:

What would have to be true for this behavior to make sense?

Maybe they have information you don’t.

Maybe the payoff you think they’re maximizing isn’t the payoff they actually care about.

Maybe they’re playing a longer game than you are.

Maybe their interpretation of your behavior is completely different from your own.

Maybe they’re responding to history you’re discounting.

And yes, maybe they’re simply making a mistake.

The point is not to explain every behavior as secretly rational.

It's to avoid using the word **irrational** as a substitute for thinking.

Learn to Separate Facts From Your Story About the Facts

Suppose someone doesn't return your call.

Fact:

They didn't return your call.

Interpretation:

They're disrespecting you.

Possible.

But not identical.

Maybe they forgot.

Maybe they're avoiding the conversation.

Maybe they're overwhelmed.

Maybe they never received the message.

Maybe they're strategically delaying.

The skill is being able to say:

Here's what I know.

Here's what I think it means.

Here's how confident I am.

That's intellectual discipline.

And it protects you from making your response to an assumption rather than to reality.

The Best Strategists Can Hold Multiple Hypotheses

Imagine a customer suddenly becomes resistant.

One interpretation:

Price objection.

Another:

They don't trust you.

Another:

They need someone else's approval.

Another:

They're interested but trying to improve their negotiating position.

Another:

They've already chosen a competitor and are being polite.

You don't need to immediately choose one.

You can hold several models.

Then ask questions or observe behavior that helps distinguish among them.

That is much more powerful than:

"I know exactly what they're doing."

Certainty feels good.

Accuracy is better.

Thinking Strategically Is Not the Same as Thinking Suspiciously

This deserves emphasis because the line can become blurry.

Suspicion asks:

What are they hiding?

Strategic thinking asks:

What possibilities should I consider?

Suspicion assumes hostile intent.

Strategic thinking tests hypotheses.

Suspicion interprets ambiguity negatively.

Strategic thinking recognizes ambiguity.

Suspicion makes trust difficult.

Strategic thinking understands why trust can be enormously valuable.

The framework should make you less naïve.

It should not make you cynical.

Sometimes the Correct Move Is Vulnerability

That may sound strange in a game-theory book.

But consider the structure.

You're in a repeated relationship.

Both sides are protecting themselves.

Neither person reveals information.

Neither takes the first cooperative step because each fears exploitation.

You could remain protected indefinitely.

Or one person can take a measured risk.

They disclose something.

They offer trust.

They make a concession.

They apologize.

That action may change the other person's beliefs.

It may shift the equilibrium.

It may reopen cooperation.

There is risk.

That's precisely why the action can carry information.

Strategic thinking isn't always about protecting yourself.

Sometimes it's understanding when taking a risk can transform the game.

Sometimes the Correct Move Is Walking Away

And sometimes the opposite is true.

You've:

- communicated clearly,
- tested your assumptions,
- enforced boundaries,
- allowed for repair,
- changed your own behavior,
- considered the other person's incentives.

The game remains destructive.

Strategic judgment may tell you:

Stop playing.

That's not defeat.

An outside option is part of the game.

A refusal to play can itself be a move.

There is no virtue in becoming brilliant at surviving a game that repeatedly damages you when a better alternative exists.

Strategy Is Often About Asking Better Questions

Think about how the questions changed throughout the book.

Instead of:

What should I say?

we asked:

What game am I in?

Instead of:

Why are they being difficult?

we asked:

What payoff would make this behavior make sense?

Instead of:

What should I do?

we asked:

What will they do if I do it?

Instead of:

Why does this keep happening?

we asked:

What equilibrium are we creating?

Instead of:

Are they telling the truth?

we asked:

What makes this claim credible?

Instead of:

Why don't they just cooperate?

we asked:

Are we in conflict, failing to coordinate, or facing a free-rider problem?

Instead of:

How do I motivate them?

we asked:

What behavior does the game reward?

Better thinking often begins with a better question.

The Eleven Questions

If you remember nothing else from the book, remember these.

1. See the Game

What game am I actually in?

2. See the Payoffs

What does each player actually value?

3. Think in Responses

Then what?

4. Manage Predictability

What should they be able to predict about me—and what becomes exploitable if they can?

5. See the Equilibrium

What pattern are we producing, and why is it stable?

6. Map Information & Beliefs

What do I know, what am I assuming, and what might they know that I don't?

7. Separate Words From Credibility

What makes this believable?

8. Commitment Changes the Game

What actually binds them—or me?

9. The Shadow of the Future

What changes because there is a next time?

10. Diagnose the Coordination Problem

Are we actually in conflict—or failing to coordinate?

11. Change the Game

Do I need a better move—or a better game?

And above them all:

What did I learn from what just happened?

That's strategic judgment.

From Conscious Thinking to Intuition

When I look back at my own life, I realize I was using pieces of these ideas long before I knew their names.

As a kid trying to figure out how to make people laugh, I was learning:

what different people responded to.

In prison, I was learning:

who the players were,
what the unwritten rules were,
which signals mattered,
which boundaries were real,
what different people valued,
and how quickly a bad read could change the game.

In sales, I was learning:

incentives,
responses,
resistance,
outside options,
credibility.

In copywriting and marketing, I was learning:

what people believed,
what information they needed,
what they valued,
and how changing a frame could change a decision.

For years, those lessons existed mostly as intuition.

Game theory gave me a language for them.

Not a perfect language.

Not a complete theory of human nature.

But a powerful set of lenses.

And once I could name the lenses, I could examine them.

Improve them.

Teach them.

And challenge them when they were wrong.
That is what I hope this book does for you.

Intuition Is Compressed Experience

People often talk about intuition as if it's almost mystical.
Sometimes what we call intuition is simply:

patterns learned so deeply that we no longer experience the reasoning consciously.

You walk into a room and something feels off.

Why?

Maybe you've unconsciously noticed:

who's speaking,
who's silent,
who looks at whom before answering,
who controls information,
who's being deferred to,
where tension appears.

An experienced negotiator hears a statement and senses:

"That isn't really the issue."

An experienced leader notices:

"These people don't disagree. They're waiting for permission."

An experienced salesperson notices:

"This objection isn't about price."

Expert judgment often begins as explicit thinking and becomes increasingly automatic through repeated exposure and feedback.

That's what we're trying to build.

But Intuition Needs Calibration

Experience alone doesn't guarantee good intuition.

You can repeat the wrong interpretation for twenty years.

If nobody gives you feedback, confidence can increase without accuracy.

That's why our loop matters:

Model → act → observe → update → adapt.

Your intuition should become more trustworthy because it has survived contact with reality.

Not because you've repeated the same story often enough to believe it.

The Goal Is Transfer

I don't want you to learn eleven ideas that work only when the examples look exactly like the chapters.

The real test is whether you can recognize the same structure in a situation you've never seen before.

Can you recognize a commitment problem in:

parenting,
business,
politics,
negotiation,
your own habits?

Can you recognize a coordination problem in:

marriage,

project management,
community organization?

Can you recognize a repeated-game problem in:

customer relationships,
leadership,
friendship?

Can you recognize a mechanism-design problem when everyone else thinks:

“We just need people to try harder”?

That’s transfer.

And that’s the difference between memorizing information and developing a mental model.

The World Will Not Label the Game for You

Nobody is going to walk into your office and say:

“Hello, Paul. I am now presenting you with an incomplete-information bargaining game containing asymmetric outside options and a credibility problem.”

They’re going to say:

“We need an answer today.”

Your child isn’t going to say:

“I have detected that your repeated concessions have created an exploitable strategy.”

They’re going to ask again.

Your employees won’t say:

“Our current compensation mechanism creates a social dilemma.”

They'll stop sharing leads.

Your partner won't say:

"We have entered a stable negative equilibrium reinforced by reciprocal expectations."

You'll just have the same argument again.

Reality gives you the symptoms.

You have to recognize the structure.

Strategic Thinking Slows the First Reaction

One of the most practical effects I hope this framework produces is simply a tiny space between:

stimulus

and

response.

Something happens.

Your first interpretation appears.

Before acting, you ask:

What game am I in?

That small pause can completely change what happens next.

Someone pressures you.

Instead of automatically resisting, you examine their incentives.

Someone rejects your proposal.

Instead of immediately arguing harder, you ask what payoff you're missing.

Someone makes a threat.

Instead of reacting to the intensity of the language, you ask whether carrying it out is credible.

A team fails.

Instead of blaming people, you inspect the system.

That pause is strategic judgment beginning to operate.

The Objective Is Not Control

There is no framework that allows you to control human beings.

And if someone promises you one, you should probably apply Chapter 7.

People surprise us.

They make mistakes.

They change.

They refuse.

They misunderstand.

They value things we didn't anticipate.

They sacrifice payoffs for principles.

They choose emotion over calculation.

They sometimes behave beautifully.

Sometimes terribly.

Strategic thinking does not give you control over them.

It gives you:

a better map of the situation you're navigating.

And a better map improves your decisions even though it cannot guarantee the destination.

You Don't Need to Win Every Game

There's one final assumption worth challenging.

Strategic thinking is often associated with:

winning.

But what does winning mean?

Chapter 2 should have made us suspicious of that question.

What is the payoff?

Sometimes winning means:

closing the deal.

Sometimes:

walking away.

Sometimes:

preserving the relationship.

Sometimes:

telling the truth even when it costs you.

Sometimes:

protecting someone else's autonomy.

Sometimes:

refusing to participate.

Sometimes:

reaching a fair agreement.

Sometimes:

discovering you were wrong.

The strategy cannot tell you what your ultimate values should be.

It can help you pursue them more intelligently.

Know What Game You Refuse to Play

Strategic judgment isn't only recognizing games.

It is deciding which games deserve your participation.

Some environments reward:

deception.

Some reward:

constant status competition.

Some reward:

betrayal.

Some require you to violate your principles to remain competitive.

You may be able to become excellent at those games.

But perhaps the wiser choice is:

not to play.

Or to change the rules.

Or to find different players.

Or to build something better.

There is strategy in refusal too.

The Game Behind the Conversation

We began the book with communication because communication is where much of strategic human interaction becomes visible.

But by now, you should see that the words themselves are only one layer.

Behind the conversation are:

- players,
- incentives,
- expectations,
- information,
- beliefs,
- signals,
- commitments,
- history,
- reputation,
- alternatives,
- rules,
- timing,
- patterns.

That's the game behind the conversation.

And once you start seeing it, many interactions that once seemed confusing become more understandable.

Not easy.

Not predictable.

But understandable.

One Last Example

Imagine someone you care about says:

“You never listen to me.”

The tactical response might be:

“What’s the best communication technique to use here?”

Maybe active listening.

Maybe mirroring.

Maybe ask a question.

Those can help.

But before reaching for a technique, see the game.

What happened before this?

Is this an isolated comment or a repeated pattern?

What do they mean by “listen”?

What payoff are they seeking?

Understanding?

Agreement?

Respect?

Attention?

Do they believe you don’t care?

Why?

What information are they using?

Have your past responses taught them that raising a concern leads to defensiveness?

Are you each anticipating the other's reaction before the conversation begins?

Have you entered an equilibrium where:

they criticize because they expect you not to listen,
and:

you defend because you expect them to criticize?

Is the issue actually disagreement—or are you failing to coordinate around what “being heard” means?

Could you change the game by changing the pattern rather than merely finding the perfect sentence?

Now you're thinking at a different level.

You may still use a communication technique.

But now the technique is serving an understanding of the situation.

That's exactly how it should be.

The Final Mental Shift

Most of us were never taught to think about interaction this way.

We were taught:

be confident,
communicate clearly,
listen,
persuade,
negotiate,
set boundaries,
build rapport.

All useful.

But those are actions.

This book has been about what comes **before** the action.

How do you understand the strategic environment in which those actions take place?

Because the exact same behavior can produce entirely different results in different games.

A concession can create trust.

Or invite exploitation.

Predictability can create security.

Or vulnerability.

A commitment can create credibility.

Or trap you.

Communication can solve coordination.

Or accomplish nothing because the real problem is incentives.

Competition can be healthy.

Cooperation can be corrupt.

A stable pattern can be excellent.

Or terrible.

There is no universal tactic.

There is only:

the game you're actually in.

So Before You Make Your Next Move

Ask:

What game am I playing?

Then look again.

Because the first game you see may not be the real one.

Ask what everyone values.

Think about how they'll respond.

Notice the pattern.

Separate facts from beliefs.

Separate words from credibility.

Look at what commitments are actually binding.

Consider tomorrow.

Determine whether you have conflict or merely failed coordination.

And when none of that seems to fix the problem, inspect the rules themselves.

Then act.

Watch what happens.

Learn.

Update.

Try again.

Over time, something changes.

You stop reacting only to the conversation in front of you.

You begin to see the structure beneath it.

And that was the goal from the beginning.

Don't just learn what to say.

Learn to see the game behind the conversation.