

THE COURSE

Practical Game Theory

See the Game Behind the Conversation

Instructor & Course-Builder Guide

PAUL MASCETTA

Contents

Module 0 — Orientation: The Game Behind the Conversation	3
Module 1 — See the Game	11
Module 2 — See the Payoffs	25
Module 3 — Think in Responses	38
Module 4 — Manage Predictability	50
Module 5 — See the Equilibrium	63
Module 6 — Map Information & Beliefs	75
Module 7 — Separate Words From Credibility	88
Module 8 — Commitment Changes the Game	100
Module 9 — The Shadow of the Future	112
Module 10 — Diagnose the Coordination Problem	125
Module 11 — Change the Game	138
Module 12 — Strategic Judgment	152

Module 0 — Orientation: The Game Behind the Conversation

The question: What if the way we normally think about communication starts too late?

Book reading	Introduction
Time	~35 min of video lessons · ~30 min of workbook · one week of observation
Tool you'll build	The Interaction Journal

Why this module matters

Most communication training starts with words: what to say, how to phrase it, which questions to ask, how to handle the objection. Those things matter. But by the time you're choosing words, you've already made a series of quieter decisions about what is happening, what the other person wants, what they know, and how they're likely to respond. If those judgments are wrong, excellent communication simply helps you execute the wrong strategy more effectively.

This course teaches the layer underneath the words. Game theory is the formal study of strategic interaction—situations where the outcome of one person's choice depends partly on the choices of others. You won't do any math. Instead, you'll learn eleven mental models that change the questions you ask before you speak.

The goal is not to make you a better manipulator. The goal is to help you see the game behind the conversation, so you have far more choices about how you play.

Learning objectives

By the end of this module you will be able to:

1. Explain, in one sentence, what “game” means in this course and why it does not mean manipulation or competition.
2. Name the eleven mental models and the question each one asks.
3. Describe the Model → Act → Observe → Update → Adapt loop that runs through the whole course.
4. State the course's ethical foundation—truth, autonomy, dignity, informed choice—and apply it as a test for any strategy.

5. Begin an Interaction Journal you will use in every module.

Lessons

Lesson 0.1 — Why most communication advice starts too late (8 min)

Teaching points

- Almost all communication training focuses on strategies and tactics: words, questions, objection handling, rapport.
- Underneath every tactic are judgments about the situation: who is involved, what each person wants, what choices they have, what they know and believe, and what they expect you to do.
- The same approach can work with one person and fail completely with another. Tactics don't explain why. The structure of the interaction does.
- The shift: from "What should I say?" to "What game am I actually in?"

Example to use: The author's own path—an only child learning to read a room through humor, a young man navigating the written and unwritten rules of federal prison, a salesperson and copywriter whose livelihood depended on understanding how people respond. In each setting, what mattered most wasn't the words, but understanding the situation surrounding the words.

On-screen takeaway: The words are only the visible layer.

Lesson 0.2 — What "game" means (and doesn't) (7 min)

Teaching points

- A game, in the game-theoretic sense, is any situation in which the outcome depends at least partly on the decisions of more than one person.
- A game can be cooperative, competitive, mixed, or a coordination problem. Families, friendships, teams, and partnerships are games in this sense.
- "Player" is a technical word, not a cynical one. People are always more than players.
- Game theory is a lens for thinking, not a guarantee that humans behave like perfectly rational machines. Throughout the course, we'll note where real behavior departs from the clean models.

Example to use: Asking your boss for a raise. Before you speak: What does your boss want? What happens if they say no? What alternatives does each of you have? Who else influences the decision? Suddenly, it isn't simply a conversation.

On-screen takeaway: Other people are decision-makers, and their decisions affect your outcomes.

Lesson 0.3 — The eleven mental models (10 min)

Teaching points

- Part I — See the structure beneath the conversation: situation → motivation → response → anticipation → pattern.
- Part II — See what isn't obvious: information, credibility, commitment.
- Part III — Think beyond the moment: the future, coordination, and the design of the game itself.
- You will not carry eleven checklists into a conversation. The aim is strategic judgment: noticing the right question at the right time.

#	Mental model	The question it asks
1	See the Game	What game am I actually in?
2	See the Payoffs	What does each player actually value?
3	Think in Responses	If I make this move, then what?
4	Manage Predictability	What am I teaching other people to predict about me?
5	See the Equilibrium	What pattern are our responses creating together?
6	Map Information & Beliefs	What do I actually know, and what am I assuming?
7	Separate Words From Credibility	What makes this message believable?
8	Commitment Changes the Game	What actually binds a promise, threat, or boundary?
9	The Shadow of the Future	What changes because there will be a next time?
10	Diagnose the Coordination Problem	Are we in conflict, or failing to coordinate?
11	Change the Game	Do I need a better move, or a better game?

On-screen takeaway: Don't just learn what to say. Learn to see the game.

Lesson 0.4 — Ethics and the learning loop (10 min)

Teaching points

- Strategic intelligence without humanity becomes manipulation. Our standard: **truth, autonomy, dignity, informed choice**.
- Understanding why someone behaves a certain way is not the same as approving of it.
- Every model you build of another person is a hypothesis. The loop: **Model → Act → Observe → Update → Adapt**.
- Surprise is information. When someone doesn't respond the way you predicted, your model of the game was incomplete.

Example to use: You assume a colleague's resistance is about status, so you offer public credit. They're still resistant. Rather than labeling them "irrational," you update: maybe the real cost to them is workload or risk. The surprise improved your model.

On-screen takeaway: Your model of the game is provisional.

Key terms

Term	Plain-language meaning
Game	Any situation in which outcomes depend on the choices of more than one person.
Player	Anyone whose choices affect the outcome, including people who aren't in the room.
Strategic interaction	An interaction in which each person's best choice depends on what others do.
Mental model	A lens that tells you what to look at; not a law that predicts exactly what people will do.
Strategic judgment	Knowing which model matters, when it matters, how much weight to give it, and when to change your mind.

Research spotlight

What game theory is, and is not. Game theory gives us a formal language for strategic interaction: players have choices and preferences, and the outcome depends on how those choices interact. Thomas Schelling, who received the 2005 Nobel Prize in economics, emphasized that most real-world interactions contain both conflicting and common interests rather than being purely competitive.

What this does NOT prove: Game theory was never created as a theory of better communication, and human beings are messier than the perfectly rational players in many

mathematical models. In this course, research illustrates and disciplines our thinking; it does not promise that any technique will work on any particular person.

The tool: The Interaction Journal

Use one page per interaction. You'll add a new column of questions in each module.

Date / setting:

Who was involved (including people not in the room):

What I wanted:

What I think they wanted:

What I said or did:

How they responded:

Did their response surprise me? What might that tell me?

If I could replay it, what question would I ask myself before speaking?

Practice exercises

Exercise A — Analyze the scenario

Dana manages a small marketing team. For three weeks, she has asked a senior designer, Marcus, to submit his drafts two days before client deadlines so she can review them. She has explained why, rephrased the request, and finally sent a firm email. Marcus keeps submitting the night before. Dana's colleague suggests she "work on her delivery." Dana suspects the problem isn't her words at all. Marcus is paid a bonus based on client satisfaction scores, not internal deadlines. The client contact loves Marcus and often requests last-minute changes directly from him.

1. Who are the players, including any who aren't in Dana's conversations?
2. Why might better wording fail here?
3. What question should Dana be asking instead of "How do I say this more clearly?"

Model answer

1. Dana, Marcus, and the client contact, who is a hidden but powerful player. The bonus system (and whoever designed it) also shapes the game.
2. Marcus's incentives reward client satisfaction, and the client rewards his last-minute responsiveness. Clear words don't change what Marcus gains by submitting late; finishing early may even cost him, because the client will request changes anyway.
3. "What game is Marcus actually in, and what makes his current behavior sensible from where he stands?" That question opens options such as involving the client in the timeline, adjusting what the bonus rewards, or agreeing on a change-request process. It's a preview of Modules 1, 2 and 11.

Exercise B — Your own game

Think of one recurring interaction that frustrates you: at work, at home, or with a customer. Write down the last thing you said in it. Then list everything you assumed before you said it: about what they wanted, what they knew, and how they'd respond. Circle any assumption you have never actually tested.

Exercise C — Partner drill / role-play

Setup (10 minutes, pairs). Partner 1 describes a recent conversation that went badly in two minutes, focusing only on what was said. Partner 2 may then ask only structural questions for five minutes: Who else was affected? What did each person want? What were their alternatives? What happened next? Switch roles.

Debrief: What did the structural questions reveal that the words didn't? Did either of you change your interpretation of the other person?

Common mistakes

- **Treating "game" as a cynical word** — In this course, a game includes cooperation, friendship and family life. Seeing the structure is about understanding, not winning at others' expense.
- **Expecting scripts** — This course will not give you seven phrases for every situation. It gives you better questions, which transfer to situations no script anticipates.
- **Overanalyzing everything** — Use as much analysis as the stakes and uncertainty justify. Most interactions don't need a strategic audit.
- **Mistaking your model for the truth** — Every read of another person is a hypothesis to test, not a fact.

Quiz

1. In this course, a “game” is best defined as:
 - A) Any competitive situation with a winner and a loser
 - B) Any situation in which outcomes depend on the choices of more than one person
 - C) A technique for gaining leverage in a conversation
 - D) A mathematical model that requires calculation
2. Why can excellent communication still produce a bad outcome?
 - A) Because people rarely listen
 - B) Because tone matters more than content
 - C) Because if your read of the situation is wrong, good words execute the wrong strategy
 - D) Because game theory says communication is irrelevant
3. Which is the course’s ethical foundation?
 - A) Win-win, always
 - B) Truth, autonomy, dignity, informed choice
 - C) Reciprocity, scarcity, authority
 - D) Whatever produces the best outcome for you
4. The learning loop that runs through the course is:
 - A) Plan → Pitch → Close
 - B) Listen → Mirror → Label
 - C) Model → Act → Observe → Update → Adapt
 - D) Analyze → Decide → Never revisit
5. Which question belongs to Mental Model 10?
 - A) What makes this message believable?
 - B) Are we in conflict, or failing to coordinate?
 - C) What does each player actually value?
 - D) What actually binds a promise?
6. A mental model in this course is best described as:
 - A) A law that predicts behavior precisely
 - B) A lens that tells you what to look at
 - C) A script for difficult conversations
 - D) A personality test
7. When someone responds in a way you didn’t predict, the course suggests you:
 - A) Conclude they are irrational
 - B) Repeat your move more forcefully
 - C) Treat the surprise as information and update your model

D) Ignore it as noise

8. How much analysis should you apply to an interaction?

- A) The full eleven-model audit every time
- B) None, because intuition is enough
- C) As much as the stakes and uncertainty justify
- D) Only when you are losing

Answer key

1. **B** — A game in the game-theoretic sense doesn't require competition; it requires interdependent choices.
2. **C** — Words are executed on top of judgments about the situation. Wrong judgments mean well-executed wrong strategy.
3. **B** — Strategy is constrained by truth, autonomy, dignity and informed choice.
4. **C** — Every model is provisional and gets tested against what actually happens.
5. **B** — Model 10 distinguishes real conflict from coordination failure.
6. **B** — The models tell you what to look at; they aren't behavioral laws.
7. **C** — Surprise reveals that your model of the game was incomplete.
8. **C** — Strategic judgment includes knowing when not to overthink.

Field assignment

This week, observe; don't change anything. Record three interactions in your Interaction Journal: one at work, one at home, and one with a stranger or customer. For each, note one thing that shaped the outcome that wasn't the words themselves, such as a person not in the room, an incentive, a deadline, or a past history.

Debrief questions for next week: 1. Which of the three interactions had the most "invisible" structure? 2. Where did you catch yourself assuming you knew what someone wanted? 3. Which of the eleven questions do you suspect you'll need most?

The question to carry

| **What did I assume before I spoke?**

Module 1 — See the Game

The question: What’s really going on here?

Book reading	Chapter 1, “See the Game”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Game Map

Why this module matters

Most communication advice begins with a deceptively simple question: What should I say? Before asking for a raise, you rehearse your words and your number. That matters, but it skips what comes first. What does your boss actually want? What happens if they say no? What alternatives does each of you have? What will your boss probably do after you speak?

Once you ask those questions, the conversation becomes a game. Not a contest against your boss, but a game in the game-theoretic sense: a situation in which the outcome depends, at least partly, on the decisions of multiple people. Your words enter a system that already holds incentives, expectations, information, constraints, alternatives, history, timing, and other decision-makers.

This module teaches the course’s first mental shift: **don’t start with the words. Start with the situation.** It isn’t a cynical view that everyone is manipulating everyone else. It’s simpler: other people are decision-makers, and their decisions affect your outcomes. Seeing the game often makes cooperation easier, because you begin to understand why another person’s behavior makes sense from where they stand.

Learning objectives

By the end of this module you will be able to:

1. **Define** a game in the game-theoretic sense and explain why a game does not have to be competitive.
2. **Identify** every player who can affect an outcome, including people who aren’t in the room.
3. **Distinguish** what a person wants from what they can actually do, and stated constraints from real ones.

4. **Map** an interaction using the nine questions of the Game Map: WHO, WHAT, OPTIONS, CONSTRAINTS, INFORMATION, TIMING, ALTERNATIVES, DEPENDENCIES, and THEN WHAT.
5. **Diagnose** whether a recurring problem is a person problem or a game problem before deciding how to respond.

Lessons

Lesson 1.1 — From “What should I say?” to “What’s really going on?” (10 min)

Teaching points

- Game theory studies strategic interaction: situations in which one decision-maker’s choices affect the outcomes available to others. The formal framework describes players, strategies, payoffs, and outcomes.
- A game doesn’t need a winner and a loser. Cooperation, coordination, friendship, partnerships, and family decisions are all games in this sense. Thomas Schelling emphasized that real interactions usually contain both common and conflicting interests.
- The starting premise isn’t “everyone is manipulating everyone.” It’s: **other people are decision-makers, and their decisions affect your outcomes.**
- Nash’s equilibrium idea captures the core point: each person’s behavior has to be understood in relation to everyone else’s.

Example to use: The raise request. Before rehearsing a word, ask: What does your boss want? What happens if they say no, or if you turn down a counteroffer? Who has the authority to decide, and who else might influence it? The “conversation” turns into a structure with several players, options, and consequences that run well past the first sentence.

On-screen takeaway: Don’t start with the words. Start with the situation.

Lesson 1.2 — Players, wants, and options (10 min)

Teaching points

- The person you’re talking to may not be the whole game. Ask: Who can actually affect the outcome? Sometimes the people who matter most aren’t in the room.
- A player is simply a decision-maker. Your spouse, customer, boss, and you are all players, and real people care about fairness, relationships, status, identity, security, and reputation, not just money.
- Interests are usually mixed. Before asking “Who’s right?”, ask: Where are our interests aligned, and where are they different?

- Look at what everyone can do. **Wanting something and being able to do something are not the same thing.** Game theory keeps available choices separate from preferences.

Example to use: Negotiating with a salesperson. You think it's you vs. the salesperson. But they may need a manager's approval, the company may have a pricing policy, another customer may want the same product, and you may have three other suppliers. The strategic situation looks very different from the surface conversation.

On-screen takeaway: Who can actually affect the outcome?

Lesson 1.3 — Constraints, alternatives, and dependence (10 min)

Teaching points

- Constraints are part of the game. "I simply can't go below \$10,000" might be a bluff, a real company minimum, or an unwillingness to go lower. These look identical on the surface and are completely different strategically. Separate what people **say** constrains them from what actually **constrains** them.
- Ask: What happens if we don't agree? Outside options shape the negotiation itself, not just what happens after it fails. But their effect depends on the game's structure. There's no rule that "the person with the better alternative always wins."
- Power often comes from the structure of **dependence and alternatives**. Ask: Who needs whom more? That isn't a persuasion tactic. It's a property of the game.

Example to use: The contractor negotiation. You want a lower price; they want theirs. Instead of only persuading, ask what happens if you both walk away. Maybe you have four other contractors and they have ten customers waiting. Maybe neither has a good alternative, or one of you badly needs the deal. Each version is a different game, even though the conversation sounds the same.

On-screen takeaway: Wanting something and being able to do something are not the same thing.

Lesson 1.4 — Timing, moving games, and the game behind the conversation (10 min)

Teaching points

- **When you act matters.** Responding on the spot, seeing the other side's offer before making yours, and choosing without knowing what they'll do are different games. Game theorists separate simultaneous-move from sequential games because timing and observability change what's possible.
- A game isn't static. Your move can change what the other person knows, believes, and expects, and which choices remain open. A broken promise, a revealed fact, or a public commitment reshapes the board.
- The stated issue may not be the whole issue. "I don't care about the money" may be about recognition. The lesson isn't "assume everyone is lying," which is paranoia. It's:

don't assume the surface description tells you everything about the strategic structure. Stay curious.

- Games can sit inside games, and you can have the perfect words and still be playing the wrong game.

Example to use: Two people argue about which business strategy to use. Underneath may be other questions: who gets final authority, who gets credit or blame, and who is setting a precedent. The visible disagreement is only one layer.

On-screen takeaway: You can have the perfect words and still be playing the wrong game.

Lesson 1.5 — Person problem or game problem? Building the Game Map (8 min)

Teaching points

- A game doesn't have to be competitive. You can view a marriage or a friendship through game theory without treating the other person as an opponent.
- Before diagnosing a **person problem**, check for a **game problem**. Behavior often comes from the structure of the interaction. That doesn't remove responsibility, but it changes where you look for a fix.
- The Game Map trains your attention. It isn't a rigid nine-step procedure. With practice, you notice these things automatically. That's what a mental model does.
- The first mental shift: from **Me → What should I say? → Them** to **Me → My move → Their response → My next move → Their next response → Outcome**.

Example to use: A manager says, "My employees aren't cooperating," and concludes they're selfish. But maybe the pay system rewards outperforming each other, nobody knows who owns the task, or speaking up has been punished before. Likewise, you think a coworker is procrastinating while they think you'll take over if they start. You both want the project done, but your expectations of each other create the problem.

On-screen takeaway: Before you diagnose a person problem, check for a game problem.

Key terms

Term	Plain-language meaning
Game	A situation in which the outcome depends, at least partly, on the decisions of multiple people. It can be cooperative, competitive, or mixed.
Player	Any decision-maker in the interaction who can affect the outcome, including people who aren't in the room.

Term	Plain-language meaning
Mixed interests	Interests that are neither fully aligned nor fully opposed.
Available actions (strategies)	What each player can actually do, as distinct from what they want to do.
Constraint	Anything that limits a player's choices, such as a policy, budget, approval requirement, or veto. A stated constraint and a real constraint aren't always the same.
Outside option / alternative	What happens to each player if the interaction produces no agreement.
Dependence	Who needs whom, and who has alternatives. Power often comes from this structure rather than from personality.
Simultaneous vs. sequential game	Choosing without seeing the other's move, versus one moving first and the other responding.
Game behind the conversation	The structure beneath the stated topic, such as trust, status, credit, or precedent.
Person problem vs. game problem	Blaming behavior on someone's character, versus recognizing that the structure of the interaction may be producing it.
Game Map	The chapter's nine-question perception tool: WHO, WHAT, OPTIONS, CONSTRAINTS, INFORMATION, TIMING, ALTERNATIVES, DEPENDENCIES, THEN WHAT.

Research spotlight

1. The “Wall Street Game” vs. the “Community Game” (Lieberman, Samuels, and Ross, 2004)

What was done: Participants, including American college students and Israeli pilots and instructors, played essentially the same Prisoner's Dilemma. The only change was the name. Some were told they were playing the “Wall Street Game” and others the “Community Game.”

What was found: The label had a substantial effect on whether people cooperated or defected, larger than observers expected. Meanwhile, judgments of participants' prior reputations did a poor job of predicting who would cooperate.

What it suggests for real conversations: People's construal of the situation, meaning what kind of game they believe they're in, can materially change how they behave inside it. Two people can face nearly identical incentives and act very differently because they read the interaction differently. How you identify the game matters.

What it does NOT prove: It doesn't show that consciously mapping a game leads to optimal decisions. It also doesn't mean the game is "only perception." There's still an underlying strategic structure, and people's interpretations of it can be accurate or inaccurate.

2. Awareness of the frame changes strategic play (Charness and Sontuoso, 2023)

What was done: In a study of "frame-dependent rationalizability," participants played coordination games built around visual objects. The researchers varied which features of the choices players were made aware of.

What was found: Changing which frames players were aware of changed their strategic behavior.

What it suggests for real conversations: People don't necessarily notice every strategically relevant feature of the situation in front of them. That's a strong reason to ask yourself on purpose: What am I failing to notice about the game?

What it does NOT prove: It was a controlled lab coordination task, not a workplace or family conversation, and it doesn't show that any particular habit of mind improves everyday outcomes.

A counterweight. Colin Camerer's review of behavioral game theory shows that framing, social preferences, and learning can push real behavior away from simple game-theoretic predictions. The models don't claim people are perfectly rational. They help you spot strategic forces you'd miss by looking only at words and personalities.

The evidence standard for this module: Representing situations in terms of players, choices, payoffs, information, and timing is basic game theory, and the evidence that construal affects play is strong. Evidence that teaching people to "see the game" improves real-world communication **has not been established**. The Game Map is a practical synthesis, not a tested intervention.

The tool: The Game Map

Use this before any important interaction. Write short answers and don't aim for perfection. The goal is to train your attention, not to complete a form.

The interaction:

WHO?

- Who are the relevant players?

-
- Who can actually affect the outcome, including anyone not in the room?
-

WHAT?

- What does each person want?
-
- Where are our interests aligned?
-
- Where are they different?
-

OPTIONS?

- What can each player actually do?
-

CONSTRAINTS?

- What limits those choices?
-
- Which constraints are stated, and which do I know are real?
-

INFORMATION?

- Who knows what?
-
- Who doesn't know what?
-

TIMING?

- Who moves first?
-
- Who responds?
-
- What can each person observe before choosing?
-

ALTERNATIVES?

- What happens if there is no agreement, for me and for them?
-

DEPENDENCIES?

- Who needs whom?
-

- Who has alternatives?
-

THEN WHAT?

- If I make this move, what is the other person likely to do next?
-

- And what would I do after that?
-

Final check

- Could there be a game behind the conversation, such as credit, trust, status, authority, or precedent?
-

- Is this a person problem or a game problem?
-

- What am I failing to notice about the game?
-

Practice exercises**Exercise A — Analyze the scenario**

Dana runs operations at a regional distribution company. For six weeks she has been pushing Marcus, the warehouse lead, to adopt a new scheduling system. Marcus keeps saying it “isn’t ready” and that his team “doesn’t have time to learn it.” Dana has started to see him as stubborn and plans to tell him at Thursday’s one-on-one that the switch is happening next Monday, no exceptions. Here’s what Dana hasn’t thought about. Marcus’s quarterly bonus is tied to on-time shipment rates, and any slowdown during the learning period will hit his numbers. The software vendor’s onboarding team is booked for the next three weeks. Two of Marcus’s most experienced pickers have quietly told him they’d rather leave than learn “another system.” Dana’s own boss, the VP, approved the system and asked her for a progress update at the end of the month. Marcus knows other warehouses in the area are hiring leads.

1. Who are the players in this game, and which ones aren’t in Thursday’s meeting?
2. Where are Dana’s and Marcus’s interests aligned, and where do they differ?
3. What constraints shape Marcus’s options, and what alternatives does each of them have if they don’t reach agreement?

4. If Dana announces “next Monday, no exceptions,” what is Marcus likely to do next? Then what?
5. Is this primarily a person problem or a game problem?

Model answer

1. Players. In the room: Dana and Marcus. Outside it: the VP, who wants an update; the two senior pickers, who could quit; the vendor’s onboarding team, whose availability sets the real start date; and local employers hiring leads, part of Marcus’s outside option. Dana has been treating this as Dana vs. Marcus. It’s a much bigger game.

2. Interests. Aligned: both want the warehouse to run well, both want to keep good people, and neither wants a failed rollout. Different: Dana needs visible progress for her VP by month’s end, while Marcus needs to protect on-time rates that drive his bonus and to keep his best pickers. Their interests are mixed, not opposed.

3. Constraints and alternatives. Marcus’s bonus structure penalizes a slowdown. The vendor can’t onboard anyone for three weeks, so “next Monday” may not even be feasible. Marcus can’t force veteran pickers to stay. If there’s no agreement, Marcus can leave for another warehouse, the pickers can leave, and Dana faces a VP update with a stalled rollout and possibly a staffing gap. Dana’s alternative to Marcus is replacing a lead mid-rollout, which is costly. Dependence runs both ways, and Marcus has more alternatives than Dana has assumed.

4. Then what? A hard Monday deadline is a commitment Dana may not be able to keep, since the vendor isn’t available. Marcus could comply on paper while shipments slip, escalate, start interviewing, or watch his pickers resign. Each makes Dana’s month-end update worse.

5. Diagnosis. Mostly a game problem. Marcus’s “stubbornness” is a reasonable response to an incentive system that punishes the change, plus real staffing risk. He still has responsibility, but the structure is driving the behavior. A better first move: ask Marcus what would make the rollout workable, check the vendor’s timeline, and raise the bonus-metric conflict with the VP. That changes the game instead of escalating inside it.

Exercise B — Your own game

Pick a live situation you’ve been thinking about mostly in terms of what to say: a conversation you’re dreading, a request, or a recurring friction.

1. Describe the situation in one sentence, as you’ve been seeing it.
2. Fill out the full Game Map. Circle any answer you’re unsure of. Those are your information gaps.
3. Name one player you hadn’t been considering, and one constraint or alternative that changes the picture.
4. Write the stated issue and one possible game behind it. What would tell you whether the stated issue is the real one?

5. Answer THEN WHAT for two different opening moves. Which leaves both of you with better options?
6. Finish this sentence: “I was about to walk into this thinking only about

. What’s really going on is

.”

Exercise C — Partner drill / role-play

Format: Pairs, about 15 minutes. One person plays the Employee and one plays the Manager. An observer is optional, or you can practice with an AI practice partner by giving it the Manager’s private brief.

Setup (both see this): The Manager has just sent the Employee a text: “We need to talk about your performance.” The meeting starts in five minutes.

Employee’s private brief: You don’t know what this is about. You have an unmentioned informal offer from another company, and you think your major project last month went unnoticed.

Manager’s private brief: You wrote that text in a hurry. You actually want your strongest performer, the Employee, to take over a struggling project. Your director is pressing you to fix it this quarter, and you can’t offer a raise until the next budget cycle.

Round 1 (3 minutes): Each player fills out a quick Game Map from their own position, noting what they know, don’t know, and assume.

Round 2 (5 minutes): Hold the conversation. The Employee’s job is to understand the game before defending themselves: ask what the Manager wants and who else is involved.

Round 3 (5 minutes), debrief:

- Employee: How much of your map was wrong? What did you assume from the wording of the text?
- Manager: How did your words change the Employee’s decision problem? What would you send instead?
- Both: How did the players outside the room (the director, the other company) and each side’s constraints shape the conversation? Where were your interests aligned?

If time allows, switch roles and rerun it with the Manager stating the real purpose up front. Notice how the game changes.

Common mistakes

- **Starting with the words.** Rehearsing the perfect sentence before understanding the situation. Words enter a system of incentives, information, and alternatives. Map the situation first.
- **Treating the person in front of you as the whole game.** Managers, policies, partners, spouses, other customers, and alternative suppliers may all affect the outcome. Ask who can actually affect it.
- **Assuming “game” means “competition.”** Most interactions mix common and conflicting interests. Seeing the game should make cooperation easier.
- **Confusing wanting with being able.** Someone may want to give you a refund or a raise and still not be able to. Look at their available actions and constraints.
- **Assuming a stated constraint is a lie, or that it’s true.** “I can’t go below \$10,000” could be a bluff, a real policy, or a preference. Ask what actually constrains them.
- **Ignoring what happens if there’s no agreement.** Alternatives and dependence shape the conversation itself, though a better alternative doesn’t automatically win.
- **Becoming paranoid about hidden motives.** Sometimes the stated issue really is the issue. Stay curious instead of suspicious.
- **Diagnosing a person problem when it’s a game problem.** Before concluding someone is selfish or lazy, check whether incentives, ownership, or history are producing the behavior.
- **Running the Game Map as a rigid script.** It trains attention. The goal is to notice these things automatically over time.

Quiz

1. In the game-theoretic sense used in this course, a “game” is best defined as:
 - A) A competition in which one person wins and the other loses
 - B) A situation in which the outcome depends, at least partly, on the decisions of multiple people
 - C) Any interaction in which someone is trying to manipulate someone else
 - D) A formal model that requires calculating an equilibrium
2. According to the chapter, what question does game theory ask before “Was that the right thing to say?”
 - A) “How do I build more rapport?”
 - B) “Which persuasion technique fits here?”
 - C) “What happens after I say it?”
 - D) “Who is right?”

3. A vendor tells you, “I’d love to give you that price, but I simply can’t go below \$10,000.” The approach the chapter recommends is to:
- A) Assume they’re bluffing and push harder
 - B) Accept the number, since they’ve stated a constraint
 - C) Ask what actually constrains them, because a bluff, a real policy, and an unwilling-but-able seller look identical on the surface
 - D) Walk away immediately to show strength
4. Two business partners both want their company to succeed but disagree about pay, reinvestment, and credit. Their interests are best described as:
- A) Fully aligned
 - B) Fully opposed
 - C) Mixed
 - D) Irrelevant to the game
5. Which statement about outside options matches the chapter?
- A) The person with the better alternative always wins.
 - B) Alternatives matter only after a negotiation fails.
 - C) Alternatives can change the negotiation itself, but their effect depends on the structure of the game.
 - D) Alternatives are a persuasion tactic you should hide.
6. A manager says, “My employees aren’t cooperating.” Their pay system rewards outperforming one another and nobody knows who owns key tasks. The chapter would most likely call this:
- A) A person problem, because the employees are selfish
 - B) A possible game problem, because the structure may be producing the behavior
 - C) A communication problem that better wording would solve
 - D) Proof that the employees bear no responsibility
7. You and a coworker both want a project finished. You think they’re procrastinating, and they think you’ll take over if they start. What does this example illustrate?
- A) Your coworker has a character flaw.
 - B) Shared goals guarantee smooth cooperation.
 - C) Expectations about each other can create a problem even when both people want the same outcome.
 - D) The situation is purely competitive.
8. Which statement best reflects the evidence on “seeing the game”?
- A) Studies show that people who use a Game Map always make better decisions.
 - B) Research shows that how people construe and notice a strategic situation can affect their behavior, but it hasn’t been shown that teaching people to see the game improves real-world communication.

- C) The Wall Street/Community Game study proves that labels don't matter.
 D) Game theory assumes people are perfectly rational, and experiments confirm it.

Answer key

1. **B** — A game is any situation whose outcome depends partly on several people's decisions. It can be cooperative, competitive, or mixed.
2. **C** — Your sentence changes the other person's decision problem, and their response changes yours. "Then what?" comes first.
3. **C** — Situations that look identical on the surface can be completely different strategically.
4. **C** — Their interests are neither completely aligned nor completely opposed, the mixture Schelling emphasized.
5. **C** — Outside options shape the interaction itself, but there's no simple rule that the better alternative always wins.
6. **B** — The incentives and unclear ownership may be producing the behavior. That doesn't erase responsibility, but don't automatically diagnose a person problem.
7. **C** — Both want the same outcome, but their expectations about each other create the impasse. Seeing that structure lets you change it.
8. **B** — Experiments (Liberman, Samuels, and Ross, 2004; Charness and Sontuoso, 2023) show construal and awareness affect strategic behavior. That teaching people to see the game improves real communication has not been established.

Field assignment

This week: map one conversation before you have it.

Choose one meaningful interaction coming up this week. Before it happens, spend 10 minutes on a Game Map instead of a script, and mark which answers are assumptions rather than facts.

During and after the conversation, observe and record: - Any player who turned out to matter that you hadn't listed - Any constraint or alternative that was different from what you assumed - One moment where the other person's response changed your next move (THEN WHAT in action) - Whether the stated topic was the whole issue, or whether there was a game behind the conversation

Debrief questions for next week: 1. What did you notice about the situation that you would have missed if you'd only prepared what to say? 2. Which assumption was most wrong, and what made the other person's behavior more reasonable from their position? 3. Was the main obstacle a person problem or a game problem, and what would change the game next time?

The question to carry

| **What's really going on here?**

Module 2 — See the Payoffs

The question: What payoff would make their behavior make sense?

Book reading	Chapter 2, “See the Payoffs”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Payoff Map

Why this module matters

In Module 1 you learned to see the game before choosing your move. But identifying the players isn’t enough. You also have to understand what each player is trying to accomplish. And here is a mistake we make all the time: we assume other people want what we want. When they don’t behave the way we expect, we conclude that they’re irrational, difficult, selfish, stubborn, or just wrong. Sometimes they are. But often they’re simply responding to a different set of incentives. They’re playing the same game, but they’re playing for a different payoff.

A payoff is simply what makes an outcome better or worse for a player. Money is one payoff. So are time, security, status, fairness, reputation, control, relationships, and identity. If you only look at the dollar amount, you can completely misunderstand the game.

This module builds one mental shift. Instead of asking, “Why would they do that?” you’ll ask, “What payoff would make that behavior make sense?” That doesn’t mean you’ll agree with them or excuse them. It means you stop treating your own preferences as the universal standard for rational behavior. And it is understanding, not manipulation: the goal is never to find someone’s hot button and push it, but to understand the incentive structure before choosing your move.

Learning objectives

By the end of this module you will be able to:

1. Define a payoff in plain language and name the non-monetary currencies that can drive a decision.
2. Distinguish opposed, aligned, and mixed interests, and locate where a disagreement conflicts and where it aligns.
3. Separate stated from revealed preferences, and what happened from what each person thinks happened.

4. Identify what each player is trying to avoid, not just what they're trying to gain.
5. Complete a Payoff Map and use it to choose a move that respects the other person's autonomy and dignity.

Lessons

Lesson 2.1 — What is a payoff? (~10 min)

Teaching points

- In game theory, a player's payoff is what they value about an outcome. The working definition: What makes this outcome better or worse for this player?
- Money is one currency among many: time, security, status, freedom, fairness, reputation, control, relationships, avoiding embarrassment, reducing risk, protecting one's identity.
- Sometimes the payoff is what the outcome represents. A person may fight over \$500 because of what conceding it means: respect, principle, or not being taken advantage of.
- The most important question: What are they actually trying to get out of this? Not what you think they should want, not what you'd want in their shoes, and not even necessarily what they said they want.

Example to use: The raise you didn't get. You want more money and assume your boss wants to save money. But your boss may also want to retain good people, stay within budget, avoid an expensive precedent, keep the team happy, and avoid explaining the raise upstairs. Your boss may even want to give you the raise but believe they can't afford it. The game shifts from "they don't want to pay me" to "We both want me to stay, but we have different constraints about how compensation is structured."

On-screen takeaway: A payoff is whatever makes an outcome better or worse for that player, and money is only one currency.

Lesson 2.2 — Opposed, aligned, and mixed interests (~10 min)

Teaching points

- **Opposed interests:** what benefits me hurts you. Zero-sum games like chess are the strong form, and they're rarer in everyday relationships than we imagine.
- **Aligned interests:** we both benefit from the same outcome. The problem is often coordination, not conflict: "How do we get our actions aligned?"
- **Mixed interests:** where most of human life happens. We both want the deal, company, or relationship to succeed, but disagree about how.
- Replace "How do I beat them?" with "Where are our interests actually aligned, and where are they opposed?"

- The pie doesn't have to be fixed. Before asking "Who gets what?" ask "Can we create more value before deciding how to divide it?" Conflict over one dimension doesn't mean conflict over everything.
- But don't turn this into a "there's always a win-win" philosophy. Sometimes payoffs really are incompatible. Misdiagnose a conflict as cooperation and you'll make bad decisions; misdiagnose cooperation as pure conflict and you'll leave value on the table.

Example to use: Buying a house. You want to pay \$500,000; the seller wants \$550,000. You conflict over price but align on completing the transaction. And in a business deal where the seller cares most about getting paid quickly while you care about paying over time, restructuring the terms can increase the total value even though you still disagree about the final allocation.

On-screen takeaway: Conflict over one dimension doesn't mean conflict over everything.

Lesson 2.3 — Hidden payoffs and why "rational" doesn't mean "selfish" (~10 min)

Teaching points

- The visible payoff may not be the complete payoff. Someone who refuses to apologize may care about the relationship but believe apologizing means admitting they were entirely at fault: relationship plus dignity plus self-image. The lesson is not "everyone has a secret hidden motive."
- Empathy and game theory meet here. Don't imagine yourself in their position with your preferences. Reconstruct their game: What do they value? What are they afraid of losing? What alternatives and constraints do they have? What would feel like winning, or losing?
- The Prisoner's Dilemma shows that payoff structure can push decent people toward a behavior. Sometimes behavior is produced by the structure of the game rather than the character of the players.
- The Ultimatum Game shows people often reject offers they see as unfair, even when rejecting leaves them with nothing. If people consistently sacrifice money to punish perceived unfairness, fairness is part of the game.
- "That person isn't being rational" often means "That person isn't maximizing the thing I assumed they cared about."
- Boundary: understanding someone's incentives doesn't mean their behavior is justified. Game theory explains why a behavior may make sense; it doesn't tell us whether it's good, moral, or wise.

Example to use: The lucrative job. You turn down a high-paying job because it would mean missing your children's lives. Financially, it looks irrational. But if your priorities include family, it makes perfect sense. The observer misunderstood your payoff function.

On-screen takeaway: Sometimes the behavior we observe is being produced by the structure of the game rather than the character of the players.

Lesson 2.4 — Gains, losses, perception, and revealed preferences (~10 min)

Teaching points

- Ask: What are they trying to avoid? Embarrassment, uncertainty, conflict, lost status, financial risk, looking incompetent, being controlled. Same upside, different perceived downside: “Let’s try it” versus “What if it fails?”
- Kahneman and Tversky’s prospect theory: choices often depend on gains and losses relative to a reference point, and losses can weigh more than equivalent gains in many contexts (not always).
- Perception is part of the game. The same deal can read as “I’m getting \$10,000” or “I’m losing \$5,000 compared with what I expected.” Distinguish what happened from what each person thinks happened.
- Stated versus revealed preferences: choices give you information beyond words. That doesn’t mean people are lying; they often don’t know their preferences until they face a real choice.
- The payoff can change. Switch from individual bonuses to a team bonus and the people don’t change; the game does. Before trying to change behavior, understand what the environment rewards.

Example to use: A manager gives two employees more responsibility. Employee A thinks, “They trust me.” Employee B thinks, “They’re dumping more work on me.” Same action, different perceived payoff, different likely response.

On-screen takeaway: The same objective outcome can produce different decisions depending on what someone perceives themselves as gaining or losing.

Lesson 2.5 — The Payoff Map and how it changes communication (~5 min)

Teaching points

- The Payoff Map adds a second layer to the game map from Module 1: what I want, what they want, where we align, where we conflict, what we’re each trying to avoid, and what I might be missing.
- The most important question on the map: Am I assuming their payoff is the same as mine?
- Before asking “How do I convince them?” ask “What payoff are they protecting?”
- Knowing what someone values tells you what they’ll likely respond to, resist, or see as a threat. That is not license to “find their hot button and push it.” That’s tactics.
- The map may help you persuade, negotiate, or cooperate, or tell you agreement isn’t possible, or that you’re solving the wrong problem.

Example to use: The curfew. A parent says, “Be home by midnight.” The teenager says, “That’s ridiculous.” Framed as parent versus teenager, it’s a standoff. Seen as safety versus independence, both positions make sense.

On-screen takeaway: Understand the incentive structure before choosing your move.

Key terms

Term	Plain-language meaning
Payoff	What makes an outcome better or worse for a particular player.
Payoff function	The full set of things a person values and how they weigh them.
Opposed interests	What benefits one player hurts the other; zero-sum is the strong form.
Aligned interests	Both benefit from the same outcome; the challenge is coordination.
Mixed interests	Shared goals on some dimensions, conflict on others.
Hidden payoff	A value such as status or dignity that drives a decision but isn’t the visible stake.
Ultimatum Game	One person proposes a split; the other accepts or rejects, and rejection leaves both with nothing.
Reference point	The baseline, often an expectation, against which an outcome feels like a gain or a loss.
Stated vs. revealed preference	What a person says they want vs. what their choices show they value.

Research spotlight

1. The Ultimatum Game: fairness is part of the payoff

What was done. One person proposes how to divide a sum of money; the responder accepts (both get the proposed amounts) or rejects (both get nothing). Güth, Schmittberger, and Schwarze’s 1982 experiment is the classic starting point. A later meta-analysis pooled 37 papers containing 75 experimental results, and Slonim and Roth (Econometrica, 1998) ran versions in Slovakia with stakes varied by a factor of 25.

What was found. A money-only model says accept anything positive, since \$1 beats \$0. People often don't. In the meta-analysis, proposers offered about 40% of the pie on average and responders rejected about 16% of offers; rejection was less common with larger offers and larger total stakes. Slonim and Roth found that raising stakes had only a small effect on inexperienced players: rejections fell somewhat, and experienced proposers gradually offered less.

What it suggests. The payoff isn't only "How much did I get?" It can also be "How was I treated?" Price only the material outcome and an insulted person's walkout will look irrational when it isn't.

What it does NOT prove. Not that people always choose fairness over money, or that money doesn't motivate people. The weight on fairness varies across people, cultures, situations, and stakes, and higher stakes reduce some of these effects without eliminating them.

2. Falk and Kosfeld: autonomy can be part of the payoff

What was done. Armin Falk and Michael Kosfeld (American Economic Review, 2006) ran a principal-agent game in which principals could impose a minimum performance requirement on agents.

What was found. Instead of performing better under control, most agents reduced performance in response to it. Many described the control as a signal of distrust and a limit on their autonomy.

What it suggests. Changing someone's material options can also change what the outcome means to them. A rule meant to secure performance can read as "I don't trust you."

What it does NOT prove. Not that oversight always backfires; the evidence is good but context-dependent. Similarly, Gneezy and Rustichini's daycare study (a late-pickup fine followed by more late pickups) is a vivid example that incentives can change social meaning, but a later replication attempt found deterrent effects from fines. Treat these as possibilities, not universal laws.

The tool: The Payoff Map

Use this map before any important interaction. It sits on top of your Game Map from Module 1.

The interaction:

Players (from your Game Map):

1. What do I want?

What outcomes matter to me? (List more than one currency.)

2. What do they want?

What outcomes appear to matter to them? (Not what they should want, not what I'd want in their shoes, not only what they said.)

3. Where do our interests align?

What can we both gain?

4. Where do they conflict?

What does one person gain at the other's expense?

5. What are we each trying to avoid?

What losses or risks matter? (Me / Them)

6. What might I be missing?

Could status, identity, fairness, autonomy, or relationships matter here?

7. Interest structure (circle one): Mostly opposed · Mostly aligned · Mixed

The most important question: Am I assuming their payoff is the same as mine?

My next move, given this map (and one that respects their autonomy and informed choice):

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small marketing agency. Her best designer, Marcus, has been there four years. Dana offers him a promotion to creative director with a 15% raise, and she expects him to be thrilled. Instead, Marcus hesitates and asks for a week to think. In a hallway conversation he mentions that the new role “means a lot more meetings” and that he’d “stop doing the actual design work.” A few days later Dana learns from a colleague that Marcus was hurt that the promotion was announced to the leadership team before it was offered to him. Dana is frustrated: “I’m offering him more money and a better title. Why is he being so difficult?” She is considering adding another 5% to close the deal. Marcus has said several times that he “doesn’t care about titles,” but last year he turned down a higher-paying offer from a larger firm.

1. What payoff is Dana assuming Marcus has, and what is the evidence that his payoff function is different?
2. List at least three currencies that may be in play for Marcus, including at least one he’s trying to avoid.
3. Is the interest structure between Dana and Marcus opposed, aligned, or mixed? Where specifically?
4. What does his past behavior (the turned-down offer) reveal compared with what he has stated?
5. Why might adding another 5% fail to solve the problem, and what would a better next move look like?

Model answer

1. Dana assumes Marcus values what she’d value: more money and a bigger title. The evidence says otherwise: he hesitated, he described the role in terms of what he’d lose, and he was hurt by how the offer was handled. “Why is he being so difficult?” signals that she may be misreading his payoff function, not that he’s irrational.
2. Likely currencies: **identity and craft** (he sees himself as a designer, and the role takes him away from “the actual design work”); **autonomy and time** (more meetings, less control over his day); **respect** (leadership heard before he did). He may be trying to avoid losing the work he’s good at or struggling in a management role he never sought. Money matters, but it doesn’t look like the binding payoff.
3. Mixed. Aligned: both want Marcus to stay and the agency’s creative work to be excellent. Opposed: Dana needs someone to carry management; Marcus may not want it.
4. Turning down a higher-paying offer is a revealed preference that backs up his stated one: money isn’t his top currency. It isn’t proof, but it’s meaningful information.
5. Another 5% doubles down on the currency that isn’t driving his decision, and could even read as “she thinks I can be bought.” A better move is to ask, not push: acknowledge the

announcement misstep, then ask what a good outcome would look like for him and what would feel like a loss. That may reveal a value-creating option, such as a senior designer track with a raise and mentoring but no management load. Or it may show the interests are genuinely incompatible, which is better to know than to pretend a win-win exists.

Exercise B — Your own game

Pick a live situation where someone's behavior frustrates or puzzles you: a coworker resisting a plan, a family member refusing a reasonable request, a client who won't commit.

1. Write down the sentence you've been thinking ("Why would they..." or "They're being..."), then rewrite it as "What payoff would make that behavior make sense?"
2. Complete the Payoff Map for this situation.
3. Answer: "What outcome would make this person feel that they got what they wanted?" and "What outcome would make them feel that they lost?"
4. Note one place where you've treated their stated preference as the whole story, and one choice of theirs that reveals something else.
5. Decide whether the interests are opposed, aligned, or mixed. If they're genuinely opposed, write that down plainly; knowing what kind of game you're in is useful too.
6. Draft one honest, open question that would test your map rather than confirm it.

Exercise C — Partner drill / role-play

Time: 15 minutes. **Format:** pairs, or one person with an AI practice partner.

Setup: Two co-owners of a small bakery are deciding what to do with this year's \$60,000 profit.

Jordan (private): You want to reinvest most of it in a second location. Growth matters a lot, and your biggest fear is the business plateauing and looking small-time to people who doubted you. You're comfortable with risk.

Riley (private): You want to distribute most of it. You're quietly helping a parent with medical bills, so financial security matters a lot and debt makes you anxious. You'd like the business to grow, just not in a way that risks your income this year.

Round 1 (5 min): Argue positions only. Try to convince each other.

Round 2 (7 min): Start over. Your only job now is to fill in the other person's Payoff Map by asking: What are you trying to gain? What are you trying to avoid? What would feel like a loss? Share only what you're comfortable sharing. Then look for a deal that creates value before dividing it.

Debrief questions:

- In Round 1, what did you assume the other person was optimizing for? How close were you?
- What payoffs surfaced in Round 2 that never came up in Round 1?
- Did you find an arrangement (for example, a smaller expansion plus a guaranteed distribution) that gave each of you more of what you cared about most? If not, was the conflict real?
- Did anyone feel pushed or “worked”? What separates understanding a payoff from exploiting it?

Common mistakes

- **Assuming they want what you want** — The mistake the chapter is built around. Ask “Am I assuming their payoff is the same as mine?” before interpreting their behavior.
- **Calling someone “irrational” too fast** — It usually means they aren’t maximizing what you assumed they cared about. Ask what payoff would make the behavior make sense.
- **Treating money as the only currency** — Status, fairness, autonomy, identity, and relationships can matter as much or more.
- **Turning hidden payoffs into conspiracy theories** — The lesson isn’t that everyone has a secret motive; it’s that the visible payoff may not be the complete payoff.
- **Assuming every interaction is zero-sum, or that there’s always a win-win** — Most interactions are mixed. Identify conflict accurately rather than inventing it or wishing it away.
- **Taking stated preferences as the whole story** — Watch choices as well as words, without concluding people are lying.
- **Confusing understanding with excusing, or with exploiting** — Seeing why a behavior makes sense doesn’t make it good, and knowing what someone values is not a hot button to push.

Quiz

1. According to the chapter, what is a payoff?
 - A) The amount of money a player receives at the end of a game
 - B) What makes an outcome better or worse for a particular player
 - C) A hidden psychological motive that people try to conceal
 - D) The reward you offer someone to change their behavior

2. Two business partners disagree about reinvesting versus distributing profits. One prioritizes growth, the other security. What best describes this?
 - A) One partner is being irrational
 - B) It's purely a disagreement about money
 - C) They are optimizing for different things
 - D) It's a zero-sum game
3. A buyer and seller disagree on price but both want the sale to close. This is best described as:
 - A) Opposed interests
 - B) Aligned interests
 - C) Mixed interests
 - D) A coordination problem only
4. In the Ultimatum Game, responders often reject offers they see as unfair, leaving both with nothing. What is the chapter's lesson?
 - A) The mathematical model is useless
 - B) People don't care about money
 - C) Fairness can be part of the actual payoff structure
 - D) Responders misunderstand the rules
5. Someone says "Price is my only concern," then rejects a cheaper option because it takes longer to deliver. What does the chapter suggest?
 - A) They are lying and can't be trusted
 - B) Their behavior reveals that price may not be their only concern
 - C) Delivery time is irrelevant to the game
 - D) You should ignore what people say entirely
6. Two employees get the same new responsibility. One feels trusted; the other feels dumped on. Which distinction does this illustrate?
 - A) Stated versus revealed preferences
 - B) What happened versus what each person thinks happened
 - C) Zero-sum versus positive-sum games
 - D) Short-term versus long-term payoffs
7. A company switches from individual bonuses to a team bonus, and two rivals start helping each other. What changed?
 - A) The people's character
 - B) Nothing; they were always cooperative
 - C) The game and its incentives
 - D) Their stated preferences

8. Which response best reflects the chapter's ethical use of payoff thinking?

- A) "Find their hot button and push it."
- B) "Once I understand their incentives, their behavior is justified."
- C) "Understand the incentive structure before choosing your move, without exploiting it."
- D) "There is always a win-win if I look hard enough."

Answer key

1. **B** — A payoff is what makes an outcome better or worse for that player, money or otherwise.
2. **C** — Neither is necessarily irrational; they optimize for growth versus security.
3. **C** — Conflicting interests over price plus aligned interests over completing the transaction make it a mixed-interest game.
4. **C** — If people consistently sacrifice money to punish perceived unfairness, fairness is part of the game. The model isn't useless; it learns the real payoffs.
5. **B** — Choices reveal what matters. That doesn't mean they're lying; people often learn their preferences by facing a real choice.
6. **B** — Same action, different perceived payoff.
7. **C** — The people didn't change; the game did.
8. **C** — The chapter rejects "push their hot button" as tactics, says understanding doesn't mean justifying, and warns against a "there's always a win-win" philosophy.

Field assignment

This week, catch yourself in one "Why would they do that?" moment at work or at home. Before you respond or judge, sketch a quick Payoff Map, even a few lines on your phone. Then, if appropriate, ask one honest, open question to test it, such as "What would a good outcome look like for you here?" or "What are you most worried about?" Record the behavior, your first interpretation, the payoffs you guessed (gain and avoid), and what you learned.

Debrief questions for next week: 1. Where was your first interpretation based on assuming their payoff was the same as yours? 2. What payoff did you miss, and which currency was it? 3. Did understanding their payoffs change your next move, or confirm that your interests were genuinely opposed?

The question to carry

What payoff would make their behavior make sense?

What are they trying to gain? What are they trying to avoid? Where are our interests aligned? Where are they genuinely opposed? What am I assuming they value that they may not?

Module 3 — Think in Responses

The question: “If I do this, what are they likely to do next? And if they do that, what will I do?”

Book reading	Chapter 3, “Think in Responses”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Response Map

Why this module matters

In Module 1 you learned to see the game. In Module 2 you learned to see the payoffs. This module is about the idea that sits at the heart of the whole book: **your move is only half of the move**. The other half is what the other person does in response, and their response depends, at least partly, on what you did.

Most of us treat communication as a straight line: I say something, they hear it, outcome. Real interaction is a chain: I act, they respond, I respond to their response. When we forget that, we judge our moves by how they feel to make (“I was just being honest”) instead of by what they cause. We confuse **intention** with **incentive**.

The shift is simple to say and hard to practice. Before an important move, don’t stop at “What should I say?” Ask what it’s likely to make them do, what you’ll do then, and, when stakes are high, where the sequence leads. Once another person is involved, your move doesn’t end when you stop talking. Your move becomes their situation.

Learning objectives

By the end of this module you will be able to:

1. **Explain** what a best response is and why the value of any move depends on the other person’s likely reply.
2. **Distinguish** between what you intend a move to produce and what the move actually makes attractive for the other person.
3. **Predict** a set of plausible responses to a planned move using the other person’s goals, information, options, and constraints, rather than your own preferences.
4. **Reason backward** from where a sequence is likely to end, and use that reasoning to test whether a threat or commitment is credible.

5. **Map** an interaction as a sequence of interdependent decisions, find its turning point, and spot escalation loops.

Lessons

Lesson 3.1 — Your move is only half of the move (~8 min)

Teaching points

- Communication isn't a straight line ("I say it → they hear it → outcome"). It's a chain: I act → they respond → I respond to their response → they respond to mine.
- A **best response** is the action that gives a player the best outcome given what the other player is doing. Nash equilibrium is built from it: each player's strategy is a best response to the others, so no one gains by changing alone.
- You don't need the math. You need two habits: "If I do this, what are they likely to do next?" and "If they do that, what will I do?"
- Restaurant example: if your friend is definitely going to the Italian place, your best response may be to go there too; if they're going to the steakhouse, it changes. Your best move is never determined in isolation.

Example to use: The employee who plans to say, "I have another job offer." Then what? The boss might offer more money, ask to see the offer, call the bluff, say they can't compete, worry you're leaving, decide you're no longer committed, or start preparing for your departure. A boss desperate to keep you and a boss already considering a replacement will respond very differently. Same sentence, different game.

On-screen takeaway: Your move is only half of the move. The other half is their response.

Lesson 3.2 — Intention vs. incentive (~8 min)

Teaching points

- Ask "What does this make the other person want to do?", not just "What do I want them to do?"
- You may want cooperation, trust, compromise, or for them to stop pushing. The test is whether your action makes that response more attractive than their alternatives.
- You can intend a response without creating the conditions that make it rational.
- People don't respond only to your words. They respond to the game those words are entering.

Example to use: A sales manager says, "I want you to collaborate more," while compensation rewards each person for outperforming colleagues. The manager sincerely wants collaboration, but each salesperson's best response may be to protect their own

leads. They heard the message perfectly; their incentives make another response more attractive.

On-screen takeaway: People respond to the game your words enter, not just to your words.

Lesson 3.3 — Predicting responses well (~10 min)

Teaching points

- Don't ask "What should they do?", "What would a reasonable person do?", or "What would I do?" Ask: "Given their goals, information, options, and constraints, what is the most likely response?"
- The projection trap: "If I were them, I'd never do that" puts your preferences into their position. Use their payoffs (Module 2).
- People aren't perfectly rational. They get emotional, miscalculate, and misread the game. Separate what someone would choose if optimizing from what they're actually likely to choose.
- **Rationalizability** as a shortcut: don't waste energy predicting moves that make no strategic sense. If Choice A is worse for you than Choice B no matter what they do, eliminate A first.

Example to use: A customer refuses your offer. Instead of "They're being unreasonable," ask what would make sense given their objectives and constraints. Maybe they distrust the delivery date, have another supplier, fear looking foolish, need their boss's approval, or are testing your flexibility. You're searching the game instead of judging the person.

On-screen takeaway: Put their preferences into their position, not yours.

Lesson 3.4 — Think backward from the end (~9 min)

Teaching points

- An angry email isn't just an email. If you reply angrily, they may escalate; then you're both defending yourselves instead of solving the problem. You've created a sequence.
- **Backward induction:** ask where the sequence eventually leads, then reason back. "If they reject, what will I do? Knowing that, what will they do? So what should I do now?"
- The Final Conversation: the same disagreement plays differently if it's your last conversation ever versus the start of twenty years working together. The best response to a move depends partly on what comes after it (more in Module 9).
- Threats: "If you do that, I'm walking away" can't be judged by how forcefully it's said. Ask, "If I actually do that, will walking away be their best response?"

Example to use: The chapter's negotiation: you offer, they accept or reject, and after a rejection you get one final response. Decide that last move first, then what they'll do knowing it, and only then choose the opening offer.

On-screen takeaway: Judge a threat by the incentives at the moment it would be carried out.

Lesson 3.5 — Loops, escalation, and the limits of prediction (~9 min)

Teaching points

- Interactions are loops: I act → you respond → I interpret your response → I act again. We respond to what we believe an action means (Module 6).
- **The escalation trap:** a bad outcome can emerge from a sequence of locally understandable responses. Loops can settle into stable but bad patterns (equilibrium, Module 5), so ask, “What pattern will my action create?”
- More thinking isn’t always better. No 47-step tree for “What do you want for dinner?” Ask, “What response could materially change the outcome?”
- Your model is a hypothesis. Use **model → act → observe → update**. When surprised, ask “What did I misunderstand about the game?” The response is data.
- They’re predicting you, too. If you always give in to pressure, people learn that pressure works (Module 4).

Example to use: Two people in a hallway step right, then left, then right again, bouncing side to side. Neither wanted it; each was responding to the other. Scale it up: someone speaks aggressively, you read disrespect and respond in kind, they read hostility and escalate, and soon both of you are acting in ways neither intended.

On-screen takeaway: A bad outcome can emerge from a sequence of locally understandable responses.

Key terms

Term	Plain-language meaning
Best response	The action that gives a player the best outcome given what the other player is doing.
Nash equilibrium	An outcome where each player’s strategy is a best response to the others, so no one benefits by changing alone.
Intention vs. incentive	The gap between the response you want and the response your action actually makes attractive.
Rationalizability (iterated elimination of dominated strategies)	Ruling out moves that make no strategic sense, such as a choice that’s worse for you no matter what the other person does.

Term	Plain-language meaning
Backward induction	Starting from the final decision in a sequence and reasoning back to what to do now.
Credible threat	A threat the person would actually want to carry out when the moment arrives, regardless of how forcefully it was stated.
Feedback loop	A cycle in which each person interprets and responds to the other's last move.
Escalation trap	A worsening pattern built from responses that each seem understandable on their own.
Model → act → observe → update	Treat your prediction as a hypothesis, compare it with what happens, and revise your understanding of the game.

Research spotlight

1. People don't naturally look far ahead, but they can learn to. Eric Johnson, Colin Camerer, Sankar Sen, and Talia Rymon studied a three-round sequential bargaining game. Under full backward induction, the first offer should have been about \$1.25; an equal split would have been \$2.50. Participants actually offered about \$2.11 on average, and the information they chose to inspect before deciding showed limited look-ahead rather than full backward induction. When participants were explicitly trained in backward induction, they learned it quickly, and their offers to computerized opponents moved close to the theoretical prediction (about \$1.22).

What it suggests: "Then what?" isn't automatic, but it appears to be teachable.

What it does NOT prove: That lab bargaining behavior mirrors everyday conversation, or that the "theoretical" offer is the wise move against real people (the trained offers were made to computer opponents). It also doesn't show that thinking further ahead always produces better outcomes.

2. People differ in how many steps ahead they think. In level-k models, a level-1 thinker asks, "What will they probably do?"; a level-2 thinker asks, "What do they think I will do?" Dale Stahl and Paul Wilson's experiments found wide variety; in one, their estimates classified roughly 24% of participants as level-1, 49% as level-2, and 27% as Nash-type. Colin Camerer, Teck-Hua Ho, and Juin-Kuan Chong's cognitive hierarchy model (Quarterly Journal of Economics) fit an average depth of roughly 1.5 steps across the games they studied. Rosemarie Nagel's "beauty contest" guessing game (American Economic Review, 1995) showed that people's choices depend on their model of everyone else's reasoning, and that they rarely jump straight to the equilibrium answer.

What it suggests: The rule isn't "think ten moves ahead." It's "think one or two responses ahead, calibrated to the actual person." Assume five layers of reasoning in someone who just wants lunch, and you can outsmart yourself.

What it does NOT prove: These figures come from specific games and participant pools; they are not universal facts about people or about anyone in your life. And McKelvey and Palfrey's centipede-game experiments add a warning: strict backward induction predicted play would stop immediately, yet participants usually continued for several rounds, which made more sense if they allowed for a chance the other player might be altruistic. Backward reasoning is only as good as the model of the other player feeding it.

The tool: The Response Map

Part 1 is the chapter's practical exercise for an interaction that went sideways. Part 2 turns the chapter's questions into a pre-move check.

Part 1 — Rewind a past interaction

Pick a recent interaction that didn't go the way you wanted. Don't judge who was right yet; just write the sequence.

1. What did I do?

2. What did they do?

3. What did I do next?

4. What did they do next?

5. (Continue if needed)

- At which point did the interaction change direction?

- What could I have predicted about their response before making my move?

Part 2 — Map an upcoming move

1. **My planned move** (what I'm about to say or do):

2. **What I want them to do** (my intention):

3. **Their side of the board.** Given their preferences, not mine:

Their goals:

What they know (and don't):

Their options:

Their constraints (boss, money, time, reputation, emotions):

4. **What does this move make them want to do?** Several plausible responses, each marked:

Response A:

(likely / possible / unlikely)

Response B:

(likely / possible / unlikely)

Response C:

(likely / possible / unlikely)

5. **Intention check:** Does my move actually make the response I want more attractive than their alternatives?

6. **If they do that, what will I do?** (One line per response above.)

7. **Where does the sequence lead?** (High stakes only.)

8. **Eliminate:** Is any option worse for me regardless of what they do?

9. **Threat check:** If a threat is involved (theirs or mine), would carrying it out be the best response when the moment comes?

10. **Pattern and relevance:** What pattern will this move create? Which response could materially change the outcome?

11. **Afterward — update:** What did they actually do? Where was my prediction off, and what did I misunderstand about the game?

Practice exercises

Exercise A — Analyze the scenario

Dana leads a five-person operations team at a logistics company. At her first team meeting she said, “If something goes wrong, tell me right away. No surprises.” Two weeks later, Marcus, a junior coordinator, admitted he had booked a shipment to the wrong warehouse, delaying a client order by two days. Dana was under pressure from her own boss that week. She raised her voice, called it “exactly the kind of careless mistake I can’t afford,” and later cited the error, by name, in the weekly team meeting. Marcus fixed the booking and apologized, and Dana felt she had set a clear standard. A month later, a much larger error surfaced only when the client called to complain; two team members had known about it for a week. Dana told her boss, “I don’t understand it. I specifically told them to come to me.”

1. What did Dana intend her reaction to produce, and what did it actually make attractive for the team?
2. Write the sequence of moves. Where did the interaction change direction?
3. Why is “I specifically told them to come to me” beside the point strategically?
4. Reasoning backward: if Dana wants early disclosure, what must the team expect, and where does that expectation come from?
5. What could Dana do now, and what responses should she anticipate?

Model answer

1. **Intention vs. incentive.** Dana intended to set a high standard, but her response raised the cost of disclosure: anger plus public embarrassment, witnessed by everyone. She didn’t want concealment, but her move made it rational.
2. **The sequence.** Dana: “Tell me right away.” → Marcus discloses. → Dana explodes and names him publicly. → The team learns disclosure is punished. → Next error: two people sit on it for a week. → The client finds out first. The turning point was Dana’s response to the disclosure, the moment her stated rule and her actual behavior pointed in opposite directions.
3. **Words vs. the game they enter.** Her team heard her perfectly; they responded to the game, which her own reaction rewrote. Today’s response became tomorrow’s incentive.
4. **Backward reasoning.** Desired behavior: early reporting. Required expectation: disclosure will be heard and worked through, not punished. Source of that expectation: Dana’s observable responses to disclosure, especially the most recent, most public one. The question isn’t “What phrase will make them open up?” but “What does my response pattern teach them to expect?”

5. **Next move.** Dana could tell the team plainly that her reaction made it harder to bring problems forward, then respond differently when bad news next arrives: thank the person, focus on the fix, handle performance concerns privately. Expect cautious skepticism or a small “test” disclosure. That test is the critical branch; her response to it is what changes the game. If reporting stays slow, she should update and ask what else makes silence attractive. The goal isn’t controlling the team; it’s an honest environment where reporting problems is a reasonable choice for them.

Exercise B — Your own game

Choose a live situation where you’ll make an important move in the next week or two: a request, hard feedback, a boundary, a negotiation, or a reply you’re tempted to send fast.

1. Complete Part 2 of the Response Map. If you catch yourself writing “If I were them, I’d...”, rewrite it from their payoffs.
2. List at least three plausible responses and mark each likely, possible, or unlikely. Replace certainty with branches.
3. For the most likely response, write what you’ll do, then what they’ll likely do next. Stop there unless the stakes are high.
4. Intention check: does your move make the response you want genuinely more attractive to them? If not, what change (timing, setting, framing, what you offer or ask) would?
5. Elimination check: cross off any option that’s worse for you no matter how they respond.
6. Afterward, complete step 11 and write one sentence on what you misunderstood about the game, if anything.

Exercise C — Partner drill / role-play

Format: Pairs (or one learner with an AI practice partner), 12–15 minutes.

Setup: Jordan, a small business owner, wants to cut a freelance designer’s monthly rate. Riley, the designer, wants to keep it.

- **Jordan (private):** You want a 15% cut and have told yourself you’ll “walk away” if Riley refuses. But replacing Riley would take at least six weeks, and your product launch is in four.
- **Riley (private):** Other clients fill about 70% of your capacity. Losing Jordan would hurt but not sink you. You suspect Jordan has a launch coming but don’t know when.

Round 1 (4 min): Jordan opens: “I need a 15% cut, or I’ll have to find someone else.” Play it out.

Round 2 (4 min): First, 60 seconds of silent prep. Jordan: “If Riley says no, will walking away really be my best response? What does my opening make Riley want to do?” Riley: “If I refuse, what will Jordan actually do at that moment?” Replay with new openings.

Debrief (5 min):

- Was Jordan's Round 1 threat credible? How did Riley decide?
- Where did either player project their own preferences onto the other?
- Did Jordan's Round 2 opening make the response Jordan wanted more attractive to Riley, or just state it more firmly?
- Ethics check: did anyone make a threat they didn't believe? What might that cost in an ongoing relationship?

Common mistakes

- **Judging a move by your intention** — "I was only being honest" may be true, but strategy asks what your move makes the other person want to do.
- **Projecting your preferences** — "If I were them, I'd never..." uses your payoffs. Use their goals, information, options, and constraints.
- **Assuming perfect rationality** — People get emotional and misread the game. Separate what's optimal for them from what's likely.
- **Judging threats by forcefulness** — Check whether carrying out the threat would be their best response when the moment arrives.
- **Overthinking** — Skip the 47-step tree. Focus on responses that could materially change the outcome.
- **Overpredicting** — Your model is a hypothesis. When surprised, ask "What did I misunderstand about the game?", not "Why are they irrational?"
- **Forgetting they're predicting you** — A predictable response, like always giving in to pressure, becomes part of their strategy.
- **Mistaking influence for control** — You can shape the environment; other people still choose their own responses.

Quiz

1. According to the chapter, what is a best response?
 - A) The most polite and respectful reply
 - B) The action that gives a player the best outcome given what the other player is doing
 - C) The first reply that comes to mind under pressure
 - D) The action that would be best if the other player didn't exist
2. A manager urges collaboration but pays bonuses only for individual results. Salespeople keep guarding their leads. The best diagnosis?
 - A) The team didn't hear the message
 - B) The team is irrational
 - C) Their incentives make another response more attractive than collaboration

- D) The manager needs to say it more forcefully
3. Which question does the chapter recommend for predicting a response?
- A) “What would a reasonable person do?”
 - B) “What would I do in their position?”
 - C) “What should they do?”
 - D) “Given their goals, information, options, and constraints, what is the most likely response?”
4. Option A leaves you worse off than Option B no matter what the other person does. What does the chapter suggest?
- A) Eliminate Option A without needing to predict them
 - B) Predict their move carefully before deciding
 - C) Choose A to keep them guessing
 - D) Ask them which option they prefer
5. Someone says, “If you do that, I’m walking away.” How should you assess the threat?
- A) By how forcefully it was said
 - B) By assuming all threats are bluffs
 - C) By whether walking away would actually be their best response when that moment arrives
 - D) By making a bigger threat in return
6. Which statement best describes the escalation trap?
- A) One person deliberately decides to make things worse
 - B) It only happens between people who dislike each other
 - C) The person who started it can always stop it
 - D) A bad outcome emerges from a sequence of locally understandable responses
7. In Johnson, Camerer, Sen, and Rymon’s bargaining study, what happened after participants were trained in backward induction?
- A) Their offers to computerized opponents moved close to the theoretical prediction
 - B) Nothing changed; they kept offering about an equal split
 - C) They refused to make offers
 - D) The study proved backward induction predicts everyday conversation
8. You predicted a colleague would welcome your proposal, but they pushed back hard. What’s the best next step?
- A) Conclude they’re being irrational
 - B) Ask, “What did I misunderstand about the game?” and update your model
 - C) Repeat the proposal until they agree
 - D) Stop making predictions about people

Answer key

1. **B** — A best response is defined relative to what the other player is doing; your best move isn't determined in isolation.
2. **C** — They may have heard perfectly. People respond to the game your words enter, and here the payoffs reward protecting leads.
3. **D** — The other options substitute someone else's preferences (yours, or an imagined "reasonable person's") for theirs.
4. **A** — The practical translation of rationalizability: if an option is worse no matter what they do, drop it before predicting anything.
5. **C** — A threat matters only if carrying it out makes sense at the point it would be carried out, not because it was stated forcefully.
6. **D** — Each person can respond understandably to the last move while the interaction gets worse. Nobody has to intend it.
7. **A** — Offers moved to about \$1.22 against computer opponents, suggesting look-ahead is teachable, not that it describes everyday behavior.
8. **B** — Model → act → observe → update. The response is data; their payoffs, information, or options may differ from what you assumed.

Field assignment

This week, run "Then what?" before three moves that matter (a request, a reply to a frustrating message, feedback, a "no"). Before each, take 60 seconds to write:

- My move:
- What I want them to do:
- Their two or three most plausible responses (with likely / possible / unlikely):
- If they do the most likely thing, what will I do?

Afterward, record what they actually did and whether your prediction was right, partly right, or wrong. Once this week, also notice a feedback loop in progress and note whether it was escalating or settling.

Debrief questions for next week: 1. How often did your prediction match the response? When it didn't, what had you misunderstood about the game? 2. Did asking "Then what?" change any move before you made it? What did you change, and what happened? 3. In the feedback loop you observed, where could one well-chosen response have changed the direction of the sequence?

The question to carry

"What is this likely to make them do? If they do that, what will I do? And if the stakes are high, where does that sequence lead?"

Module 4 — Manage Predictability

The question: What does my predictability enable the other player to do?

Book reading	Chapter 4, “Manage Predictability”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Predictability Audit

Why this module matters

In Module 3 you learned to think in responses: before you make a move, ask what it is likely to make the other person do. This module turns that question around. What happens when the other person can predict you? If someone knows exactly how you’re going to respond, they can begin planning around you. Your behavior becomes part of their strategy.

Think of the coworker who knows you always say yes to favors, or the counterpart who knows you always concede after their first counteroffer. You never announced a rule, but they learned one: Ask, and he’ll say yes. Push, and he’ll move. You’ve taught them the game.

But this is not a module about becoming erratic, deceptive, or hard to read. Predictability is also what makes trust, coordination, and credible consequences possible. The skill is knowing when each is strategically valuable. The shift you’ll practice is adding one question: **“What are they learning about me from what I repeatedly do?”**

Learning objectives

By the end of this module you will be able to:

1. Identify the patterns in your own behavior that other people have learned to predict, including rules you never intended to create.
2. Distinguish situations where predictability builds trust and coordination from situations where it creates exploitation and strategic vulnerability.
3. Explain the mixed-strategy idea in plain language, and why variety is not the same as unpredictability.
4. Distinguish withholding information you are not obligated to disclose from deliberately providing false information.

5. Choose deliberately among four moves — be predictable, protect your predictability, vary your behavior, or be deliberately consistent — based on what your predictability enables the other player to do.

Lessons

Lesson 4.1 — The Habit They Learn (8 min)

Teaching points

- If someone can reliably predict your action, they can choose a response that takes advantage of it. They no longer wonder “Should I push?” — your pattern tells them.
- Your patterns can create rules even when you never intended to create them. Letting something go once is a choice; letting it go every time is a policy.
- Frame the ethics up front: this chapter is not about lying, manipulating, or becoming erratic so nobody knows what you’ll do.

Example to use: Someone repeatedly interrupts you in meetings. You let it go three times. Now they’ve learned something: Interrupting you has no cost. You never consciously decided to establish that rule, but your repeated behavior established one anyway.

On-screen takeaway: Your patterns can create rules even when you never intended to create them.

Lesson 4.2 — Rock-Paper-Scissors and the Value of Uncertainty (9 min)

Teaching points

- In formal game theory, a **mixed strategy** is one where a player chooses among actions according to probabilities rather than choosing the same action every time.
- If you always throw rock, your opponent learns to throw paper. They don’t even need your exact next move; a tendency is enough to adjust.
- A good mixed strategy is not “do something random.” It is often a carefully chosen probability distribution that prevents the opponent from exploiting your pattern. You don’t need to calculate probabilities to grasp the idea: sometimes uncertainty itself has strategic value.
- Caution: the mental model is not “Randomize everything.” Randomization helps only where another player gains by correctly anticipating you.

Example to use: Play it out live. You and a partner are playing rock-paper-scissors, and you favor rock 70 percent of the time. Your partner only needs to notice the tilt and lean toward paper. That is what strategic predictability means.

On-screen takeaway: If they know your pattern, your pattern becomes a weakness.

Lesson 4.3 — The Predictability Tax: When Your Pattern Becomes Information (10 min)

Teaching points

- Your actions can reveal information you didn't intend to reveal. That hidden cost is the **predictability tax**.
- Every action communicates. Always answering emails immediately signals availability; always threatening to quit but never doing it teaches people your threats aren't credible.
- A bluff that once worked stops working when others learn you bluff a lot. The words didn't change; the other players' beliefs did.
- Giving away your bottom line changes the information structure: you've told them how far they need to push. You can decline without lying: "I'm not prepared to disclose my bottom line."

Example to use: A negotiator always makes three concessions: \$100, then \$90, then \$80. The other side learns the pattern. They don't need to know the true minimum; they just need to know you'll move again. So they push, you move, they push again. By the third round, your pattern has effectively told them: There's more room.

On-screen takeaway: Your behavior creates information — sometimes more than your words do.

Lesson 4.4 — Predictability Is Not the Enemy (9 min)

Teaching points

- Coordination requires predictability. At a stoplight, you want everyone to know red means stop. Conventions and norms work because everyone can predict everyone else.
- Predictability can create trust and coordination, or exploitation and strategic vulnerability. The challenge is knowing which environment you're in.
- Don't become unpredictable in the wrong way: tolerating an interruption one day and exploding the next isn't sophistication. It makes you difficult to trust and destroys reliable expectations.
- When rebuilding trust, become more predictable. If you say you'll call, call. Over time they learn: his words predict his actions.
- A consistently enforced consequence ("falsify customer records and you will be terminated") changes behavior before a violation occurs.

Example to use: Two managers. One says, "I really don't like firing people," and employees learn he never actually fires anyone — so some begin taking risks they otherwise wouldn't. The other is known for this: if you violate this rule, there will be consequences. That predictability improves cooperation. Same trait, opposite strategic effect.

On-screen takeaway: Predictability isn't inherently good or bad. Its value depends on what it allows other players to do.

Lesson 4.5 — Managing Predictability: Secrecy, Surprise, and the Feedback Loop (9 min)

Teaching points

- Distinguish **manipulative unpredictability** from **strategic unpredictability**. Varying launch timing so a competitor can't optimize against a fixed pattern is competing, not deceiving.
- Secrecy is not deception: withholding information you aren't obligated to disclose differs from deliberately providing false information to induce a decision.
- Don't romanticize surprise. Ask: Who benefits from knowing what I'm going to do? If "both of us," predictability helps. If "they can use it to take advantage of me," some uncertainty may be valuable.
- Add learning to the response chain: action → response → learning → adaptation. Your behavior today can change the game tomorrow.
- The Prisoner's Dilemma paradox: in a repeated game, if you know I'll always retaliate against defection, that predictability can support cooperation. The same predictability that makes you exploitable in one game can make you trustworthy in another.
- The four moves: **be predictable** (when trust, coordination, credibility, or cooperation are valuable), **protect your predictability** (when revealing your pattern gives another player leverage), **vary your behavior** (when someone could exploit a fixed pattern), **be deliberately consistent** (when you want someone to understand a consequence is real).

Example to use: A competitor who always responds to a price cut with another price cut can be trapped into a predictable escalation. Contrast that with a negotiator asked, "What's the lowest price you'll accept?" who answers, "I'm not prepared to disclose my bottom line." One is trapped by a fixed rule; the other protects information without saying anything false.

On-screen takeaway: Don't confuse unpredictability with strategic intelligence. The skill is managing predictability.

Key terms

Term	Plain-language meaning
Strategic predictability	When someone knows enough about your pattern to adjust and gain an advantage, even without knowing your exact next move.

Term	Plain-language meaning
Mixed strategy	Choosing among actions by carefully chosen probabilities so an opponent can't exploit your pattern. Not "doing something random."
Predictability tax	The hidden cost of predictable behavior: your actions reveal information you didn't intend to reveal, such as your limits or how far someone needs to push.
Unintended rule	A rule other people learn from your repeated behavior even though you never consciously chose it ("Interrupting her has no cost").
Strategic unpredictability	Varying your behavior so another player can't simply optimize against a fixed pattern — legitimate in competition and compatible with honesty.
Manipulative unpredictability	Unpredictability used to deceive or to destabilize another person, rather than to prevent legitimate exploitation.
Reputation	Accumulated evidence about your future behavior. Every repeated pattern teaches the other player something about you.
Feedback loop	Action → response → learning → adaptation. Today's behavior changes tomorrow's game.

Research spotlight

1. Penalty kicks: when being predictable gets punished. Chiappori, Levitt, and Groseclose (2002) analyzed 459 penalty kicks from the top French and Italian professional soccer leagues. Penalty kicks are a clean strategic setting: opposing goals, known tendencies, high stakes. The researchers found that the testable predictions of mixed-strategy theory fit the data surprisingly well. A later analysis using a different statistical method reached broadly consistent results.

What it suggests for real conversations: When another person's gain depends on correctly anticipating you, your predictability changes the payoff of your action. A kick to the left is worth less if the goalkeeper expects it — and a concession is worth less if the other side knows it's coming.

What it does NOT prove: Professional soccer is close to a pure zero-sum contest. It does not show that every conversation works like a penalty kick, that "unpredictable people win," or that you should randomize your communication. Nor does it show that people

perfectly randomize; related research on baseball pitches (Bhattacharya and Howard) found professionals only approximate the mixed-strategy benchmark, more closely as stakes rise.

2. Trust games: when being predictable creates value. Bracht and Feltovich ran a trust-game experiment in which players could see one previous action of the person they were interacting with. Simply making past behavior observable produced substantially higher cooperation and efficiency than the basic trust game, while cheap-talk messages alone had little effect. Past cooperative behavior predicted current cooperative behavior and influenced whether others were willing to trust.

What it suggests for real conversations: When cooperation depends on trust, being readable is an asset, and what you've done outweighs what you say — which is why the book says to rebuild trust by becoming more predictable.

What it does NOT prove: It is a simple laboratory game. It does not show that more transparency is always better or that you should disclose everything. Observable track records support cooperation where trust matters — a context-dependent result.

Taken together: behavior others can anticipate is a liability in the first setting and an asset in the second — which is why the chapter asks what your predictability enables, rather than offering a universal rule.

The tool: The Predictability Audit

Use this on one recurring interaction where a pattern may be shaping the game.

Step 1 — Name the interaction

The person or group:

The recurring situation (meeting, negotiation, household routine, request):

Step 2 — Find the pattern

What do I repeatedly do in this situation?

What have my past actions taught this person to expect from me?

Have I taught this person how to respond to me? If so, what is the rule they've learned (finish the sentence "If I , **she/he will** ")?

Did I choose this rule, or did my repeated behavior establish it?

Step 3 — Calculate the predictability tax

What does this pattern reveal that I didn't intend to reveal (my limits, availability, whether my threats or boundaries are real)?

Step 4 — Asset or liability?

Who benefits from knowing what I'm going to do? (Circle one: both of us / mainly them / mainly me)

Is this pattern helping me? Helping them? Helping both of us?

Could they exploit this pattern? How?

Check any that apply. This pattern currently: - ☐ creates trust, cooperation, or a credible commitment - ☐ helps others coordinate with me - ☐ shows people exactly how to pressure me - ☐ reveals threats, concessions, or unenforced boundaries

Step 5 — Choose your move

Would greater consistency create trust — or greater flexibility create strategic room?

Which of the four moves fits? (Circle one) - **Be predictable** — trust, coordination, credibility, or cooperation are valuable here. - **Protect my predictability** — revealing my pattern gives them leverage. - **Vary my behavior** — they could otherwise exploit a fixed pattern. - **Be deliberately consistent** — I want them to understand a consequence is real.

The specific behavior I'll change or reinforce:

Step 6 — Ethics check

Am I protecting legitimate information or providing false information?

Could my change destabilize someone or make me harder to trust where reliability matters?

Practice exercises

Exercise A — Analyze the scenario

Dana is a senior designer at a marketing agency. For two years, whenever account manager Rob brings her an “urgent” client revision on Friday afternoon, she stays late and finishes it. Lately the Friday requests have multiplied, and several had been sitting in his inbox since Tuesday. Dana has complained in meetings that “people need to plan better” and once told Rob she “wouldn’t keep doing this,” then stayed late anyway. Dana also reviews junior designers’ work every Thursday morning, and the juniors have built their weeks around it. Frustrated, she is considering becoming “harder to read” across the board so nobody takes her for granted.

1. What rule has Rob learned from Dana’s behavior, and did she ever choose to create it?
2. What is Dana’s predictability tax here? What information has her pattern revealed?
3. Which of Dana’s patterns is an asset and which is a liability? Why?
4. Is “becoming harder to read across the board” a good plan? What would you recommend instead, using the four moves?
5. How can Dana change the pattern without deception?

Model answer

1. **The learned rule:** “If I bring Dana an urgent revision on Friday, she’ll stay late and do it.” Dana never chose this rule; her repeated behavior established it. Rob has adapted — the feedback loop of action, response, learning, adaptation has made Friday requests more frequent.
2. **The predictability tax:** Her pattern revealed that her evenings are flexible and her complaints carry no consequence. “I won’t keep doing this,” followed by doing it anyway, taught Rob her warnings aren’t credible. Her behavior communicated more than her words.
3. **Asset vs. liability:** The Thursday review is an asset — the stoplight kind of predictability that helps the juniors coordinate. Both sides benefit from knowing what she’ll do. The Friday pattern is a liability: it shows exactly how to pressure her and reveals a boundary that’s stated but unenforced.
4. **Not across the board.** That confuses unpredictability with strategic intelligence and would destroy expectations her team depends on. Instead: **be predictable** about the Thursday review, and **be deliberately consistent** about a clear Friday boundary — for example, “Revisions that arrive after 2 p.m. Friday will be done Monday morning unless the client deadline truly requires otherwise.” Then enforce it every time. Staying late only sometimes would be the wrong kind of unpredictability; it would keep Rob testing her.
5. **Without deception:** No fake evening plans needed. She states the rule honestly and lets consistent follow-through teach the new pattern. She needn’t explain her whole schedule

— withholding what she doesn't owe is not lying. Over time Rob learns: "Her words predict her actions."

Exercise B — Your own game

Pick one live situation where someone has likely learned to predict you — a coworker, negotiation partner, client, or family member. Then:

1. Complete Steps 1-3 of the Predictability Audit. Write the rule they've learned in one sentence, starting with "If I , **they will** ."
2. In Step 4, list one way this pattern helps both of you and one way it could be exploited. No exploitation risk? This may be a pattern to protect.
3. In Step 5, choose one of the four moves. Write the exact words or behavior you'll use the next three times this comes up.
4. Run the Step 6 ethics check. If your plan involves saying anything false, rewrite it so you are withholding rather than misrepresenting.

Exercise C — Partner drill / role-play

Format: Pairs, 15 minutes (or with an AI practice partner). Two rounds.

Round 1 — The concession pattern (6 min)

Setup: A freelance photographer (Seller) is quoting a corporate event. A procurement manager (Buyer) wants the lowest price.

Seller's private card: Your opening quote is \$4,000. Your true minimum is \$3,000. Follow a fixed rule: every time the Buyer rejects your number, drop \$300. If asked your minimum, tell the truth.

Buyer's private card: You have budget up to \$4,000 but want to pay as little as possible. Figure out the Seller's pattern and use it. Ask once, "What's the lowest you'll accept?"

Play for four minutes. Then the Buyer states the rule they think the Seller was following.

Round 2 — Managing predictability (6 min)

Same roles, same private numbers, but the Seller now has no fixed rule. The Seller must (a) decline to disclose the bottom line honestly, (b) tie any movement to a reason or a change in scope rather than to rejection alone, and (c) say nothing false. The Buyer again tries to find a pattern.

Debrief (3 min)

- In Round 1, how many rejections did it take the Buyer to learn the pattern? What did the Seller's behavior reveal that their words didn't?
- In Round 2, how did "I'm not prepared to disclose my bottom line" change the Buyer's decision-making?

- Did the Seller ever say anything untrue? Where was the line between protecting information and deceiving?

Common mistakes

- **Treating “keep them guessing” as the lesson.** Game theory doesn’t say that. Where trust matters, excessive unpredictability is destructive. Ask what your predictability enables the other player to do first.
- **Becoming unpredictable in the wrong way.** Tolerating something one day and exploding the next doesn’t make you strategic; it makes you difficult to trust. Where you need a consequence to be real, be deliberately consistent.
- **Confusing variety with unpredictability.** A rigid sequence of different moves (rock, then paper, then scissors) is still a pattern. If the other player can learn it, it is still exploitable.
- **Stating a boundary but not enforcing it.** A boundary that’s predictable but unenforced teaches the opposite of what you said. Your behavior creates more information than your words.
- **Answering every question about your limits.** Giving away your bottom line tells the other side how far they need to push. You can decline honestly; protecting strategically valuable information isn’t deception.
- **Crossing from secrecy into deception.** Withholding what you’re not obligated to disclose is legitimate. Deliberately providing false information to induce someone into a decision is not.
- **Romanticizing surprise.** Surprise is useful only when the uncertainty it creates is valuable. A spouse or customer gains nothing from wondering whether you’ll keep your word.

Quiz

1. According to the chapter, what is the core principle of managing predictability?
 - A) Never let people know what you’re going to do.
 - B) Always be consistent so people trust you.
 - C) Don’t allow your patterns to become so predictable that other people can consistently use them against you.
 - D) Randomize your behavior in every competitive situation.
2. In rock-paper-scissors, you know your opponent chooses rock 70 percent of the time. What does this illustrate?
 - A) You need to know their exact next move to gain an advantage.
 - B) Knowing someone’s pattern is enough to adjust your behavior and gain an advantage.

- C) Mixed strategies only work in games of pure chance.
 - D) Predictability matters only when the other player is irrational.
3. A negotiator always concedes from \$100 to \$90 to \$80. What is the “predictability tax” here?
- A) The cost of making three offers instead of one.
 - B) The pattern reveals there’s more room, so the other side keeps pushing.
 - C) The negotiator must pay a fee for each round.
 - D) The other side will refuse to negotiate.
4. Someone interrupts you in three meetings and you let it go each time. What has most likely happened?
- A) Nothing; one-off choices don’t create patterns.
 - B) You’ve established a rule — interrupting you has no cost — even though you never intended to.
 - C) You’ve used a mixed strategy effectively.
 - D) You’ve built trust that will reduce future interruptions.
5. Which situation is the best example of predictability creating value?
- A) A competitor always matches your price cuts.
 - B) A parent who always gives in after a third request.
 - C) A manager who says rule violations will have consequences and consistently enforces it.
 - D) A poker player who bluffs frequently.
6. A buyer asks, “What’s the lowest price you’ll accept?” Which response best fits the chapter?
- A) Give a number well above your real minimum and call it your floor.
 - B) Give your true minimum to show good faith.
 - C) “I’m not prepared to disclose my bottom line.”
 - D) Change the subject without acknowledging the question.
7. You’re trying to rebuild trust with a friend after missing several commitments. What does the chapter recommend?
- A) Become less predictable so they stop taking you for granted.
 - B) Become more predictable: if you say you’ll call, call.
 - C) Make fewer statements so there’s less to be held to.
 - D) Vary your follow-through so they can’t exploit you.
8. In a repeated Prisoner’s Dilemma, you’re known to always retaliate against defection. According to the chapter, this predictability:
- A) Always makes you exploitable.
 - B) Is irrelevant because the game is repeated.

- C) Can support cooperation — showing that the same predictability that's exploitable in one game can make you trustworthy in another.
- D) Should be replaced with random retaliation.

Answer key

1. **C** — Sometimes that means being unpredictable, sometimes extremely predictable. A and D overcorrect; B ignores exploitability.
2. **B** — The opponent doesn't need your exact move; knowing enough about your pattern to adjust creates the advantage. That's what strategic predictability means.
3. **B** — The pattern reveals information you didn't intend to share. The other side doesn't need your true minimum, only the knowledge that you'll move again.
4. **B** — Your patterns can create rules even when you never intended to create them. Once people learn the rule, they can act accordingly.
5. **C** — Known, enforced consequences let people anticipate outcomes and change behavior before a violation. The others are exploitable patterns.
6. **C** — Declining protects valuable information without deception. A is false; B hands over your boundary.
7. **B** — When rebuilding trust, become more predictable so the other person learns that your words predict your actions.
8. **C** — Predictable retaliation can support cooperation in repeated play, which is why you must understand the structure of the game first.

Field assignment

This week: run a pattern log. Choose one recurring interaction — a weekly meeting, a regular request from a colleague, a household routine, a client relationship. Each time it happens, write down three things: what you did, how the other person responded, and what they might have learned about you from that exchange.

By midweek, complete the Predictability Audit, choose one of the four moves, and carry it out on purpose for the rest of the week. Stay honest throughout — no false statements.

Record: - The rule you think the other person had learned before the week started. - Any sign that they noticed a change (testing, surprise, adjustment, relief). - Whether the change made the interaction better for you, for them, or for both of you.

Debrief questions for next week: 1. What had your past actions taught this person to expect from you — and how accurate was your guess once you watched closely? 2. Did they adapt to your new pattern? What does that tell you about the feedback loop of action, response, learning, and adaptation? 3. Where in your life would greater consistency create trust, and where would greater flexibility create strategic room?

The question to carry

Have I taught this person how to respond to me? What have my past actions taught this person to expect from me?

Module 5 — See the Equilibrium

The question: What equilibrium have we created?

Book reading	Chapter 5, “See the Equilibrium”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Equilibrium Audit

Why this module matters

Imagine two people who have been arguing for months. Neither likes it, yet neither changes. You stop calling because “they never call me.” They don’t call because “I always have to be the one who reaches out.” Nobody sat down and decided, “Let’s create a miserable relationship.” Yet here you are, in a stable pattern.

You’ve learned to see the game, see the payoffs, and think in responses. This module adds **equilibrium**. Sometimes people aren’t doing what they want; they’re doing what makes sense given what everyone else is doing. And the most important thing to understand: **stable doesn’t mean good**. A bad situation can persist without anyone needing to be irrational.

The mental shift is from moves to patterns. Instead of asking, “Why did they say that?” you start asking, “Why do we keep ending up here?” Instead of asking, “Who’s to blame?” you ask, “What would have to change for the pattern to move?” Sometimes the problem isn’t that you haven’t found the right words. Sometimes you’ve both become very good at playing the same bad game.

Learning objectives

By the end of this module you will be able to:

1. Define a Nash equilibrium in plain language and explain why “stable” is not the same as “good,” “fair,” or “optimal.”
2. Walk through the Prisoner’s Dilemma and show how individually rational responses can produce an outcome both players dislike.
3. Map a recurring interaction as a loop of mutual responses rather than a series of isolated moves.
4. Explain why “just change” can backfire, and identify what would have to change for a different pattern to become attractive.

5. Distinguish between using equilibrium as a mental model and using it as a prediction, including the problem of multiple equilibria.

Lessons

Lesson 5.1 — Stable Doesn't Mean Good (10 min)

Teaching points

- A Nash equilibrium is a combination of strategies in which each player's strategy is a best response to the others'. Simply: given what everyone else is doing, no one has a reason to change on their own.
- Equilibrium does not mean everyone is happy, that everyone got the best possible outcome, or even that anyone thinks the situation is fair. It means the existing pattern is stable because changing unilaterally doesn't look attractive.
- Everyday language hears "equilibrium" as balance or harmony. In game theory, an equilibrium can be terrible for everyone.
- The Prisoner's Dilemma shows it: if the other person cooperates, you'd do better defecting; if they defect, you'd better defect too. Both reason the same way, both defect, and both end up worse off than mutual cooperation.

Example to use: Walk the Prisoner's Dilemma in three questions. "Which outcome would both players prefer?" (Both cooperate.) "What should each player do if the other's choice is fixed?" (Defect.) "So where do they end up?" (Both defect.) Land the realization: everyone doing what makes sense locally can produce an outcome everyone dislikes globally.

On-screen takeaway: A bad situation can persist without anyone needing to be irrational.

Lesson 5.2 — Look for the Pattern, Not Just the Move (10 min)

Teaching points

- Chapter 3 asked, "What will they do after I move?" This chapter asks something deeper: "What pattern are our responses creating?"
- One interaction looks like "I push, you resist." Repeated, it becomes "I push, you resist, I push harder, you resist harder." Now you're looking at a system, and systems develop stable patterns.
- Equilibrium traps are everywhere: rival companies cutting prices because each fears being undercut; coworkers working unsustainable hours because each fears being taken advantage of.
- In relationships, complaint responds to withdrawal and withdrawal responds to complaint. The loop is self-reinforcing, which doesn't make it healthy.
- Most communication advice asks, "What should I say differently?" Game theory adds: "What equilibrium have we created?"

Example to use: The “Can you help me with this?” loop from the chapter. You ask; they’re too busy; “You’re always too busy”; “Because you always ask at the last minute”; “Because you never tell me when you’re available”; “Because you don’t plan ahead.” Skip who’s right and name the pattern: you ask late, they refuse, you criticize, they defend, you feel unsupported, you ask late again. The interaction itself is producing the outcome.

On-screen takeaway: Don’t ask “Why did they say that?” Ask “Why do we keep ending up here?”

Lesson 5.3 — Unwritten, Self-Reinforcing, and Sometimes Invisible (8 min)

Teaching points

- An equilibrium doesn’t require an agreement. Many of the strongest social patterns are unwritten: who speaks first, who pays, who apologizes, who gives in, who gets the last word.
- Once a pattern develops, expectations adjust to it. People respond not only to incentives but to expectations about what others will do.
- Some equilibria are invisible: everyone hates the two-hour meeting, but nobody wants to be the one who says, “I’m leaving after 30 minutes.”
- Repetition turns a pattern into a habit. After enough rounds, “We need to talk” triggers “Here comes another criticism” before anything is said. Nobody chooses the pattern anymore; they’ve learned it. The game starts playing you.

Example to use: The new employee in a company where everyone stays late. At first: “Why is everyone still here?” After a few weeks: “I guess this is normal.” After a few months: “If I leave at five, people will think I’m not committed.” Now their own behavior reinforces the very norm they first questioned. The equilibrium has become cultural.

On-screen takeaway: The game starts playing you.

Lesson 5.4 — Why “Just Change” Can Be Bad Advice (9 min)

Teaching points

- Equilibrium is about unilateral change. If I change while you don’t, I’m worse off, and vice versa, so neither of us moves alone, whether the pattern is beneficial, harmful, unfair, or merely acceptable.
- “Just stop participating in the argument” sounds reasonable. But if the partner reads silence as abandonment, they may escalate. A unilateral change can initially make things worse.
- A company that raises prices alone during a price war may simply lose customers. Good intentions alone rarely break an equilibrium; sometimes the players need to change their expectations about what the others will do.

- So you have to change something that affects the strategic structure: your response, the sequence, the incentives, the information, a commitment, a third party, expectations, or the available choices.
- Equilibrium is the bridge between the individual models: it explains why the pieces you've learned lock together into persistent patterns.

Example to use: The company where everyone waits for senior management before deciding. Management asks, “Why doesn’t anyone take initiative?” Employees think, “Because whenever we do, management overrides us.” Each side’s behavior reinforces the other. “Take more initiative” doesn’t touch what holds the pattern in place.

On-screen takeaway: Ask “What would have to change for the pattern to move?” instead of “Who’s to blame?”

Lesson 5.5 — Many Equilibria, and Why It’s a Lens, Not a Crystal Ball (8 min)

Teaching points

- Some games have several equilibria. If two people each do best by matching the other, both “A/A” and “B/B” can be stable. The problem is getting to the same one: a coordination problem.
- Knowing the equilibria doesn’t tell you which one people will land on (the equilibrium-selection problem). Salient cues help, but research after Schelling found even small payoff asymmetries can disrupt coordination around an otherwise obvious choice.
- Equilibrium isn’t prediction. People make mistakes, have limited information, use rules of thumb, and learn over time. They can sometimes learn their way into an equilibrium (model, act, observe, update), but convergence isn’t universal.
- A credibility rule: say “Nash equilibrium” for formal games you’ve actually specified; for everyday patterns, say “equilibrium lens.” Not everything that persists is an equilibrium. Habit, coercion, inertia, or rules can also keep things going. Test it: would changing alone actually make the player worse off?

Example to use: Schelling’s meeting-in-New-York thought experiment. You’ve agreed to meet someone tomorrow but set no place or time. People frequently converged on noon at Grand Central Station, though nothing made it mathematically unique. It worked as a focal point: people can coordinate through context and shared expectations without talking.

On-screen takeaway: Equilibrium is a mental model, not a crystal ball.

Key terms

Term	Plain-language meaning
Nash equilibrium	A set of strategies where each player’s choice is a best response to the others’, so no one has a reason to change on their own.

Term	Plain-language meaning
Unilateral change	One player changing while everyone else keeps doing what they were doing. Equilibrium is defined by the lack of reward for this.
Prisoner's Dilemma	The classic game where each player's best response is to defect, so both defect, even though both would prefer mutual cooperation.
Equilibrium trap	A stable pattern that the people inside it dislike but that no one sees a good way to leave alone.
Self-reinforcing pattern	A loop in which each person's response provokes the response that sustains it, such as complaining and withdrawing.
Coordination problem	The challenge of getting everyone to the same equilibrium when several are possible.
Focal point	Something that stands out and helps people coordinate their expectations without explicit communication, such as noon at Grand Central.
Equilibrium-selection problem	Knowing the possible stable outcomes without knowing which one people will land on.
Equilibrium lens	Using equilibrium thinking to ask why an everyday pattern persists, without claiming it is a formally proven Nash equilibrium.

Research spotlight

1. Cooperation in one-shot Prisoner's Dilemmas (Charness, Rigotti, and Rustichini).

- What was done: The researchers ran clean, one-shot Prisoner's Dilemma experiments with 348 participants and varied how much players earned from mutual cooperation.
- What was found: Cooperation ranged from 23% to 60% depending on the mutual-cooperation payoff. Participants' beliefs about whether others would cooperate also shifted as the payoff structure changed.
- What it suggests for real conversations: Structure matters: changing the payoffs changed behavior, so you shift a pattern by changing what makes it stable. But the simple money-only prediction ("everyone defects") didn't hold. Beliefs, norms, social preferences, and errors matter too.
- What it does NOT prove: It doesn't show that people always cooperate, that any particular relationship is a Prisoner's Dilemma, or that Nash equilibrium is either a

reliable predictor or useless. It is a useful structure to think with and an incomplete description of real people.

2. Cooperation in repeated Prisoner's Dilemmas (Dal Bó and Fréchette, *American Economic Review*; and their later survey in the *Journal of Economic Literature*).

- What was done: Participants played indefinitely repeated Prisoner's Dilemma games with no known end. Theory says mutual cooperation can become sustainable in repeated play if players value the future enough.
- What was found: Cooperation can emerge, but under conditions more demanding than the bare theoretical requirement. The survey concludes that being supportable in equilibrium is not enough; cooperation must also be robust to strategic uncertainty (the risk the other person won't go along).
- What it suggests for real conversations: Knowing a better equilibrium exists doesn't mean people will find it. Two partners who both want openness may stay guarded because whoever opens up first absorbs the risk. The practical question: how do we move to a better pattern without one person carrying all the risk?
- What it does NOT prove: These are lab games with defined payoffs, not marriages or teams. They don't tell you when your own relationship will shift, and they don't say cooperation is always right. Some competitive patterns exist for good reasons.

A note on honesty: that equilibria are mutual best responses, and can be bad for everyone, is foundational. Applying the idea to everyday relationships is the book's own synthesis, a lens rather than a proof. A pattern can "resemble an equilibrium trap"; your marriage isn't "literally in a Nash equilibrium" unless you've modeled it.

The tool: The Equilibrium Audit

Use this when you're stuck in a recurring interaction. Work through it in writing, at a calm moment.

The situation (one line):

1. What keeps happening? Forget the individual incident. Describe the repeated pattern as a loop (my move, their response, my response...).

2. What does my behavior cause them to do?

3. What does their behavior cause me to do?

4. Why might neither of us want to change unilaterally? If I changed alone, what would I risk? If they changed alone, what would they risk? Me:

Them:

5. Is this pattern actually good for either of us? What does each of us get from maintaining it? What does it cost each of us?

6. Are there alternative stable patterns? Describe at least one other pattern that could hold steady if we both settled into it.

Equilibrium check: Is this really equilibrium-like, or is it held in place by something else (habit, a rule, coercion, lack of awareness)? Would changing alone actually make someone worse off?

The most important question: What would have to change for a different pattern to become attractive? Check any levers that apply, then describe one: - ☐ My response - ☐ The sequence (who moves first, when) - ☐ The incentives - ☐ New information - ☐ A commitment - ☐ A third party - ☐ Expectations about what the other will do - ☐ The available choices

My first experiment:

Practice exercises

Exercise A — Analyze the scenario

Priya and Marcus co-lead a product team. Every Friday, Priya sends Marcus a long status email flagging risks she sees in his part of the project. Marcus has started replying with a short “Noted, thanks” and handling the issues on his own without updating her. Because she hears nothing, Priya assumes the risks aren’t being handled, so the next Friday’s email is longer and copies their director. Marcus reads the cc as Priya trying to make him look bad, so he shares even less and starts sending his own updates straight to the director. Both say they want a real partnership, and both say the other started it. Last month Marcus tried skipping his separate update to the director “to be the bigger person,” and the director asked him in a meeting why he’d gone quiet, which embarrassed him.

1. Describe the pattern as a loop of mutual responses.
2. Why does each person’s current behavior make sense given the other’s?

3. Why did Marcus's unilateral change backfire, and what does that show about this pattern?
4. Is this a formal Nash equilibrium? How should you describe it honestly?
5. What would have to change for a different pattern to become attractive to both?

Model answer

1. **The loop:** Priya flags risks, Marcus goes quiet, Priya escalates (longer emails, cc to the director), Marcus feels attacked and shares less, Priya sees even less and escalates further. It's the chapter's "I push, you resist, I push harder" in a work setting.
2. **Why each response makes sense:** Given Marcus's silence, escalating is Priya's way to make sure risks get handled and she isn't blamed. Given Priya's escalation, guarding information and protecting his reputation is Marcus's. Neither is irrational. Both want a partnership; the pattern persists anyway. Stable doesn't mean good.
3. **Why "just change" backfired:** Marcus changed while everyone else kept their expectations, so his silence read as disengagement and he paid for moving alone. That's the heart of equilibrium: if I change while you don't, I'm worse off.
4. **Honest description:** With no formal strategies and payoffs, it isn't a proven Nash equilibrium. It "resembles an equilibrium trap," and the equilibrium check supports that: changing alone did leave Marcus worse off.
5. **What would have to change:** Sequence and information: replace the one-way email with a 15-minute joint risk review, so Marcus's handling is visible and Priya needn't escalate. Commitment: neither copies the director on the other's area without talking first. Third party: they brief the director together, so a quieter channel isn't read as disengagement. Each now expects the other's cooperation, so neither carries the whole risk of going first. "Marcus should communicate more" doesn't do this; the change has to touch the structure.

Exercise B — Your own game

Choose a recurring interaction that frustrates you (the same argument, a standoff with a coworker, a stalled negotiation), ideally one where you're tempted to say, "It's their fault."

1. Fill in the Equilibrium Audit. Write the pattern as a loop at least five steps long.
2. For question 4, name the exact cost you fear if you change alone, then the cost you think they fear. Mark your guess about theirs "untested."
3. Run the equilibrium check: is it really the cost of changing alone holding this in place?
4. Name one alternative stable pattern. Pick one lever and design a small, honest experiment, ideally one that lowers the other person's risk of changing too.
5. Complete: "Regardless of who started it, what keeps it going now is ____." Keep this page for the field assignment.

Exercise C — Partner drill / role-play

Time: 15 minutes. **Players:** two people plus an optional observer (or one person with an AI practice partner).

Setup: Two roommates. The kitchen is always messy. Each player can “Clean” or “Wait.”

Player A privately knows/wants: You want a clean kitchen, but you cleaned three times last month and your roommate never did. “If I clean and they don’t, I’m the doormat.”

Player B privately knows/wants: You want a clean kitchen too, but whenever you’ve cleaned, your roommate redid it “the right way,” which felt like criticism. “If I clean, I’ll get corrected. Better to wait.”

Round 1 (4 min): Play the usual conversation in character, each making the move that seems sensible from their private view. The observer writes down the loop.

Round 2 (4 min): Player A changes unilaterally (cleans without comment) while Player B keeps the same beliefs. Notice what happens.

Round 3 (5 min): Step out of character for 60 seconds and share “what I fear if I change alone.” Then replay, looking for a structural change: a schedule (sequence), a shared standard (information), a simultaneous commitment, or a different split of chores (available choices).

Debrief questions: 1. In Round 1, did either player behave irrationally given what they believed? What kept the pattern stable? 2. What happened to the unilateral change in Round 2? Did it shift the other player’s expectations, or get misread? 3. Which lever in Round 3 made a different pattern attractive to both? Why did it work better than good intentions?

Common mistakes

- **Treating “stable” as “good” or “chosen.”** A repeating pattern may simply mean nobody sees a better unilateral move. Ask what keeps it stable, not who secretly prefers it.
- **Hunting for the villain.** Ask, “What interaction is producing this pattern?” This doesn’t erase responsibility; it adds the question of what keeps the pattern going now.
- **Fixing the next sentence instead of the pattern.** If the loop is self-reinforcing, even a better line gets pulled back into it. Look at the system, not just the move.
- **“Just change.”** A one-sided change can make things worse at first if the other person reads it through the old pattern. Change something structural, especially expectations about what the other will do.
- **Calling everything a Nash equilibrium.** Persistence can come from habit, coercion, inertia, or rules. Test whether changing alone would actually hurt the person, and call it an equilibrium lens.

- **Treating equilibrium as prediction.** It doesn't tell you which stable outcome people will reach or exactly what they'll do. Model, act, observe, update.

Quiz

1. In game theory, a Nash equilibrium is best described as:
 - A) The outcome that gives everyone the highest possible payoff
 - B) A combination of strategies in which each player's strategy is a best response to the others'
 - C) A fair agreement that all players have explicitly accepted
 - D) The outcome most people will choose in the real world
2. Which statement about equilibrium is true according to the chapter?
 - A) If a pattern is stable, it must be good for at least one person
 - B) An equilibrium requires an explicit agreement
 - C) A pattern can be stable even if everyone involved dislikes it
 - D) Equilibria only occur when someone acts irrationally
3. In the Prisoner's Dilemma, why do both players end up defecting?
 - A) They don't realize mutual cooperation would be better
 - B) Regardless of what the other does, defecting looks like each player's best response
 - C) One player tricks the other into defecting
 - D) The rules forbid cooperation
4. Dana and her brother haven't called each other in months. Dana says, "He never calls me." Her brother says, "I'm always the one who has to reach out." Which question best reflects this module's approach?
 - A) "Which of them is being unreasonable?"
 - B) "What should Dana say on the next call?"
 - C) "What pattern are their responses creating, and why is it stable?"
 - D) "How can Dana get her brother to admit he's wrong?"
5. A friend says, "My partner and I fight constantly. I'm just going to stop engaging." Based on the chapter, what's the most important caution?
 - A) Stopping engagement always resolves conflict
 - B) A unilateral change may initially make things worse if the partner reads silence as abandonment
 - C) The friend should engage even more forcefully instead
 - D) Nothing can change once an equilibrium exists
6. What is the "equilibrium-selection problem"?
 - A) Choosing which player is responsible for the bad pattern

- B) Knowing the possible stable outcomes without knowing which one people will land on
 - C) Selecting the best possible outcome for yourself
 - D) Deciding whether to use game theory at all
7. A new hire at a firm where everyone stays until 7 p.m. starts staying late too, worried that leaving at five would look uncommitted. This best illustrates:
- A) An equilibrium sustained by expectations about what others will do
 - B) A written company policy
 - C) A one-shot Prisoner's Dilemma with a unique cooperative outcome
 - D) Proof that everyone at the firm prefers working late
8. The Charness, Rigotti, and Rustichini experiments found one-shot Prisoner's Dilemma cooperation ranging from 23% to 60% depending on the payoffs. What is the most accurate takeaway?
- A) Nash equilibrium perfectly predicts human behavior
 - B) Game structure is irrelevant to how people behave
 - C) Payoff structure shapes behavior, but the simple money-only Nash prediction doesn't fully describe real people
 - D) People always cooperate when given the chance

Answer key

1. **B** — Mutual best responses, so no one gains by changing alone. Nothing about best, fair, agreed, or predicted.
2. **C** — Stable doesn't mean good. A Nash equilibrium can be terrible for everyone, and it doesn't require any agreement.
3. **B** — Whatever the other does, defecting looks better. They know cooperation would be better; the structure pushes them apart anyway.
4. **C** — Each sibling's behavior responds to the other's. The chapter shifts from blame and wording to the pattern.
5. **B** — "Just change" can be bad advice: a one-sided move can be misread and trigger escalation.
6. **B** — Knowing the equilibria doesn't tell you which one people select; context, focal points, and expectations come into play.
7. **A** — Nobody wrote the rule. The new hire responds to expectations, then reinforces the norm: a cultural equilibrium.
8. **C** — Payoffs changed cooperation rates, so structure matters, but many people cooperated against the money-only prediction.

Field assignment

This week: catch one equilibrium in the wild and run one experiment.

1. **Observe (days 1 to 3):** Take your Exercise B situation or a new one. Each time it happens, jot down the loop: their move, your response, their response. Don't fix anything yet. Note when you reacted to what you expected them to do rather than what they did.
2. **Audit (day 4):** Update the Equilibrium Audit from your notes, including the equilibrium check.
3. **Experiment (days 5 to 7):** Pick one lever and try one small, honest change, preferably one that doesn't leave the other person carrying all the risk of going first. For example, name the pattern without blame: "I think we're in a loop where I ask late and you say no, and we both end up frustrated. Could we try a weekly check-in?" Don't use this to corner or pressure anyone. The aim is a better pattern both of you would choose freely.
4. **Record:** What happened, and what did you learn about what keeps the pattern stable?

Debrief questions for next week: 1. When you mapped the loop, what did you see that you hadn't seen when you looked only at individual incidents? 2. What made changing alone risky, for you and, as best you could tell, for them? Were your guesses about their fears confirmed or wrong? 3. Did your experiment change the other person's expectations? If not, what would you change about the structure next time?

The question to carry

What equilibrium have we created? Why is it stable? What does each of us get from maintaining it? What makes changing it risky? Are there other equilibria available? And what would have to change for us to move toward one of them?

Module 6 — Map Information & Beliefs

The question: “What do they know—or believe—that I don’t?”

Book reading	Chapter 6, “Map Information & Beliefs”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Information Map

Why this module matters

You and another person can be standing in exactly the same situation and still be playing very different games, because you have different information. You may know something they don’t. They may know something you don’t. You may both know something, but one of you may not know that the other knows it. And sometimes what matters most isn’t what is actually true. It’s what each person believes is true.

Take the used car. The seller knows the car has a serious mechanical problem; you don’t. You think you’re negotiating over “What is this car worth?” The seller knows you’re really negotiating over “How much can I get you to pay before you discover the problem?” Same car, same conversation, different information, and therefore a very different game. Reverse it: if you’ve already been approved for financing elsewhere and the seller doesn’t know, your alternatives look different to you than they do to the seller. The information structure changes the game.

The previous modules taught you to see the game, the payoffs, the responses, the patterns others learn about you, and the equilibrium those patterns create. This module goes one level deeper: to the beliefs underneath the behavior. Communication is, in large part, an attempt to change what another person knows or believes. Before you can communicate well, you need an honest map of who knows what, who believes what, and what you are merely assuming.

Learning objectives

By the end of this module you will be able to:

1. Distinguish what is actually true, what each player knows, what each player believes, and what you are assuming.
2. Identify information asymmetries in everyday interactions and ask what each side knows that the other doesn’t.

3. Evaluate a statement or signal by asking what the speaker gains, what it costs them, and why you should believe it.
4. Recognize higher-order beliefs and the difference between shared information and common knowledge.
5. Treat your interpretations of other people's behavior as revisable hypotheses, and name the evidence that would change your mind.

Lessons

Lesson 6.1 — Information Is Not the Same as Truth (~10 min)

Teaching points

- There is what is actually true, and there is what each player believes is true. They aren't necessarily the same.
- People make decisions based on the game as they perceive it, not the game as it objectively is.
- Beliefs about what other people will do can change what you do, and that can help create the very outcome people feared.
- This is not "perception is reality." Truth determines whether a belief is accurate, and inaccurate beliefs can collapse when reality produces contradictory evidence. The precise claim: perception is part of the causal structure of strategic interaction.

Example to use: The bank run. You believe your bank may fail; you don't know whether it's true, but you know that if enough depositors withdraw, it could fail. So you withdraw. Others see you, get nervous, and withdraw too. The outcome is shaped not only by the bank's actual condition but by what depositors believe other depositors will do. Pair it with the company doing well while employees believe layoffs are coming: they start job hunting, stop taking risks, and become less committed, which may then affect the company.

On-screen takeaway: People make decisions based on the game as they perceive it.

Lesson 6.2 — The Information Gap (~10 min)

Teaching points

- Information asymmetry: one player knows something relevant that another doesn't. It is everywhere: doctor and patient, mechanic and customer, lawyer and client, employee and manager, child and parent.
- Information has strategic value. That's why people ask questions, investigate, stay silent, reveal selectively, observe, and test one another. They're trying to improve their model of the game.
- "Just ask" has a problem: the other person gets to decide what to tell you, and they may have an incentive to tell you what benefits them.

- In formal terms, players can have private information about their “type.” In plain terms, a type is something important about me that you don’t know: skilled or inexperienced, patient or desperate, willing to walk away or not. You have to infer it, and I know you’re trying to.

Example to use: The job interview. The candidate knows how capable and hardworking they really are; the employer has to infer it from résumés, references, work samples, and answers. But the employer also knows the candidate has an incentive to present themselves favorably, so claims aren’t automatically treated as facts. The employer has to separate information from self-interested persuasion.

On-screen takeaway: How do you make decisions when someone else knows something you don’t?

Lesson 6.3 — Signals and Cheap Talk (~10 min)

Teaching points

- A signal is an action or piece of information that helps another player infer something they can’t directly observe. Michael Spence’s job-market model showed how education can function as a signal when acquiring it costs different types differently.
- A signal is evidence, not proof. Its value depends on how informative it actually is.
- Costly signals can be more informative because different types of players face different costs or benefits from sending them. The everyday question: what does this person have to lose by making this claim?
- Cheap talk is communication that doesn’t directly change the material payoffs. It isn’t useless. It can carry real information, especially when the speaker’s interests are aligned with the listener’s. The key question is: why should I believe you?
- The same words can carry different information depending on the speaker, incentives, history, context, alternatives, and consequences.

Example to use: “I’m extremely confident” costs almost nothing to say. Putting \$10,000 of your own money behind the prediction is different. Similarly, “I really value our relationship” is a cheap statement, while repeated sacrifices, kept commitments, and showing up when it’s hard are far more informative. And when two salespeople each say “This is our best price,” the one with a reputation for transparency is telling you more than the one known for exaggeration.

On-screen takeaway: A signal is evidence, not proof.

Lesson 6.4 — Beliefs About Beliefs and Common Knowledge (~8 min)

Teaching points

- Higher-order beliefs (belief hierarchies): you have beliefs about what the other person believes about you, and so on. “I believe you think I’m bluffing.”

- Sometimes you're not responding to what someone knows. You're responding to what you think they know about what you know.
- A bluff works not because it's true but because the other player believes it enough to act differently.
- Mistaken beliefs can sustain a stable pattern. If an equilibrium rests on incorrect beliefs, changing the information may change the equilibrium.
- Common knowledge: we both know something, and we both know that we both know it, and so on. Coordination can fail not from lack of information but from lack of shared knowledge that the information is shared.

Example to use: Two departments. A believes "B will never cooperate," so A stops sharing information. B sees this, concludes "They don't trust us," and stops sharing too. A says, "See? We were right." Neither started with malice; the belief changed behavior and the behavior reinforced the belief. For common knowledge, use the room of ten people: if everyone hears "At noon we all leave the building," coordination is easy. If only half hear it, those who heard don't know whether everyone else did.

On-screen takeaway: Beliefs can influence actions, and actions can reinforce beliefs.

Lesson 6.5 — Hypotheses, Not Facts (~7 min)

Teaching points

- Don't confuse uncertainty with deception. You may not know what someone is thinking because they're withholding it, because they don't know themselves, because the information isn't available, because you haven't asked, or because you've misread ambiguous evidence.
- The healthier question: what information do I actually have, and how confident should I be in the conclusions I'm drawing from it?
- Treat your first interpretation as a hypothesis. What evidence would support it? What would contradict it? What else could explain the behavior?
- Humans are good at forming an interpretation and then hunting for confirmation. Your model of the other player should be revisable.

Example to use: You text someone, "Can we talk?" No reply. You know very little, but your mind generates explanations: angry, busy, didn't see it, avoiding you, something happened. That belief drives your next move (another text, withdrawal, anger), and now your belief about their behavior is influencing their future behavior. Uncertainty can create conflict even when neither person intends it.

On-screen takeaway: Your model of the other player should be revisable.

Key terms

Term	Plain-language meaning
Information structure	Who knows what in an interaction. Change it and you can change the game.
Information asymmetry	One player knows something relevant that another player doesn't.
Incomplete information	Uncertainty about other players' payoffs or types.
Type	Something important about a player that others don't directly observe and have to infer (skill, patience, desperation, trustworthiness, willingness to walk away).
Signal	An action or piece of information that helps another player infer something they don't directly observe. Evidence, not proof.
Costly signal	A signal that would cost the sender something, often more informative because different types face different costs for sending it.
Cheap talk	Communication that doesn't directly change the material payoffs. It can still be informative, especially when interests are aligned.
Higher-order beliefs	Beliefs about what others believe, including what they believe about your beliefs. Also called belief hierarchies.
Common knowledge	We both know something, we both know that we both know it, and so on. Especially important for coordination.
Self-fulfilling belief	A belief that changes behavior in a way that produces the evidence that seems to confirm it.

Research spotlight

Tadelis and Zettelmeyer, wholesale auto auctions (American Economic Review).

The researchers ran a large-scale field experiment in wholesale car auctions, randomly disclosing quality information about some vehicles. Both good and bad news could increase competition and revenue, because buyers differed in which cars suited them, and better information improved the matching between buyers and vehicles. The cars didn't change. The buyers didn't change. The information changed, and so the game changed.

What it suggests for real conversations: Changing the information structure can change behavior and outcomes, sometimes in ways people don't expect. Disclosing an unflattering fact isn't automatically a loss; it can help the right counterpart find you.

What it does NOT prove: It does not show that more information always improves outcomes. The effect depends on the game. It's evidence from a specific market setting, not a universal rule that disclosure always pays.

Costa-Gomes and Weizsäcker, stated beliefs and play (Review of Economic Studies). Participants played 14 one-shot strategic games and separately reported what they expected their opponents to do. If people simply chose the best response to their own stated beliefs, the two should line up. They didn't: on average, participants failed to best respond to their own stated beliefs in almost half the games, and their stated beliefs often reflected deeper strategic reasoning than their actual choices.

What it suggests for real conversations: Beliefs matter, but people don't always translate beliefs into consistent action. Knowing what someone believes is useful, not a crystal ball. That includes you: your own stated beliefs may not be what is actually driving your moves.

What it does NOT prove: It doesn't show that beliefs are irrelevant, or that people are random. And it doesn't support the claim that correctly identifying someone's belief lets you reliably predict their behavior. A related methodological note in the chapter's research (Rutström and Wilcox) found that asking people for their beliefs with incentives could alter early play in one repeated game under some conditions, a reminder that even measuring beliefs can affect the game.

The tool: The Information Map

Use this before any important interaction, or after one that confused you. The first seven questions are the chapter's Information Map. The sorting table and the final checks help you keep facts, beliefs, and assumptions in separate columns.

The situation:

The other player(s):

Part 1 — The seven questions

1. **What do I know?** (Separate facts from assumptions.)

2. **What don't I know?** (Identify the uncertainty.)

3. **What do they know that I don't?** (Look for information asymmetry.)

4. **What do I know that they don't?** (Recognize your own informational advantage.)

5. **What do I think they believe?** (Model their beliefs.)

6. **What do they think I believe?** (If the stakes are high, this matters too.)

7. **What evidence could change my mind?** (Protect yourself from your own assumptions.)

Part 2 — Sort it

Layer	Question	My answer
Facts	What do I actually know happened? (Observable, verifiable.)	
Unknowns	What relevant information am I missing?	
Beliefs	What do I currently think is true?	
Assumptions about them	What am I assuming they know or believe?	

For each belief, mark your confidence: **I know / I have good evidence / I suspect / I'm guessing.**

Part 3 — Shared or common?

- Is the key information actually in their information set? What evidence do I have that they know it?
-

- Do we both know that we both know it, or am I assuming it's shared?
-

Part 4 — Messages I'm relying on

- What did they tell me that matters?
-

- What do they gain if I believe it? What would it cost them if it were false?
-

- How could I verify it?
-

-
- **Why should I believe it?**
-

Part 5 — Decision-relevant information

- What single piece of information would change my decision, or theirs?
-
- How can I get it (ask, observe, verify, share)?
-

Practice exercises

Exercise A — Analyze the scenario

Dana manages a five-person analytics team. On Friday, the quarterly report Jordan owned went out two days late, and Dana's director asked her about it in front of other managers. Dana had told the team on Monday that "the director is watching this one closely," but she said it in a quick hallway exchange that Jordan wasn't part of. Jordan had sent a message mid-week saying "Still waiting on finance's numbers," and Dana, skimming it on her phone, read it as a routine status update. On Friday afternoon, Dana writes to Jordan: "We need to talk Monday about ownership." Jordan, who stayed late Thursday rebuilding the model after finance finally delivered, reads the message as a sign that Dana thinks he's careless. Over the weekend, he tells a friend, "She already knows finance was late. She just wants someone to blame." Dana, meanwhile, believes Jordan didn't flag the risk because he doesn't take the director's priorities seriously.

1. Sort what Dana knows (facts) from what she believes and what she's assuming.
2. What information asymmetries exist in each direction?
3. Identify at least one higher-order belief and one place where Dana treated shared information as common knowledge.
4. How could the beliefs on both sides become self-fulfilling?
5. What information would change either person's decision, and how should Dana open Monday's conversation?

Model answer

1. **Dana's facts:** the report was two days late; the director raised it publicly; Jordan sent a mid-week message mentioning finance's numbers. **Her beliefs:** Jordan didn't flag the risk; he doesn't take the director's priorities seriously. **Her assumptions:** that Jordan knew the director was watching closely, and that his message was routine. Her confidence in "he doesn't take it seriously" sits at "I suspect," but she's acting as if it's "I know."

2. **Jordan knows what Dana doesn't:** finance delivered late, he worked late Thursday, and his mid-week message was meant as a warning. **Dana knows what Jordan doesn't:** the director's heightened attention and the public embarrassment on Friday. Each is reacting to a different game.
3. **Higher-order belief:** Jordan believes Dana already knows finance was late and is choosing to blame him anyway. He's responding not to what Dana knows but to what he thinks she knows. **Common knowledge failure:** Dana said "the director is watching" in a hallway Jordan wasn't part of. The information existed on the team, but it wasn't shared knowledge that everyone had it. She assumed it was in Jordan's information set without evidence.
4. If Dana walks in treating her hypothesis as fact, Jordan will likely become defensive and guarded, which Dana may read as confirmation that he doesn't care. Jordan, expecting blame, may stop sending early warnings ("What's the point?"), which confirms Dana's belief that he doesn't flag risks. This is the two-departments pattern in miniature.
5. Decision-relevant information: when finance actually delivered, what Jordan's mid-week message was meant to signal, and whether Jordan knew about the director's attention. Dana should open by naming facts and asking, not concluding: "The report went out two days late and the director raised it with me. I'm not sure I have the full picture. Walk me through what happened with finance's numbers, and tell me if you'd heard the director was watching this one." She should also share what she knows, since Jordan can't respond to information he never had. Neither of them needs to deceive or pressure anyone; both need a more accurate map.

Exercise B — Your own game

Pick a live situation where someone's behavior is confusing, frustrating, or worrying you: a coworker who has gone quiet, a family member who reacted oddly, a client who stopped responding, a negotiation that stalled.

1. Write one sentence describing what happened, using only observable facts. No adjectives about motives.
2. Write your current interpretation. Label it: "Hypothesis."
3. Work through all five parts of the Information Map.
4. List at least three other explanations for the behavior. Use the chapter's list: are they withholding, do they not know themselves, is the information unavailable, have you not asked, or have you misread ambiguous evidence?
5. For your hypothesis, write one piece of evidence that would support it and one that would contradict it.
6. Identify one thing you know that they may not, and decide whether sharing it would change the game for the better.
7. Decide on a single next step that gathers decision-relevant information honestly, such as a direct question or a check-in, rather than a move based on an untested belief.

Exercise C — Partner drill / role-play

Time: 15 minutes. **Players:** two people (or one person plus an AI practice partner), plus an optional observer.

Setup: A small business owner (Player 1) is considering hiring a freelance designer (Player 2) for a three-month website project.

Player 1 privately knows: The budget is firm at the quoted amount, but the deadline is more flexible than the posting suggests. A previous freelancer missed deadlines badly, so Player 1 is nervous about reliability. Player 1 has one other candidate, who is cheaper but less experienced.

Player 2 privately knows: They're strong at design but have never built a site this size. They have another offer that starts in six weeks and would overlap. They really want this client for their portfolio.

Round 1 (5 min): Have the conversation. Each player tries to learn what matters and decides what to share. No lying about material facts; you may keep things private.

Round 2 (5 min): Stop. Each player silently fills in Part 1 of the Information Map: what I know, what I think they know, what I think they believe about me. Then compare maps out loud. Where were you right? Where were you wrong?

Round 3 (5 min): Replay the last part of the conversation, this time asking the one question whose answer would most change your decision.

Debrief questions

- Which of your beliefs about the other person turned out to be assumptions?
- Which claims did you treat as strong evidence, and which as cheap talk? What made the difference?
- What did you keep private, and was that different from deceiving? Did the other person have what they needed for an informed choice?
- What single piece of information changed the game most when it came out?

Common mistakes

- **Promoting a hypothesis to a fact.** Forming an interpretation is natural. The error is treating it as settled before you have evidence. Label it a hypothesis, name what would contradict it, and update as information arrives.
- **Assuming “they know.”** “He knows why I’m upset” is often an assumption. A person cannot respond to information they never had. Ask what evidence you have that it’s in their information set.
- **Taking statements at face value, or dismissing them entirely.** Words aren’t automatically facts, and cheap talk isn’t automatically worthless. Evaluate the speaker, incentives, history, and context, and ask why you should believe it.

- **Treating a signal as proof.** A degree, a confident tone, or a single gesture is evidence. Its weight depends on how informative it is and what it cost to send.
- **Assuming uncertainty means deception.** Not knowing what someone thinks doesn't mean they're hiding something. They may not know themselves, the information may not exist yet, or you may not have asked.
- **Confusing shared information with common knowledge.** Telling half the room isn't the same as everyone knowing that everyone knows. If coordination matters, make the information public and mutual.
- **Saying "perception is reality."** Beliefs drive behavior, but truth still decides whether beliefs are accurate, and false beliefs can collapse when reality pushes back.
- **Gathering more information instead of the right information.** Ask what would actually change your decision, and go after that.

Quiz

1. According to the chapter, what do people base their decisions on?
 - A) The objective facts of the situation
 - B) The game as they perceive it
 - C) Whatever the most powerful player tells them
 - D) Their payoffs alone, regardless of beliefs
2. In the bank-run example, what helps produce the outcome people fear?
 - A) The bank's hidden accounting fraud
 - B) Government regulation
 - C) What depositors believe other depositors will do
 - D) A single large withdrawal by an insider
3. A coworker says, "Trust me, this vendor is the best option." Which question from the chapter best tests the claim?
 - A) "Is my coworker a nice person?"
 - B) "What does this person have to lose by making this claim?"
 - C) "How confidently did they say it?"
 - D) "Have I heard this vendor's name before?"
4. Which statement about cheap talk matches the chapter?
 - A) It is always worthless because it is costless.
 - B) It is the same thing as lying.
 - C) It can convey meaningful information, especially when interests are aligned.
 - D) It only exists in formal games, not in everyday life.
5. "I think my manager believes I think she's micromanaging me." This is an example of:
 - A) A costly signal

- B) A higher-order belief
 - C) Common knowledge
 - D) Information asymmetry about type
6. A team lead tells three of six members that Friday's meeting is canceled. The other three show up. What was missing?
- A) Payoffs
 - B) A costly signal
 - C) Common knowledge that the information was shared
 - D) A focal point based on reputation
7. Your partner hasn't replied to "Can we talk?" for three hours. What does the chapter recommend first?
- A) Assume they're avoiding you and respond in kind
 - B) Send several more messages until they answer
 - C) Treat "they're angry" as a hypothesis and ask what evidence would support or contradict it
 - D) Assume they're deliberately withholding information
8. What does the Costa-Gomes and Weizsäcker experiment suggest?
- A) People always best respond to their stated beliefs.
 - B) Stated beliefs are meaningless.
 - C) People often fail to best respond to their own stated beliefs, so knowing a belief doesn't guarantee predicting behavior.
 - D) More information always leads to better choices.

Answer key

1. **B** — "People make decisions based on the game as they perceive it." Their perception may or may not match the facts.
2. **C** — The outcome depends partly on what depositors believe others will do, and a belief can help create the failure it predicts.
3. **B** — The chapter's everyday signaling question. If the claim costs the speaker nothing, it's weaker evidence.
4. **C** — Cheap talk doesn't change material payoffs, but it can still carry information, particularly when the sender's and receiver's interests are aligned.
5. **B** — It's a belief about someone's belief about your belief: a belief hierarchy.
6. **C** — Like the room where only half hear the noon announcement, the problem wasn't lack of information but lack of shared knowledge that it was shared.
7. **C** — Treat your first interpretation as a hypothesis and keep your model revisable. Uncertainty isn't the same as deception.

8. **C** — Participants failed to best respond to their own stated beliefs in almost half the games. Beliefs matter, but they don't translate perfectly into action.

Field assignment

This week, catch yourself in one moment of confusion about someone's behavior: an unanswered message, a cool response in a meeting, an unexpected "no." Before you react, take two minutes and run the Information Map, at least Part 1 and the confidence labels in Part 2. Then take one honest step to get decision-relevant information: ask a direct question, share something you know that they might not, or verify a claim you were relying on.

Record: the situation; your first interpretation; your confidence label before; what you learned; and whether your belief survived, changed, or collapsed.

Debrief questions for next week 1. How far apart were your first interpretation and what you eventually learned? 2. Did you discover any information the other person lacked, or any you lacked? What changed once it was shared? 3. Where else in your life might a belief, yours or theirs, be sustaining a pattern that new information could change?

The question to carry

"What do they know—or believe—that I don't?"

Then ask: What do I actually know? What am I assuming? What might they be assuming? What information would change either of our decisions? And when someone tells you something important, don't stop at "What did they say?" Ask: **"Why should I believe it?"**

Module 7 — Separate Words From Credibility

The question: What makes this message credible?

Book reading	Chapter 7, “Separate Words From Credibility”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Credibility Test

Why this module matters

In Module 6 you learned that people decide based on what they know, what they believe, and what they think others believe. That creates a new problem. Suppose I know something you don’t, and I tell you. Should you believe me? That depends, because I’m not simply a source of information. I’m also a player in the game, with interests and incentives, and what I want you to believe may affect what you decide.

Consider “Trust me. I’ll take care of this.” From a business partner who has kept every major commitment for fifteen years, it means one thing. From a stranger you’re about to wire \$50,000 to, it means something else. Same words, completely different credibility. Credibility isn’t contained inside the sentence. It comes from the game surrounding it: history, incentives, evidence, knowledge, reputation, consequences, verification, and what the speaker stands to gain or lose.

This module trains one habit: separating what someone says from what makes it credible. It replaces “Do I believe them?” with “What makes this believable?” And it turns the lens on you as a speaker. When you’re telling the truth and still not being believed, the problem is often not your wording. The other person may simply have no good reason to believe you yet.

Learning objectives

By the end of this module you will be able to:

1. **Distinguish** how persuasive a message is from how credible the underlying claim is.
2. **Evaluate** a claim along separate dimensions: honesty, knowledge, and incentives.
3. **Apply** the “Would they say this anyway?” test to spot statements that carry little information on their own.

4. **Explain** why a costly signal is informative only when it is easier or more worthwhile for someone whose claim is true.
5. **Design** ways to make your own truthful claims easier for others to verify.

Lessons

Lesson 7.1 — The Game Around the Sentence (9 min)

Teaching points

- Credibility lives in the surrounding game, not in the words. Persuasion and credibility are different: a message can be persuasive without being credible, and credible without being persuasive.
- The formal setup: an informed **sender** communicates, a **receiver** decides, and the sender cares what the receiver decides. That makes communication strategic.
- Crawford and Sobel's model shows that how much information can be credibly transmitted depends partly on how closely the players' interests are aligned. So don't evaluate a message without considering the sender's incentives.
- Avoid paranoia. "They benefit from saying it, therefore it's false" is not strategic thinking.

Example to use: A stranger tells you the train station is two blocks east. Their payoff barely depends on which way you walk, so the words are informative. Now the same person is selling you their car and says the best offer they've had is \$20,000. The answer affects how much they receive. That doesn't mean they're lying. It means the statement alone gives you less reason to update than it would if their incentives were neutral or aligned.

On-screen takeaway: They benefit from my believing this, so the statement by itself is weaker evidence. Not false. Weaker.

Lesson 7.2 — Cheap Talk Isn't Worthless Talk (8 min)

Teaching points

- **Cheap talk** is costless, nonbinding communication. Saying it doesn't directly change the material structure of the game.
- "If talk is cheap, talk is meaningless" is wrong. Farrell and Rabin argued both extremes fail: cheap talk doesn't automatically produce truth-telling, but it isn't necessarily meaningless.
- In experiments, communication was largely truthful until the incentive to exaggerate rose. Credibility is partly a property of incentives.

Example to use: You can see the GPS and I can't. You say "Turn left." It costs you nothing and is highly informative because we both want to arrive. Now we're racing to get there

first. Same words, different strategic meaning. I have to ask, “What do you gain if I believe you?” Communication didn’t change; the game did.

On-screen takeaway: Words can be cheap and still useful.

Lesson 7.3 — Honest, Informed, and Aligned Are Different Things (9 min)

Teaching points

- Don’t stop at “Are they honest?” Also ask: Are they informed (in a position to know)? What are their incentives (what do they gain or lose depending on what I believe)?
- **Expertise without alignment:** a salesperson for a complex financial product knows far more than you and earns a large commission if you buy. Don’t ignore them; separate “How much do they know?” from “How much does their payoff depend on my decision?”
- **Alignment without expertise:** your brother wants the best for you but knows nothing about your medical decision. Ask, “Trust him to do what?”
- **Truthfulness versus accuracy:** a false statement isn’t automatically a lie. Someone can be truthful but wrong.
- **Credibility is not confidence.** Certainty and eye contact are not evidence, and confidence can itself be strategic. Don’t confuse the delivery of a message with the evidence behind it.

Example to use: A friend sincerely tells you an investment is “guaranteed to go up,” while knowing almost nothing about investing. Honesty isn’t accuracy. Contrast the salesperson: high knowledge, uncertain alignment. Two different credibility problems that one “Do I trust them?” blurs together.

On-screen takeaway: Don’t ask “Do I trust them?” Ask “Trust them to do what?”

Lesson 7.4 — Signals, Evidence, and “Would They Say This Anyway?” (11 min)

Teaching points

- Ask: **Would they say this anyway?** If nearly everyone benefits from saying it whether or not it’s true, the words don’t distinguish between cases. “I’m extremely hardworking” is an example; a demanding work sample is not.
- Evidence that is harder for different types of people to produce equally moves us from cheap talk to **signaling**. Spence showed how behavior can carry information that words alone do not.
- Cost alone doesn’t make a signal credible. A terrible company can spend \$20 million on advertising. Ask why the signal would be easier, cheaper, or more worthwhile for someone whose claim is true.
- Actions can be strategic too. Treat behavior as **evidence, not proof**.
- Verification (inspections, audits, references, independent reviews) changes the information environment. Reputation is a summary of previous rounds of the game, and credibility can collapse faster than it builds.

Example to use: Two contractors both say, “We stand behind our work.” Contractor A wants full payment upfront, offers no references or warranty, and disappears after the job. Contractor B has hundreds of verifiable reviews, a written warranty, twenty years in the community, and fixes problems at their own expense. B has created evidence that would be hard to maintain if the claim were repeatedly false. Pair it with the manager who says “Challenge me,” then sidelines the employee who did. The observed consequence of disagreement now outweighs the invitation.

On-screen takeaway: Why would this signal be easier for someone whose claim is true?

Lesson 7.5 — The Speaker’s Side: Make Truth Easier to Verify (8 min)

Teaching points

- A persuasion failure may be a credibility failure. If the listener understood you and still doesn’t believe you, the question isn’t “How do I say this better?” It’s “How do I make this claim more credible?”
- “I know I’m honest” isn’t evidence for them. They may lack your information, have been misled before, or see that you benefit from being believed.
- When you’re telling the truth, make it easier to evaluate: evidence, demonstrations, references, guarantees, documentation, and consistency between words and behavior.
- Signaling asks, “What does my action tell you about me?” Commitment asks, “How does my action change my future options?” Module 8 covers commitment.
- Ethics: never manufacture the appearance of truth. Fake reviews, fabricated evidence, and staged signals are deception. And don’t become cynical. You can trust and still understand the game.

Example to use: You’re selling your car and know the transmission is fine. The buyer wants a mechanic to inspect it, and you feel insulted. From their side, the seller knows more and benefits from the sale. The inspection isn’t an accusation; it’s a response to an information asymmetry. Similarly, a client who doesn’t believe your on-time promise after three late projects won’t be moved by a better story, but milestones they can check might help.

On-screen takeaway: Ask “What reason do they have to believe me?”

Key terms

Term	Plain-language meaning
Credibility	How much reason a claim gives you to update, given the speaker’s knowledge, incentives, evidence, behavior, and history. Not the same as persuasiveness.

Term	Plain-language meaning
Sender / receiver	The informed player who communicates, and the player who then decides. Communication is strategic when the sender cares what the receiver decides.
Cheap talk	Costless, nonbinding communication. It can still carry real information when interests are aligned.
“Would they say this anyway?”	A test of whether a statement distinguishes true from false cases. If everyone would say it regardless, the words alone carry little information.
Signal	An observable action that conveys information about something hidden, especially one that is easier or more worthwhile for someone whose claim is true.
Verification	Inspections, audits, certifications, references, or independent reviews that change the information environment around a claim.
Reputation	A summary of information from previous rounds of the game. It accumulates and can collapse.
Truthfulness vs. accuracy	Do they report what they believe? Is what they believe correct?
Commitment	An action that changes what you can rationally do later, as opposed to a signal, which tells others something about you.

Research spotlight

1. Cheap talk: Crawford and Sobel; Farrell and Rabin; Cox and Stoddard

What was done. Crawford and Sobel modeled a better-informed sender communicating with a receiver whose later action affects both. Farrell and Rabin reviewed the cheap-talk literature. Cox and Stoddard ran an experiment in which players received private information about whether contributing to a public good would be valuable and could send costless messages first. The group gained from honesty, but individuals sometimes had an incentive to exaggerate.

What was found. In Crawford and Sobel’s model, how much information can be credibly passed depends on how aligned the players’ interests are, and informative communication can still occur when interests partly conflict. Farrell and Rabin argued that cheap talk is

neither automatically truthful nor automatically meaningless. In Cox and Stoddard's experiment, communication was largely truthful and improved efficiency. When the incentive to exaggerate increased, truthfulness declined and the benefits shrank.

What it suggests. The people didn't change; the payoff for exaggerating did, and communication changed with it. When you hear a claim, ask how much the speaker gains from moving your decision.

What it does NOT prove. It does not show that people always lie when interests conflict, that costless words are meaningless, or that aligned interests guarantee truth.

2. Signaling: Spence; Kübler, Müller, and Normann

What was done. Michael Spence's signaling work, recognized in the 2001 Nobel Prize in economics, showed how education can, under certain conditions, help employers distinguish among workers whose productivity they can't observe. Kübler, Müller, and Normann tested versions of this job-market signaling game in the lab.

What was found. More productive worker types invested in education more often, and employers paid higher wages to those who invested. But separation was incomplete, and wages didn't converge perfectly to the theoretical equilibrium.

What it suggests. Behavior that is harder to produce for someone whose claim is false (a work sample, a warranty, a verifiable track record) tells you more than self-description.

What it does NOT prove. It does not show that observable actions perfectly reveal hidden qualities, or that a costly action proves a claim true. Signals contain information; they don't perfectly reveal reality.

The tool: The Credibility Test

Use this for high-stakes claims, not for directions to the restroom.

The claim (in their words):

What decision of mine does it affect?

Part 1 — The eight questions

1. **Do they know?** Are they in a position to possess this information?

2. **What do they want me to believe?**

3. **What happens if I believe them?** How does my response affect their payoff?

4. **Would they say this even if it were false?** If yes, the words alone may not be highly informative.

5. **What evidence supports the claim?** Assertion:

Verification:

6. **What does their behavior suggest?** (Evidence, not proof.)

7. **What does their history suggest?**

8. **What would make this claim costly to fake?**

Part 2 — Separate the dimensions (high / medium / low / unknown)

- Honesty: _____ because

- Knowledge: _____ because

- Incentive alignment: _____ because

Part 3 — Persuasive versus credible

- How convincing did it feel (confidence, charisma, emotion)?

- How credible is the underlying claim?

- Did it change my mind, or give me good reason to change my mind?

- What verification is available, and is it worth getting?

Part 4 — Speaker's side (when you're the one making the claim)

- What reason does the other person have to believe me?

- Is this a message problem or a credibility problem?

- How can I make my truthful claim easier to evaluate?

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small marketing agency and is choosing a web developer for a large client project. Marcus is polished and energetic. He says, “I’ve never missed a deadline, and I’m the best React developer in this city.” He never hesitates, has a beautiful slide deck, and offers 10 percent off if Dana signs this week. Asked for references, he says past clients are “under NDA, but trust me, they were thrilled.” Priya is quieter and a bit awkward. She says, “I think I can hit your timeline, but the payment integration is the risky part, and I’d want a week to scope it.” She sends three live sites she built, names two former clients Dana can call, and proposes milestone payments tied to working deliverables. Dana’s business partner, who knows Marcus socially, says, “Marcus is a great guy. He’d never lie to you.”

1. Apply “Would they say this anyway?” to Marcus’s two main claims.
2. Which parts of each pitch are cheap talk, and which are signals or verification?
3. Separate honesty, knowledge, and incentives for the partner’s endorsement.
4. What should Dana do next, without becoming cynical?

Model answer

1. Almost any developer chasing this contract, reliable or not, benefits from saying “never missed a deadline” and “best in the city.” The words don’t distinguish dependable from undependable. That doesn’t make Marcus a liar; it means the claims alone give Dana little reason to update.
2. Marcus’s pitch is mostly cheap talk, delivered with confidence. The deck cost effort, but a weak developer could make one just as polished, so it doesn’t separate types. The signing discount is pressure, not evidence. Priya offers things harder to produce if her claims were false: live sites, callable references, and milestone pay. Milestone pay is far more attractive to someone who expects to deliver. Her volunteered risk also works against her short-term interest, which makes her other statements more believable. Dana shouldn’t read Priya’s awkwardness as weakness or Marcus’s polish as accuracy.
3. The partner is probably honest. Knowledge is uncertain: knowing Marcus socially says little about his deadline record or skill. Incentives are mostly aligned with Dana, though friendship may color the view. “He’d never lie” addresses honesty, not accuracy. Marcus could be sincere and still overestimate himself. Trust the partner to do what?
4. Ask Marcus for verification that doesn’t breach an NDA: a public portfolio, a work sample, or a general reference. Offer both candidates the same milestone structure and see how each responds. Actually call Priya’s references. This isn’t distrust of Marcus, who may be

excellent. It's a reasonable response to an information asymmetry where both candidates benefit from being believed.

Exercise B — Your own game

Pick a live situation where a claim matters to a decision you're about to make: a vendor's promise, a colleague's estimate, a contractor's quote, or advice from family.

1. Write the claim in the speaker's exact words and run Parts 1 through 3 of the Credibility Test.
2. Circle the weakest link: knowledge, incentives, evidence, behavior, or history.
3. Identify one piece of verification you could get this week, and decide whether the stakes justify it.

Then flip it. Pick a situation where you're telling the truth but aren't being believed, or soon need to be.

1. Complete Part 4. Honestly: would someone in your position say this even if it were false?
2. Write two ways to make your claim easier to evaluate, at least one the other person can check without relying on your word.
3. Check the ethics: does each step make the truth easier to see, or only make you look credible? Keep only the first kind.

Exercise C — Partner drill / role-play

Setup (15 minutes, pairs or threes). One **seller**, one **buyer**, an optional **observer**. The seller is selling a used bicycle at a fixed price.

Private information. The seller secretly draws a card. Card 1: "The bike is in excellent condition." Card 2: "The frame has a hairline crack. You needn't volunteer it, but you may not state a direct falsehood if asked."

Round 1 (4 min). The seller makes claims; the buyer asks questions but may not request verification. The buyer then guesses the card and rates confidence from 1 to 10.

Round 2 (4 min). New card draw. The buyer may now request verification: an inspection, a return window, a call to the previous owner. A Card 2 seller must honor any guarantee they offer.

Debrief (5-7 min).

- In Round 1, did the buyer's guess track the truth, or the seller's confidence?
- Which statements would the seller have made regardless of the card?
- In Round 2, which offers were easier for a Card 1 seller? Why did those carry more information?
- Card 1 sellers: how did it feel to be doubted while telling the truth? What helped?

AI version. Ask an AI to play a seller who secretly picks a card and reveals it only at the end.

Common mistakes

- **Treating “they benefit from saying it” as “it’s false.”** That’s paranoia. The correct conclusion is that the statement alone is weaker evidence.
- **Treating cheap talk as meaningless.** When interests align, costless words like directions or “Turn left” are highly informative.
- **Reducing credibility to “Are they honest?”** Ask about honesty, knowledge, and incentives separately, and ask “Trust them to do what?”
- **Confusing confidence with evidence.** Smooth delivery can be strategic, and an awkward, nervous person can be completely correct.
- **Assuming any costly signal is credible.** A \$20 million ad campaign doesn’t prove a product is good.
- **Answering a credibility problem with more rhetoric.** If they understood you and still don’t believe you, make the claim more credible instead of more eloquent.
- **Being offended by verification.** An inspection is a response to an information asymmetry, not an accusation.
- **Becoming cynical.** Most people cooperate and tell the truth. Don’t confuse trusting someone with failing to think, but keep trusting.

Quiz

1. According to the chapter, where does credibility primarily come from?
 - A) The precise wording of the sentence
 - B) The game surrounding the sentence: history, incentives, evidence, reputation, and verification
 - C) How confidently the message is delivered
 - D) Whether you like the speaker
2. A seller says the highest offer they’ve received is \$20,000. Which conclusion best reflects the chapter?
 - A) They benefit from saying it, so it’s false.
 - B) Treat it exactly like directions to the train station.
 - C) They benefit from your believing it, so the statement alone gives you less reason to update.
 - D) Ignore price claims entirely.
3. What did Farrell and Rabin argue about cheap talk?
 - A) Costless messages are always meaningless.

- B) Costless messages always produce truth-telling.
 - C) Both extremes are wrong; under the right incentives, cheap talk can carry substantial information.
 - D) Cheap talk works only between family members.
4. In Cox and Stoddard's experiment, what happened when the incentive to exaggerate increased?
- A) Truthfulness declined, and the benefits of communication shrank.
 - B) Truthfulness stayed the same.
 - C) Players stopped sending messages.
 - D) Efficiency improved.
5. A candidate says, "I'm extremely hardworking." Why is this weak evidence on its own?
- A) Candidates who say this are usually lying.
 - B) Almost any candidate benefits from saying it, so it doesn't distinguish between types.
 - C) Hard work can't be measured.
 - D) Self-descriptions are never useful.
6. Which question best captures what makes a signal credible?
- A) How expensive was it?
 - B) How dramatic was it?
 - C) Why would it be easier or more worthwhile for someone whose claim is true than for someone whose claim is false?
 - D) Did the person sound sincere?
7. You explained your timeline brilliantly, but the client still doubts you. You were late on the last three projects. What does the chapter suggest?
- A) Build a more persuasive story.
 - B) Treat it as a credibility problem: ask how to make the claim more credible.
 - C) Insist you're telling the truth.
 - D) Offer a discount to change the subject.
8. What distinguishes signaling from commitment?
- A) Signaling uses words; commitment uses actions.
 - B) Signaling asks what an action tells you about me; commitment asks how an action changes my future options.
 - C) Signaling is ethical; commitment is manipulative.
 - D) They are the same idea.

Answer key

1. **B** — The same "Trust me" means different things from a fifteen-year partner and a stranger asking for a wire. Credibility comes from the game around the sentence.

2. **C** — Their incentive doesn't make the claim false; it makes the statement weaker evidence than it would be if their interests were neutral.
3. **C** — Farrell and Rabin rejected both extremes. Words can be cheap and still useful.
4. **A** — Communication was largely truthful until the payoff for exaggerating rose. Credibility is partly a property of incentives.
5. **B** — The "Would they say this anyway?" test. A demanding work sample or track record is harder to fake and more informative.
6. **C** — Cost alone doesn't create credibility. A signal informs when its structure helps separate true claims from false ones.
7. **B** — The client understood you; they just don't believe you. Evidence or verification addresses the real problem, not more rhetoric.
8. **B** — One action can do both, but they're different ideas. Module 8 is about commitment.

Field assignment

This week: run the Credibility Test on three real claims, and make one of your own claims easier to verify.

1. Find three high-stakes claims made to you this week, at work, in a purchase, or at home. Run the eight questions on each and note whether the claim felt more or less credible than the evidence supported.
2. Catch one "Would they say this anyway?" statement and note what evidence would actually distinguish the cases.
3. Before making one truthful claim of your own (a status update, a commitment, a recommendation), add one piece of evidence the listener can check without relying on your word. Record how they responded.

Debrief questions for next week

1. Where did persuasiveness and credibility point in different directions, and which one did you act on?
2. When you added verification to your own claim, how did the other person's response change?
3. Where did you slide toward cynicism, and how did you return to "trust and still understand the game"?

The question to carry

What makes this message credible?

Module 8 — Commitment Changes the Game

The question: What makes this commitment real?

Book reading	Chapter 8, “Commitment Changes the Game”
Time	~45 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Commitment Test

Why this module matters

In the last module we asked what makes a message credible. Sometimes the answer is evidence, reputation, aligned incentives, or a signal that’s hard to fake. This module adds another answer: you can change the game so that doing something different later becomes difficult, costly, or impossible. That is commitment.

It leads to one of the most counterintuitive ideas in game theory: sometimes having fewer options can make you stronger. Usually more flexibility means more power. But when another person is deciding based partly on what they expect you to do later, your freedom to change your mind can become a weakness. “One hundred thousand is my absolute minimum” means little if you can still take ninety. Deliberately and visibly limiting that freedom can change what the other person rationally chooses to do now.

But a credible commitment can also lock you into a terrible outcome. So the goal isn’t “commit harder.” It’s to understand what a commitment does to the game, and to ask a sharper question when someone tells you “I have no choice”: what actually binds them?

Learning objectives

By the end of this module you will be able to:

1. Distinguish an intention, a promise, and a strategic commitment, and explain why only the third changes the future game.
2. Evaluate a threat, promise, boundary, deadline, or “final” position by asking whether carrying it out will still make sense when the time comes.
3. Identify what actually binds a person who claims “my hands are tied,” and separate real constraints from words.

4. Predict how a commitment will change the other player's best response, including the risk of impasse when both sides commit.
5. Distinguish strategic commitment from signaling, from stubbornness, and from the sunk-cost fallacy, and apply the ethical standard of truth, autonomy, dignity, and informed choice.

Lessons

Lesson 8.1 — Fewer options, more strength (10 min)

Teaching points

- “This is my absolute minimum” is only a statement. If you can still accept less later, the other side can test you, and when you move they learn your words didn't bind you.
- When a real, visible constraint makes a lower number impossible (a binding agreement, another party's enforceable claim, delegation to someone without authority), your future choice set changes. They no longer need to ask, “Will he cave?”
- Thomas Schelling made this paradox central to strategic commitment: a bargaining player can sometimes strengthen its position by openly making its own alternatives worse.
- A commitment influences behavior today by changing what another player expects will happen tomorrow. It isn't that you sound more serious; the future game has changed.
- The burning-the-bridge image illustrates it, but the lesson is not “burn every bridge.”

Example to use: The \$100,000 negotiation. I say \$100,000 is my floor, you offer \$85,000, I say no, you wait, and I come back at \$95,000. My “absolute minimum” wasn't absolute. Now suppose I genuinely cannot accept less because of a binding agreement. Pushing below \$100,000 may no longer make sense for you. I have fewer options, and a stronger position.

On-screen takeaway: Flexibility has value, but flexibility can also undermine credibility.

Lesson 8.2 — Intention, promise, commitment (8 min)

Teaching points

- Intention: I plan to do something. Promise: I tell you I plan to do something. Strategic commitment: I take an action that changes the incentives or options surrounding what I can do later.
- Signaling asks, “What does this action tell the other player about me?” Commitment asks, “How does this action change what I can or will do later?” One action can do both.
- A commitment usually has to be observable to affect the other player.

- Commitment devices work against your future self: the alarm across the room, the prepaid session, the tempting food removed from the house.

Example to use: “I’m definitely waking up at 5:00” is intention. “I promise I’ll meet you at the gym at 5:30” is a promise. Arranging for your training partner to pick you up at 5:15, with a prepaid session you lose if you skip, changes the game your future self will face.

On-screen takeaway: A hidden commitment constrains you. An observable commitment can constrain you and change them.

Lesson 8.3 — Threats, promises, and boundaries (10 min)

Teaching points

- Severity is not credibility. The question is: will carrying it out still make sense when the time comes?
- Threats (“If you do X, I will do Y”) and promises (“If you do X, I will reward you with Y”) share one problem: will they believe you’ll follow through? A contract, published criteria, or money already in an incentive pool makes the future action harder to reverse.
- Don’t threaten what you won’t do. Every unenforced threat, and every broken promise, teaches the other player that your words don’t predict your behavior.
- A boundary is not “I demand that you behave differently.” It is “If this occurs, here is what I will do,” and it matters only when your behavior makes it credible.

Example to use: A parent threatens six months of grounding for an uncleaned room. The child knows it’s never happening, because enforcing it would be miserable for the parent too. Compare “If the room isn’t cleaned by six, there’s no gaming tonight,” consistently enforced. Less dramatic, more strategic effect. Then the boundary: you say you’ll end the conversation if the disrespect continues, and you stay, and stay again. Calmly saying “We’ll continue this when we can talk respectfully” and leaving, consistently, changes the game.

On-screen takeaway: Use a consequence you can actually enforce.

Lesson 8.4 — “I can’t” versus “I don’t want to” (7 min)

Teaching points

- “I’d rather not” invites pressure, because preferences can change. A truthful “My authority stops at this number” is different, because your available actions have changed. This is why delegation can create commitment.
- People know this, so they sometimes pretend to be constrained. Ask what actually binds them: a contract, rule, third-party authority, legal restriction, public commitment, financial or reputational cost, irreversible investment? Or merely words?
- A credible deadline can reduce uncertainty. “This offer expires Friday at noon, and then the asset goes to another buyer,” if true, can make agreement easier. Commitment can provide clarity, not just pressure.

Example to use: Asked for another 10 percent discount, “I’d rather not” invites a push. Truthfully explaining that anything beyond your number requires board approval changes their decision problem, not because the board is magical but because your options really are different. Contrast “corporate won’t let me,” when corporate may be sitting ten feet away.

On-screen takeaway: When someone says “I have no choice,” ask: what actually binds them?

Lesson 8.5 — When commitment becomes a trap (10 min)

Teaching points

- A credible commitment can lock you into a terrible outcome. Two countries that each publicly commit (“If you cross this line, we retaliate” / “We will cross the line”) may have removed every path to de-escalation.
- Incompatible commitments create no deal. Evaluate a commitment not by how strongly it binds you, but by the game it creates.
- The less you know about the future, the more flexibility is worth. Public commitment creates discipline and rigidity at the same time.
- Commitment is not the sunk-cost fallacy, and it is not toughness. Refusing to adapt with no strategic benefit is just rigidity.
- A commitment is itself a move: they may concede, walk away, commit back, or stop negotiating.
- Ethics: deliberately creating a crisis so the other person “has no choice” can become coercive. Contracts, boundaries, transparent deadlines, and public promises help people coordinate and trust each other; destroying someone’s meaningful ability to choose does not.

Example to use: A business owner announces, “We will never raise our prices.” Customers love it, and years later costs double. Or a manager declares that anyone who makes a certain mistake is automatically fired, and then the best employee makes an honest mistake under unusual circumstances. Now the manager must break the commitment or carry out a punishment that no longer makes sense.

On-screen takeaway: A credible commitment can lock you into a terrible outcome.

Key terms

Term	Plain-language meaning
Intention	“I plan to do something.” A private plan that changes nothing about the game.

Term	Plain-language meaning
Promise	“I tell you that I plan to do something.” Communication that may or may not be backed by anything.
Strategic commitment	An action that changes the incentives or options surrounding what you can do later, so that doing something else becomes difficult, costly, or impossible.
Credible threat or promise	“If you do X, I will do Y” (or “reward you with Y”), where following through will still make sense, or be hard to avoid, when the moment arrives.
Observable commitment	A commitment the other player knows about or can credibly infer. Only these can change their present decision.
Delegation	Transferring authority to someone whose choices are more restricted, turning “I don’t want to” into a truthful “I can’t.”
Boundary	“If this occurs, here is what I will do.” A commitment that matters only when your behavior makes it credible.
Commitment device	An arrangement that constrains your future self’s options, such as a prepaid session.
Sunk-cost fallacy	Letting yesterday’s irreversible cost distort today’s decision when that cost should no longer matter to the choice.

Research spotlight

1. Schelling’s commitment insight (theory). The Nobel committee summarized one of Thomas Schelling’s key insights this way: a bargaining player can sometimes strengthen its position by openly making its own alternatives worse. The committee’s materials also noted applications to delegating decision power.

- What it suggests: Ask how your future options look to the other side. Changing them can change their best response today.
- What it does NOT prove: That commitment always creates power, that you should never leave an escape route, or that the most extreme commitment wins.

2. Commitment in the bargaining lab (Swope, Cadigan, and Schmitt). In laboratory bargaining experiments where players could credibly commit to a position, Kurtis Swope,

John Cadigan, and Pamela Schmitt found that commitment could increase the committer's payoff, particularly where others had incentives to hold out for better terms. But it could also increase the probability of bargaining failure and reduce overall efficiency. The same research notes that credible ultimatums can reduce uncertainty and remove incentives for strategic delay in some settings.

- What it suggests: Commitment may strengthen your position while making agreement harder: two sides of one mechanism.
- What it does NOT prove: That commitment helps in every negotiation, or that a “final offer” works because it sounds firm. These are specific, structured lab games.

Also cited in the chapter. A 2026 experiment on precommitment before sequential ultimatum bargaining found people used commitment strategically but less aggressively than theory predicts, with reciprocity and equal opportunity apparently mattering. A field experiment by Emily Breza, Martin Kanz, and Leora Klapper found factory workers were more likely to open a commitment savings account when the employer would observe it, and managers were more willing to invest in workers who held one. Neither shows that people use commitment exactly as the model predicts.

Bottom line from the book: The evidence does not support “commit harder and you’ll win.” It supports this: changing your future options can change another player’s current incentives, but commitment creates tradeoffs between credibility, bargaining power, flexibility, and the possibility of agreement.

The tool: The Commitment Test

Use this before making a threat, promise, deadline, boundary, or “final” position, and to evaluate one someone else makes to you.

The situation

The commitment I’m considering (or facing), and whether it’s an intention, promise, or strategic commitment:

Part 1 — What makes this commitment real?

What actually binds it? (contract / rule / third-party authority / legal restriction / public commitment / financial consequence / reputational cost / irreversible investment / merely words)

Can the other player observe or verify it?

Part 2 — What am I giving up, and why?

What future option am I giving up? (Be precise.)

What strategic benefit does the constraint create?

Part 3 — Threats and promises

“If you

, I will

.” Threat or promise?

When the moment arrives, will carrying it out still make sense? Is there a smaller consequence I would enforce every time?

Part 4 — Boundaries

“If this occurs, here is what I will do”:

What I will actually do, calmly and consistently, when it’s crossed:

Part 5 — Their response

How will their best response change (concede, walk away, commit back, find alternatives, retaliate, stop)?

Can this produce an impasse? What if they refuse to move?

Part 6 — The costs of commitment

What happens if new information appears? What flexibility will I wish I still had?

Is this strategy, or stubbornness and sunk cost?

Part 7 — Ethics

Does it preserve truth, autonomy, dignity, and informed choice?

Decision: Commit as planned / commit to something smaller or more enforceable / stay flexible. Why?

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small design studio and is renewing a contract with her largest client. The client's procurement lead, Marcus, asks for a 15 percent cut. Dana says, "Our rates are final; we can't go lower," then agrees to 5 percent off when he pushes. Now he wants another 5 percent, and Dana says this time she means it. Marcus mentions that his authority caps the budget at a set number and anything above it goes to his CFO. At home, Dana has told her teenage son three times this month that if he misses curfew again, he loses the car "for the rest of the year," and each time he kept the keys. She is also drafting a public post promising her studio will "never charge clients for revisions," though her biggest project next quarter has a scope that's hard to predict.

1. Why does Dana's "we can't go lower" carry little weight with Marcus, and what would make a price position real?
2. Analyze the curfew threat. Why isn't it working, and what would a more credible version look like?
3. Evaluate the planned "never charge for revisions" announcement using the costs of commitment.
4. How should Dana treat Marcus's claim about his budget cap?
5. What is one ethical line Dana should keep in mind as she firms up her positions?

Model answer

1. **The fake "can't."** Dana said "can't" and then moved, so Marcus learned her words don't bind her. "This time I mean it" is still only words. A real position needs something Marcus can see that makes reversal costly: published rates she applies to every client, or a partner who must approve exceptions. If nothing like that exists, an honest "I'd rather not, and here's why" beats a fake constraint.
2. **The curfew threat.** Severity isn't credibility. Taking the car for a year would be miserable for Dana too, so her son predicts she'll back down, and three unenforced threats have taught him her words don't predict her behavior. A credible version is smaller and enforced every time: "If you're home after curfew, you don't have the car next weekend."
3. **The revisions pledge.** Going public raises the cost of reversal, which builds credibility. But Dana's next big project has an unpredictable scope. If revisions balloon, she faces the commitment trap: break a public promise or absorb costs that no longer make sense. A narrower version ("two rounds of revisions included") keeps much of the credibility without locking in so much uncertainty.

4. **Marcus's cap.** Delegated authority can be a real commitment, and if it is, pushing above it accomplishes nothing. But "my hands are tied" can be fake. Dana should ask, respectfully, what actually binds him: Is the cap documented? Can it go to the CFO? Is there room on scope or timing instead?
5. **Ethics.** Dana can use transparent deadlines, published rates, and enforceable boundaries. She should not invent constraints or engineer a crisis that leaves Marcus no meaningful choice.

Exercise B — Your own game

Pick one live situation where you are about to state, or recently stated, a threat, promise, deadline, boundary, or "final" offer: a negotiation, a rule with your kids, a boundary with a relative, a public goal at work.

1. Write the statement exactly and classify it: intention, promise, or strategic commitment.
2. Work through all seven parts of the Commitment Test.
3. Rewrite the statement so it is something you will actually enforce, observable to the other person, and flexible enough to survive the most likely change in circumstances.
4. Flip it: take a "final" position or "I have no choice" someone recently gave you. What actually binds them? If the answer is "nothing except that they said so," what would you do differently?

Exercise C — Partner drill / role-play

Format: Pairs, 12–15 minutes, plus a 5-minute debrief. Works with an AI practice partner too.

Setup: A freelance photographer (Player A) and an event planner (Player B) are negotiating the fee for a wedding shoot. Run it in three short rounds of three to four minutes each.

Private instructions, Player A (photographer): You'd prefer \$3,000 and would accept \$2,400. Round 1: no real constraint. Round 2: a genuine, verifiable constraint: another booking on the same date at \$2,600, which you must confirm by tomorrow at noon. Share it if you choose. Round 3: decide whether to make a firm "final" statement and whether to back it with anything.

Private instructions, Player B (planner): Your client's budget is \$2,800, but you'd love to come in under it. Round 1: test A's stated minimums. Round 2: before reacting to any constraint, ask what actually binds it. Round 3: you may make your own firm commitment, such as truthfully stating your client's cap.

Rules: No inventing constraints that don't exist on your instruction card. Truth, autonomy, dignity, and informed choice apply.

Debrief questions: 1. In Round 1, did A's "minimum" hold? What did B learn when it moved? 2. In Round 2, did the observable deadline change B's best response? Did it reduce uncertainty, create pressure, or both? 3. In Round 3, did two commitments collide and threaten an impasse? Where did fairness or reciprocity shape the outcome?

Common mistakes

- **Confusing severity with credibility** — A huge threat you won't carry out does less than a small one you will. Use a consequence you can actually enforce.
- **Threatening what you won't do** — Every unenforced threat teaches the other player that your words don't predict your behavior. If you won't follow through, don't create the threat.
- **Keeping the commitment hidden** — A private constraint may shape your behavior, but it can't change their present decision unless they know about it or can credibly infer it.
- **Believing every "my hands are tied"** — Sometimes it's true, and sometimes corporate is ten feet away. Ask what actually binds them.
- **Judging a commitment by how strongly it binds you** — Two credible, incompatible commitments produce no deal, and under deep uncertainty flexibility may be worth more than credibility. Evaluate a commitment by the game it creates.
- **Mistaking stubbornness or sunk costs for commitment** — Refusing to adapt with no strategic benefit is rigidity, and "we've invested too much to quit" is the sunk-cost fallacy. Know when a past commitment has become a liability.

Quiz

1. In the book's framework, which of these is a strategic commitment rather than an intention or a promise?
 - A) "I'm definitely waking up at 5:00 tomorrow."
 - B) "I promise I'll meet you at the gym at 5:30."
 - C) Arranging for your partner to pick you up at 5:15 and prepaying a session you lose if you don't attend.
 - D) Telling yourself you're truly committed to getting in shape.
2. Schelling's paradox of commitment holds that:
 - A) More options always produce more bargaining power.
 - B) A player can sometimes strengthen its position by openly making its own alternatives worse.
 - C) The player who commits most aggressively always wins.
 - D) Commitments work best when kept secret.

3. A parent threatens six months of grounding over an uncleaned room. According to the chapter, the main problem is that:
 - A) The threat isn't severe enough.
 - B) Carrying it out would be costly for the parent too, so the child expects them to back down.
 - C) Threats never work with children.
 - D) The parent didn't state it publicly.
4. A seller tells you, "I'd love to go lower, but corporate won't let me." The most useful question to ask is:
 - A) "How committed are you, really?"
 - B) "What actually binds you: a rule, an authority, a contract, or just words?"
 - C) "Why don't you just quit?"
 - D) "Can you make it your final offer?"
5. Why does a commitment usually need to be observable?
 - A) Hidden commitments are unethical by definition.
 - B) A hidden commitment can constrain you, but it can't change the other player's present decision unless they know or can infer it.
 - C) Observable commitments are always irreversible.
 - D) Only public commitments affect your own future behavior.
6. In the Swope, Cadigan, and Schmitt bargaining experiments, credible commitment:
 - A) Increased the committer's payoff and also increased the probability of bargaining failure.
 - B) Had no effect on outcomes.
 - C) Always improved efficiency for both parties.
 - D) Only worked when both sides committed at once.
7. A team keeps funding a failing project because "we've already spent too much to stop." This is best described as:
 - A) A strategic commitment.
 - B) A credible promise.
 - C) The sunk-cost fallacy.
 - D) A boundary.
8. You tell a relative, "If you keep yelling, I'll end the conversation," and you've stayed the last three times. What does the chapter suggest?
 - A) Make the consequence more dramatic so it gets noticed.
 - B) Stop setting boundaries, since they don't work.
 - C) Calmly say "We'll continue this when we can talk respectfully" and leave, consistently, so your behavior makes the boundary credible.
 - D) Announce the boundary publicly to raise its cost.

Answer key

1. **C** — Only C changes the payoff your future self will face. A and D are intentions; B is a promise.
2. **B** — The Nobel committee summarized Schelling’s insight as strengthening your position by openly making your own alternatives worse. It does not claim extreme commitment always wins.
3. **B** — Severity is not credibility. The question is whether carrying out the threat will still make sense when the time comes.
4. **B** — A claimed constraint is only useful if it’s credible. Ask whether anything real binds them, or merely words.
5. **B** — A hidden commitment constrains you; an observable commitment can constrain you and change them.
6. **A** — Commitment could raise the committer’s payoff, but it could also make bargaining failure more likely and reduce efficiency. Those are two sides of the same mechanism.
7. **C** — Strategic commitment deliberately uses today’s action to change tomorrow’s incentives. The sunk-cost fallacy lets an unrecoverable past cost distort today’s decision.
8. **C** — A boundary becomes meaningful only when your behavior makes it credible. A more dramatic threat you won’t keep would make things worse.

Field assignment

This week: audit your commitments. Each day, notice one threat, promise, deadline, boundary, or “final” position, yours or someone else’s. Record the exact words, whether it was an intention, promise, or strategic commitment, what actually bound it (or “words only”), and whether it was followed through.

Then pick one boundary or consequence you have stated but not enforced. Replace it with a smaller, clear consequence you will actually carry out, state it once calmly, and follow through the next time the condition occurs. Record what happens.

Debrief questions for next week: 1. How many of the “final” positions or “I can’t” statements you heard were backed by something real? How could you tell? 2. When you enforced a smaller consequence, how did the other person’s behavior or expectations change, if at all? 3. Where did you notice a commitment, yours or someone else’s, that had become a trap because circumstances changed? What flexibility would have helped?

The question to carry**What makes this commitment real?**

What option am I giving up? How does giving it up change their decision? What flexibility will I wish I still had if circumstances change? What happens if they refuse to move?

Module 9 — The Shadow of the Future

The question: What changes because there is a next time?

Book reading	Chapter 9, “The Shadow of the Future”
Time	~55 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Repeated-Game Audit

Why this module matters

Imagine buying something from a street vendor in a city you will never visit again. Now imagine buying the same thing from the store around the corner from your house, a store you expect to visit every week for the next twenty years. Same transaction, same product, same price. But not necessarily the same game.

That is the shadow of the future. Once there is a tomorrow, today’s move doesn’t only determine today’s payoff. It changes what the other person expects from you, how they treat you next time, whether they trust you, and whether they cooperate with you again. Trust matters. Reputation matters. Retaliation matters. Repair matters. The future has entered the payoff.

This module trains a shift in horizon. Earlier in the course you learned to ask, “Then what?” Now you’ll ask, “What happens next time?” A choice can be excellent in a one-shot game and terrible in a repeated one. Most of the important parts of life (your marriage, your children, your closest friends, your team, your business partners) are repeated games. You are always teaching the people in them what to expect from you, whether you intend to or not.

Learning objectives

By the end of this module you will be able to:

1. Distinguish a one-time transaction from a repeated relationship, and estimate how long the shadow of the future is in a given interaction.
2. Explain why repeated interaction makes reciprocity and conditional cooperation strategically possible, without claiming it guarantees cooperation.
3. Identify what a specific response is teaching another person to expect next time, including unintended lessons.

4. Design responses that discourage exploitation while leaving a path back to cooperation after honest mistakes (boundaries plus repair).
5. Recognize when “thinking long term” has become a reason to tolerate endless losses, and when changing the rules or exiting is the long-term strategy.

Lessons

Lesson 9.1 — Tomorrow changes today (~12 min)

Teaching points

- In the one-shot Prisoner’s Dilemma, defection is individually attractive, so players can end up at defect/defect even though cooperate/cooperate would leave both better off.
- Change one thing: the players will meet again tomorrow, and the day after. The immediate payoff hasn’t changed. The continuation of the relationship has, and that can completely change the strategy.
- Today’s move can change the value of tomorrow’s game. A gain today comes bundled with its future consequences.
- Cooperation is not only a personality trait. It can also be a strategic response to the structure of the relationship, and that doesn’t make it fake.
- In a repeated game, today becomes history. Memory creates consequences.

Example to use: Two business partners have \$10,000 on the table. One discovers a loophole that would let them take an extra \$1,000 at the other’s expense. If it’s their last interaction, the choice looks like \$1,000 versus \$0. If they expect millions in business over the next ten years, that \$1,000 comes bundled with less trust, tighter contracts, fewer opportunities, retaliation, or the end of the partnership.

On-screen takeaway: Today’s move can change the value of tomorrow’s game.

Lesson 9.2 — Reciprocity, reputation, and trust (~12 min)

Teaching points

- Axelrod’s iterated Prisoner’s Dilemma tournaments: a simple strategy, Tit for Tat, won. It cooperated first, then copied the other player’s previous move. It let cooperation continue while still responding to exploitation.
- The lesson is not “Tit for Tat is the best strategy for life.” Its success depended on the environment. The broader lesson: once people expect a future together, reciprocity becomes strategically possible.
- Reputation is a prediction. It helps others answer, “What is this person likely to do next?” It can be wrong or unfair, but it shapes expectations.
- Trust lowers the cost of interaction. Demonstrated reliability lets contracts, deposits, verification, and monitoring fall away for both sides.

- The most valuable reputation is often simply: “This person is safe to play repeated games with.” Squeezing every round teaches the opposite: every interaction requires protection.

Example to use: The neighborhood restaurant versus the highway restaurant. Mistreat a local customer and they may never return, tell neighbors, and take their family with them. The highway restaurant won’t necessarily behave badly, but structurally the local one has a longer direct shadow of the future. Online ratings and seller histories extend that shadow to strangers: the platform has created memory.

On-screen takeaway: Reputation is accumulated information about future behavior.

Lesson 9.3 — What are you teaching them? (~11 min)

Teaching points

- Every important repeated interaction contains a learning process: what happens when I tell you the truth, disagree, say no, or make a mistake?
- Your words and your behavior can send different messages. When they conflict, behavior usually wins.
- Winning the round can lose the game. Sometimes strategic thinking requires resisting a short-term win.
- Judge some strategies across the relationship, not the incident: leaving \$500 on the table, enforcing an uncomfortable boundary, or admitting a mistake can cost today and pay for years.

Example to use: A manager tells employees, “Bring me bad news early.” An employee reports a serious problem, and the manager humiliates and blames them. Next month, nobody reports anything. The words said “report problems.” The behavior said “reporting problems is dangerous.” Pair it with the spouse who wins an argument by using something their partner once confided, and teaches them that vulnerability is dangerous.

On-screen takeaway: What am I teaching this person to expect from me next time?

Lesson 9.4 — Boundaries, noise, and repair (~12 min)

Teaching points

- The long game is not “always be nice.” Unconditional cooperation with someone who repeatedly exploits you can make exploitation profitable. Tit for Tat was cooperative but not passive.
- Boundaries can protect cooperation: “I want to keep cooperating with you, but this behavior changes what I will do next.” Conditional cooperation is often more durable than unconditional cooperation.

- But punishment creates its own problem. Real relationships contain noise: we mishear, forget, and misread motives. An accidental betrayal punished as if intentional can become a feud.
- Repair (apology, clarification, restitution, forgiveness, renewed commitment) is not just a soft skill. In a noisy repeated game, it stops errors from becoming permanent retaliation loops.
- Forgiveness is not forgetting. It means: do not allow every deviation to make future cooperation impossible.
- History changes meaning. The same action, a cancelled dinner, means “something must have come up” from someone reliable for ten years and “here we go again” from someone who cancelled the last six times.

Example to use: Two families who have hated each other for years. Each side describes a response to a response to something years earlier. The retaliation has become the game.

On-screen takeaway: Cooperation does not require defenselessness, and protection does not require a feud.

Lesson 9.5 — The limits of the long game (~8 min)

Teaching points

- The end of the game changes behavior. When someone knows Friday is their final day, some future consequences disappear. Values may still hold, but certain incentives have changed.
- Cooperation does not automatically emerge. The future creates possibilities; it doesn’t guarantee outcomes.
- Cooperation is not the same thing as goodness. Cartels and corrupt organizations cooperate too. Game theory tells us how patterns become sustainable; ethics tells us whether they’re worth sustaining: truth, autonomy, dignity, informed choice.
- Reputation can trap you (“I never back down,” “I’m always the nice one”). Ask: what reputation do I want to be bound by?
- Don’t use “the long term” to justify endless losses. If someone repeatedly exploits cooperation, change your strategy, change the rules, or leave the game. The shadow of the future matters only if the future game is worth playing.

Example to use: A business relationship that keeps producing broken promises, escalating monitoring, and conflict. Instead of optimizing each new round, ask, “Why am I still playing this repeated game?” The existence of a future does not obligate you to preserve the relationship.

On-screen takeaway: Sometimes exit is the long-term strategy.

Key terms

Term	Plain-language meaning
Shadow of the future	The influence that expected future interaction has on today's choices. The longer and more valuable the future, the longer the shadow.
One-shot game	An interaction both sides expect to happen only once, so today's move carries no future consequences between them.
Repeated game	The same strategic interaction played again over time, so today's behavior can shape tomorrow's treatment.
Tit for Tat	A strategy that cooperates first, then copies whatever the other player did last round. Simple and successful in Axelrod's tournaments, but not universally best.
Reciprocity	Rewarding cooperation and responding to exploitation over time; made strategically possible by a shared future.
Conditional cooperation	Cooperating as long as the other side does, with a credible response if they don't. Often more durable than unconditional cooperation.
Reputation	Accumulated information about how someone is likely to behave next; a form of strategic predictability.
Noise	Misunderstandings, errors, and imperfect observation that make it hard to tell mistakes from exploitation.
Repair	A route back to cooperation after a mistake: apology, clarification, restitution, forgiveness, renewed commitment.
End-game effect	The change in incentives when players believe there will be no next round.
Strategic uncertainty	Uncertainty about what the other player is going to do; cooperation tends to flourish when it is robust to it.

Research spotlight

1. Dal Bó's test of the shadow of the future. Pedro Dal Bó had participants play repeated Prisoner's Dilemma games in the laboratory while he varied the probability that another round would occur. A longer expected future significantly reduced opportunistic behavior, broadly in the direction repeated-game theory predicts.

What it suggests for real conversations: when future interaction becomes more likely, the future consequences of present behavior become more important, and that can make cooperation more attractive.

What it does not prove: it does not show that repeated interaction automatically produces cooperation. Later work by Dal Bó and Guillaume Fréchette found that cooperation can prevail as subjects gain experience, but the conditions were more demanding than cooperation simply being theoretically possible. Their survey of the experimental literature found high cooperation tends to emerge when cooperative strategies are robust to strategic uncertainty. These are laboratory games; they don't predict any particular person.

2. Noise and the value of leniency. Masaki Aoyagi, V. Bhaskar, and Guillaume Fréchette compared repeated Prisoner's Dilemma games under perfect monitoring, where players could see exactly what happened, and under imperfect (noisy) monitoring. Subjects sustained cooperation in all conditions, but under noisy monitoring they used strategies that were more complex and more lenient, which helped limit the damage from observational errors.

What it suggests: when misunderstandings are possible, a completely unforgiving strategy is fragile. Some tolerance and a path to repair can be strategically valuable.

What it does not prove: it is not evidence for "always forgive." Leniency helps in noisy settings, but exploitation is still possible, and a too-forgiving strategy can make exploitation profitable. Similarly, Axelrod's tournaments do not show that Tit for Tat is the best strategy everywhere. Its success depended on the mix of strategies, the payoffs, and how errors were handled.

The tool: The Repeated-Game Audit

Use this when an interaction is likely to continue. Don't evaluate it as though it ends today.

The interaction:

The move I'm tempted to make:

Step 1 — Is there a next round? Is there likely to be a next round? (yes / no / unsure)

How valuable is that future interaction to me? To them?

Is an endpoint visible (last day, final contract, move away)?

Step 2 — Transaction or relationship? Is this a transaction that ends when the exchange ends, or a relationship that creates another round?

Step 3 — Teaching and learning What is today's move teaching the other person to expect from me?

What am I learning about them?

What does their history tell me (reliable for years, or "here we go again")?

Step 4 — Reputation What reputation are we each building?

Is that reputation useful predictability, or exploitable predictability?

What reputation do I want to be bound by?

Step 5 — Incentives What behavior am I rewarding?

What behavior am I making costly?

Step 6 — Repair If a mistake occurs, is there a path back to cooperation?

How would I tell a mistake from exploitation?

Step 7 — The hundred-round test What happens if we keep playing this way for the next hundred rounds?

My move, revised:

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small marketing agency. For four years her biggest client has been a regional furniture retailer, and her main contact there is Marcus. Last month Marcus signed off on an ad with the wrong sale dates, and the client lost a weekend of sales. Dana's contract technically lets her bill for the rework, and her account lead, Priya, wants to bill the full \$4,000 and send Marcus an email quoting his approval line by line "so it never happens again." Meanwhile Dana also has a freelance designer, Theo, who has missed three of his last four deadlines, each time with an apology and a promise, and Dana has absorbed the overtime "to keep the relationship good." Theo just asked for a raise. Separately, Dana learned that Marcus is leaving the retailer in two months, and his replacement is unknown.

1. Which of these relationships has the longer shadow of the future, and does Marcus's departure shorten it?
2. What would Priya's approach win in this round, and what might it teach the client to expect next time?
3. What has Dana been teaching Theo to expect, and what does the chapter suggest she do?
4. How does noise apply to the Marcus situation, and how should Dana distinguish mistake from exploitation in each case?
5. Run the hundred-round test on Dana's current approach with Theo.

Model answer

1. **The shadow.** The relationship with the retailer has the long shadow. Marcus's departure shortens his personal horizon, but not the account's. How Dana handles this becomes history Marcus's replacement inherits. The game has an endpoint for one player, not for the relationship.
2. **Winning the round.** Priya's approach wins the round: \$4,000 and a documented record that the error wasn't the agency's. But it may teach the client that every mistake with this agency produces a bill and an itemized accusation, so every interaction requires protection. Dana could still protect herself (clarify how approvals work going forward, perhaps split the rework cost) without turning an honest error into a punishment.
3. **Theo.** Dana has been cooperating unconditionally with someone who repeatedly defects. She's been teaching Theo that missed deadlines cost him nothing, because she absorbs the cost. Thinking long term doesn't mean refusing to protect yourself. A conditional-cooperation boundary might sound like, "I want to keep working with you, but missed deadlines change what I can do next. Before we talk about rates, I need the next three projects on time." If the pattern continues, the long-run data are telling her something, and the long-term strategy may be to change the rules or leave the game.
4. **Noise.** Marcus's error looks like a one-time mistake inside four years of reliability: same action, different history. Treating it as exploitation risks starting a retaliation loop in which

the client reads the bill as hostility and responds in kind. Theo's case is different. One missed deadline might be noise; three of four, each followed by a new promise, is a pattern. Apologies without changed behavior are not repair.

5. **Hundred rounds.** If Dana keeps absorbing Theo's misses for a hundred rounds: more overtime, growing resentment, rising rates for unreliable work, and a likely eventual collapse that ends the relationship anyway, on worse terms. A clear boundary now gives it a better chance, or ends it before more value is lost.

Exercise B — Your own game

Choose a relationship you expect to continue (coworker, partner, child, client) and one recurring moment in it, such as how mistakes or disagreements are handled.

1. Write down the last three times this moment happened and how you responded.
2. Complete the Repeated-Game Audit for the next time it happens.
3. For each of your three past responses, finish this sentence: "From this, they probably learned that

."

4. Where do your words ("tell me the truth") and your history disagree?
5. Decide on one change: a behavior you'll start rewarding, a behavior you'll make more costly, or a repair route you'll offer. Write the exact sentence you might say.

Exercise C — Partner drill / role-play

Format: pairs, 15 minutes. (An AI practice partner can play Player B.)

Setup: Two roommates, or two coworkers who share a project, have been splitting tasks for a year. Last week Player B didn't do a task they had agreed to do, and Player A had to cover it.

Player A privately knows: You're annoyed. This is the second time in a month. You aren't sure whether B forgot or is starting to take advantage of you. You want the arrangement to continue.

Player B privately knows: You genuinely misread a message and thought A had swapped tasks with you. The first miss a few weeks ago was real carelessness on your part.

Round 1 (4 min): Player A responds as if this is a one-shot game: maximize being right today. Player B responds naturally.

Round 2 (4 min): Reset. Player A uses the audit (mistake or exploitation? what will this teach? is there a path back?) and aims to set a clear expectation without starting a feud. Player B responds naturally.

Debrief (5-7 min):

- In Round 1, what did B learn to expect from A? Did B start retaliating, even subtly?
- In Round 2, did A find a way to address the earlier, real miss without punishing the honest mistake?
- What would the hundred-round version of each conversation look like?

Common mistakes

- **Treating a relationship like a transaction.** Optimizing every round can make the game itself less valuable. Give the future a place in the decision.
- **Assuming “long term” means “always be nice.”** Unconditional cooperation can make exploitation profitable. Use conditional cooperation and credible boundaries.
- **Treating every apparent defection as intentional.** In a noisy world, rigid retaliation turns mistakes into feuds. Separate mistakes from exploitation, and build in a path to repair.
- **Confusing forgiveness with forgetting.** It means not letting every deviation make future cooperation impossible.
- **Using “the long term” to justify endless losses.** Repeated exploitation is data. Change your strategy, change the rules, or leave the game.
- **Treating cooperation as automatically good.** Cartels and corrupt groups cooperate too. Game theory explains what patterns can persist; ethics decides whether they deserve to.

Quiz

1. According to the chapter, what changes when a one-shot Prisoner’s Dilemma becomes a repeated game?
 - A) The immediate payoffs of cooperating and defecting change
 - B) The continuation of the relationship enters the calculation, so today’s move can affect future treatment
 - C) Defection becomes impossible
 - D) Players become more generous by nature
2. What did Pedro Dal Bó’s experiments vary, and what did they find?
 - A) The size of the payoffs; larger payoffs increased cooperation
 - B) Whether players could talk; communication eliminated defection
 - C) The probability of another round; a longer expected future significantly reduced opportunistic behavior
 - D) Players’ personalities; cooperative people always cooperated

3. What is the most accurate lesson from Axelrod's tournaments?
 - A) Tit for Tat is the universally best strategy for life
 - B) Complex strategies always beat simple ones
 - C) Once people expect a future together, reciprocity becomes strategically possible
 - D) Cooperation never pays against selfish players
4. A manager says "bring me bad news early" but berates the first person who does. Next month nobody reports a problem. What best explains this?
 - A) The employees are responding to the history of the game, not the manager's words
 - B) The employees lack integrity
 - C) The manager's instruction was unclear
 - D) Bad news is rare in well-run teams
5. Aoyagi, Bhaskar, and Fréchette found that under noisy monitoring, players who sustained cooperation tended to use strategies that were:
 - A) Simpler and harsher
 - B) Identical to those under perfect monitoring
 - C) Purely random
 - D) More complex and more lenient
6. Your friend has cancelled plans with you six times in a row, each time with a heartfelt apology and a promise to make it up. Following the chapter, what is the best response?
 - A) Keep accepting, because thinking long term means always cooperating
 - B) Treat the pattern as information; set a clear expectation or change the terms, and be willing to step back if it continues
 - C) Retaliate by cancelling on them six times
 - D) Ignore history and evaluate only the most recent cancellation
7. Why does the chapter say "cooperation is not the same thing as goodness"?
 - A) Because cooperation is always a form of manipulation
 - B) Because repeated games make cooperation impossible
 - C) Because cartels, corrupt organizations, and people deceiving a third party can sustain cooperation too
 - D) Because selfish motives make cooperation fake
8. A neighborhood restaurant has a longer direct shadow of the future than a highway restaurant mainly because:
 - A) Its customers are likely to return and to talk with neighbors who might come
 - B) Its owners are more ethical
 - C) It charges higher prices
 - D) Highway restaurants are not subject to reviews

Answer key

1. **B** — The immediate payoff hasn't necessarily changed; the continuation of the relationship has, and today's move can now shape trust, punishment, and future cooperation.
2. **C** — Dal Bó varied the probability of another round, and a longer expected future significantly reduced opportunistic behavior.
3. **C** — The chapter explicitly rejects "Tit for Tat is universally best". The broader lesson is that a shared future makes reciprocity possible.
4. **A** — The manager's words said "report problems," but the behavior said "reporting is dangerous." Today's response teaches tomorrow's expectations.
5. **D** — Under imperfect monitoring, players shifted to more complex and more lenient strategies, which helped limit damage from observational errors.
6. **B** — Don't use "the long term" to justify endless losses. Six repeats is a pattern, not noise; apology without changed behavior isn't repair.
7. **C** — Game theory explains which patterns are sustainable; ethics (truth, autonomy, dignity, informed choice) tells us whether they're worth sustaining.
8. **A** — Structurally, mistreating a local customer can cost repeat visits and word of mouth among neighbors.

Field assignment

This week, pick one ongoing relationship and pay attention to what you are teaching. Each day, note one moment where someone told you the truth, disagreed, made a mistake, or kept or broke their word. Record:

- What they did.
- How you responded.
- What your response likely taught them to expect next time.
- Whether it was a mistake or a pattern, and whether there was a path back to cooperation.

At least once this week, deliberately choose the move that makes tomorrow's game better over the one that gives you the most today. Note what happened.

Debrief questions for next week: 1. Where did your words and your behavior teach different lessons? 2. Did you find a relationship where you've been cooperating unconditionally with repeated exploitation, or punishing honest mistakes as if they were exploitation? What would conditional cooperation look like there? 3. Running the hundred-round test on your most important relationship this week, where does the current pattern lead?

The question to carry

| **What changes because there is a next time?**

Module 10 — Diagnose the Coordination Problem

The question: Are we actually in conflict—or are we failing to coordinate?

Book reading	Chapter 10, “Diagnose the Coordination Problem”
Time	~40 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Coordination Diagnostic

Why this module matters

Sometimes two people end up with a bad outcome even though neither of them wanted it. Not because one betrayed the other. Not because their interests are opposed. Not because someone failed to persuade. Sometimes the problem is much simpler: they failed to coordinate.

That distinction matters. Misdiagnose a coordination problem as a conflict, and you can spend enormous energy persuading someone who already wants essentially what you want. Misdiagnose a free-rider problem as coordination, and you can write the clearest plan in the world and watch it fail, because everyone knows exactly what to do and still has a reason not to do it.

This module trains one mental shift. Instead of asking “Why won’t they cooperate?”, you separate three possibilities: **conflict** (we genuinely want different outcomes), **coordination failure** (we could both benefit from compatible actions, but we aren’t aligned on what each of us will do), and **social dilemma** (we all like the collective outcome, but each of us has an incentive to let others carry the cost). They look remarkably similar from the outside. They are not the same game, and they do not have the same solution.

Learning objectives

By the end of this module you will be able to:

1. Distinguish conflict, coordination failure, and social dilemma using the chapter’s simple test.

2. Identify focal points and conventions that help a group select one workable option among several.
3. Explain the difference between everyone knowing something and everyone knowing that everyone knows it.
4. Locate where coordination is likely to break, especially at handoffs, ambiguous deadlines, and unowned tasks.
5. Match an intervention (shared plan, public commitment, clear owner, standard, deadline, focal point, or information about others' intentions) to the problem you have actually diagnosed.

Lessons

Lesson 10.1 — Not Every Disagreement Is a Conflict (~8 min)

Teaching points

- A coordination game is a situation in which players benefit when their actions fit together. In a pure coordination game, which outcome you land on may matter less than everyone choosing the same one.
- Driving on the right and on the left both work. The disaster is half the people choosing each.
- So the question is often not “Which option is objectively better?” but “Which option can we all reliably choose together?”
- Real life is often mixed. In the Battle of the Sexes structure (the name is dated; the structure is useful), two people want to spend the evening together but prefer different events. They need to coordinate, and they disagree over which equilibrium to coordinate on. Founders, co-parents, and departments live here.
- “We all want the same thing” doesn’t settle the strategic problem. Ask: where exactly do our interests align, and where do they diverge?

Example to use: The dinner problem. We both prefer eating together to eating alone; I slightly prefer the Italian restaurant, you slightly prefer the steakhouse. What matters most is ending up at the same restaurant. Contrast it with one available reservation that we both want for separate events: same word, “disagreement,” very different game.

On-screen takeaway: Not every disagreement is a conflict of interests.

Lesson 10.2 — Focal Points and Conventions (~9 min)

Teaching points

- Thomas Schelling asked: you must meet someone in New York City tomorrow and cannot communicate. Where do you go? People look for the obvious: a landmark, a central station, a conventional time. They ask, “What place will the other person think I will think is obvious?” These are **focal points**.

- Players don't need to prefer the focal point. They only need to believe the other person will recognize it too. Asked to match a positive whole number, many people pick 1. It isn't mathematically superior; it stands out.
- When several solutions would work, the easiest one for everyone to recognize may beat the theoretically perfect one. A slightly inferior standard used by everyone can be worth more than a superior standard used by half the group. Quality still matters; coordination itself also has value.
- Conventions save strategic energy. Without them, groups re-solve the same game daily: Who runs the meeting? Who follows up? A convention says, "This is what we do here."
- Fifty-fifty can be a coordination device as well as a fairness preference: it is extremely easy for both sides to anticipate.
- Sometimes the best choice is simply to pick something: alternate, rotate, flip a coin. The rule just needs to be a mutually understood way of selecting one equilibrium.
- Bad conventions can be stable too. Everyone may follow a rule only because everyone expects everyone else to.

Example to use: The project-management system. System A is slightly better; System B is already used by everyone. You could spend weeks proving A's features, but if half the organization stays on B, you've created a coordination failure.

On-screen takeaway: The easiest solution for everyone to recognize may beat the theoretically perfect one.

Lesson 10.3 — Strategic Uncertainty (~8 min)

Teaching points

- Ordinary uncertainty asks, "Will it rain?" **Strategic uncertainty** asks, "What are the other players going to do?"
- When switching alone is costly, everyone waits for everyone else. Management concludes "employees resist change," when perhaps each employee was willing to change, conditional on everyone else changing too.
- Sometimes people aren't unwilling to cooperate. They're unwilling to cooperate alone.
- A shared goal tells people where they want to end up. It doesn't tell them how everyone will get there.
- Cheap talk, weak as proof (Chapter 7), can be powerful as coordination. Different jobs.
- Before saying "We need better communication," ask what it should accomplish: information, credibility, persuasion, bargaining, or coordination? Two people at different airport terminals can fix their problem in thirty seconds. Two people arguing over a \$100,000 payment can talk for hours and only clarify the disagreement.

Example to use: Five departments each prefer the new reporting procedure if all five use it, but switching first means extra work, so each waits. When management publicly announces "Every department switches Monday" and everyone can see the migration plan,

the calculation changes, not because anyone was persuaded, but because uncertainty about the others dropped.

On-screen takeaway: Sometimes people aren't unwilling to cooperate. They're unwilling to cooperate alone.

Lesson 10.4 — Common Knowledge, Owners, and Handoffs (~8 min)

Teaching points

- There's a difference between everyone knowing something and everyone knowing that everyone knows it. The stronger condition is **common knowledge**.
- "We're moving the meeting to 3:00," said privately, leaves you wondering whether others heard. Said to the whole room, nothing about the meeting changes, but the mutual awareness does. Public announcements, dashboards, visible schedules, and rituals can create shared expectations, not just transmit information.
- "Everyone knows what to do" is often false. The question isn't "Have I said it?" but "Do we share the same expectation about what happens next?" Two people both hear "Thursday"; one means finished, the other means draft review.
- "Paul owns this by 3 PM" changes the game: Paul knows, everyone knows, and Paul knows everyone knows.
- Coordination often fails at the handoff. Inspect not just each person's performance but the connection between tasks.
- Roles, deadlines, agendas, and chains of command are coordination technology. Leadership often means creating a shared coordination structure ("You call them. Everyone meets back here at six") rather than persuading.

Example to use: The project nobody owned. Everyone agrees to follow up with a customer today. Tomorrow: "Did anyone call them?" Silence. Sales thought customer service would; customer service thought sales owned it; the owner thought the sales manager had assigned it. Before judging character, ask: was there actually a mutually understood assignment?

On-screen takeaway: Sometimes coordination requires making expectations mutually visible.

Lesson 10.5 — Coordination or Social Dilemma? (~7 min)

Teaching points

- Shared goals don't always mean coordination. Four roommates want a clean apartment, but each would rather someone else cleaned. That's a free-rider, public-goods problem: a **social dilemma**.
- The simple test: "If I knew for certain that everyone else was going to cooperate, would I want to cooperate too?" Yes, absolutely: likely coordination. "I'd still be better off letting everyone else do the work": social dilemma.

- The fix changes with the diagnosis. Coordination calls for clearer expectations, shared information, a focal point, a standard. Free riding may require changing incentives, monitoring, reputation, repeated interaction, commitments, or the rules.
- Focal points have limits. When players strongly disagree about which outcome they prefer, salience may not be enough. A focal point can solve uncertainty; it cannot erase incompatible interests.
- Coordination is morally neutral; groups can coordinate around harm. Keep the foundation: truth, autonomy, dignity, informed choice.

Example to use: If three roommates will definitely clean, the fourth may think, “Great. Then I don’t have to.” They coordinated just fine. The individual incentive is the problem.

On-screen takeaway: Diagnose first. The same symptom can come from three different games.

Key terms

Term	Plain-language meaning
Coordination game	A situation in which players benefit when their actions fit together.
Coordination failure	We could both benefit from compatible actions, but we aren’t aligned on what each of us will do.
Conflict	We genuinely want different, incompatible outcomes.
Social dilemma / free-rider problem	We all like the collective outcome, but each of us has an incentive to let others carry more of the cost.
Mixed-motive game	Players want to coordinate but disagree about which outcome to coordinate on (e.g., Battle of the Sexes).
Focal point	An option that stands out so each player expects the other to recognize it too.
Convention	A shared rule that solves a repeated coordination problem: “This is what we do here.”
Strategic uncertainty	Uncertainty about what the other players will do.
Common knowledge	Everyone knows something, and everyone knows that everyone knows it.

Research spotlight

1. Focal points: Schelling; Mehta, Starmer, and Sugden

What was done. Schelling posed matching puzzles without communication, such as choosing a meeting place in New York or picking the same positive integer. Judith Mehta, Chris Starmer, and Robert Sugden later tested focal points experimentally in pure coordination games, where any option works as long as both players match.

What was found. People coordinated far above random chance by using shared culture, context, and salience; the Nobel committee specifically highlighted this part of Schelling's work. Mehta, Starmer, and Sugden found evidence that people use salience rather than choosing randomly among equivalent options.

What it suggests. When a group has several workable options, the one everyone recognizes most easily (the existing standard, the obvious default) has real strategic value.

What it does NOT prove. It doesn't show the most salient option is the best one; it may just be easiest to coordinate around. And salience has limits: experiments with asymmetric payoffs, where players prefer different outcomes, find lower coordination than pure coordination settings, and later studies (such as Parravano and Poulsen; Rojo Arjona, Sitzia, and Zheng) found that higher stakes or stronger salience did not reliably help once conflicts of interest were built in.

2. Communication as coordination: Cooper et al.; the Senegal cooperatives

What was done. Russell Cooper, Douglas DeJong, Robert Forsythe, and Thomas Ross ran lab coordination games where players could send nonbinding messages first. In the field, Aflagah, Bernard, and Viceisza worked with 79 groundnut cooperatives in Senegal, running coordination games with 839 farmers and then a natural field experiment with roughly 2,800 members. They asked members how much they intended to sell through the cooperative and revealed the group's aggregate intentions.

What was found. Nonbinding communication often substantially improved coordination: one-way communication increased selection of a Pareto-superior equilibrium in some conditions, and in another type of game two-way communication consistently produced the Pareto-dominant one. In Senegal, treated farmers contributed more than twice as much toward collective selling as the control group on average, and the effect was larger in bigger groups. No new incentive was added; what changed was information about what everyone else intended to do.

What it suggests. When people already want an outcome but doubt others will act, making intentions visible can move behavior.

What it does NOT prove. Communication did not eliminate every coordination failure. Neither result supports "talk more and every group problem disappears." Communication helps when it resolves the uncertainty actually blocking coordination; it won't dissolve a

real conflict or a free-rider incentive. And shared goals alone don't guarantee action: the farmers already saw the benefits and still held back.

The tool: The Coordination Diagnostic

Use this when a group or relationship keeps producing an outcome nobody seems to want.

The bad outcome that keeps happening (no character judgments):

1. Do we actually want incompatible things? If yes, you may have a conflict or bargaining problem. If partly, note: Where we align:

Where we diverge:

2. Would I take the desired action if I knew the others would too? My honest answer:

What I suspect others would honestly answer:

(Yes → likely coordination. "I'd rather others bear the cost" → likely social dilemma; clarity alone may not fix it.)

3. My diagnosis (circle one): Conflict / Coordination failure / Social dilemma / Mixed
Why:

4. What am I uncertain about? (their actions, intentions, knowledge, commitment)

5. What are they uncertain about? (Your intentions may be part of the problem.)

6. What expectation needs to become mutually clear? Who acts?

When?

According to what standard?

Where is the handoff?

7. Is there a focal point or convention we can coordinate around?

8. Does everyone know that everyone knows? Would making it public change the game?

9. Choose the intervention:

- ☐ A shared plan
- ☐ A public commitment
- ☐ A clear owner
- ☐ A standard
- ☐ A deadline
- ☐ A focal point
- ☐ Information about what everyone else intends to do

My one sentence that makes the next move obvious to everyone:

Ethics check: Does this respect truth, autonomy, dignity, and informed choice for everyone affected?

Practice exercises

Exercise A — Analyze the scenario

Dana runs a twelve-person marketing team at a regional credit union. Three months ago, leadership approved a shared content calendar to replace the team's spreadsheets and email threads. Dana's kickoff email said, "Everyone should start using the new calendar when you can." Today only two people use it. The others say they like the tool but keep their spreadsheets "because that's where everyone else's dates are." Last week two staff members scheduled competing promotional emails for the same morning. A member newsletter also went out without its compliance review: the writer assumed the designer had sent it for sign-off, and the designer assumed the writer had. Dana's director says, "Your team doesn't take ownership and resists change. Motivate them." Dana is drafting a speech about accountability.

1. Is the calendar problem primarily conflict, coordination failure, or social dilemma? Apply the simple test.
2. What is blocking adoption, and why didn't the kickoff email solve it?

3. Where exactly did the newsletter process break?
4. Is a motivational speech the right intervention? What would you do instead?
5. Could any part of this turn out not to be coordination?

Model answer

1. **Mostly coordination failure.** “If I knew everyone else would use the calendar, would I?” The team’s own words (“that’s where everyone else’s dates are”) suggest yes. They aren’t opposed to the tool; they’re unwilling to switch alone, because a calendar one person uses is worse than spreadsheets everyone checks.
2. **Strategic uncertainty.** “When you can” gave no shared date, no plan, and no visible sign anyone else was switching. Each person read it privately and couldn’t tell whether the others would act, so everyone waited. From outside, that looks like resistance to change.
3. **The handoff, with no clear owner.** Nobody opposed the review. Each held a reasonable but incompatible expectation about who sends it. The individual tasks were done; the connection between them failed.
4. **No, motivation isn’t the bottleneck.** Better: announce in a team meeting, so everyone sees everyone hear it, “Starting Monday the 3rd, the shared calendar is the only calendar; spreadsheets are retired.” Name an owner for keeping it current, define what “scheduled” means, and review the calendar briefly at each weekly meeting so people can see others using it. The calendar becomes the focal point, which also prevents colliding sends. For the newsletter: “The writer submits for compliance review; nothing is released without sign-off.”
5. **Possibly.** If, after expectations are public and clear, some people still skip the calendar because updating it is extra work that mostly helps others, that looks like a social dilemma. Then clarity may not be enough; Dana may need to build updating into the workflow or make it visible. Diagnose first, and re-diagnose if the fix doesn’t hold.

Exercise B — Your own game

Pick a live situation where people “can’t seem to get on the same page”: a household chore, a team handoff, a family schedule, a volunteer group.

1. Describe the bad outcome in one sentence, without blaming anyone’s character.
2. Work through the Coordination Diagnostic.
3. On step 2, answer honestly for yourself first. If your answer is “I’d rather someone else did it,” say so; it changes the diagnosis.
4. Write the one sentence that makes the next move obvious, in the style of “Paul owns this by 3 PM.”
5. Decide how you’ll make it mutually visible: in front of everyone, in a shared document, on a visible schedule.

6. If it's something you renegotiate every week, draft a convention everyone can know before the next conflict arises.

Exercise C — Partner drill / role-play

Format: Pairs or small groups, 12–15 minutes; an AI practice partner can play Player B.

Setup: A shared office kitchen keeps getting left messy. Player A facilitates a short conversation to fix it. Player B privately draws one card:

- **Card 1 — Coordination:** “I’d happily clean on a set day, but there’s no schedule and I never know if someone already did it.”
- **Card 2 — Social dilemma:** “If others clean, I’ll enjoy it and not bother.”
- **Card 3 — Mixed:** “I support a rotation, but I want Fridays excluded since I’m remote then; others want everyone equal on every day.”

Rules: Player A has 7 minutes to diagnose the card using questions, not accusations, including some version of the simple test and at least one question about what B is uncertain about. Player B answers honestly in character without naming the card. A then proposes one intervention.

Debrief (5 minutes):

- Which card did A guess, and what evidence led there?
- Did the intervention fit? (Card 1: schedule, owner, public visibility. Card 2: something that changes incentives or accountability. Card 3: a fair selection rule among options people disagree about.)
- How would the wrong card’s intervention have failed?
- Swap roles and draw again.

Common mistakes

- **Treating every disagreement as conflict** — You end up persuading someone who already wants what you want. First ask whether you actually want incompatible things.
- **Calling coordination failure a character flaw** — Before “Nobody takes initiative,” ask whether there was a mutually understood assignment, and whether people were simply unwilling to act alone.
- **Assuming a shared goal settles it** — A shared goal tells people where they want to end up, not how everyone gets there. And it can still hide a free-rider problem; run the simple test.
- **Saying “everyone knows what to do”** — Ask instead, “Do we actually share the same expectation about what happens next?”

- **Telling people privately what needs to be public** — Private messages leave each person wondering about the others. Say it in front of everyone when mutual awareness matters.
- **Over-optimizing among acceptable options** — When failing to coordinate costs more than the difference among options, stop optimizing and coordinate.
- **Prescribing “better communication” without a diagnosis** — Ask what communication must accomplish. Talking won’t dissolve a genuine conflict.
- **Assuming a stable convention is a good one** — Everyone may follow it only because everyone expects everyone else to. Sometimes the system needs redesigning (Module 11).

Quiz

- Two friends prefer eating together to eating alone; one slightly prefers Italian, the other steak. What matters most?
 - The person with the stronger preference gets their choice
 - Ending up at the same restaurant
 - Finding the objectively better restaurant
 - Proving whose preference is more reasonable
- Which question best separates coordination from a social dilemma?
 - “Do we all want the same final outcome?”
 - “Who has more power in this group?”
 - “If I knew for certain everyone else would cooperate, would I want to cooperate too?”
 - “Have we communicated clearly enough?”
- Asked to match another person’s positive whole number, many choose 1. Why does this work?
 - 1 is mathematically optimal for matching
 - It stands out, and you can imagine the other person thinking the same
 - People are told 1 is the default
 - Smaller numbers pay more
- Employees all like a new system if everyone uses it, but switching alone makes work harder, so nobody switches. The barrier is best described as:
 - Low motivation
 - A genuine conflict of interests
 - Strategic uncertainty
 - A free-rider problem

5. What does telling five employees together “You five are responsible” add over telling each privately?
 - A) A financial incentive
 - B) Mutual awareness: each knows the others know who is responsible
 - C) Stronger arguments for the change
 - D) Nothing; the information is identical
6. Which claim do the Cooper et al. and Senegal studies support?
 - A) Talking more makes every group problem disappear
 - B) Nonbinding messages are useless because they aren’t credible
 - C) Communication can substantially improve coordination when it resolves the blocking uncertainty, but it doesn’t eliminate every failure
 - D) Revealing intentions works only in small groups
7. Two co-founders both want the company to succeed but each wants their own strategy adopted. This is:
 - A) Pure coordination
 - B) A mixed-motive game with coordination and conflict
 - C) A pure free-rider problem
 - D) Not a strategic problem
8. A team agrees “someone should call the customer today,” and nobody does; sales and service each assumed the other would. The best intervention is:
 - A) A talk about taking initiative
 - B) A penalty for everyone
 - C) “Jordan owns the call by 3 PM,” said in front of the team
 - D) More debate about customer ownership in general

Answer key

1. **B** — It’s a coordination problem: going to different restaurants is a loss for both, so matching matters more than first choices.
2. **C** — The chapter’s simple test. A can’t distinguish them; the free-riding roommates also share the goal.
3. **B** — 1 is a focal point. Players only need to believe the other will recognize it too.
4. **C** — People are willing to change conditional on others changing; the barrier is uncertainty about others, not motivation.
5. **B** — The public statement makes the expectation mutually visible, moving the group toward common knowledge.
6. **C** — Both show communication can be powerful for coordination (the Senegal effect was larger in bigger groups), but neither supports “talk more and everything is fixed.”
7. **B** — Like the Battle of the Sexes: aligned on coordinating, in conflict over which outcome.

8. **C** — Nobody opposed the call; there was no mutually understood assignment. A clear owner, deadline, and public statement remove the ambiguity.

Field assignment

This week, pick one recurring friction point you've been treating as a conflict or a motivation problem: a task that falls through the cracks, a failing handoff, a schedule you renegotiate every week.

1. Run the Coordination Diagnostic before saying anything, and commit to a diagnosis.
2. If it's coordination, make one expectation mutually clear (owner, time, standard) where everyone involved can see everyone else receive it, or propose a simple convention.
3. If it's a social dilemma or conflict, don't force a coordination fix; note what kind of intervention it really needs.
4. Record what you said, where, what happened over the following days, and whether the problem recurred.

Debrief questions for next week:

- Did your initial assumption match your final diagnosis? What changed your mind?
- Did a public, explicit expectation produce a different result than past private reminders?
- If the problem returned despite clear, mutually known expectations, what does that suggest: a hidden incentive, a real conflict, or a game that needs redesigning?

The question to carry

Are we actually in conflict—or are we failing to coordinate? If it's coordination: What expectation needs to become mutually clear? Make the next move obvious to everyone.

Module 11 — Change the Game

The question: Do I need a better move—or a better game?

Book reading	Chapter 11, “Change the Game”
Time	~55 min of video lessons · ~60 min of workbook · one week of field practice
Tool you’ll build	The Game-Change Audit

Why this module matters

For ten modules you have learned to see more clearly inside a strategic interaction. Those questions help you make better moves. But eventually you reach a more powerful question: do you need a better move, or a better game? You can become extraordinarily skilled at playing a game that should never have been designed that way in the first place.

Picture a sales organization where management keeps saying, “We want everyone to collaborate,” while publicly ranking salespeople every month and paying bonuses only to the top performers. People guard information, management holds a teamwork workshop, and nothing changes, because the game is saying something different from the speech. Management says “cooperate.” The incentive system says “your coworkers are your competitors.” At that point the strategic question is no longer “How can we persuade these people to collaborate?” It is “What behavior does our system actually reward?”

This module shifts you from fixing people to examining systems, and from reacting to bad outcomes to redesigning the structure producing them. Because designers usually hold more power than players, every lever here is tied back to the book’s foundation: truth, autonomy, dignity, and informed choice.

Learning objectives

By the end of this module you will be able to:

1. **Distinguish** a “better move” problem from a “better game” problem by asking what behavior the current structure actually rewards, rather than what it claims to reward.
2. **Explain** mechanism design and incentive compatibility in plain language, using the “one cuts, the other chooses” example.
3. **Predict** how a metric or incentive could be gamed, and how effort may shift away from valuable work that is harder to measure.

4. **Identify** the design levers in a recurring problem: rules, payoffs, information, timing, options, authority, players, observability, defaults, feedback, friction, and exit.
5. **Design** a small, ethical, testable change to a game you are part of, and plan how you will observe and update it.

Lessons

Lesson 11.1 — From playing the game to designing it (~12 min)

Teaching points

- Most of game theory asks: given the game, what will the players do? Mechanism design reverses the question: given the outcome we want, how should we design the game?
- Leonid Hurwicz, Eric Maskin, and Roger Myerson received the 2007 Nobel Prize in Economic Sciences for laying the foundations of mechanism design. You don't need the mathematics; the mental shift is what matters.
- Don't fix the player before you inspect the game. We tend to explain behavior through character ("they're lazy," "they're dishonest," "they're gaming it"). Strategic thinking adds one more question: what does the game make attractive?
- Behavior emerges from the interaction between people and structure. Sometimes perfectly ordinary people behave differently when the rules change.
- Incentives speak louder than mission statements. Say "tell us the truth" but retaliate against bad news, and the reward wins. Incentives don't explain everything; people have values. But if you repeatedly pay for a behavior, expect more of it.

Example to use: The cake. Two children need to divide a cake and each wants the larger piece. Instead of dividing it perfectly and refereeing the complaints, the parent changes the mechanism: one child cuts, the other chooses first. The cutter now has every reason to make the pieces equal. The parent didn't change anyone's personality or give a lecture on fairness; the rules made the children's self-interest produce the fair outcome.

On-screen takeaway: Instead of constantly fighting incentives, sometimes redesign them.

Lesson 11.2 — Incentives work, which is exactly why they are dangerous (~12 min)

Teaching points

- Changing an incentive can change behavior substantially (see the Pakistan study in the Research spotlight).
- Designing incentives is dangerous precisely because incentives work. Once a measurement is tied to money, status, or punishment, people optimize against the scoreboard, and the measure may stop representing what you cared about.
- The metric is not the goal. You want learning; you measure test scores. You want a healthy relationship; you measure who apologized first.

- The principal-agent problem: an agent acting for a principal has somewhat different incentives and often knows more about their own actions.
- Multitask incentives: strongly reward the easiest task to measure and effort can drift away from valuable work that is harder to see.
- The key question: “What exactly am I rewarding?” Not what you intend to reward, but what behavior actually improves the metric.

Example to use: The customer-service department. Management says solving the customer’s problem is the top priority, but representatives are evaluated mainly on average call time. A complicated customer becomes strategically expensive, so people rush, transfer calls, and end conversations early. Then management asks why nobody takes ownership. That isn’t a communication problem; it’s a game-design problem.

On-screen takeaway: How could a smart player maximize this payoff without giving me what I really want?

Lesson 11.3 — The levers: how a game can be changed (~14 min)

Teaching points

- **Rules:** a clear sales-to-customer-success handoff removes one source of strategic ambiguity.
- **Payoffs:** change the compensation, the payment schedule, or add a refundable deposit. Not every behavior should be bought, but payoffs matter.
- **Information:** who knows what, and when. Dashboards, transparent criteria, independent verification.
- **Timing:** sequential instead of simultaneous moves, immediate instead of annual feedback, staged payments. Sometimes the problem is the sequence.
- **Options:** trial periods, phased implementation, partial agreements. New options can alter the equilibrium.
- **Who decides and players:** assigned decision rights, an appraiser, a mediator. Who gets to decide is part of the game.
- **Observability, defaults, feedback:** what happens when nobody chooses is itself a rule; shorter feedback makes a game easier to learn.
- **Removing an incentive:** sometimes simplify, reduce measurement, and restore discretion.
- **Exit:** when you can’t redesign a destructive game, you can stop playing, but exit has payoffs too.
- **Good rules can reduce the need for strategy** by making the desired behavior the easy behavior.

Example to use: A couple stops renegotiating chores every Saturday and creates a system: you handle these, I handle those, we rotate this one. A consultant facing scope

creep introduces a defined scope, a change-request process, and a fee for extra work, protecting the relationship from resentment. And the person who says “I need more discipline” turns off notifications instead of fighting a willpower contest every five minutes.

On-screen takeaway: Sometimes the highest form of strategy is removing the need for strategy.

Lesson 11.4 — Ethical design, humility, and the system that talks back (~12 min)

Teaching points

- Designers are players too: set rule, players respond, learn, adapt, system evolves. Ask “then what?”
- Design for the behavior people are likely to have. Don’t require saints; don’t design for monsters.
- Don’t confuse control with design: design, observe, test, update, adapt. When results surprise you, ask “What did I misunderstand about the game?” Start with the smallest change that tests your hypothesis.
- Friction is part of the game. The fact that you can make something strategically difficult does not mean you should. A dark pattern is still game design, serving the wrong ethical objective.
- Ethical design asks: Does the player understand the game? Do they retain meaningful choice? Are important facts hidden? Does it treat people as decision-makers? Sometimes rules should protect the weaker player, as disclosure and inspection rights do in the used-car problem.

Example to use: Boston school choice. Boston’s old assignment system let families benefit from strategically ranking schools instead of listing their true preferences. Sophisticated parents could game it; parents who simply reported what they wanted could be disadvantaged. In 2005 Boston replaced it with a deferred-acceptance mechanism designed to remove much of that incentive. You could tell parents “don’t game the system,” or you could ask “why does the system make gaming strategically useful?” The second question is mechanism design.

On-screen takeaway: Create a game in which good outcomes emerge without violating the people playing it.

Key terms

Term	Plain-language meaning
Mechanism design	Starting from the outcome you want and asking what rules would make that outcome more likely, instead of taking the rules as given.

Term	Plain-language meaning
Incentive compatibility	A design in which following the desired strategy lines up with each participant's own incentives, so people don't have to fight their self-interest to do the right thing.
Strategy-proof mechanism	A system designed so that truthfully reporting your preferences is a good strategy, reducing the need to outsmart the rules.
Metric vs. goal	The measurement (test scores, call time, clicks) is not necessarily the outcome you care about (learning, solved problems, good content). Once payoffs attach to a metric, people may play the metric.
Principal-agent problem	A principal wants one thing; the agent acting on their behalf has somewhat different incentives and often knows more about their own actions.
Multitask incentives	When a job has several valuable tasks, strongly rewarding the most measurable one can pull effort away from the others.
Default	What happens when nobody chooses. It is itself a rule, and it shapes outcomes.
Friction	How easy or hard each choice is made to be. Friction is part of the design and carries ethical weight.
Exit	Stopping play when the game is destructive and you can't redesign it. A strategic choice with its own payoffs and costs.

Research spotlight

1. Property-tax inspectors in Punjab, Pakistan (Khan, Khwaja, and Olken)

What was done: In a field experiment, 525 property-tax inspectors took part in a system where performance rankings determined the order in which inspectors could choose their future job postings. Instead of telling bureaucrats to work harder, the design connected performance to something many employees actually valued: where they would be assigned to work.

What was found: The mechanism increased annual tax-revenue growth by roughly 30 to 41 percent. The research dossier notes that inspectors the researchers' model predicted would face stronger incentives under the new system improved more.

What it suggests for real conversations: Before giving another speech, ask what people genuinely value and whether the system connects the behavior you want to it.

What it does NOT prove: Not that any incentive redesign will improve outcomes, or that this design transfers unchanged to your team or family. Nor that responses will take the form the designer intended.

2. Multitask incentives: Holmström and Milgrom's theory, illustrated by the Kenyan teacher-incentive study (Glewwe, Ilias, and Kremer)

What was done: Holmström and Milgrom's multitask principal-agent analysis formalized a problem: strong incentives on one measurable activity can distort effort away from other valuable but poorly measured activities. Glewwe, Ilias, and Kremer studied a program in Kenyan primary schools that rewarded teachers based on students' test performance.

What was found: Scores on the incentivized exams improved. But researchers found little evidence that teacher attendance, homework assignment, or classroom pedagogy improved. Teachers increased test-preparation sessions, gains were concentrated on the exams linked to the incentive rather than unrelated tests, and the improvements did not persist as broader long-run learning gains.

What it suggests for real conversations: Whenever you attach a payoff to a number (calls made, bugs fixed, chores done), ask what behavior actually moves it, and what harder-to-measure work might be quietly abandoned.

What it does NOT prove: Not that incentives "don't work." The chapter's dossier notes teacher-incentive experiments in rural India and Tanzania with different results; in Tanzania a multiple-threshold design produced greater learning gains at lower cost than a rank-based one. The design of the incentive changes the game. Nor is there always one correct rule: goals like fairness, simplicity, and strategy-proofness can trade off.

The tool: The Game-Change Audit

Use this when the same bad outcome keeps returning, after better arguments, better wording, more motivation, and more pressure haven't fixed it.

Step 1 — Name the recurring problem

The outcome that keeps returning:

What I (or we) have already tried, inside the current game:

My current character explanation (“they’re lazy / selfish / gaming it”):

Step 2 — What does the game actually reward?

What behavior is the current game rewarding? (Not what the policy claims to reward. What actually pays?)

What behavior is it punishing? Could the desirable action currently be costly?

What metric are people optimizing? Is it actually the outcome we care about?

What does my stated value say, and what does the reward say?

Step 3 — Inspect the levers

Lever	How it works now	One possible change
Rules: what players can and can't do		
Payoffs: what pays, what costs		
Information: who knows what, when		
Who decides		
Timing: who moves when		
Options: what choices exist		
Players: could a third party help?		
Observability and feedback		
Default: what happens if nobody acts		

What is currently easy?

What is currently hard? (Friction is part of the design.)

Is there an incentive or rule I should remove rather than add?

Step 4 — Stress-test the redesign

How could a smart player game this system, including my new version of it?

What unintended behavior might the new rule create? (Think one response further: then what?)

What harder-to-measure work might lose effort if I reward the measurable part?

Step 5 — Ethics check

Does the player understand the game?

Do they retain meaningful choice?

Are important facts hidden, or does the system exploit confusion?

Does it treat people as decision-makers rather than objects to be manipulated?

Step 6 — Smallest testable change

What specific mechanism appears to be producing the unwanted behavior?

The smallest change capable of testing that hypothesis:

What I'll observe, and when:

If exit is on the table: what would I lose, and what future would I give up?

The central question: Can I redesign the game so that the behavior I want becomes a more natural response to the structure itself?

Practice exercises

Exercise A — Analyze the scenario

Dana manages a twelve-person software support team. A leaderboard ranks engineers by “tickets closed per week,” and the top three each quarter get a bonus and first choice of vacation weeks. Dana keeps saying customer satisfaction and knowledge-sharing matter most, and has run two mentoring workshops. Yet senior engineers rarely help juniors, complex tickets sit untouched until someone is assigned them, and customers increasingly reopen tickets marked “resolved” within days. Reopened tickets are counted as new tickets. Escalations to engineering require Dana’s personal sign-off, and she is often in meetings until late afternoon. Last month, two engineers were caught splitting single customer issues into several tickets. Dana’s boss has suggested firing them and “hiring people with better values.”

1. What behavior is Dana’s current game rewarding, and what is it punishing?
2. What metric is the team optimizing, and how does it differ from the goal?
3. Which design levers could Dana change? Name at least three.
4. How could a smart player game your proposed redesign?
5. Is “hire people with better values” the right diagnosis? What would you test first?

Model answer

1. **Rewarded and punished.** The game rewards closing many tickets quickly and not helping others, since mentoring or untangling a hard ticket lowers your count while coworkers are your rivals. It punishes complex work, mentoring, and thorough fixes. The workshops are speeches; the leaderboard is the game.
2. **Metric versus goal.** The metric is tickets closed; the goal is solved problems and a stronger team. Because reopened tickets count as new, a shallow fix earns credit twice, and ticket-splitting is people playing the metric. It is also a multitask problem: mentoring and durable fixes are hard to measure, so effort drains away from them.
3. **Levers.** Payoffs: a team-level reward, or one weighted toward problems that stay resolved. Rules: count a reopened ticket against the original resolution. Observability: publish reopen rates. Who decides: let senior engineers escalate within defined limits so tickets don’t wait on Dana’s calendar. Removing an incentive: drop the public leaderboard and restore discretion.
4. **Gaming the redesign.** Reward low reopen rates and engineers might discourage customers from reopening. Reward a team target and some may free ride. Reward satisfaction scores and people might survey only happy customers. That is stress-testing, not accusation.
5. **Diagnosis.** Firing people fixes the player before inspecting the game. Accountability can still matter, but a new hire faces the same payoffs. Dana should name the mechanism (individual ranking on a count that double-credits shallow fixes), make the smallest

testable change (stop counting reopens as new tickets for one quarter), explain it openly to the team, then observe, update, and adapt.

Exercise B — Your own game

Pick a problem in your own life that keeps coming back despite your best efforts: a recurring argument at home, a work process that keeps producing the same failure, a client pattern, or a habit you keep “trying harder” to change.

1. List every “better move” you have already tried: the conversations, the wording, the willpower.
2. Work the problem through all six steps of the Game-Change Audit.
3. Circle the single lever most responsible for the behavior.
4. Design the smallest change that would test your hypothesis. If you are both designer and player, design for the behavior you are likely to have, not the behavior you wish you had.
5. Run the ethics questions. If others are affected, can you design the change with them rather than impose it?
6. If you lack authority to redesign anything, ask: why am I trying to become better at this game? List what exit would cost before deciding.

Exercise C — Partner drill / role-play

Format: pairs or small groups, 15 minutes. Works with an AI practice partner too.

Setup: A small family business, a bakery, has a recurring problem. Every Friday the two co-owners, siblings, argue about who gets to approve the weekend order, and the order goes in late about half the time. Every order decision requires both of them to agree.

Roles:

- **Sibling A (front of house).** Privately: you think your sibling over-orders and you discard unsold goods. You lack the sales numbers.
- **Sibling B (kitchen and ordering).** Privately: you have the sales data, believe running out on Saturday costs more than waste, and feel micromanaged.
- **Designer (optional third player).** Your job is not to decide who is right about quantities but to help them redesign the game.

Round 1 (5 min): The siblings argue Friday’s order under the current rules. The Designer notes the rules, information gaps, timing, and who decides.

Round 2 (7 min): The Designer leads the group through the audit’s lever table. Propose at least two redesigns (for example, B decides within a set range, a shared sales-and-waste sheet, or a standing Wednesday review), then stress-test: how could either sibling game the new rule?

Debrief (3 min): 1. In Round 1, which disagreement was really about payoffs, and which was about missing information? 2. Which lever changed the most in your redesign? Who decides, information, or timing? 3. How might the new arrangement be gamed, and what would you watch for in the first month? 4. Did the redesign treat both siblings as informed decision-makers, or did it quietly favor one?

Common mistakes

- **Giving another speech when the structure rewards the opposite behavior.** Before another workshop, ask what the system actually pays for.
- **Blaming character before inspecting the game.** Ordinary people often respond predictably to a badly built system. Ask what the game makes attractive, without assuming structure explains everything.
- **Confusing the metric with the goal.** Ask what behavior actually improves the number, and what hard-to-measure work may lose effort.
- **Assuming “set rule, get outcome.”** Players respond, learn, and adapt. Treat your design as a hypothesis.
- **Designing for saints or for monsters.** Design for ordinary human beings, not only saints and not only villains.
- **Always adding rules and rewards.** Sometimes the accumulated rules are the problem; remove the distorting incentive.
- **Redesigning everything at once.** Systems are interconnected. Make the smallest testable change, then observe and update.
- **Treating design as control or manipulation.** Friction, defaults, and hidden terms are powerful; respect truth, autonomy, dignity, and informed choice.
- **Ignoring exit, or treating it as automatically right.** Not every game deserves another move, but leaving has costs too.

Quiz

1. How does mechanism design differ from the game theory used in earlier modules?
 - A) It focuses on predicting a single opponent’s next move.
 - B) It starts from the outcome you want and asks how to design the rules to make it more likely.
 - C) It assumes players are always irrational.
 - D) It replaces incentives with persuasion.
2. In the “one cuts, the other chooses” cake example, why does the cutter make the pieces equal?
 - A) The parent lectures them on fairness first.
 - B) The cutter has become less selfish.

- C) If one piece is clearly larger, the other child will take it, so self-interest favors equal pieces.
- D) The parent inspects the pieces afterward.
3. A company says “tell us the truth” but penalizes people who report bad news. According to the chapter, what should you expect?
- A) The stated value will win because people have integrity.
- B) The rewarded behavior will tend to win.
- C) Communication training will resolve it.
- D) Behavior will be unaffected by the penalty.
4. What did the Kenyan teacher-incentive study illustrate?
- A) Incentives never change behavior.
- B) Incentives tied to test scores produced lasting, broad learning gains.
- C) Incentives raised scores on the incentivized exams, mainly through test preparation, without clear gains in attendance, homework, or pedagogy.
- D) Teachers ignored the incentive entirely.
5. A manager rewards her team solely on the number of leads generated, then complains that lead quality has dropped. What is the best diagnosis?
- A) The team missed the point and needs a firmer talk.
- B) The scoreboard told them what the point was; the metric is not the goal.
- C) The team has poor values.
- D) Lead quality cannot be influenced by design.
6. A project automatically continues unless someone explicitly stops it. Which lever is this?
- A) The default: what happens when nobody chooses.
- B) The players.
- C) Exit.
- D) Incentive compatibility.
7. You introduce a productivity metric. Productivity rises, but quality collapses. What is the chapter’s recommended response?
- A) “Why aren’t people behaving according to my design?”
- B) Add many more rules at once.
- C) “What did I misunderstand about the game?” and then observe, update, and adapt.
- D) Abandon all measurement permanently.
8. A subscription takes one click to start and twenty to cancel. From the chapter’s ethical standpoint, what is the key point?
- A) Since cancelling is technically possible, the design is neutral.
- B) Friction is part of the game, and being able to make a choice strategically difficult does not mean you should.

- C) Friction only matters in government systems.
- D) Ethical design means removing all choices.

Answer key

1. **B** — Mechanism design reverses the usual question: instead of “given the game, what will players do?”, it asks “given the outcome we want, how should we design the game?”
2. **C** — The mechanism uses the cutter’s preference for the bigger piece. Because the other child chooses first, self-interest produces fairness. That is incentive compatibility.
3. **B** — When the stated value and the rewarded behavior conflict, don’t be surprised when the reward wins. People can resist bad incentives, but repeated payoffs produce more of the paid-for behavior.
4. **C** — Scores on the linked exams rose, largely through test preparation, with little evidence of broader teaching improvement and no persistent long-run learning gains. The incentive worked, just not entirely as intended.
5. **B** — People optimize against the scoreboard. If pay depends entirely on lead count, lead count is what tends to rise.
6. **A** — Defaults structure what happens when no player acts. What happens when nobody chooses is itself a rule.
7. **C** — The system talks back. Unexpected results are data, and the right response is strategic humility: ask what you misunderstood, then make small, testable adjustments.
8. **B** — Friction shapes behavior. Ethical mechanism design asks whether players understand the game and retain meaningful choice; a dark pattern is still game design, serving the wrong objective.

Field assignment

This week: run one small, ethical game redesign.

Choose one recurring problem you have the authority to influence: a household routine, a team process, a client arrangement, or your own work environment. Complete the Game-Change Audit, identify the mechanism you believe is producing the unwanted behavior, and implement the smallest change that tests it. If others are affected, explain the new game clearly, or better, design it with them.

Each day, record in a few lines: - What behavior did you observe under the new structure? - Did anyone respond in a way you didn’t predict, including ways of gaming the new rule? - What became easier, and what became harder?

Debrief questions for next week: 1. What did the system “talk back” with, and what did it reveal that you had misunderstood about the game? 2. Before this week, how much of the problem had you attributed to someone’s character? How has that changed? 3. What is your next smallest adjustment, or have you concluded that the better move is to leave this game?

The question to carry

Do I need a better move—or a better game?

Module 12 — Strategic Judgment

The question: What did I learn from what just happened?

Book reading	Conclusion, “Strategic Judgment”
Time	~60 min of video lessons · ~75 min of workbook · one week of field practice, plus the Capstone project (two to three weeks)
Tool you’ll build	The Strategic Judgment Worksheet (the Eleven Questions) and the Model → Act → Observe → Update → Adapt Practice Log

Why this module matters

You now have eleven lenses. The risk at this point is not that you know too little. It is that you try to use everything, everywhere, all at once. The book is blunt about this: the eleven models are not eleven steps. You are not supposed to walk into a conversation and think, “Step one: identify the game. Step two: identify the payoffs.” An experienced driver doesn’t recite “check mirror, calculate distance, adjust steering angle.” They have developed perception. The goal of this course was never recitation. It was recognition.

Strategic judgment is the skill that sits above the models. It means choosing the two or three lenses that matter right now, knowing when a situation doesn’t need analysis at all (“Sure” sometimes just means “Sure”), keeping “Will this work?” separate from “Should I do it?”, and treating your read of any situation as provisional. Your first model will often be wrong. The person you thought cared about money cares about status. The person you thought was angry is scared. The threat you thought was a bluff was a commitment.

So this module adds one rule over all eleven models: **your model of the game is provisional, so update it from play.** The loop that holds everything together is Model → Act → Observe → Update → Adapt. That is what strategic thinking is. You won’t predict perfectly or be right every time. What you can do is learn faster from what happens. The module ends with the Capstone, where you put the whole framework to work on one real interaction that matters to you.

Learning objectives

By the end of this module you will be able to:

1. **Select** the two or three lenses that matter most in a given interaction, and justify why the others can wait.
2. **Distinguish** situations that warrant strategic analysis from ordinary interactions where more analysis would only add suspicion and confusion.
3. **Separate** facts from your story about the facts, state your confidence, and hold several competing hypotheses at once.
4. **Apply** the Model → Act → Observe → Update → Adapt loop, treating surprise and prediction errors as information.
5. **Evaluate** any planned move against both “Will this work?” and “Should I do it?”, using truth, autonomy, dignity, and informed choice.

Lessons

Lesson 12.1 — Recognition, not recitation (12 min)

Teaching points

- The eleven models are lenses, not a procedure. At first you’ll ask the questions consciously. Over time you’ll start noticing the structure almost automatically.
- Recognition sounds like short internal prompts: “What payoff are they protecting?” “Would they actually want to carry that out?” “This is a pattern, not an incident.” “What makes that credible?” “They’re not opposed. They don’t know what the other is going to do.” “The incentive system rewards the exact opposite behavior.”
- The world will not label the game for you. You’ll get symptoms, like “We need an answer today,” a child asking again, employees who stop sharing leads, or the same argument with your partner. Your job is to recognize the structure underneath.

Example to use: The book’s list of disguised games. Nobody walks into your office and announces “an incomplete-information bargaining game containing asymmetric outside options and a credibility problem.” They say “We need an answer today.” Your partner won’t say “We have entered a stable negative equilibrium reinforced by reciprocal expectations.” You’ll just have the same argument again.

On-screen takeaway: Reality gives you the symptoms. You have to recognize the structure.

Lesson 12.2 — Choosing the right lens, and knowing when to put it down (14 min)

Teaching points

- Walk through the book’s “final offer by tomorrow” example. Every one of the eleven lenses produces a useful question. Who has authority? What happens to them if I say

no? Have I taught them that artificial pressure makes me move? What makes “final offer” believable? Do I need to accept their yes-or-no structure at all?

- You probably don’t need all eleven. Judgment means recognizing which two or three matter most right now.
- More analysis is not always better. Use the framework when the stakes matter, something doesn’t make sense, a pattern keeps repeating, incentives are unclear, trust is uncertain, multiple people affect the outcome, or your normal way of thinking isn’t working.
- Strategic thinking should reduce confusion. If it makes ordinary life needlessly suspicious and complicated, you’re using the tool badly.
- Strategic thinking is not suspicious thinking. Suspicion asks “What are they hiding?” and assumes hostile intent. Strategic thinking asks “What possibilities should I consider?” and tests hypotheses. The framework should make you less naive. It should not make you cynical.

Example to use: The friend who texts “Sure.” Then the overanalysis: “What strategic signal does the absence of an exclamation point communicate about their latent payoff structure?” Don’t do that. Sometimes “Sure” means “Sure.”

On-screen takeaway: Which two or three lenses matter most right now?

Lesson 12.3 — Your model is provisional: the learning loop (14 min)

Teaching points

- Everything depends on your interpretation of the game, and your interpretation can be wrong. Don’t fall in love with your diagnosis. Test it.
- The loop: **Model** (form your best understanding), **Act** (make a move), **Observe** (what actually happens?), **Update** (where did reality differ from your prediction?), **Adapt** (change your strategy, or change your model). Then repeat.
- Your mistakes are information. If you explain your reasoning clearly and they understand perfectly but still refuse, maybe it was never an information problem. Maybe your interests are genuinely opposed.
- Surprise is a gift. They didn’t cooperate, so ask what you misunderstood. They stayed, so ask what payoff you underestimated. They followed through, so ask what made the commitment stronger than you realized. The new rule made things worse, so ask what response the design created that you failed to anticipate.
- Don’t label someone irrational too quickly. Ask, “What would have to be true for this behavior to make sense?” They might have information you don’t, a different payoff, a longer game, a different reading of your behavior, or history you’re discounting. And yes, they might simply be making a mistake.

Example to use: The refused request. You predict, “If I refuse, they’ll become angry.” You refuse. They say, “No problem. I completely understand.” Good. You were wrong, and

that's valuable. Maybe you overestimated their dependency, misread their payoff, or believed the relationship was more fragile than it was.

On-screen takeaway: Model → act → observe → update → adapt.

Lesson 12.4 — Facts, stories, and multiple hypotheses (10 min)

Teaching points

- Separate the fact from the interpretation. Fact: they didn't return your call. Interpretation: they're disrespecting you. That's possible, but not the same thing. Maybe they forgot, are overwhelmed, never got the message, are avoiding the conversation, or are delaying on purpose.
- The discipline has three parts: "Here's what I know. Here's what I think it means. Here's how confident I am." This keeps you responding to reality instead of to an assumption.
- Strong strategists hold several hypotheses at once, then ask questions or watch for behavior that helps tell them apart.
- Certainty feels good. Accuracy is better.

Example to use: The customer who suddenly becomes resistant. It could be a price objection, a lack of trust, a need for someone else's approval, a play to improve their negotiating position, or a polite cover for having already chosen a competitor. You don't need to choose one immediately.

On-screen takeaway: Here's what I know. Here's what I think it means. Here's how confident I am.

Lesson 12.5 — Should I do it? Ethics, vulnerability, walking away (10 min)

Teaching points

- Truth, autonomy, dignity, and informed choice are not decorative. They constrain how this framework should be used.
- A strategy can "work" and still be wrong. You may be able to exploit someone's ignorance, press on an insecurity, or hide information so they choose what you prefer. Doing so means they did not make an informed choice. Keep two questions apart, "Will this work?" and "Should I do it?", and never collapse them into one.
- Cooperation isn't automatically good. People can cooperate to build a hospital or to commit a fraud. The ethical question is "What are they cooperating toward?"
- People are not pieces on a board. Calling someone a "player" reminds you that they are a decision-maker too, with agency and a different vantage point. That should make you more respectful of their agency, not less.
- Sometimes the correct move is vulnerability. In a repeated relationship where both sides protect themselves, a measured risk (a disclosure, a concession, an apology) can change beliefs and shift the equilibrium. The risk is exactly why the action carries information.

- Sometimes the correct move is walking away. If you have communicated clearly, tested assumptions, enforced boundaries, allowed for repair, and the game is still destructive, then stop playing. An outside option is part of the game. Know which games you refuse to play.

Example to use: “You never listen to me.” The tactical reflex is to pick a technique, like active listening or mirroring. See the game first. Is this a repeated pattern? What payoff are they seeking: understanding, agreement, respect, or attention? Have your past responses taught them that raising a concern leads to defensiveness? Are you in an equilibrium where they criticize because they expect you not to listen, and you defend because you expect them to criticize? Is this a coordination failure about what “being heard” means? You may still use a technique, but now it serves an understanding of the situation.

On-screen takeaway: Will this work? Should I do it? Never collapse them into one.

Key terms

Term	Plain-language meaning
Strategic judgment	Choosing which lenses matter in this situation, acting on your best model, and learning from what happens.
Recognition	Noticing the structure of an interaction almost automatically, as opposed to reciting steps.
Provisional model	Your current best interpretation of the game, held as a hypothesis to test, never as a conclusion to defend.
Model → Act → Observe → Update → Adapt	The learning loop: form a model, make a move, watch what happens, find where reality differed from your prediction, then change your strategy or your model.
Prediction error / surprise	The gap between what you expected and what happened. It’s the most useful signal that your model has a hole in it.
Multiple hypotheses	Holding several plausible explanations at once and gathering evidence that tells them apart.
Transfer	Recognizing a familiar structure (commitment, coordination, repeated game, mechanism design) in a situation you’ve never seen before.

Term	Plain-language meaning
Calibration	Making intuition trustworthy by checking it against outcomes, so confidence rises only as accuracy does.

Research spotlight

The Conclusion introduces no new experiments. Its evidence base comes from the research audit behind the whole framework. Two findings from that audit explain why this module exists.

1. Learning in games. The full-framework audit compared the eleven models with the topics covered in university game-theory curricula. Those curricula include learning in games (evolutionary and non-equilibrium learning) as well as equilibrium analysis. The dossier notes that researchers such as Fudenberg and Levine have argued that learning models improve game theory’s descriptive power, and that experimental evidence shows behavior can change substantially through experience rather than starting at equilibrium.

- What it suggests: Real players don’t arrive already knowing the solution. They adjust from feedback, and so should you. That is why the audit made “Model → act → observe → update → adapt” a meta-rule for the whole book instead of a twelfth model.
- What it does NOT prove: It doesn’t show that any particular person will learn quickly, or learn the right lesson. Experience without feedback can entrench a wrong story. Learning models also describe patterns across many players and rounds, so they don’t predict what one individual will do in your next conversation.

2. The audit’s completeness test. The audit asked whether any major area of game theory (bargaining, coalitions, dominance, social dilemmas, auctions, voting, networks) needed its own model. It applied one test: does the concept give an ordinary person a genuinely different and broadly useful way to see strategic human interaction? The answer was to keep eleven models and fold the other ideas in as supporting concepts. Examples include outside options and coalitions under See the Game, dominance under Think in Responses, and the coordination versus free-rider distinction under Diagnose the Coordination Problem.

- What it suggests: The framework is meant as “the smallest useful latticework of game-theoretic lenses for human interaction,” not as all of game theory. Choosing lenses well matters more than owning more of them.
- What it does NOT prove: Eleven is a design judgment, not an empirical discovery. The audit doesn’t claim the models are exhaustive, and it doesn’t claim that a situation fits neatly into one of them.

The tool: The Strategic Judgment Worksheet

Use this when the stakes matter, something doesn't make sense, a pattern keeps repeating, incentives are unclear, trust is uncertain, or several people affect the outcome. Don't use it for "Sure."

Part 1 — The situation

What happened (facts only):

My first interpretation (my story):

How confident am I in that story (0–100%)?

What is my immediate impulse, the thing I'd say or do with no pause?

Part 2 — The Eleven Questions (scan all, mark the 2-3 that matter most)

✓	Model	Question	My answer
☐	1. See the Game	What game am I actually in? (Players, actions, timing, outside options, who can combine with whom)	
☐	2. See the Payoffs	What does each player actually value? Are our interests opposed, aligned, or mixed?	
☐	3. Think in Responses	Then what? Is there an option I shouldn't choose regardless of what they do?	
☐	4. Manage Predictability	What should they be able to predict about me, and what becomes exploitable if they can?	
☐	5. See the Equilibrium	What pattern are we producing, and why is it stable?	

✓	Model	Question	My answer
☐	6. Map Information & Beliefs	What do I know, what am I assuming, and what might they know that I don't?	_____
☐	7. Separate Words From Credibility	What makes this believable?	_____
☐	8. Commitment Changes the Game	What actually binds them, or me?	_____
☐	9. The Shadow of the Future	What changes because there is a next time?	_____
☐	10. Diagnose the Coordination Problem	Are we actually in conflict, or failing to coordinate? (If everyone else cooperated, would I still have an incentive not to?)	_____
☐	11. Change the Game	Do I need a better move, or a better game?	_____

The 2-3 lenses that matter most, and why:

The lenses I'm deliberately setting aside, and why:

Part 3 — Competing hypotheses

Hypothesis about what's going on	Evidence for	What would tell this apart from the others?
H1: _____	_____	_____
H2: _____	_____	_____
H3: _____	_____	_____

What would have to be true for their behavior to make sense?

Part 4 — The move

My planned move:

What I predict will happen:

Will this work?

Should I do it? Check each: - Truth: Am I saying only what I believe is true and not creating a false impression?

- Autonomy: Does the other person stay free to choose, including to say no?

- Dignity: Would I be comfortable if they saw my full reasoning?

- Informed choice: Do they have the information they'd need to decide well?

What are we cooperating toward?

Is the right move vulnerability, walking away, or refusing this game?

The Practice Log: Model → Act → Observe → Update → Adapt

Keep one entry per meaningful interaction. Over weeks, this log is how your intuition gets calibrated.

Date	Model (my read + key lens)	Act (what I did)	Prediction	Observe (what actually happened)	Update (where reality differed, and why)	Adapt (new strategy or new model)
------	----------------------------	------------------	------------	----------------------------------	--	-----------------------------------

Surprise prompts for the Update column: - They didn't cooperate. What did I misunderstand?

- They stayed. What payoff did I underestimate?

- They followed through. What made the commitment stronger than I realized?

- My rule or design backfired. What response did it create that I didn't anticipate?

Practice exercises

Exercise A — Analyze the scenario

Dana is the operations director at a regional logistics company. For the third quarter in a row, the sales team has promised clients delivery windows her warehouse can't meet, and her team absorbs the overtime and the angry calls. On Monday, the VP of Sales, Marcus, emails her and copies the CEO: "We just landed the Harlan account. They need three-day turnaround starting next month. This is non-negotiable. The CEO is aware." Dana's first impulse is to reply-all with a detailed list of every missed promise from the last nine months. She knows Marcus's team is paid commission on signed contracts, not on retained or profitable accounts. She doesn't know whether the CEO actually approved three-day turnaround or was simply told about the deal. She'll work with Marcus for years, and the Harlan account would be good for the company if it could be served reliably. Two of her shift leads have privately told her they're close to quitting.

1. Which two or three lenses matter most here, and which can Dana set aside for now?
2. Separate the facts from Dana's likely story. What competing hypotheses should she hold about Marcus and the CEO?
3. What is wrong with Dana's first impulse, judged by both "Will this work?" and "Should I do it?"
4. Design a move and state a prediction. Then describe what Dana should do in the Observe and Update steps if her prediction turns out wrong.

Model answer

1. The lenses that matter. Three stand out.

- **See the Payoffs / Change the Game.** Sales is paid on signed contracts, not on deliverable ones. Over-promising is the rational response to that incentive, so this is a pattern the system produces, not one bad act by Marcus. A better email won't fix it. The durable fix sits at the level of the game, for example a sign-off rule for delivery terms or commission tied partly to retention.
- **Map Information & Beliefs / Separate Words From Credibility.** "Non-negotiable" and "The CEO is aware" are claims. What makes them credible? "Aware" is not the same as "approved three-day turnaround knowing warehouse capacity." Dana doesn't know what the CEO knows, and the CEO may not know about the overtime and the shift leads who are about to quit.
- **The Shadow of the Future.** Dana and Marcus will play this game for years, and her response sets a precedent. If she quietly absorbs this one, she teaches the organization that operations will always bend (which also touches Manage Predictability).

She can set aside, for now, Diagnose the Coordination Problem (the underlying interests aren't really aligned while the incentive stays as it is) and a detailed Think in Responses tree for Marcus alone. Marcus may not be the whole game. The CEO is the decisive player.

2. Facts vs. story. Facts: three quarters of missed promises, the commission structure, the email, the CC, the shift leads' warnings. Story: "Marcus is deliberately throwing my team under the bus." Competing hypotheses: (H1) Marcus knows the windows are unrealistic and is using the CEO's name as pressure. (H2) Marcus genuinely doesn't understand warehouse constraints. (H3) The CEO did approve three-day turnaround but without capacity data. (H4) The CEO approved it and plans to add resources Dana hasn't heard about. A question to the CEO will distinguish H1/H2 from H3/H4.

3. The reply-all list. Will it work? Probably not. It turns a design problem into a public fight, invites Marcus to defend himself, and makes Dana look like an obstacle in front of the CEO. Then what? Marcus escalates and the pattern hardens. Should she do it? Her facts may be true, but publicly shaming a colleague is a dignity problem, and the "gotcha" framing aims to win the exchange rather than to give the CEO the information needed for an informed choice.

4. A move. Dana replies to both, briefly and without blame: "Glad we landed Harlan, and I want us to keep them. To commit to three-day turnaround reliably, here's what it would take [staffing, shift, cost], and here's what we can guarantee today [four days, with three-day by a set date]. Can the three of us take 20 minutes this week to decide which we commit to?" She also asks the CEO separately for a broader conversation about how delivery terms get approved. She doesn't need to name the commission plan yet, but she brings data on the three quarters of missed windows and overtime cost.

Prediction: the CEO, given capacity information, picks a staged commitment, and Marcus accepts because keeping Harlan protects his future commissions.

If she's wrong. Say the CEO replies, "Just make three days work." Dana should update rather than label anyone irrational. Maybe the CEO has information she doesn't, like a strategic reason Harlan matters more than its margin. Maybe her capacity numbers weren't credible to someone outside operations. Her adaptations could be to ask what's behind the urgency, to make the cost concrete and verifiable, or to tie her agreement to a specific commitment of resources. If over several rounds the game stays destructive (promises without resources, staff leaving), strategic judgment may tell her that her own outside options deserve a serious look.

Exercise B — Your own game

Pick one live situation from your own life that meets at least two of the book's criteria: the stakes matter, something doesn't make sense, a pattern keeps repeating, incentives are unclear, trust is uncertain, multiple people affect the outcome, or your normal approach isn't working.

1. Write the facts in one paragraph with no adjectives or motives. Then write your story in a second paragraph. Put a confidence percentage on the story.

2. Scan all eleven questions in the Strategic Judgment Worksheet for about 60 seconds each. Then circle no more than three. For each lens you set aside, write one sentence explaining why.
3. Write at least three competing hypotheses. For each, name one question you could ask or one behavior you could watch for that would separate it from the others.
4. Draft your move and your prediction. Run the four ethics checks. If any one fails, redesign the move before going further.
5. Ask yourself whether there is a version of this where the best move is vulnerability, or one where it is walking away. Write one sentence on each.
6. Enter the situation as the first row of your Practice Log. You'll return to it in the Field assignment and may choose it for your Capstone.

Exercise C — Partner drill / role-play

“Two Lenses, Then Update” (15 minutes, pairs or trios, or with an AI practice partner)

Setup. A small nonprofit's development director (Player A) is meeting a longtime major donor (Player B). The donor has given \$50,000 a year for six years. This year the donor says, “I'm thinking of pausing my gift. I'll decide by Friday.”

Player A privately knows: The organization recently changed its program focus. The board chair told A that the donor “always comes around.” A has no information about the donor's finances.

Player B privately knows (choose one card at random and don't reveal it):

- Card 1: B is upset that nobody told them about the program change before it was announced publicly. B feels taken for granted and wants to be consulted.
- Card 2: B's business had a bad year and B is embarrassed to say so. B would like to give a smaller amount without losing face.
- Card 3: B's adult daughter has started her own nonprofit and B feels torn. B is open to a joint project.

Round 1 (5 min). Before speaking, Player A names out loud the two lenses they'll lead with and one hypothesis they're testing. A then conducts the conversation. A's goal is to tell the hypotheses apart, not to secure the gift.

Round 2 (3 min). Pause. A states what they observed, what surprised them, and how they're updating. B says nothing yet.

Round 3 (4 min). A adapts and continues. A must stay within truth, autonomy, dignity, and informed choice. That means no inventing urgency, no guilt, and no misrepresenting the program change.

Debrief (3 min).

- Which card did A think B held, and when? What evidence moved the estimate?

- Did the board chair's "always comes around" belief help or hurt? (It's a story, not a fact.)
- Did A respect B's ability to say no? Would B feel comfortable if they saw A's full reasoning?
- Which lens turned out to matter most, and was it one A chose at the start?

Common mistakes

- **Running all eleven models every time.** The models are lenses, not a checklist. Full analysis of everything is slow and noisy. Instead, pick the two or three that matter most right now.
- **Analyzing interactions that don't need it.** Turning "Sure" into a signaling puzzle makes ordinary life needlessly suspicious and complicated. Save the framework for high stakes, recurring patterns, unclear incentives, and uncertain trust.
- **Falling in love with your diagnosis.** Your first model will often be wrong. Treat it as provisional and look for evidence that would disprove it.
- **Labeling people irrational.** That word often replaces thinking. Ask what would have to be true for the behavior to make sense, while allowing that sometimes people simply make mistakes.
- **Confusing strategic with suspicious.** Suspicion assumes hostile intent and reads ambiguity negatively. Strategic thinking recognizes ambiguity and tests hypotheses. Become less naive without becoming cynical.
- **Collapsing "Will this work?" into "Should I do it?"** A strategy can work and still violate truth, autonomy, dignity, or informed choice. Ask both questions every time.
- **Trusting uncalibrated intuition.** Experience without feedback can produce confident error. Use the Practice Log so your intuition earns its trust.

Quiz

1. According to the Conclusion, how should the eleven models be used in a live interaction?
 - A) As a fixed sequence, from See the Game through Change the Game
 - B) As lenses, choosing the two or three that matter most right now
 - C) Only in formal negotiations
 - D) All eleven, written out, before every important conversation
2. Which of these is NOT one of the book's signals that a situation warrants strategic analysis?
 - A) A pattern keeps repeating
 - B) Trust is uncertain
 - C) A friend's text message is short
 - D) Multiple people affect the outcome

3. A manager is sure a team member's missed deadlines come from laziness. He introduces a public "late list," and the deadlines get worse. Using the loop, what is the best next step?
 - A) Make the late list more prominent
 - B) Conclude the employee is irrational
 - C) Ask what response the new rule created that he failed to anticipate, and revise his model
 - D) Stop tracking deadlines entirely
4. What best distinguishes strategic thinking from suspicion?
 - A) Strategic thinking assumes people usually lie
 - B) Suspicion tests hypotheses, and strategic thinking doesn't
 - C) Strategic thinking considers possibilities and tests them, while suspicion assumes hostile intent
 - D) There is no meaningful difference
5. You explain your reasoning clearly. The other person restates it accurately and still refuses. What does the book suggest you consider?
 - A) Explaining it again more slowly
 - B) That it may never have been an information problem, and your interests may be genuinely opposed
 - C) That they are acting irrationally
 - D) That you should hide some information next time
6. A sales rep learns that a prospect is anxious about falling behind competitors. She considers inventing a story that two of the prospect's rivals just signed. Which statement reflects the book's position?
 - A) If it closes the deal, it's good strategy
 - B) It's fine as long as the product is genuinely good
 - C) It may "work," but it violates truth and informed choice, and "Should I do it?" must be asked separately
 - D) It's acceptable in a one-shot interaction because there's no shadow of the future
7. Two longtime co-founders each withhold their real concerns about the company's direction because each fears the other will use it against them. Nothing changes for months. What move does the Conclusion suggest may sometimes be correct here?
 - A) A measured act of vulnerability, such as disclosing a concern first, which can change beliefs and shift the equilibrium
 - B) Withholding even more, to protect your position
 - C) Issuing an ultimatum
 - D) Waiting for the other person to go first, however long it takes

8. What sits “above all” eleven questions in the book’s framework?
- A) “Who has the most power?”
 - B) “What should I say?”
 - C) “What did I learn from what just happened?”
 - D) “How do I win?”

Answer key

1. **B** — The models are “not eleven steps.” Judgment means recognizing which two or three matter most right now.
2. **C** — A short text usually needs no analysis. Sometimes “Sure” means “Sure.” The others are on the book’s list of situations where the framework helps.
3. **C** — The design backfired, so this is a surprise to learn from. Model → act → observe → update → adapt means checking where reality differed from the prediction, possibly a mistaken payoff or a mechanism that rewards hiding problems, and then changing the model or the strategy.
4. **C** — Suspicion asks “What are they hiding?” and interprets ambiguity negatively. Strategic thinking asks “What possibilities should I consider?” and recognizes ambiguity.
5. **B** — The book uses this exact case. Full understanding followed by refusal suggests the problem wasn’t missing information, so update the model.
6. **C** — A strategy can work and still be wrong. Fabricating competitor behavior violates truth and prevents an informed choice. Never collapse “Will this work?” and “Should I do it?” into one question.
7. **A** — When both sides protect themselves in a repeated relationship, a measured risk can carry information precisely because it’s risky, and it may reopen cooperation.
8. **C** — After the eleven questions, the book adds: “And above them all: What did I learn from what just happened?”

Field assignment

This week: run the loop five times.

Choose five interactions this week that clear the bar for analysis. At least one should involve someone you’ll deal with again, and at least one should involve more than two people. For each one:

1. Before acting, take the pause between stimulus and response. Write your story, your confidence, and the one or two lenses you’re leading with.
2. Make a specific prediction: “If I do X, they’ll do Y.”
3. Act, then log in the Practice Log what actually happened.
4. Score your prediction as right, partly right, or wrong, and fill in the Update and Adapt columns.

Also record one ordinary interaction where you noticed yourself starting to over-analyze and chose not to.

Debrief questions for next week 1. How many of your five predictions were right? In the ones that were wrong, was the error mostly about payoffs, information, credibility, or the structure of the game? 2. Which surprise taught you the most, and how did you change your model because of it? 3. Did any move tempt you toward something that would “work” but fail the truth, autonomy, dignity, or informed-choice check? What did you do instead?

Capstone project

Purpose

This is the course-completion project. Over two to three weeks, you’ll take one real, high-stakes interaction from your own life through the full strategic-judgment cycle. You’ll see the game, choose the lenses that matter, design an ethical move, act, and learn from what happens. The Capstone does not test whether you get the outcome you wanted. People surprise us, and no framework gives you control over them. It tests whether you saw the game clearly, acted with integrity, and learned faster from what happened.

Choosing your interaction

Choose something real, current, and consequential where you’ll actually make a move within the project window. Good candidates include a compensation or contract negotiation, a recurring conflict with a partner, family member, or colleague, a team that keeps failing the same way, a boundary you need to set, a decision about whether to stay in or leave a situation, or a proposal you need others to adopt.

It must meet at least three of the book’s criteria: high stakes, something doesn’t make sense, a recurring pattern, unclear incentives, uncertain trust, multiple people affecting the outcome, or your normal approach not working.

Avoid situations involving physical danger or ongoing abuse. In those cases, safety and appropriate professional support come before any analysis.

Deliverables

1. **Situation brief** (250–400 words). Facts, kept strictly separate from your story, with a confidence rating, plus your stakes and what “winning” would actually mean to you.
2. **Completed Strategic Judgment Worksheet.** All eleven questions scanned, with the 2–4 lenses you chose analyzed in depth and a written rationale for the lenses you set aside.
3. **Hypothesis table.** At least three competing explanations of the other players’ behavior, with a plan to tell them apart.

4. **Move design** (300–500 words). Your move, the alternatives you rejected and why, your explicit prediction, and the full ethics check. Include whether a game-changing move (new options, new rules, new players, or refusing to play) is available.
5. **Practice Log**. At least three Model → Act → Observe → Update → Adapt entries covering the interaction and its follow-ups.
6. **Debrief** (600–900 words). Observe, Update, Adapt. What happened, where reality differed from your prediction and why, what you'd do differently, and what you'll do next. End with one paragraph on which lens you'll now notice faster in other situations (transfer).

Step-by-step process

1. **Pick and frame (days 1-2)**. Choose the interaction. Write the situation brief. Name what you'd consider a good outcome, including outcomes other than “winning,” such as preserving the relationship, telling the truth at a cost, or walking away.
2. **Scan wide (day 3)**. Answer all eleven questions briefly, without editing. This is the only time you deliberately run the whole set.
3. **Choose the lens (day 3)**. Circle the two to four lenses that explain the most. For the others, write one line each on why they matter less here. Over-analysis counts against you, not for you.
4. **Go deep (days 4-5)**. Fully analyze the chosen lenses. Map players (including people not in the room), payoffs, likely responses, the stabilizing pattern, information gaps, credibility, commitments, and future rounds, as relevant.
5. **Hold hypotheses (day 5)**. Build the hypothesis table. Where you can, design a low-risk question or observation that would discriminate among them before your main move.
6. **Design the move (day 6)**. Draft your move and two alternatives. Ask “Will this work?” and “Should I do it?” separately. Consider vulnerability, consider walking away, and consider changing the game. Write a specific, checkable prediction.
7. **Act (days 7-14)**. Make the move. Log what happened as soon as you can, with facts first and interpretation second.
8. **Observe, update, adapt (days 7-21)**. Compare outcome to prediction. Use the surprise prompts. Revise your model, make at least one follow-up move if appropriate, and log it.
9. **Write the debrief (final days)**. Be honest about what you got wrong. A clean prediction that came true teaches less than a surprise you learned from.

Scoring rubric

Criterion	Developing	Proficient	Strong
1. Seeing the game	Describes the conversation but not the structure. Players, options, or outside options are missing. Facts and story are blended.	Identifies the main players, options, and payoffs. Separates facts from story.	Maps the full structure, including players not in the room, outside options, coalitions, and timing. Separates facts from story with explicit confidence levels.
2. Lens selection	Applies every model mechanically, or relies on just one without justification.	Chooses 2–4 relevant lenses and explains the choice.	Chooses lenses with a sharp rationale, explains why the others matter less, and shows the restraint of “more analysis isn’t always better.”
3. Hypotheses and prediction	One explanation, held with certainty. Prediction vague or missing.	Several hypotheses and a specific prediction.	Several well-differentiated hypotheses, a deliberate way to tell them apart, and a specific, checkable prediction made before acting.
4. Move design	Move is a script or tactic disconnected from the analysis.	Move follows from the analysis and considers likely responses.	Move follows from the analysis, weighs alternatives (including vulnerability, walking away, or changing the game), and accounts for future rounds and precedent.

Criterion	Developing	Proficient	Strong
5. Ethics: truth, autonomy, dignity, informed choice	Ethics not considered, or the move relies on deception, pressure, or withheld information.	Move passes all four checks, and “Will this work?” and “Should I do it?” are addressed separately.	Move passes all four checks with specific reasoning, considers what any cooperation is directed toward, and treats the other person as a decision-maker whose agency is respected even when it costs a payoff.
6. Observe / Update / Adapt	Reports the outcome only, or explains away surprises. Calls others “irrational.”	Compares outcome to prediction, identifies at least one model error, and states an adaptation.	Treats surprise as information, asks what would have to be true for others’ behavior to make sense, revises the model specifically, follows up, and names a transferable lesson.

A passing Capstone is Proficient or better on all six criteria. Ethics (criterion 5) must be at least Proficient regardless of other scores. A strategy that works but fails that check is incomplete.

The question to carry

What did I learn from what just happened?