

THE COURSE

Practical Game Theory

See the Game Behind the Conversation

Student Workbook

PAUL MASCETTA

Contents

Module 0 — Orientation: The Game Behind the Conversation	3
Module 1 — See the Game	7
Module 2 — See the Payoffs	14
Module 3 — Think in Responses	20
Module 4 — Manage Predictability	27
Module 5 — See the Equilibrium	33
Module 6 — Map Information & Beliefs	39
Module 7 — Separate Words From Credibility	46
Module 8 — Commitment Changes the Game	52
Module 9 — The Shadow of the Future	58
Module 10 — Diagnose the Coordination Problem	64
Module 11 — Change the Game	70
Module 12 — Strategic Judgment	77

Module 0 — Orientation: The Game Behind the Conversation

The question: What if the way we normally think about communication starts too late?

The tool: The Interaction Journal

Use one page per interaction. You'll add a new column of questions in each module.

Date / setting:

Who was involved (including people not in the room):

What I wanted:

What I think they wanted:

What I said or did:

How they responded:

Did their response surprise me? What might that tell me?

If I could replay it, what question would I ask myself before speaking?

Practice exercises

Exercise A — Analyze the scenario

Dana manages a small marketing team. For three weeks, she has asked a senior designer, Marcus, to submit his drafts two days before client deadlines so she can review them. She has explained why, rephrased the request, and finally sent a firm email. Marcus keeps submitting the night before. Dana's colleague suggests she "work on her delivery." Dana suspects the problem isn't her words at all. Marcus is paid a bonus based on client satisfaction scores, not internal deadlines. The client contact loves Marcus and often requests last-minute changes directly from him.

1. Who are the players, including any who aren't in Dana's conversations?
2. Why might better wording fail here?
3. What question should Dana be asking instead of "How do I say this more clearly?"

Exercise B — Your own game

Think of one recurring interaction that frustrates you: at work, at home, or with a customer. Write down the last thing you said in it. Then list everything you assumed before you said it: about what they wanted, what they knew, and how they'd respond. Circle any assumption you have never actually tested.

Exercise C — Partner drill / role-play

Setup (10 minutes, pairs). Partner 1 describes a recent conversation that went badly in two minutes, focusing only on what was said. Partner 2 may then ask only structural questions for five minutes: Who else was affected? What did each person want? What were their alternatives? What happened next? Switch roles.

Debrief: What did the structural questions reveal that the words didn't? Did either of you change your interpretation of the other person?

Quiz

1. In this course, a "game" is best defined as:
 - A) Any competitive situation with a winner and a loser
 - B) Any situation in which outcomes depend on the choices of more than one person
 - C) A technique for gaining leverage in a conversation
 - D) A mathematical model that requires calculation
2. Why can excellent communication still produce a bad outcome?
 - A) Because people rarely listen

- B) Because tone matters more than content
 - C) Because if your read of the situation is wrong, good words execute the wrong strategy
 - D) Because game theory says communication is irrelevant
3. Which is the course's ethical foundation?
- A) Win-win, always
 - B) Truth, autonomy, dignity, informed choice
 - C) Reciprocity, scarcity, authority
 - D) Whatever produces the best outcome for you
4. The learning loop that runs through the course is:
- A) Plan → Pitch → Close
 - B) Listen → Mirror → Label
 - C) Model → Act → Observe → Update → Adapt
 - D) Analyze → Decide → Never revisit
5. Which question belongs to Mental Model 10?
- A) What makes this message believable?
 - B) Are we in conflict, or failing to coordinate?
 - C) What does each player actually value?
 - D) What actually binds a promise?
6. A mental model in this course is best described as:
- A) A law that predicts behavior precisely
 - B) A lens that tells you what to look at
 - C) A script for difficult conversations
 - D) A personality test
7. When someone responds in a way you didn't predict, the course suggests you:
- A) Conclude they are irrational
 - B) Repeat your move more forcefully
 - C) Treat the surprise as information and update your model
 - D) Ignore it as noise
8. How much analysis should you apply to an interaction?
- A) The full eleven-model audit every time
 - B) None, because intuition is enough
 - C) As much as the stakes and uncertainty justify
 - D) Only when you are losing

Field assignment

This week, observe; don't change anything. Record three interactions in your Interaction Journal: one at work, one at home, and one with a stranger or customer. For each, note one thing that shaped the outcome that wasn't the words themselves, such as a person not in the room, an incentive, a deadline, or a past history.

Debrief questions for next week: 1. Which of the three interactions had the most "invisible" structure? 2. Where did you catch yourself assuming you knew what someone wanted? 3. Which of the eleven questions do you suspect you'll need most?

The question to carry

| What did I assume before I spoke?

My notes

Module 1 — See the Game

The question: What's really going on here?

The tool: The Game Map

Use this before any important interaction. Write short answers and don't aim for perfection. The goal is to train your attention, not to complete a form.

The interaction:

WHO?

- Who are the relevant players?

- Who can actually affect the outcome, including anyone not in the room?

WHAT?

- What does each person want?

- Where are our interests aligned?

- Where are they different?

OPTIONS?

- What can each player actually do?

CONSTRAINTS?

- What limits those choices?

- Which constraints are stated, and which do I know are real?

INFORMATION?

- Who knows what?

- Who doesn't know what?

TIMING?

- Who moves first?

- Who responds?

- What can each person observe before choosing?

ALTERNATIVES?

- What happens if there is no agreement, for me and for them?

DEPENDENCIES?

- Who needs whom?

- Who has alternatives?

THEN WHAT?

- If I make this move, what is the other person likely to do next?

- And what would I do after that?

Final check

- Could there be a game behind the conversation, such as credit, trust, status, authority, or precedent?

- Is this a person problem or a game problem?

- What am I failing to notice about the game?

Practice exercises**Exercise A — Analyze the scenario**

Dana runs operations at a regional distribution company. For six weeks she has been pushing Marcus, the warehouse lead, to adopt a new scheduling system. Marcus keeps saying it “isn’t ready” and that his team “doesn’t have time to learn it.” Dana has started to see him as stubborn and plans to tell him at Thursday’s one-on-one that the switch is happening next Monday, no exceptions. Here’s what Dana hasn’t thought about. Marcus’s quarterly bonus is tied to on-time shipment rates, and any slowdown during the learning period will hit his numbers. The software vendor’s onboarding team is booked for the next three weeks. Two of Marcus’s most experienced pickers have quietly told him they’d rather leave than learn “another system.” Dana’s own boss, the VP, approved the system and asked her for a progress update at the end of the month. Marcus knows other warehouses in the area are hiring leads.

1. Who are the players in this game, and which ones aren’t in Thursday’s meeting?
2. Where are Dana’s and Marcus’s interests aligned, and where do they differ?

3. What constraints shape Marcus's options, and what alternatives does each of them have if they don't reach agreement?
4. If Dana announces "next Monday, no exceptions," what is Marcus likely to do next? Then what?
5. Is this primarily a person problem or a game problem?

Exercise B — Your own game

Pick a live situation you've been thinking about mostly in terms of what to say: a conversation you're dreading, a request, or a recurring friction.

1. Describe the situation in one sentence, as you've been seeing it.
2. Fill out the full Game Map. Circle any answer you're unsure of. Those are your information gaps.
3. Name one player you hadn't been considering, and one constraint or alternative that changes the picture.
4. Write the stated issue and one possible game behind it. What would tell you whether the stated issue is the real one?
5. Answer THEN WHAT for two different opening moves. Which leaves both of you with better options?
6. Finish this sentence: "I was about to walk into this thinking only about

. What's really going on is

."

Exercise C — Partner drill / role-play

Format: Pairs, about 15 minutes. One person plays the Employee and one plays the Manager. An observer is optional, or you can practice with an AI practice partner by giving it the Manager's private brief.

Setup (both see this): The Manager has just sent the Employee a text: "We need to talk about your performance." The meeting starts in five minutes.

Employee's private brief: You don't know what this is about. You have an unmentioned informal offer from another company, and you think your major project last month went unnoticed.

Manager's private brief: You wrote that text in a hurry. You actually want your strongest performer, the Employee, to take over a struggling project. Your director is pressing you to fix it this quarter, and you can't offer a raise until the next budget cycle.

Round 1 (3 minutes): Each player fills out a quick Game Map from their own position, noting what they know, don't know, and assume.

Round 2 (5 minutes): Hold the conversation. The Employee's job is to understand the game before defending themselves: ask what the Manager wants and who else is involved.

Round 3 (5 minutes), debrief:

- Employee: How much of your map was wrong? What did you assume from the wording of the text?
- Manager: How did your words change the Employee's decision problem? What would you send instead?
- Both: How did the players outside the room (the director, the other company) and each side's constraints shape the conversation? Where were your interests aligned?

If time allows, switch roles and rerun it with the Manager stating the real purpose up front. Notice how the game changes.

Quiz

1. In the game-theoretic sense used in this course, a "game" is best defined as:
 - A) A competition in which one person wins and the other loses
 - B) A situation in which the outcome depends, at least partly, on the decisions of multiple people
 - C) Any interaction in which someone is trying to manipulate someone else
 - D) A formal model that requires calculating an equilibrium
2. According to the chapter, what question does game theory ask before "Was that the right thing to say?"
 - A) "How do I build more rapport?"
 - B) "Which persuasion technique fits here?"
 - C) "What happens after I say it?"
 - D) "Who is right?"
3. A vendor tells you, "I'd love to give you that price, but I simply can't go below \$10,000." The approach the chapter recommends is to:
 - A) Assume they're bluffing and push harder
 - B) Accept the number, since they've stated a constraint
 - C) Ask what actually constrains them, because a bluff, a real policy, and an unwilling-but-able seller look identical on the surface
 - D) Walk away immediately to show strength
4. Two business partners both want their company to succeed but disagree about pay, reinvestment, and credit. Their interests are best described as:
 - A) Fully aligned

- B) Fully opposed
 - C) Mixed
 - D) Irrelevant to the game
5. Which statement about outside options matches the chapter?
- A) The person with the better alternative always wins.
 - B) Alternatives matter only after a negotiation fails.
 - C) Alternatives can change the negotiation itself, but their effect depends on the structure of the game.
 - D) Alternatives are a persuasion tactic you should hide.
6. A manager says, “My employees aren’t cooperating.” Their pay system rewards outperforming one another and nobody knows who owns key tasks. The chapter would most likely call this:
- A) A person problem, because the employees are selfish
 - B) A possible game problem, because the structure may be producing the behavior
 - C) A communication problem that better wording would solve
 - D) Proof that the employees bear no responsibility
7. You and a coworker both want a project finished. You think they’re procrastinating, and they think you’ll take over if they start. What does this example illustrate?
- A) Your coworker has a character flaw.
 - B) Shared goals guarantee smooth cooperation.
 - C) Expectations about each other can create a problem even when both people want the same outcome.
 - D) The situation is purely competitive.
8. Which statement best reflects the evidence on “seeing the game”?
- A) Studies show that people who use a Game Map always make better decisions.
 - B) Research shows that how people construe and notice a strategic situation can affect their behavior, but it hasn’t been shown that teaching people to see the game improves real-world communication.
 - C) The Wall Street/Community Game study proves that labels don’t matter.
 - D) Game theory assumes people are perfectly rational, and experiments confirm it.

Field assignment

This week: map one conversation before you have it.

Choose one meaningful interaction coming up this week. Before it happens, spend 10 minutes on a Game Map instead of a script, and mark which answers are assumptions rather than facts.

During and after the conversation, observe and record: - Any player who turned out to matter that you hadn't listed - Any constraint or alternative that was different from what you assumed - One moment where the other person's response changed your next move (THEN WHAT in action) - Whether the stated topic was the whole issue, or whether there was a game behind the conversation

Debrief questions for next week: 1. What did you notice about the situation that you would have missed if you'd only prepared what to say? 2. Which assumption was most wrong, and what made the other person's behavior more reasonable from their position? 3. Was the main obstacle a person problem or a game problem, and what would change the game next time?

The question to carry

| What's really going on here?

My notes

Module 2 — See the Payoffs

The question: What payoff would make their behavior make sense?

The tool: The Payoff Map

Use this map before any important interaction. It sits on top of your Game Map from Module 1.

The interaction:

Players (from your Game Map):

1. What do I want?

What outcomes matter to me? (List more than one currency.)

2. What do they want?

What outcomes appear to matter to them? (Not what they should want, not what I'd want in their shoes, not only what they said.)

3. Where do our interests align?

What can we both gain?

4. Where do they conflict?

What does one person gain at the other's expense?

5. What are we each trying to avoid?

What losses or risks matter? (Me / Them)

6. What might I be missing?

Could status, identity, fairness, autonomy, or relationships matter here?

7. Interest structure (circle one): Mostly opposed · Mostly aligned · Mixed

The most important question: Am I assuming their payoff is the same as mine?

My next move, given this map (and one that respects their autonomy and informed choice):

Practice exercises**Exercise A — Analyze the scenario**

Dana runs a small marketing agency. Her best designer, Marcus, has been there four years. Dana offers him a promotion to creative director with a 15% raise, and she expects him to be thrilled. Instead, Marcus hesitates and asks for a week to think. In a hallway conversation he mentions that the new role “means a lot more meetings” and that he’d “stop doing the actual design work.” A few days later Dana learns from a colleague that Marcus was hurt that the promotion was announced to the leadership team before it was offered to him. Dana is frustrated: “I’m offering him more money and a better title. Why is he being so difficult?” She is considering adding another 5% to close the deal. Marcus has

said several times that he “doesn’t care about titles,” but last year he turned down a higher-paying offer from a larger firm.

1. What payoff is Dana assuming Marcus has, and what is the evidence that his payoff function is different?
2. List at least three currencies that may be in play for Marcus, including at least one he’s trying to avoid.
3. Is the interest structure between Dana and Marcus opposed, aligned, or mixed? Where specifically?
4. What does his past behavior (the turned-down offer) reveal compared with what he has stated?
5. Why might adding another 5% fail to solve the problem, and what would a better next move look like?

Exercise B — Your own game

Pick a live situation where someone’s behavior frustrates or puzzles you: a coworker resisting a plan, a family member refusing a reasonable request, a client who won’t commit.

1. Write down the sentence you’ve been thinking (“Why would they...” or “They’re being...”), then rewrite it as “What payoff would make that behavior make sense?”
2. Complete the Payoff Map for this situation.
3. Answer: “What outcome would make this person feel that they got what they wanted?” and “What outcome would make them feel that they lost?”
4. Note one place where you’ve treated their stated preference as the whole story, and one choice of theirs that reveals something else.
5. Decide whether the interests are opposed, aligned, or mixed. If they’re genuinely opposed, write that down plainly; knowing what kind of game you’re in is useful too.
6. Draft one honest, open question that would test your map rather than confirm it.

Exercise C — Partner drill / role-play

Time: 15 minutes. **Format:** pairs, or one person with an AI practice partner.

Setup: Two co-owners of a small bakery are deciding what to do with this year’s \$60,000 profit.

Jordan (private): You want to reinvest most of it in a second location. Growth matters a lot, and your biggest fear is the business plateauing and looking small-time to people who doubted you. You’re comfortable with risk.

Riley (private): You want to distribute most of it. You’re quietly helping a parent with medical bills, so financial security matters a lot and debt makes you anxious. You’d like the business to grow, just not in a way that risks your income this year.

Round 1 (5 min): Argue positions only. Try to convince each other.

Round 2 (7 min): Start over. Your only job now is to fill in the other person's Payoff Map by asking: What are you trying to gain? What are you trying to avoid? What would feel like a loss? Share only what you're comfortable sharing. Then look for a deal that creates value before dividing it.

Debrief questions:

- In Round 1, what did you assume the other person was optimizing for? How close were you?
- What payoffs surfaced in Round 2 that never came up in Round 1?
- Did you find an arrangement (for example, a smaller expansion plus a guaranteed distribution) that gave each of you more of what you cared about most? If not, was the conflict real?
- Did anyone feel pushed or "worked"? What separates understanding a payoff from exploiting it?

Quiz

1. According to the chapter, what is a payoff?
 - A) The amount of money a player receives at the end of a game
 - B) What makes an outcome better or worse for a particular player
 - C) A hidden psychological motive that people try to conceal
 - D) The reward you offer someone to change their behavior
2. Two business partners disagree about reinvesting versus distributing profits. One prioritizes growth, the other security. What best describes this?
 - A) One partner is being irrational
 - B) It's purely a disagreement about money
 - C) They are optimizing for different things
 - D) It's a zero-sum game
3. A buyer and seller disagree on price but both want the sale to close. This is best described as:
 - A) Opposed interests
 - B) Aligned interests
 - C) Mixed interests
 - D) A coordination problem only
4. In the Ultimatum Game, responders often reject offers they see as unfair, leaving both with nothing. What is the chapter's lesson?
 - A) The mathematical model is useless
 - B) People don't care about money

- C) Fairness can be part of the actual payoff structure
 - D) Responders misunderstand the rules
5. Someone says “Price is my only concern,” then rejects a cheaper option because it takes longer to deliver. What does the chapter suggest?
- A) They are lying and can’t be trusted
 - B) Their behavior reveals that price may not be their only concern
 - C) Delivery time is irrelevant to the game
 - D) You should ignore what people say entirely
6. Two employees get the same new responsibility. One feels trusted; the other feels dumped on. Which distinction does this illustrate?
- A) Stated versus revealed preferences
 - B) What happened versus what each person thinks happened
 - C) Zero-sum versus positive-sum games
 - D) Short-term versus long-term payoffs
7. A company switches from individual bonuses to a team bonus, and two rivals start helping each other. What changed?
- A) The people’s character
 - B) Nothing; they were always cooperative
 - C) The game and its incentives
 - D) Their stated preferences
8. Which response best reflects the chapter’s ethical use of payoff thinking?
- A) “Find their hot button and push it.”
 - B) “Once I understand their incentives, their behavior is justified.”
 - C) “Understand the incentive structure before choosing your move, without exploiting it.”
 - D) “There is always a win-win if I look hard enough.”

Field assignment

This week, catch yourself in one “Why would they do that?” moment at work or at home. Before you respond or judge, sketch a quick Payoff Map, even a few lines on your phone. Then, if appropriate, ask one honest, open question to test it, such as “What would a good outcome look like for you here?” or “What are you most worried about?” Record the behavior, your first interpretation, the payoffs you guessed (gain and avoid), and what you learned.

Debrief questions for next week: 1. Where was your first interpretation based on assuming their payoff was the same as yours? 2. What payoff did you miss, and which

currency was it? 3. Did understanding their payoffs change your next move, or confirm that your interests were genuinely opposed?

The question to carry

What payoff would make their behavior make sense?

What are they trying to gain? What are they trying to avoid? Where are our interests aligned? Where are they genuinely opposed? What am I assuming they value that they may not?

My notes

Module 3 — Think in Responses

The question: “If I do this, what are they likely to do next? And if they do that, what will I do?”

The tool: The Response Map

Part 1 is the chapter’s practical exercise for an interaction that went sideways. Part 2 turns the chapter’s questions into a pre-move check.

Part 1 — Rewind a past interaction

Pick a recent interaction that didn’t go the way you wanted. Don’t judge who was right yet; just write the sequence.

1. What did I do?

2. What did they do?

3. What did I do next?

4. What did they do next?

5. (Continue if needed)

- At which point did the interaction change direction?

- What could I have predicted about their response before making my move?

Part 2 — Map an upcoming move

1. **My planned move** (what I'm about to say or do):

2. **What I want them to do** (my intention):

3. **Their side of the board.** Given their preferences, not mine:

Their goals:

What they know (and don't):

Their options:

Their constraints (boss, money, time, reputation, emotions):

4. **What does this move make them want to do?** Several plausible responses, each marked:

Response A:

(likely / possible / unlikely)

Response B:

(likely / possible / unlikely)

Response C:

(likely / possible / unlikely)

5. **Intention check:** Does my move actually make the response I want more attractive than their alternatives?

6. **If they do that, what will I do?** (One line per response above.)

7. **Where does the sequence lead?** (High stakes only.)

8. **Eliminate:** Is any option worse for me regardless of what they do?

9. **Threat check:** If a threat is involved (theirs or mine), would carrying it out be the best response when the moment comes?

10. **Pattern and relevance:** What pattern will this move create? Which response could materially change the outcome?

11. **Afterward — update:** What did they actually do? Where was my prediction off, and what did I misunderstand about the game?

Practice exercises

Exercise A — Analyze the scenario

Dana leads a five-person operations team at a logistics company. At her first team meeting she said, “If something goes wrong, tell me right away. No surprises.” Two weeks later, Marcus, a junior coordinator, admitted he had booked a shipment to the wrong warehouse, delaying a client order by two days. Dana was under pressure from her own boss that week. She raised her voice, called it “exactly the kind of careless mistake I can’t afford,” and later cited the error, by name, in the weekly team meeting. Marcus fixed the booking and apologized, and Dana felt she had set a clear standard. A month later, a much larger error surfaced only when the client called to complain; two team members had known

about it for a week. Dana told her boss, “I don’t understand it. I specifically told them to come to me.”

1. What did Dana intend her reaction to produce, and what did it actually make attractive for the team?
2. Write the sequence of moves. Where did the interaction change direction?
3. Why is “I specifically told them to come to me” beside the point strategically?
4. Reasoning backward: if Dana wants early disclosure, what must the team expect, and where does that expectation come from?
5. What could Dana do now, and what responses should she anticipate?

Exercise B — Your own game

Choose a live situation where you’ll make an important move in the next week or two: a request, hard feedback, a boundary, a negotiation, or a reply you’re tempted to send fast.

1. Complete Part 2 of the Response Map. If you catch yourself writing “If I were them, I’d...”, rewrite it from their payoffs.
2. List at least three plausible responses and mark each likely, possible, or unlikely. Replace certainty with branches.
3. For the most likely response, write what you’ll do, then what they’ll likely do next. Stop there unless the stakes are high.
4. Intention check: does your move make the response you want genuinely more attractive to them? If not, what change (timing, setting, framing, what you offer or ask) would?
5. Elimination check: cross off any option that’s worse for you no matter how they respond.
6. Afterward, complete step 11 and write one sentence on what you misunderstood about the game, if anything.

Exercise C — Partner drill / role-play

Format: Pairs (or one learner with an AI practice partner), 12–15 minutes.

Setup: Jordan, a small business owner, wants to cut a freelance designer’s monthly rate. Riley, the designer, wants to keep it.

- **Jordan (private):** You want a 15% cut and have told yourself you’ll “walk away” if Riley refuses. But replacing Riley would take at least six weeks, and your product launch is in four.
- **Riley (private):** Other clients fill about 70% of your capacity. Losing Jordan would hurt but not sink you. You suspect Jordan has a launch coming but don’t know when.

Round 1 (4 min): Jordan opens: “I need a 15% cut, or I’ll have to find someone else.” Play it out.

Round 2 (4 min): First, 60 seconds of silent prep. Jordan: “If Riley says no, will walking away really be my best response? What does my opening make Riley want to do?” Riley: “If I refuse, what will Jordan actually do at that moment?” Replay with new openings.

Debrief (5 min):

- Was Jordan’s Round 1 threat credible? How did Riley decide?
- Where did either player project their own preferences onto the other?
- Did Jordan’s Round 2 opening make the response Jordan wanted more attractive to Riley, or just state it more firmly?
- Ethics check: did anyone make a threat they didn’t believe? What might that cost in an ongoing relationship?

Quiz

1. According to the chapter, what is a best response?
 - A) The most polite and respectful reply
 - B) The action that gives a player the best outcome given what the other player is doing
 - C) The first reply that comes to mind under pressure
 - D) The action that would be best if the other player didn’t exist
2. A manager urges collaboration but pays bonuses only for individual results. Salespeople keep guarding their leads. The best diagnosis?
 - A) The team didn’t hear the message
 - B) The team is irrational
 - C) Their incentives make another response more attractive than collaboration
 - D) The manager needs to say it more forcefully
3. Which question does the chapter recommend for predicting a response?
 - A) “What would a reasonable person do?”
 - B) “What would I do in their position?”
 - C) “What should they do?”
 - D) “Given their goals, information, options, and constraints, what is the most likely response?”
4. Option A leaves you worse off than Option B no matter what the other person does. What does the chapter suggest?
 - A) Eliminate Option A without needing to predict them
 - B) Predict their move carefully before deciding
 - C) Choose A to keep them guessing
 - D) Ask them which option they prefer

5. Someone says, “If you do that, I’m walking away.” How should you assess the threat?
 - A) By how forcefully it was said
 - B) By assuming all threats are bluffs
 - C) By whether walking away would actually be their best response when that moment arrives
 - D) By making a bigger threat in return
6. Which statement best describes the escalation trap?
 - A) One person deliberately decides to make things worse
 - B) It only happens between people who dislike each other
 - C) The person who started it can always stop it
 - D) A bad outcome emerges from a sequence of locally understandable responses
7. In Johnson, Camerer, Sen, and Rymon’s bargaining study, what happened after participants were trained in backward induction?
 - A) Their offers to computerized opponents moved close to the theoretical prediction
 - B) Nothing changed; they kept offering about an equal split
 - C) They refused to make offers
 - D) The study proved backward induction predicts everyday conversation
8. You predicted a colleague would welcome your proposal, but they pushed back hard. What’s the best next step?
 - A) Conclude they’re being irrational
 - B) Ask, “What did I misunderstand about the game?” and update your model
 - C) Repeat the proposal until they agree
 - D) Stop making predictions about people

Field assignment

This week, run “Then what?” before three moves that matter (a request, a reply to a frustrating message, feedback, a “no”). Before each, take 60 seconds to write:

- My move:
- What I want them to do:
- Their two or three most plausible responses (with likely / possible / unlikely):
- If they do the most likely thing, what will I do?

Afterward, record what they actually did and whether your prediction was right, partly right, or wrong. Once this week, also notice a feedback loop in progress and note whether it was escalating or settling.

Debrief questions for next week: 1. How often did your prediction match the response? When it didn’t, what had you misunderstood about the game? 2. Did asking “Then what?” change any move before you made it? What did you change, and what happened? 3. In the

feedback loop you observed, where could one well-chosen response have changed the direction of the sequence?

The question to carry

“What is this likely to make them do? If they do that, what will I do? And if the stakes are high, where does that sequence lead?”

My notes

Module 4 — Manage Predictability

The question: What does my predictability enable the other player to do?

The tool: The Predictability Audit

Use this on one recurring interaction where a pattern may be shaping the game.

Step 1 — Name the interaction

The person or group:

The recurring situation (meeting, negotiation, household routine, request):

Step 2 — Find the pattern

What do I repeatedly do in this situation?

What have my past actions taught this person to expect from me?

Have I taught this person how to respond to me? If so, what is the rule they've learned (finish the sentence "If I , **she/he will** ")?

Did I choose this rule, or did my repeated behavior establish it?

Step 3 — Calculate the predictability tax

What does this pattern reveal that I didn't intend to reveal (my limits, availability, whether my threats or boundaries are real)?

Step 4 — Asset or liability?

Who benefits from knowing what I'm going to do? (Circle one: both of us / mainly them / mainly me)

Is this pattern helping me? Helping them? Helping both of us?

Could they exploit this pattern? How?

Check any that apply. This pattern currently: - ☐ creates trust, cooperation, or a credible commitment - ☐ helps others coordinate with me - ☐ shows people exactly how to pressure me - ☐ reveals threats, concessions, or unenforced boundaries

Step 5 — Choose your move

Would greater consistency create trust — or greater flexibility create strategic room?

Which of the four moves fits? (Circle one) - **Be predictable** — trust, coordination, credibility, or cooperation are valuable here. - **Protect my predictability** — revealing my pattern gives them leverage. - **Vary my behavior** — they could otherwise exploit a fixed pattern. - **Be deliberately consistent** — I want them to understand a consequence is real.

The specific behavior I'll change or reinforce:

Step 6 — Ethics check

Am I protecting legitimate information or providing false information?

Could my change destabilize someone or make me harder to trust where reliability matters?

Practice exercises

Exercise A — Analyze the scenario

Dana is a senior designer at a marketing agency. For two years, whenever account manager Rob brings her an “urgent” client revision on Friday afternoon, she stays late and finishes it. Lately the Friday requests have multiplied, and several had been sitting in his inbox since Tuesday. Dana has complained in meetings that “people need to plan better” and once told Rob she “wouldn’t keep doing this,” then stayed late anyway. Dana also reviews junior designers’ work every Thursday morning, and the juniors have built their weeks around it. Frustrated, she is considering becoming “harder to read” across the board so nobody takes her for granted.

1. What rule has Rob learned from Dana’s behavior, and did she ever choose to create it?
2. What is Dana’s predictability tax here? What information has her pattern revealed?
3. Which of Dana’s patterns is an asset and which is a liability? Why?
4. Is “becoming harder to read across the board” a good plan? What would you recommend instead, using the four moves?
5. How can Dana change the pattern without deception?

Exercise B — Your own game

Pick one live situation where someone has likely learned to predict you — a coworker, negotiation partner, client, or family member. Then:

1. Complete Steps 1–3 of the Predictability Audit. Write the rule they’ve learned in one sentence, starting with “If I , **they will** .”
2. In Step 4, list one way this pattern helps both of you and one way it could be exploited. No exploitation risk? This may be a pattern to protect.
3. In Step 5, choose one of the four moves. Write the exact words or behavior you’ll use the next three times this comes up.
4. Run the Step 6 ethics check. If your plan involves saying anything false, rewrite it so you are withholding rather than misrepresenting.

Exercise C — Partner drill / role-play

Format: Pairs, 15 minutes (or with an AI practice partner). Two rounds.

Round 1 — The concession pattern (6 min)

Setup: A freelance photographer (Seller) is quoting a corporate event. A procurement manager (Buyer) wants the lowest price.

Seller’s private card: Your opening quote is \$4,000. Your true minimum is \$3,000. Follow a fixed rule: every time the Buyer rejects your number, drop \$300. If asked your minimum, tell the truth.

Buyer's private card: You have budget up to \$4,000 but want to pay as little as possible. Figure out the Seller's pattern and use it. Ask once, "What's the lowest you'll accept?"

Play for four minutes. Then the Buyer states the rule they think the Seller was following.

Round 2 — Managing predictability (6 min)

Same roles, same private numbers, but the Seller now has no fixed rule. The Seller must (a) decline to disclose the bottom line honestly, (b) tie any movement to a reason or a change in scope rather than to rejection alone, and (c) say nothing false. The Buyer again tries to find a pattern.

Debrief (3 min)

- In Round 1, how many rejections did it take the Buyer to learn the pattern? What did the Seller's behavior reveal that their words didn't?
- In Round 2, how did "I'm not prepared to disclose my bottom line" change the Buyer's decision-making?
- Did the Seller ever say anything untrue? Where was the line between protecting information and deceiving?

Quiz

1. According to the chapter, what is the core principle of managing predictability?
 - A) Never let people know what you're going to do.
 - B) Always be consistent so people trust you.
 - C) Don't allow your patterns to become so predictable that other people can consistently use them against you.
 - D) Randomize your behavior in every competitive situation.
2. In rock-paper-scissors, you know your opponent chooses rock 70 percent of the time. What does this illustrate?
 - A) You need to know their exact next move to gain an advantage.
 - B) Knowing someone's pattern is enough to adjust your behavior and gain an advantage.
 - C) Mixed strategies only work in games of pure chance.
 - D) Predictability matters only when the other player is irrational.
3. A negotiator always concedes from \$100 to \$90 to \$80. What is the "predictability tax" here?
 - A) The cost of making three offers instead of one.
 - B) The pattern reveals there's more room, so the other side keeps pushing.
 - C) The negotiator must pay a fee for each round.
 - D) The other side will refuse to negotiate.

4. Someone interrupts you in three meetings and you let it go each time. What has most likely happened?
 - A) Nothing; one-off choices don't create patterns.
 - B) You've established a rule — interrupting you has no cost — even though you never intended to.
 - C) You've used a mixed strategy effectively.
 - D) You've built trust that will reduce future interruptions.
5. Which situation is the best example of predictability creating value?
 - A) A competitor always matches your price cuts.
 - B) A parent who always gives in after a third request.
 - C) A manager who says rule violations will have consequences and consistently enforces it.
 - D) A poker player who bluffs frequently.
6. A buyer asks, "What's the lowest price you'll accept?" Which response best fits the chapter?
 - A) Give a number well above your real minimum and call it your floor.
 - B) Give your true minimum to show good faith.
 - C) "I'm not prepared to disclose my bottom line."
 - D) Change the subject without acknowledging the question.
7. You're trying to rebuild trust with a friend after missing several commitments. What does the chapter recommend?
 - A) Become less predictable so they stop taking you for granted.
 - B) Become more predictable: if you say you'll call, call.
 - C) Make fewer statements so there's less to be held to.
 - D) Vary your follow-through so they can't exploit you.
8. In a repeated Prisoner's Dilemma, you're known to always retaliate against defection. According to the chapter, this predictability:
 - A) Always makes you exploitable.
 - B) Is irrelevant because the game is repeated.
 - C) Can support cooperation — showing that the same predictability that's exploitable in one game can make you trustworthy in another.
 - D) Should be replaced with random retaliation.

Field assignment

This week: run a pattern log. Choose one recurring interaction — a weekly meeting, a regular request from a colleague, a household routine, a client relationship. Each time it happens, write down three things: what you did, how the other person responded, and what they might have learned about you from that exchange.

By midweek, complete the Predictability Audit, choose one of the four moves, and carry it out on purpose for the rest of the week. Stay honest throughout — no false statements.

Record: - The rule you think the other person had learned before the week started. - Any sign that they noticed a change (testing, surprise, adjustment, relief). - Whether the change made the interaction better for you, for them, or for both of you.

Debrief questions for next week: 1. What had your past actions taught this person to expect from you — and how accurate was your guess once you watched closely? 2. Did they adapt to your new pattern? What does that tell you about the feedback loop of action, response, learning, and adaptation? 3. Where in your life would greater consistency create trust, and where would greater flexibility create strategic room?

The question to carry

Have I taught this person how to respond to me? What have my past actions taught this person to expect from me?

My notes

Module 5 — See the Equilibrium

The question: What equilibrium have we created?

The tool: The Equilibrium Audit

Use this when you're stuck in a recurring interaction. Work through it in writing, at a calm moment.

The situation (one line):

1. What keeps happening? Forget the individual incident. Describe the repeated pattern as a loop (my move, their response, my response...).

2. What does my behavior cause them to do?

3. What does their behavior cause me to do?

4. Why might neither of us want to change unilaterally? If I changed alone, what would I risk? If they changed alone, what would they risk? Me:

Them:

5. Is this pattern actually good for either of us? What does each of us get from maintaining it? What does it cost each of us?

6. Are there alternative stable patterns? Describe at least one other pattern that could hold steady if we both settled into it.

Equilibrium check: Is this really equilibrium-like, or is it held in place by something else (habit, a rule, coercion, lack of awareness)? Would changing alone actually make someone worse off?

The most important question: What would have to change for a different pattern to become attractive? Check any levers that apply, then describe one: - ☐ My response - ☐ The sequence (who moves first, when) - ☐ The incentives - ☐ New information - ☐ A commitment - ☐ A third party - ☐ Expectations about what the other will do - ☐ The available choices

My first experiment:

Practice exercises

Exercise A — Analyze the scenario

Priya and Marcus co-lead a product team. Every Friday, Priya sends Marcus a long status email flagging risks she sees in his part of the project. Marcus has started replying with a short “Noted, thanks” and handling the issues on his own without updating her. Because she hears nothing, Priya assumes the risks aren’t being handled, so the next Friday’s email is longer and copies their director. Marcus reads the cc as Priya trying to make him look bad, so he shares even less and starts sending his own updates straight to the director. Both say they want a real partnership, and both say the other started it. Last month Marcus tried skipping his separate update to the director “to be the bigger person,” and the director asked him in a meeting why he’d gone quiet, which embarrassed him.

1. Describe the pattern as a loop of mutual responses.
2. Why does each person’s current behavior make sense given the other’s?
3. Why did Marcus’s unilateral change backfire, and what does that show about this pattern?
4. Is this a formal Nash equilibrium? How should you describe it honestly?
5. What would have to change for a different pattern to become attractive to both?

Exercise B — Your own game

Choose a recurring interaction that frustrates you (the same argument, a standoff with a coworker, a stalled negotiation), ideally one where you're tempted to say, "It's their fault."

1. Fill in the Equilibrium Audit. Write the pattern as a loop at least five steps long.
2. For question 4, name the exact cost you fear if you change alone, then the cost you think they fear. Mark your guess about theirs "untested."
3. Run the equilibrium check: is it really the cost of changing alone holding this in place?
4. Name one alternative stable pattern. Pick one lever and design a small, honest experiment, ideally one that lowers the other person's risk of changing too.
5. Complete: "Regardless of who started it, what keeps it going now is ____." Keep this page for the field assignment.

Exercise C — Partner drill / role-play

Time: 15 minutes. **Players:** two people plus an optional observer (or one person with an AI practice partner).

Setup: Two roommates. The kitchen is always messy. Each player can "Clean" or "Wait."

Player A privately knows/wants: You want a clean kitchen, but you cleaned three times last month and your roommate never did. "If I clean and they don't, I'm the doormat."

Player B privately knows/wants: You want a clean kitchen too, but whenever you've cleaned, your roommate redid it "the right way," which felt like criticism. "If I clean, I'll get corrected. Better to wait."

Round 1 (4 min): Play the usual conversation in character, each making the move that seems sensible from their private view. The observer writes down the loop.

Round 2 (4 min): Player A changes unilaterally (cleans without comment) while Player B keeps the same beliefs. Notice what happens.

Round 3 (5 min): Step out of character for 60 seconds and share "what I fear if I change alone." Then replay, looking for a structural change: a schedule (sequence), a shared standard (information), a simultaneous commitment, or a different split of chores (available choices).

Debrief questions: 1. In Round 1, did either player behave irrationally given what they believed? What kept the pattern stable? 2. What happened to the unilateral change in Round 2? Did it shift the other player's expectations, or get misread? 3. Which lever in Round 3 made a different pattern attractive to both? Why did it work better than good intentions?

Quiz

1. In game theory, a Nash equilibrium is best described as:
 - A) The outcome that gives everyone the highest possible payoff
 - B) A combination of strategies in which each player's strategy is a best response to the others'
 - C) A fair agreement that all players have explicitly accepted
 - D) The outcome most people will choose in the real world
2. Which statement about equilibrium is true according to the chapter?
 - A) If a pattern is stable, it must be good for at least one person
 - B) An equilibrium requires an explicit agreement
 - C) A pattern can be stable even if everyone involved dislikes it
 - D) Equilibria only occur when someone acts irrationally
3. In the Prisoner's Dilemma, why do both players end up defecting?
 - A) They don't realize mutual cooperation would be better
 - B) Regardless of what the other does, defecting looks like each player's best response
 - C) One player tricks the other into defecting
 - D) The rules forbid cooperation
4. Dana and her brother haven't called each other in months. Dana says, "He never calls me." Her brother says, "I'm always the one who has to reach out." Which question best reflects this module's approach?
 - A) "Which of them is being unreasonable?"
 - B) "What should Dana say on the next call?"
 - C) "What pattern are their responses creating, and why is it stable?"
 - D) "How can Dana get her brother to admit he's wrong?"
5. A friend says, "My partner and I fight constantly. I'm just going to stop engaging." Based on the chapter, what's the most important caution?
 - A) Stopping engagement always resolves conflict
 - B) A unilateral change may initially make things worse if the partner reads silence as abandonment
 - C) The friend should engage even more forcefully instead
 - D) Nothing can change once an equilibrium exists
6. What is the "equilibrium-selection problem"?
 - A) Choosing which player is responsible for the bad pattern
 - B) Knowing the possible stable outcomes without knowing which one people will land on
 - C) Selecting the best possible outcome for yourself
 - D) Deciding whether to use game theory at all

7. A new hire at a firm where everyone stays until 7 p.m. starts staying late too, worried that leaving at five would look uncommitted. This best illustrates:
 - A) An equilibrium sustained by expectations about what others will do
 - B) A written company policy
 - C) A one-shot Prisoner's Dilemma with a unique cooperative outcome
 - D) Proof that everyone at the firm prefers working late
8. The Charness, Rigotti, and Rustichini experiments found one-shot Prisoner's Dilemma cooperation ranging from 23% to 60% depending on the payoffs. What is the most accurate takeaway?
 - A) Nash equilibrium perfectly predicts human behavior
 - B) Game structure is irrelevant to how people behave
 - C) Payoff structure shapes behavior, but the simple money-only Nash prediction doesn't fully describe real people
 - D) People always cooperate when given the chance

Field assignment

This week: catch one equilibrium in the wild and run one experiment.

1. **Observe (days 1 to 3):** Take your Exercise B situation or a new one. Each time it happens, jot down the loop: their move, your response, their response. Don't fix anything yet. Note when you reacted to what you expected them to do rather than what they did.
2. **Audit (day 4):** Update the Equilibrium Audit from your notes, including the equilibrium check.
3. **Experiment (days 5 to 7):** Pick one lever and try one small, honest change, preferably one that doesn't leave the other person carrying all the risk of going first. For example, name the pattern without blame: "I think we're in a loop where I ask late and you say no, and we both end up frustrated. Could we try a weekly check-in?" Don't use this to corner or pressure anyone. The aim is a better pattern both of you would choose freely.
4. **Record:** What happened, and what did you learn about what keeps the pattern stable?

Debrief questions for next week: 1. When you mapped the loop, what did you see that you hadn't seen when you looked only at individual incidents? 2. What made changing alone risky, for you and, as best you could tell, for them? Were your guesses about their fears confirmed or wrong? 3. Did your experiment change the other person's expectations? If not, what would you change about the structure next time?

The question to carry

What equilibrium have we created? Why is it stable? What does each of us get from maintaining it? What makes changing it risky? Are there other equilibria available? And what would have to change for us to move toward one of them?

My notes

Module 6 — Map Information & Beliefs

The question: “What do they know—or believe—that I don’t?”

The tool: The Information Map

Use this before any important interaction, or after one that confused you. The first seven questions are the chapter’s Information Map. The sorting table and the final checks help you keep facts, beliefs, and assumptions in separate columns.

The situation:

The other player(s):

Part 1 — The seven questions

1. **What do I know?** (Separate facts from assumptions.)

2. **What don’t I know?** (Identify the uncertainty.)

3. **What do they know that I don’t?** (Look for information asymmetry.)

4. **What do I know that they don’t?** (Recognize your own informational advantage.)

5. **What do I think they believe?** (Model their beliefs.)

6. **What do they think I believe?** (If the stakes are high, this matters too.)

7. What evidence could change my mind? (Protect yourself from your own assumptions.)

Part 2 — Sort it

Layer	Question	My answer
Facts	What do I actually know happened? (Observable, verifiable.)	
Unknowns	What relevant information am I missing?	
Beliefs	What do I currently think is true?	
Assumptions about them	What am I assuming they know or believe?	

For each belief, mark your confidence: **I know** / **I have good evidence** / **I suspect** / **I'm guessing**.

Part 3 — Shared or common?

- Is the key information actually in their information set? What evidence do I have that they know it?

- Do we both know that we both know it, or am I assuming it's shared?

Part 4 — Messages I'm relying on

- What did they tell me that matters?

- What do they gain if I believe it? What would it cost them if it were false?

- How could I verify it?

- **Why should I believe it?**

Part 5 — Decision-relevant information

- What single piece of information would change my decision, or theirs?

- How can I get it (ask, observe, verify, share)?

Practice exercises

Exercise A — Analyze the scenario

Dana manages a five-person analytics team. On Friday, the quarterly report Jordan owned went out two days late, and Dana's director asked her about it in front of other managers. Dana had told the team on Monday that "the director is watching this one closely," but she said it in a quick hallway exchange that Jordan wasn't part of. Jordan had sent a message mid-week saying "Still waiting on finance's numbers," and Dana, skimming it on her phone, read it as a routine status update. On Friday afternoon, Dana writes to Jordan: "We need to talk Monday about ownership." Jordan, who stayed late Thursday rebuilding the model after finance finally delivered, reads the message as a sign that Dana thinks he's careless. Over the weekend, he tells a friend, "She already knows finance was late. She just wants someone to blame." Dana, meanwhile, believes Jordan didn't flag the risk because he doesn't take the director's priorities seriously.

1. Sort what Dana knows (facts) from what she believes and what she's assuming.
2. What information asymmetries exist in each direction?
3. Identify at least one higher-order belief and one place where Dana treated shared information as common knowledge.
4. How could the beliefs on both sides become self-fulfilling?
5. What information would change either person's decision, and how should Dana open Monday's conversation?

Exercise B — Your own game

Pick a live situation where someone's behavior is confusing, frustrating, or worrying you: a coworker who has gone quiet, a family member who reacted oddly, a client who stopped responding, a negotiation that stalled.

1. Write one sentence describing what happened, using only observable facts. No adjectives about motives.
2. Write your current interpretation. Label it: "Hypothesis."
3. Work through all five parts of the Information Map.
4. List at least three other explanations for the behavior. Use the chapter's list: are they withholding, do they not know themselves, is the information unavailable, have you not asked, or have you misread ambiguous evidence?
5. For your hypothesis, write one piece of evidence that would support it and one that would contradict it.
6. Identify one thing you know that they may not, and decide whether sharing it would change the game for the better.
7. Decide on a single next step that gathers decision-relevant information honestly, such as a direct question or a check-in, rather than a move based on an untested belief.

Exercise C — Partner drill / role-play

Time: 15 minutes. **Players:** two people (or one person plus an AI practice partner), plus an optional observer.

Setup: A small business owner (Player 1) is considering hiring a freelance designer (Player 2) for a three-month website project.

Player 1 privately knows: The budget is firm at the quoted amount, but the deadline is more flexible than the posting suggests. A previous freelancer missed deadlines badly, so Player 1 is nervous about reliability. Player 1 has one other candidate, who is cheaper but less experienced.

Player 2 privately knows: They're strong at design but have never built a site this size. They have another offer that starts in six weeks and would overlap. They really want this client for their portfolio.

Round 1 (5 min): Have the conversation. Each player tries to learn what matters and decides what to share. No lying about material facts; you may keep things private.

Round 2 (5 min): Stop. Each player silently fills in Part 1 of the Information Map: what I know, what I think they know, what I think they believe about me. Then compare maps out loud. Where were you right? Where were you wrong?

Round 3 (5 min): Replay the last part of the conversation, this time asking the one question whose answer would most change your decision.

Debrief questions

- Which of your beliefs about the other person turned out to be assumptions?
- Which claims did you treat as strong evidence, and which as cheap talk? What made the difference?
- What did you keep private, and was that different from deceiving? Did the other person have what they needed for an informed choice?
- What single piece of information changed the game most when it came out?

Quiz

1. According to the chapter, what do people base their decisions on?
 - A) The objective facts of the situation
 - B) The game as they perceive it
 - C) Whatever the most powerful player tells them
 - D) Their payoffs alone, regardless of beliefs
2. In the bank-run example, what helps produce the outcome people fear?
 - A) The bank's hidden accounting fraud
 - B) Government regulation
 - C) What depositors believe other depositors will do
 - D) A single large withdrawal by an insider
3. A coworker says, "Trust me, this vendor is the best option." Which question from the chapter best tests the claim?
 - A) "Is my coworker a nice person?"
 - B) "What does this person have to lose by making this claim?"
 - C) "How confidently did they say it?"
 - D) "Have I heard this vendor's name before?"
4. Which statement about cheap talk matches the chapter?
 - A) It is always worthless because it is costless.
 - B) It is the same thing as lying.
 - C) It can convey meaningful information, especially when interests are aligned.
 - D) It only exists in formal games, not in everyday life.
5. "I think my manager believes I think she's micromanaging me." This is an example of:
 - A) A costly signal
 - B) A higher-order belief
 - C) Common knowledge
 - D) Information asymmetry about type

6. A team lead tells three of six members that Friday's meeting is canceled. The other three show up. What was missing?
 - A) Payoffs
 - B) A costly signal
 - C) Common knowledge that the information was shared
 - D) A focal point based on reputation
7. Your partner hasn't replied to "Can we talk?" for three hours. What does the chapter recommend first?
 - A) Assume they're avoiding you and respond in kind
 - B) Send several more messages until they answer
 - C) Treat "they're angry" as a hypothesis and ask what evidence would support or contradict it
 - D) Assume they're deliberately withholding information
8. What does the Costa-Gomes and Weizsäcker experiment suggest?
 - A) People always best respond to their stated beliefs.
 - B) Stated beliefs are meaningless.
 - C) People often fail to best respond to their own stated beliefs, so knowing a belief doesn't guarantee predicting behavior.
 - D) More information always leads to better choices.

Field assignment

This week, catch yourself in one moment of confusion about someone's behavior: an unanswered message, a cool response in a meeting, an unexpected "no." Before you react, take two minutes and run the Information Map, at least Part 1 and the confidence labels in Part 2. Then take one honest step to get decision-relevant information: ask a direct question, share something you know that they might not, or verify a claim you were relying on.

Record: the situation; your first interpretation; your confidence label before; what you learned; and whether your belief survived, changed, or collapsed.

Debrief questions for next week 1. How far apart were your first interpretation and what you eventually learned? 2. Did you discover any information the other person lacked, or any you lacked? What changed once it was shared? 3. Where else in your life might a belief, yours or theirs, be sustaining a pattern that new information could change?

The question to carry

"What do they know—or believe—that I don't?"

Then ask: What do I actually know? What am I assuming? What might they be assuming? What information would change either of our decisions? And when someone tells you something important, don't stop at "What did they say?" Ask: **"Why should I believe it?"**

My notes

Module 7 — Separate Words From Credibility

The question: What makes this message credible?

The tool: The Credibility Test

Use this for high-stakes claims, not for directions to the restroom.

The claim (in their words):

What decision of mine does it affect?

Part 1 — The eight questions

1. **Do they know?** Are they in a position to possess this information?

2. **What do they want me to believe?**

3. **What happens if I believe them?** How does my response affect their payoff?

4. **Would they say this even if it were false?** If yes, the words alone may not be highly informative.

5. **What evidence supports the claim?** Assertion:

Verification:

6. What does their behavior suggest? (Evidence, not proof.)

7. What does their history suggest?

8. What would make this claim costly to fake?

Part 2 — Separate the dimensions (high / medium / low / unknown)

- Honesty: _____ because

- Knowledge: _____ because

- Incentive alignment: _____ because

Part 3 — Persuasive versus credible

- How convincing did it feel (confidence, charisma, emotion)?

- How credible is the underlying claim?

- Did it change my mind, or give me good reason to change my mind?

- What verification is available, and is it worth getting?

Part 4 — Speaker's side (when you're the one making the claim)

- What reason does the other person have to believe me?

-
- Is this a message problem or a credibility problem?
-

-
- How can I make my truthful claim easier to evaluate?
-
-

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small marketing agency and is choosing a web developer for a large client project. Marcus is polished and energetic. He says, “I’ve never missed a deadline, and I’m the best React developer in this city.” He never hesitates, has a beautiful slide deck, and offers 10 percent off if Dana signs this week. Asked for references, he says past clients are “under NDA, but trust me, they were thrilled.” Priya is quieter and a bit awkward. She says, “I think I can hit your timeline, but the payment integration is the risky part, and I’d want a week to scope it.” She sends three live sites she built, names two former clients Dana can call, and proposes milestone payments tied to working deliverables. Dana’s business partner, who knows Marcus socially, says, “Marcus is a great guy. He’d never lie to you.”

1. Apply “Would they say this anyway?” to Marcus’s two main claims.
2. Which parts of each pitch are cheap talk, and which are signals or verification?
3. Separate honesty, knowledge, and incentives for the partner’s endorsement.
4. What should Dana do next, without becoming cynical?

Exercise B — Your own game

Pick a live situation where a claim matters to a decision you’re about to make: a vendor’s promise, a colleague’s estimate, a contractor’s quote, or advice from family.

1. Write the claim in the speaker’s exact words and run Parts 1 through 3 of the Credibility Test.
2. Circle the weakest link: knowledge, incentives, evidence, behavior, or history.
3. Identify one piece of verification you could get this week, and decide whether the stakes justify it.

Then flip it. Pick a situation where you’re telling the truth but aren’t being believed, or soon need to be.

1. Complete Part 4. Honestly: would someone in your position say this even if it were false?

2. Write two ways to make your claim easier to evaluate, at least one the other person can check without relying on your word.
3. Check the ethics: does each step make the truth easier to see, or only make you look credible? Keep only the first kind.

Exercise C — Partner drill / role-play

Setup (15 minutes, pairs or threes). One **seller**, one **buyer**, an optional **observer**. The seller is selling a used bicycle at a fixed price.

Private information. The seller secretly draws a card. Card 1: “The bike is in excellent condition.” Card 2: “The frame has a hairline crack. You needn’t volunteer it, but you may not state a direct falsehood if asked.”

Round 1 (4 min). The seller makes claims; the buyer asks questions but may not request verification. The buyer then guesses the card and rates confidence from 1 to 10.

Round 2 (4 min). New card draw. The buyer may now request verification: an inspection, a return window, a call to the previous owner. A Card 2 seller must honor any guarantee they offer.

Debrief (5-7 min).

- In Round 1, did the buyer’s guess track the truth, or the seller’s confidence?
- Which statements would the seller have made regardless of the card?
- In Round 2, which offers were easier for a Card 1 seller? Why did those carry more information?
- Card 1 sellers: how did it feel to be doubted while telling the truth? What helped?

AI version. Ask an AI to play a seller who secretly picks a card and reveals it only at the end.

Quiz

1. According to the chapter, where does credibility primarily come from?
 - A) The precise wording of the sentence
 - B) The game surrounding the sentence: history, incentives, evidence, reputation, and verification
 - C) How confidently the message is delivered
 - D) Whether you like the speaker
2. A seller says the highest offer they’ve received is \$20,000. Which conclusion best reflects the chapter?
 - A) They benefit from saying it, so it’s false.
 - B) Treat it exactly like directions to the train station.

- C) They benefit from your believing it, so the statement alone gives you less reason to update.
- D) Ignore price claims entirely.
3. What did Farrell and Rabin argue about cheap talk?
- A) Costless messages are always meaningless.
- B) Costless messages always produce truth-telling.
- C) Both extremes are wrong; under the right incentives, cheap talk can carry substantial information.
- D) Cheap talk works only between family members.
4. In Cox and Stoddard's experiment, what happened when the incentive to exaggerate increased?
- A) Truthfulness declined, and the benefits of communication shrank.
- B) Truthfulness stayed the same.
- C) Players stopped sending messages.
- D) Efficiency improved.
5. A candidate says, "I'm extremely hardworking." Why is this weak evidence on its own?
- A) Candidates who say this are usually lying.
- B) Almost any candidate benefits from saying it, so it doesn't distinguish between types.
- C) Hard work can't be measured.
- D) Self-descriptions are never useful.
6. Which question best captures what makes a signal credible?
- A) How expensive was it?
- B) How dramatic was it?
- C) Why would it be easier or more worthwhile for someone whose claim is true than for someone whose claim is false?
- D) Did the person sound sincere?
7. You explained your timeline brilliantly, but the client still doubts you. You were late on the last three projects. What does the chapter suggest?
- A) Build a more persuasive story.
- B) Treat it as a credibility problem: ask how to make the claim more credible.
- C) Insist you're telling the truth.
- D) Offer a discount to change the subject.
8. What distinguishes signaling from commitment?
- A) Signaling uses words; commitment uses actions.
- B) Signaling asks what an action tells you about me; commitment asks how an action changes my future options.
- C) Signaling is ethical; commitment is manipulative.

D) They are the same idea.

Field assignment

This week: run the Credibility Test on three real claims, and make one of your own claims easier to verify.

1. Find three high-stakes claims made to you this week, at work, in a purchase, or at home. Run the eight questions on each and note whether the claim felt more or less credible than the evidence supported.
2. Catch one “Would they say this anyway?” statement and note what evidence would actually distinguish the cases.
3. Before making one truthful claim of your own (a status update, a commitment, a recommendation), add one piece of evidence the listener can check without relying on your word. Record how they responded.

Debrief questions for next week

1. Where did persuasiveness and credibility point in different directions, and which one did you act on?
2. When you added verification to your own claim, how did the other person’s response change?
3. Where did you slide toward cynicism, and how did you return to “trust and still understand the game”?

The question to carry

What makes this message credible?

My notes

Module 8 — Commitment Changes the Game

The question: What makes this commitment real?

The tool: The Commitment Test

Use this before making a threat, promise, deadline, boundary, or “final” position, and to evaluate one someone else makes to you.

The situation

The commitment I’m considering (or facing), and whether it’s an intention, promise, or strategic commitment:

Part 1 — What makes this commitment real?

What actually binds it? (contract / rule / third-party authority / legal restriction / public commitment / financial consequence / reputational cost / irreversible investment / merely words)

Can the other player observe or verify it?

Part 2 — What am I giving up, and why?

What future option am I giving up? (Be precise.)

What strategic benefit does the constraint create?

Part 3 — Threats and promises

“If you

, I will

.” Threat or promise?

When the moment arrives, will carrying it out still make sense? Is there a smaller consequence I would enforce every time?

Part 4 — Boundaries

“If this occurs, here is what I will do”:

What I will actually do, calmly and consistently, when it’s crossed:

Part 5 — Their response

How will their best response change (concede, walk away, commit back, find alternatives, retaliate, stop)?

Can this produce an impasse? What if they refuse to move?

Part 6 — The costs of commitment

What happens if new information appears? What flexibility will I wish I still had?

Is this strategy, or stubbornness and sunk cost?

Part 7 — Ethics

Does it preserve truth, autonomy, dignity, and informed choice?

Decision: Commit as planned / commit to something smaller or more enforceable / stay flexible. Why?

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small design studio and is renewing a contract with her largest client. The client's procurement lead, Marcus, asks for a 15 percent cut. Dana says, "Our rates are final; we can't go lower," then agrees to 5 percent off when he pushes. Now he wants another 5 percent, and Dana says this time she means it. Marcus mentions that his authority caps the budget at a set number and anything above it goes to his CFO. At home, Dana has told her teenage son three times this month that if he misses curfew again, he loses the car "for the rest of the year," and each time he kept the keys. She is also drafting a public post promising her studio will "never charge clients for revisions," though her biggest project next quarter has a scope that's hard to predict.

1. Why does Dana's "we can't go lower" carry little weight with Marcus, and what would make a price position real?
2. Analyze the curfew threat. Why isn't it working, and what would a more credible version look like?
3. Evaluate the planned "never charge for revisions" announcement using the costs of commitment.
4. How should Dana treat Marcus's claim about his budget cap?
5. What is one ethical line Dana should keep in mind as she firms up her positions?

Exercise B — Your own game

Pick one live situation where you are about to state, or recently stated, a threat, promise, deadline, boundary, or "final" offer: a negotiation, a rule with your kids, a boundary with a relative, a public goal at work.

1. Write the statement exactly and classify it: intention, promise, or strategic commitment.
2. Work through all seven parts of the Commitment Test.
3. Rewrite the statement so it is something you will actually enforce, observable to the other person, and flexible enough to survive the most likely change in circumstances.

4. Flip it: take a “final” position or “I have no choice” someone recently gave you. What actually binds them? If the answer is “nothing except that they said so,” what would you do differently?

Exercise C — Partner drill / role-play

Format: Pairs, 12–15 minutes, plus a 5-minute debrief. Works with an AI practice partner too.

Setup: A freelance photographer (Player A) and an event planner (Player B) are negotiating the fee for a wedding shoot. Run it in three short rounds of three to four minutes each.

Private instructions, Player A (photographer): You’d prefer \$3,000 and would accept \$2,400. Round 1: no real constraint. Round 2: a genuine, verifiable constraint: another booking on the same date at \$2,600, which you must confirm by tomorrow at noon. Share it if you choose. Round 3: decide whether to make a firm “final” statement and whether to back it with anything.

Private instructions, Player B (planner): Your client’s budget is \$2,800, but you’d love to come in under it. Round 1: test A’s stated minimums. Round 2: before reacting to any constraint, ask what actually binds it. Round 3: you may make your own firm commitment, such as truthfully stating your client’s cap.

Rules: No inventing constraints that don’t exist on your instruction card. Truth, autonomy, dignity, and informed choice apply.

Debrief questions: 1. In Round 1, did A’s “minimum” hold? What did B learn when it moved? 2. In Round 2, did the observable deadline change B’s best response? Did it reduce uncertainty, create pressure, or both? 3. In Round 3, did two commitments collide and threaten an impasse? Where did fairness or reciprocity shape the outcome?

Quiz

- In the book’s framework, which of these is a strategic commitment rather than an intention or a promise?
 - “I’m definitely waking up at 5:00 tomorrow.”
 - “I promise I’ll meet you at the gym at 5:30.”
 - Arranging for your partner to pick you up at 5:15 and prepaying a session you lose if you don’t attend.
 - Telling yourself you’re truly committed to getting in shape.
- Schelling’s paradox of commitment holds that:
 - More options always produce more bargaining power.
 - A player can sometimes strengthen its position by openly making its own alternatives worse.

- C) The player who commits most aggressively always wins.
 - D) Commitments work best when kept secret.
3. A parent threatens six months of grounding over an uncleaned room. According to the chapter, the main problem is that:
- A) The threat isn't severe enough.
 - B) Carrying it out would be costly for the parent too, so the child expects them to back down.
 - C) Threats never work with children.
 - D) The parent didn't state it publicly.
4. A seller tells you, "I'd love to go lower, but corporate won't let me." The most useful question to ask is:
- A) "How committed are you, really?"
 - B) "What actually binds you: a rule, an authority, a contract, or just words?"
 - C) "Why don't you just quit?"
 - D) "Can you make it your final offer?"
5. Why does a commitment usually need to be observable?
- A) Hidden commitments are unethical by definition.
 - B) A hidden commitment can constrain you, but it can't change the other player's present decision unless they know or can infer it.
 - C) Observable commitments are always irreversible.
 - D) Only public commitments affect your own future behavior.
6. In the Swope, Cadigan, and Schmitt bargaining experiments, credible commitment:
- A) Increased the committer's payoff and also increased the probability of bargaining failure.
 - B) Had no effect on outcomes.
 - C) Always improved efficiency for both parties.
 - D) Only worked when both sides committed at once.
7. A team keeps funding a failing project because "we've already spent too much to stop." This is best described as:
- A) A strategic commitment.
 - B) A credible promise.
 - C) The sunk-cost fallacy.
 - D) A boundary.
8. You tell a relative, "If you keep yelling, I'll end the conversation," and you've stayed the last three times. What does the chapter suggest?
- A) Make the consequence more dramatic so it gets noticed.
 - B) Stop setting boundaries, since they don't work.

- C) Calmly say “We’ll continue this when we can talk respectfully” and leave, consistently, so your behavior makes the boundary credible.
- D) Announce the boundary publicly to raise its cost.

Field assignment

This week: audit your commitments. Each day, notice one threat, promise, deadline, boundary, or “final” position, yours or someone else’s. Record the exact words, whether it was an intention, promise, or strategic commitment, what actually bound it (or “words only”), and whether it was followed through.

Then pick one boundary or consequence you have stated but not enforced. Replace it with a smaller, clear consequence you will actually carry out, state it once calmly, and follow through the next time the condition occurs. Record what happens.

Debrief questions for next week: 1. How many of the “final” positions or “I can’t” statements you heard were backed by something real? How could you tell? 2. When you enforced a smaller consequence, how did the other person’s behavior or expectations change, if at all? 3. Where did you notice a commitment, yours or someone else’s, that had become a trap because circumstances changed? What flexibility would have helped?

The question to carry

What makes this commitment real?

What option am I giving up? How does giving it up change their decision? What flexibility will I wish I still had if circumstances change? What happens if they refuse to move?

My notes

Module 9 — The Shadow of the Future

The question: What changes because there is a next time?

The tool: The Repeated-Game Audit

Use this when an interaction is likely to continue. Don't evaluate it as though it ends today.

The interaction:

The move I'm tempted to make:

Step 1 — Is there a next round? Is there likely to be a next round? (yes / no / unsure)

How valuable is that future interaction to me? To them?

Is an endpoint visible (last day, final contract, move away)?

Step 2 — Transaction or relationship? Is this a transaction that ends when the exchange ends, or a relationship that creates another round?

Step 3 — Teaching and learning What is today's move teaching the other person to expect from me?

What am I learning about them?

What does their history tell me (reliable for years, or "here we go again")?

Step 4 — Reputation What reputation are we each building?

Is that reputation useful predictability, or exploitable predictability?

What reputation do I want to be bound by?

Step 5 — Incentives What behavior am I rewarding?

What behavior am I making costly?

Step 6 — Repair If a mistake occurs, is there a path back to cooperation?

How would I tell a mistake from exploitation?

Step 7 — The hundred-round test What happens if we keep playing this way for the next hundred rounds?

My move, revised:

Practice exercises

Exercise A — Analyze the scenario

Dana runs a small marketing agency. For four years her biggest client has been a regional furniture retailer, and her main contact there is Marcus. Last month Marcus signed off on

an ad with the wrong sale dates, and the client lost a weekend of sales. Dana's contract technically lets her bill for the rework, and her account lead, Priya, wants to bill the full \$4,000 and send Marcus an email quoting his approval line by line "so it never happens again." Meanwhile Dana also has a freelance designer, Theo, who has missed three of his last four deadlines, each time with an apology and a promise, and Dana has absorbed the overtime "to keep the relationship good." Theo just asked for a raise. Separately, Dana learned that Marcus is leaving the retailer in two months, and his replacement is unknown.

1. Which of these relationships has the longer shadow of the future, and does Marcus's departure shorten it?
2. What would Priya's approach win in this round, and what might it teach the client to expect next time?
3. What has Dana been teaching Theo to expect, and what does the chapter suggest she do?
4. How does noise apply to the Marcus situation, and how should Dana distinguish mistake from exploitation in each case?
5. Run the hundred-round test on Dana's current approach with Theo.

Exercise B — Your own game

Choose a relationship you expect to continue (coworker, partner, child, client) and one recurring moment in it, such as how mistakes or disagreements are handled.

1. Write down the last three times this moment happened and how you responded.
2. Complete the Repeated-Game Audit for the next time it happens.
3. For each of your three past responses, finish this sentence: "From this, they probably learned that

."

4. Where do your words ("tell me the truth") and your history disagree?
5. Decide on one change: a behavior you'll start rewarding, a behavior you'll make more costly, or a repair route you'll offer. Write the exact sentence you might say.

Exercise C — Partner drill / role-play

Format: pairs, 15 minutes. (An AI practice partner can play Player B.)

Setup: Two roommates, or two coworkers who share a project, have been splitting tasks for a year. Last week Player B didn't do a task they had agreed to do, and Player A had to cover it.

Player A privately knows: You're annoyed. This is the second time in a month. You aren't sure whether B forgot or is starting to take advantage of you. You want the arrangement to continue.

Player B privately knows: You genuinely misread a message and thought A had swapped tasks with you. The first miss a few weeks ago was real carelessness on your part.

Round 1 (4 min): Player A responds as if this is a one-shot game: maximize being right today. Player B responds naturally.

Round 2 (4 min): Reset. Player A uses the audit (mistake or exploitation? what will this teach? is there a path back?) and aims to set a clear expectation without starting a feud. Player B responds naturally.

Debrief (5-7 min):

- In Round 1, what did B learn to expect from A? Did B start retaliating, even subtly?
- In Round 2, did A find a way to address the earlier, real miss without punishing the honest mistake?
- What would the hundred-round version of each conversation look like?

Quiz

1. According to the chapter, what changes when a one-shot Prisoner's Dilemma becomes a repeated game?
 - A) The immediate payoffs of cooperating and defecting change
 - B) The continuation of the relationship enters the calculation, so today's move can affect future treatment
 - C) Defection becomes impossible
 - D) Players become more generous by nature
2. What did Pedro Dal Bó's experiments vary, and what did they find?
 - A) The size of the payoffs; larger payoffs increased cooperation
 - B) Whether players could talk; communication eliminated defection
 - C) The probability of another round; a longer expected future significantly reduced opportunistic behavior
 - D) Players' personalities; cooperative people always cooperated
3. What is the most accurate lesson from Axelrod's tournaments?
 - A) Tit for Tat is the universally best strategy for life
 - B) Complex strategies always beat simple ones
 - C) Once people expect a future together, reciprocity becomes strategically possible
 - D) Cooperation never pays against selfish players
4. A manager says "bring me bad news early" but berates the first person who does. Next month nobody reports a problem. What best explains this?
 - A) The employees are responding to the history of the game, not the manager's words
 - B) The employees lack integrity

- C) The manager's instruction was unclear
D) Bad news is rare in well-run teams
5. Aoyagi, Bhaskar, and Fréchette found that under noisy monitoring, players who sustained cooperation tended to use strategies that were:
A) Simpler and harsher
B) Identical to those under perfect monitoring
C) Purely random
D) More complex and more lenient
6. Your friend has cancelled plans with you six times in a row, each time with a heartfelt apology and a promise to make it up. Following the chapter, what is the best response?
A) Keep accepting, because thinking long term means always cooperating
B) Treat the pattern as information; set a clear expectation or change the terms, and be willing to step back if it continues
C) Retaliate by cancelling on them six times
D) Ignore history and evaluate only the most recent cancellation
7. Why does the chapter say "cooperation is not the same thing as goodness"?
A) Because cooperation is always a form of manipulation
B) Because repeated games make cooperation impossible
C) Because cartels, corrupt organizations, and people deceiving a third party can sustain cooperation too
D) Because selfish motives make cooperation fake
8. A neighborhood restaurant has a longer direct shadow of the future than a highway restaurant mainly because:
A) Its customers are likely to return and to talk with neighbors who might come
B) Its owners are more ethical
C) It charges higher prices
D) Highway restaurants are not subject to reviews

Field assignment

This week, pick one ongoing relationship and pay attention to what you are teaching. Each day, note one moment where someone told you the truth, disagreed, made a mistake, or kept or broke their word. Record:

- What they did.
- How you responded.
- What your response likely taught them to expect next time.
- Whether it was a mistake or a pattern, and whether there was a path back to cooperation.

At least once this week, deliberately choose the move that makes tomorrow's game better over the one that gives you the most today. Note what happened.

Debrief questions for next week: 1. Where did your words and your behavior teach different lessons? 2. Did you find a relationship where you've been cooperating unconditionally with repeated exploitation, or punishing honest mistakes as if they were exploitation? What would conditional cooperation look like there? 3. Running the hundred-round test on your most important relationship this week, where does the current pattern lead?

The question to carry

| What changes because there is a next time?

My notes

Module 10 — Diagnose the Coordination Problem

The question: Are we actually in conflict—or are we failing to coordinate?

The tool: The Coordination Diagnostic

Use this when a group or relationship keeps producing an outcome nobody seems to want.

The bad outcome that keeps happening (no character judgments):

1. Do we actually want incompatible things? If yes, you may have a conflict or bargaining problem. If partly, note: Where we align:

Where we diverge:

2. Would I take the desired action if I knew the others would too? My honest answer:

What I suspect others would honestly answer:

(Yes → likely coordination. “I’d rather others bear the cost” → likely social dilemma; clarity alone may not fix it.)

3. My diagnosis (circle one): Conflict / Coordination failure / Social dilemma / Mixed
Why:

4. What am I uncertain about? (their actions, intentions, knowledge, commitment)

5. What are they uncertain about? (Your intentions may be part of the problem.)

6. What expectation needs to become mutually clear? Who acts?

When?

According to what standard?

Where is the handoff?

7. Is there a focal point or convention we can coordinate around?

8. Does everyone know that everyone knows? Would making it public change the game?

9. Choose the intervention:

- ☐ A shared plan
- ☐ A public commitment
- ☐ A clear owner
- ☐ A standard
- ☐ A deadline
- ☐ A focal point
- ☐ Information about what everyone else intends to do

My one sentence that makes the next move obvious to everyone:

Ethics check: Does this respect truth, autonomy, dignity, and informed choice for everyone affected?

Practice exercises

Exercise A — Analyze the scenario

Dana runs a twelve-person marketing team at a regional credit union. Three months ago, leadership approved a shared content calendar to replace the team's spreadsheets and email threads. Dana's kickoff email said, "Everyone should start using the new calendar when you can." Today only two people use it. The others say they like the tool but keep their spreadsheets "because that's where everyone else's dates are." Last week two staff members scheduled competing promotional emails for the same morning. A member newsletter also went out without its compliance review: the writer assumed the designer had sent it for sign-off, and the designer assumed the writer had. Dana's director says, "Your team doesn't take ownership and resists change. Motivate them." Dana is drafting a speech about accountability.

1. Is the calendar problem primarily conflict, coordination failure, or social dilemma? Apply the simple test.
2. What is blocking adoption, and why didn't the kickoff email solve it?
3. Where exactly did the newsletter process break?
4. Is a motivational speech the right intervention? What would you do instead?
5. Could any part of this turn out not to be coordination?

Exercise B — Your own game

Pick a live situation where people "can't seem to get on the same page": a household chore, a team handoff, a family schedule, a volunteer group.

1. Describe the bad outcome in one sentence, without blaming anyone's character.
2. Work through the Coordination Diagnostic.
3. On step 2, answer honestly for yourself first. If your answer is "I'd rather someone else did it," say so; it changes the diagnosis.
4. Write the one sentence that makes the next move obvious, in the style of "Paul owns this by 3 PM."
5. Decide how you'll make it mutually visible: in front of everyone, in a shared document, on a visible schedule.

6. If it's something you renegotiate every week, draft a convention everyone can know before the next conflict arises.

Exercise C — Partner drill / role-play

Format: Pairs or small groups, 12–15 minutes; an AI practice partner can play Player B.

Setup: A shared office kitchen keeps getting left messy. Player A facilitates a short conversation to fix it. Player B privately draws one card:

- **Card 1 — Coordination:** “I’d happily clean on a set day, but there’s no schedule and I never know if someone already did it.”
- **Card 2 — Social dilemma:** “If others clean, I’ll enjoy it and not bother.”
- **Card 3 — Mixed:** “I support a rotation, but I want Fridays excluded since I’m remote then; others want everyone equal on every day.”

Rules: Player A has 7 minutes to diagnose the card using questions, not accusations, including some version of the simple test and at least one question about what B is uncertain about. Player B answers honestly in character without naming the card. A then proposes one intervention.

Debrief (5 minutes):

- Which card did A guess, and what evidence led there?
- Did the intervention fit? (Card 1: schedule, owner, public visibility. Card 2: something that changes incentives or accountability. Card 3: a fair selection rule among options people disagree about.)
- How would the wrong card’s intervention have failed?
- Swap roles and draw again.

Quiz

1. Two friends prefer eating together to eating alone; one slightly prefers Italian, the other steak. What matters most?
 - A) The person with the stronger preference gets their choice
 - B) Ending up at the same restaurant
 - C) Finding the objectively better restaurant
 - D) Proving whose preference is more reasonable
2. Which question best separates coordination from a social dilemma?
 - A) “Do we all want the same final outcome?”
 - B) “Who has more power in this group?”
 - C) “If I knew for certain everyone else would cooperate, would I want to cooperate too?”
 - D) “Have we communicated clearly enough?”

3. Asked to match another person's positive whole number, many choose 1. Why does this work?
 - A) 1 is mathematically optimal for matching
 - B) It stands out, and you can imagine the other person thinking the same
 - C) People are told 1 is the default
 - D) Smaller numbers pay more
4. Employees all like a new system if everyone uses it, but switching alone makes work harder, so nobody switches. The barrier is best described as:
 - A) Low motivation
 - B) A genuine conflict of interests
 - C) Strategic uncertainty
 - D) A free-rider problem
5. What does telling five employees together "You five are responsible" add over telling each privately?
 - A) A financial incentive
 - B) Mutual awareness: each knows the others know who is responsible
 - C) Stronger arguments for the change
 - D) Nothing; the information is identical
6. Which claim do the Cooper et al. and Senegal studies support?
 - A) Talking more makes every group problem disappear
 - B) Nonbinding messages are useless because they aren't credible
 - C) Communication can substantially improve coordination when it resolves the blocking uncertainty, but it doesn't eliminate every failure
 - D) Revealing intentions works only in small groups
7. Two co-founders both want the company to succeed but each wants their own strategy adopted. This is:
 - A) Pure coordination
 - B) A mixed-motive game with coordination and conflict
 - C) A pure free-rider problem
 - D) Not a strategic problem
8. A team agrees "someone should call the customer today," and nobody does; sales and service each assumed the other would. The best intervention is:
 - A) A talk about taking initiative
 - B) A penalty for everyone
 - C) "Jordan owns the call by 3 PM," said in front of the team
 - D) More debate about customer ownership in general

Field assignment

This week, pick one recurring friction point you've been treating as a conflict or a motivation problem: a task that falls through the cracks, a failing handoff, a schedule you renegotiate every week.

1. Run the Coordination Diagnostic before saying anything, and commit to a diagnosis.
2. If it's coordination, make one expectation mutually clear (owner, time, standard) where everyone involved can see everyone else receive it, or propose a simple convention.
3. If it's a social dilemma or conflict, don't force a coordination fix; note what kind of intervention it really needs.
4. Record what you said, where, what happened over the following days, and whether the problem recurred.

Debrief questions for next week:

- Did your initial assumption match your final diagnosis? What changed your mind?
- Did a public, explicit expectation produce a different result than past private reminders?
- If the problem returned despite clear, mutually known expectations, what does that suggest: a hidden incentive, a real conflict, or a game that needs redesigning?

The question to carry

Are we actually in conflict—or are we failing to coordinate? If it's coordination: What expectation needs to become mutually clear? Make the next move obvious to everyone.

My notes

Module 11 — Change the Game

The question: Do I need a better move—or a better game?

The tool: The Game-Change Audit

Use this when the same bad outcome keeps returning, after better arguments, better wording, more motivation, and more pressure haven't fixed it.

Step 1 — Name the recurring problem

The outcome that keeps returning:

What I (or we) have already tried, inside the current game:

My current character explanation ("they're lazy / selfish / gaming it"):

Step 2 — What does the game actually reward?

What behavior is the current game rewarding? (Not what the policy claims to reward. What actually pays?)

What behavior is it punishing? Could the desirable action currently be costly?

What metric are people optimizing? Is it actually the outcome we care about?

What does my stated value say, and what does the reward say?

Step 3 — Inspect the levers

Lever	How it works now	One possible change
Rules: what players can and can't do		
Payoffs: what pays, what costs		
Information: who knows what, when		
Who decides		
Timing: who moves when		
Options: what choices exist		
Players: could a third party help?		
Observability and feedback		
Default: what happens if nobody acts		

What is currently easy?

What is currently hard? (Friction is part of the design.)

Is there an incentive or rule I should remove rather than add?

Step 4 — Stress-test the redesign

How could a smart player game this system, including my new version of it?

What unintended behavior might the new rule create? (Think one response further: then what?)

What harder-to-measure work might lose effort if I reward the measurable part?

Step 5 — Ethics check

Does the player understand the game?

Do they retain meaningful choice?

Are important facts hidden, or does the system exploit confusion?

Does it treat people as decision-makers rather than objects to be manipulated?

Step 6 — Smallest testable change

What specific mechanism appears to be producing the unwanted behavior?

The smallest change capable of testing that hypothesis:

What I'll observe, and when:

If exit is on the table: what would I lose, and what future would I give up?

The central question: Can I redesign the game so that the behavior I want becomes a more natural response to the structure itself?

Practice exercises

Exercise A — Analyze the scenario

Dana manages a twelve-person software support team. A leaderboard ranks engineers by “tickets closed per week,” and the top three each quarter get a bonus and first choice of vacation weeks. Dana keeps saying customer satisfaction and knowledge-sharing matter most, and has run two mentoring workshops. Yet senior engineers rarely help juniors, complex tickets sit untouched until someone is assigned them, and customers increasingly reopen tickets marked “resolved” within days. Reopened tickets are counted as new tickets. Escalations to engineering require Dana’s personal sign-off, and she is often in meetings until late afternoon. Last month, two engineers were caught splitting single customer issues into several tickets. Dana’s boss has suggested firing them and “hiring people with better values.”

1. What behavior is Dana’s current game rewarding, and what is it punishing?
2. What metric is the team optimizing, and how does it differ from the goal?
3. Which design levers could Dana change? Name at least three.
4. How could a smart player game your proposed redesign?
5. Is “hire people with better values” the right diagnosis? What would you test first?

Exercise B — Your own game

Pick a problem in your own life that keeps coming back despite your best efforts: a recurring argument at home, a work process that keeps producing the same failure, a client pattern, or a habit you keep “trying harder” to change.

1. List every “better move” you have already tried: the conversations, the wording, the willpower.
2. Work the problem through all six steps of the Game-Change Audit.
3. Circle the single lever most responsible for the behavior.
4. Design the smallest change that would test your hypothesis. If you are both designer and player, design for the behavior you are likely to have, not the behavior you wish you had.
5. Run the ethics questions. If others are affected, can you design the change with them rather than impose it?
6. If you lack authority to redesign anything, ask: why am I trying to become better at this game? List what exit would cost before deciding.

Exercise C — Partner drill / role-play

Format: pairs or small groups, 15 minutes. Works with an AI practice partner too.

Setup: A small family business, a bakery, has a recurring problem. Every Friday the two co-owners, siblings, argue about who gets to approve the weekend order, and the order goes in late about half the time. Every order decision requires both of them to agree.

Roles:

- **Sibling A (front of house).** Privately: you think your sibling over-orders and you discard unsold goods. You lack the sales numbers.
- **Sibling B (kitchen and ordering).** Privately: you have the sales data, believe running out on Saturday costs more than waste, and feel micromanaged.
- **Designer (optional third player).** Your job is not to decide who is right about quantities but to help them redesign the game.

Round 1 (5 min): The siblings argue Friday's order under the current rules. The Designer notes the rules, information gaps, timing, and who decides.

Round 2 (7 min): The Designer leads the group through the audit's lever table. Propose at least two redesigns (for example, B decides within a set range, a shared sales-and-waste sheet, or a standing Wednesday review), then stress-test: how could either sibling game the new rule?

Debrief (3 min): 1. In Round 1, which disagreement was really about payoffs, and which was about missing information? 2. Which lever changed the most in your redesign? Who decides, information, or timing? 3. How might the new arrangement be gamed, and what would you watch for in the first month? 4. Did the redesign treat both siblings as informed decision-makers, or did it quietly favor one?

Quiz

1. How does mechanism design differ from the game theory used in earlier modules?
 - A) It focuses on predicting a single opponent's next move.
 - B) It starts from the outcome you want and asks how to design the rules to make it more likely.
 - C) It assumes players are always irrational.
 - D) It replaces incentives with persuasion.
2. In the "one cuts, the other chooses" cake example, why does the cutter make the pieces equal?
 - A) The parent lectures them on fairness first.
 - B) The cutter has become less selfish.
 - C) If one piece is clearly larger, the other child will take it, so self-interest favors equal pieces.

- D) The parent inspects the pieces afterward.
3. A company says “tell us the truth” but penalizes people who report bad news. According to the chapter, what should you expect?
- A) The stated value will win because people have integrity.
 - B) The rewarded behavior will tend to win.
 - C) Communication training will resolve it.
 - D) Behavior will be unaffected by the penalty.
4. What did the Kenyan teacher-incentive study illustrate?
- A) Incentives never change behavior.
 - B) Incentives tied to test scores produced lasting, broad learning gains.
 - C) Incentives raised scores on the incentivized exams, mainly through test preparation, without clear gains in attendance, homework, or pedagogy.
 - D) Teachers ignored the incentive entirely.
5. A manager rewards her team solely on the number of leads generated, then complains that lead quality has dropped. What is the best diagnosis?
- A) The team missed the point and needs a firmer talk.
 - B) The scoreboard told them what the point was; the metric is not the goal.
 - C) The team has poor values.
 - D) Lead quality cannot be influenced by design.
6. A project automatically continues unless someone explicitly stops it. Which lever is this?
- A) The default: what happens when nobody chooses.
 - B) The players.
 - C) Exit.
 - D) Incentive compatibility.
7. You introduce a productivity metric. Productivity rises, but quality collapses. What is the chapter’s recommended response?
- A) “Why aren’t people behaving according to my design?”
 - B) Add many more rules at once.
 - C) “What did I misunderstand about the game?” and then observe, update, and adapt.
 - D) Abandon all measurement permanently.
8. A subscription takes one click to start and twenty to cancel. From the chapter’s ethical standpoint, what is the key point?
- A) Since cancelling is technically possible, the design is neutral.
 - B) Friction is part of the game, and being able to make a choice strategically difficult does not mean you should.
 - C) Friction only matters in government systems.
 - D) Ethical design means removing all choices.

Field assignment

This week: run one small, ethical game redesign.

Choose one recurring problem you have the authority to influence: a household routine, a team process, a client arrangement, or your own work environment. Complete the Game-Change Audit, identify the mechanism you believe is producing the unwanted behavior, and implement the smallest change that tests it. If others are affected, explain the new game clearly, or better, design it with them.

Each day, record in a few lines: - What behavior did you observe under the new structure? - Did anyone respond in a way you didn't predict, including ways of gaming the new rule? - What became easier, and what became harder?

Debrief questions for next week: 1. What did the system “talk back” with, and what did it reveal that you had misunderstood about the game? 2. Before this week, how much of the problem had you attributed to someone's character? How has that changed? 3. What is your next smallest adjustment, or have you concluded that the better move is to leave this game?

The question to carry

Do I need a better move—or a better game?

My notes

Module 12 — Strategic Judgment

The question: What did I learn from what just happened?

The tool: The Strategic Judgment Worksheet

Use this when the stakes matter, something doesn't make sense, a pattern keeps repeating, incentives are unclear, trust is uncertain, or several people affect the outcome. Don't use it for "Sure."

Part 1 — The situation

What happened (facts only):

My first interpretation (my story):

How confident am I in that story (0-100%)?

What is my immediate impulse, the thing I'd say or do with no pause?

Part 2 — The Eleven Questions (scan all, mark the 2-3 that matter most)

✓	Model	Question	My answer
☐	1. See the Game	What game am I actually in? (Players, actions, timing, outside options, who can combine with whom)	
☐	2. See the Payoffs	What does each player actually value? Are our interests opposed, aligned, or mixed?	

✓	Model	Question	My answer
☐	3. Think in Responses	Then what? Is there an option I shouldn't choose regardless of what they do?	_____ _____
☐	4. Manage Predictability	What should they be able to predict about me, and what becomes exploitable if they can?	_____ _____
☐	5. See the Equilibrium	What pattern are we producing, and why is it stable?	_____ _____
☐	6. Map Information & Beliefs	What do I know, what am I assuming, and what might they know that I don't?	_____ _____
☐	7. Separate Words From Credibility	What makes this believable?	_____ _____
☐	8. Commitment Changes the Game	What actually binds them, or me?	_____ _____
☐	9. The Shadow of the Future	What changes because there is a next time?	_____ _____
☐	10. Diagnose the Coordination Problem	Are we actually in conflict, or failing to coordinate? (If everyone else cooperated, would I still have an incentive not to?)	_____ _____
☐	11. Change the Game	Do I need a better move, or a better game?	_____ _____

The 2-3 lenses that matter most, and why:

The lenses I'm deliberately setting aside, and why:

Part 3 — Competing hypotheses

Hypothesis about what's going on	Evidence for	What would tell this apart from the others?
H1: _____	_____	_____
H2: _____	_____	_____
H3: _____	_____	_____

What would have to be true for their behavior to make sense?

Part 4 — The move

My planned move:

What I predict will happen:

Will this work?

Should I do it? Check each: - Truth: Am I saying only what I believe is true and not creating a false impression?

- Autonomy: Does the other person stay free to choose, including to say no?

- Dignity: Would I be comfortable if they saw my full reasoning?

- Informed choice: Do they have the information they'd need to decide well?

What are we cooperating toward?

Is the right move vulnerability, walking away, or refusing this game?

The Practice Log: Model → Act → Observe → Update → Adapt

Keep one entry per meaningful interaction. Over weeks, this log is how your intuition gets calibrated.

Date	Model (my read + key lens)	Act (what I did)	Prediction	Observe (what actually happened)	Update (where reality differed, and why)	Adapt (new strategy or new model)

Surprise prompts for the Update column: - They didn't cooperate. What did I misunderstand?

- They stayed. What payoff did I underestimate?

- They followed through. What made the commitment stronger than I realized?

- My rule or design backfired. What response did it create that I didn't anticipate?

Practice exercises

Exercise A — Analyze the scenario

Dana is the operations director at a regional logistics company. For the third quarter in a row, the sales team has promised clients delivery windows her warehouse can't meet, and her team absorbs the overtime and the angry calls. On Monday, the VP of Sales, Marcus, emails her and copies the CEO: "We just landed the Harlan account. They need three-day

turnaround starting next month. This is non-negotiable. The CEO is aware.” Dana’s first impulse is to reply-all with a detailed list of every missed promise from the last nine months. She knows Marcus’s team is paid commission on signed contracts, not on retained or profitable accounts. She doesn’t know whether the CEO actually approved three-day turnaround or was simply told about the deal. She’ll work with Marcus for years, and the Harlan account would be good for the company if it could be served reliably. Two of her shift leads have privately told her they’re close to quitting.

1. Which two or three lenses matter most here, and which can Dana set aside for now?
2. Separate the facts from Dana’s likely story. What competing hypotheses should she hold about Marcus and the CEO?
3. What is wrong with Dana’s first impulse, judged by both “Will this work?” and “Should I do it?”
4. Design a move and state a prediction. Then describe what Dana should do in the Observe and Update steps if her prediction turns out wrong.

Exercise B — Your own game

Pick one live situation from your own life that meets at least two of the book’s criteria: the stakes matter, something doesn’t make sense, a pattern keeps repeating, incentives are unclear, trust is uncertain, multiple people affect the outcome, or your normal approach isn’t working.

1. Write the facts in one paragraph with no adjectives or motives. Then write your story in a second paragraph. Put a confidence percentage on the story.
2. Scan all eleven questions in the Strategic Judgment Worksheet for about 60 seconds each. Then circle no more than three. For each lens you set aside, write one sentence explaining why.
3. Write at least three competing hypotheses. For each, name one question you could ask or one behavior you could watch for that would separate it from the others.
4. Draft your move and your prediction. Run the four ethics checks. If any one fails, redesign the move before going further.
5. Ask yourself whether there is a version of this where the best move is vulnerability, or one where it is walking away. Write one sentence on each.
6. Enter the situation as the first row of your Practice Log. You’ll return to it in the Field assignment and may choose it for your Capstone.

Exercise C — Partner drill / role-play

“Two Lenses, Then Update” (15 minutes, pairs or trios, or with an AI practice partner)

Setup. A small nonprofit’s development director (Player A) is meeting a longtime major donor (Player B). The donor has given \$50,000 a year for six years. This year the donor says, “I’m thinking of pausing my gift. I’ll decide by Friday.”

Player A privately knows: The organization recently changed its program focus. The board chair told A that the donor “always comes around.” A has no information about the donor’s finances.

Player B privately knows (choose one card at random and don’t reveal it):

- Card 1: B is upset that nobody told them about the program change before it was announced publicly. B feels taken for granted and wants to be consulted.
- Card 2: B’s business had a bad year and B is embarrassed to say so. B would like to give a smaller amount without losing face.
- Card 3: B’s adult daughter has started her own nonprofit and B feels torn. B is open to a joint project.

Round 1 (5 min). Before speaking, Player A names out loud the two lenses they’ll lead with and one hypothesis they’re testing. A then conducts the conversation. A’s goal is to tell the hypotheses apart, not to secure the gift.

Round 2 (3 min). Pause. A states what they observed, what surprised them, and how they’re updating. B says nothing yet.

Round 3 (4 min). A adapts and continues. A must stay within truth, autonomy, dignity, and informed choice. That means no inventing urgency, no guilt, and no misrepresenting the program change.

Debrief (3 min).

- Which card did A think B held, and when? What evidence moved the estimate?
- Did the board chair’s “always comes around” belief help or hurt? (It’s a story, not a fact.)
- Did A respect B’s ability to say no? Would B feel comfortable if they saw A’s full reasoning?
- Which lens turned out to matter most, and was it one A chose at the start?

Quiz

1. According to the Conclusion, how should the eleven models be used in a live interaction?
 - A) As a fixed sequence, from See the Game through Change the Game
 - B) As lenses, choosing the two or three that matter most right now
 - C) Only in formal negotiations
 - D) All eleven, written out, before every important conversation
2. Which of these is NOT one of the book’s signals that a situation warrants strategic analysis?
 - A) A pattern keeps repeating
 - B) Trust is uncertain

- C) A friend's text message is short
 - D) Multiple people affect the outcome
3. A manager is sure a team member's missed deadlines come from laziness. He introduces a public "late list," and the deadlines get worse. Using the loop, what is the best next step?
- A) Make the late list more prominent
 - B) Conclude the employee is irrational
 - C) Ask what response the new rule created that he failed to anticipate, and revise his model
 - D) Stop tracking deadlines entirely
4. What best distinguishes strategic thinking from suspicion?
- A) Strategic thinking assumes people usually lie
 - B) Suspicion tests hypotheses, and strategic thinking doesn't
 - C) Strategic thinking considers possibilities and tests them, while suspicion assumes hostile intent
 - D) There is no meaningful difference
5. You explain your reasoning clearly. The other person restates it accurately and still refuses. What does the book suggest you consider?
- A) Explaining it again more slowly
 - B) That it may never have been an information problem, and your interests may be genuinely opposed
 - C) That they are acting irrationally
 - D) That you should hide some information next time
6. A sales rep learns that a prospect is anxious about falling behind competitors. She considers inventing a story that two of the prospect's rivals just signed. Which statement reflects the book's position?
- A) If it closes the deal, it's good strategy
 - B) It's fine as long as the product is genuinely good
 - C) It may "work," but it violates truth and informed choice, and "Should I do it?" must be asked separately
 - D) It's acceptable in a one-shot interaction because there's no shadow of the future
7. Two longtime co-founders each withhold their real concerns about the company's direction because each fears the other will use it against them. Nothing changes for months. What move does the Conclusion suggest may sometimes be correct here?
- A) A measured act of vulnerability, such as disclosing a concern first, which can change beliefs and shift the equilibrium
 - B) Withholding even more, to protect your position
 - C) Issuing an ultimatum

D) Waiting for the other person to go first, however long it takes

8. What sits “above all” eleven questions in the book’s framework?

- A) “Who has the most power?”
- B) “What should I say?”
- C) “What did I learn from what just happened?”
- D) “How do I win?”

Field assignment

This week: run the loop five times.

Choose five interactions this week that clear the bar for analysis. At least one should involve someone you’ll deal with again, and at least one should involve more than two people. For each one:

1. Before acting, take the pause between stimulus and response. Write your story, your confidence, and the one or two lenses you’re leading with.
2. Make a specific prediction: “If I do X, they’ll do Y.”
3. Act, then log in the Practice Log what actually happened.
4. Score your prediction as right, partly right, or wrong, and fill in the Update and Adapt columns.

Also record one ordinary interaction where you noticed yourself starting to over-analyze and chose not to.

Debrief questions for next week 1. How many of your five predictions were right? In the ones that were wrong, was the error mostly about payoffs, information, credibility, or the structure of the game? 2. Which surprise taught you the most, and how did you change your model because of it? 3. Did any move tempt you toward something that would “work” but fail the truth, autonomy, dignity, or informed-choice check? What did you do instead?

Capstone project

Purpose

This is the course-completion project. Over two to three weeks, you’ll take one real, high-stakes interaction from your own life through the full strategic-judgment cycle. You’ll see the game, choose the lenses that matter, design an ethical move, act, and learn from what happens. The Capstone does not test whether you get the outcome you wanted. People surprise us, and no framework gives you control over them. It tests whether you saw the game clearly, acted with integrity, and learned faster from what happened.

Choosing your interaction

Choose something real, current, and consequential where you'll actually make a move within the project window. Good candidates include a compensation or contract negotiation, a recurring conflict with a partner, family member, or colleague, a team that keeps failing the same way, a boundary you need to set, a decision about whether to stay in or leave a situation, or a proposal you need others to adopt.

It must meet at least three of the book's criteria: high stakes, something doesn't make sense, a recurring pattern, unclear incentives, uncertain trust, multiple people affecting the outcome, or your normal approach not working.

Avoid situations involving physical danger or ongoing abuse. In those cases, safety and appropriate professional support come before any analysis.

Deliverables

1. **Situation brief** (250–400 words). Facts, kept strictly separate from your story, with a confidence rating, plus your stakes and what “winning” would actually mean to you.
2. **Completed Strategic Judgment Worksheet**. All eleven questions scanned, with the 2–4 lenses you chose analyzed in depth and a written rationale for the lenses you set aside.
3. **Hypothesis table**. At least three competing explanations of the other players' behavior, with a plan to tell them apart.
4. **Move design** (300–500 words). Your move, the alternatives you rejected and why, your explicit prediction, and the full ethics check. Include whether a game-changing move (new options, new rules, new players, or refusing to play) is available.
5. **Practice Log**. At least three Model → Act → Observe → Update → Adapt entries covering the interaction and its follow-ups.
6. **Debrief** (600–900 words). Observe, Update, Adapt. What happened, where reality differed from your prediction and why, what you'd do differently, and what you'll do next. End with one paragraph on which lens you'll now notice faster in other situations (transfer).

Step-by-step process

1. **Pick and frame (days 1–2)**. Choose the interaction. Write the situation brief. Name what you'd consider a good outcome, including outcomes other than “winning,” such as preserving the relationship, telling the truth at a cost, or walking away.
2. **Scan wide (day 3)**. Answer all eleven questions briefly, without editing. This is the only time you deliberately run the whole set.
3. **Choose the lens (day 3)**. Circle the two to four lenses that explain the most. For the others, write one line each on why they matter less here. Over-analysis counts against you, not for you.

4. **Go deep (days 4-5).** Fully analyze the chosen lenses. Map players (including people not in the room), payoffs, likely responses, the stabilizing pattern, information gaps, credibility, commitments, and future rounds, as relevant.
5. **Hold hypotheses (day 5).** Build the hypothesis table. Where you can, design a low-risk question or observation that would discriminate among them before your main move.
6. **Design the move (day 6).** Draft your move and two alternatives. Ask “Will this work?” and “Should I do it?” separately. Consider vulnerability, consider walking away, and consider changing the game. Write a specific, checkable prediction.
7. **Act (days 7-14).** Make the move. Log what happened as soon as you can, with facts first and interpretation second.
8. **Observe, update, adapt (days 7-21).** Compare outcome to prediction. Use the surprise prompts. Revise your model, make at least one follow-up move if appropriate, and log it.
9. **Write the debrief (final days).** Be honest about what you got wrong. A clean prediction that came true teaches less than a surprise you learned from.

Scoring rubric

Criterion	Developing	Proficient	Strong
1. Seeing the game	Describes the conversation but not the structure. Players, options, or outside options are missing. Facts and story are blended.	Identifies the main players, options, and payoffs. Separates facts from story.	Maps the full structure, including players not in the room, outside options, coalitions, and timing. Separates facts from story with explicit confidence levels.
2. Lens selection	Applies every model mechanically, or relies on just one without justification.	Chooses 2-4 relevant lenses and explains the choice.	Chooses lenses with a sharp rationale, explains why the others matter less, and shows the restraint of “more analysis isn’t always better.”

Criterion	Developing	Proficient	Strong
3. Hypotheses and prediction	One explanation, held with certainty. Prediction vague or missing.	Several hypotheses and a specific prediction.	Several well-differentiated hypotheses, a deliberate way to tell them apart, and a specific, checkable prediction made before acting.
4. Move design	Move is a script or tactic disconnected from the analysis.	Move follows from the analysis and considers likely responses.	Move follows from the analysis, weighs alternatives (including vulnerability, walking away, or changing the game), and accounts for future rounds and precedent.
5. Ethics: truth, autonomy, dignity, informed choice	Ethics not considered, or the move relies on deception, pressure, or withheld information.	Move passes all four checks, and “Will this work?” and “Should I do it?” are addressed separately.	Move passes all four checks with specific reasoning, considers what any cooperation is directed toward, and treats the other person as a decision-maker whose agency is respected even when it costs a payoff.
6. Observe / Update / Adapt	Reports the outcome only, or explains away surprises. Calls others “irrational.”	Compares outcome to prediction, identifies at least one model error, and states an adaptation.	Treats surprise as information, asks what would have to be true for others’ behavior to make sense, revises the model specifically, follows up, and names a transferable lesson.

A passing Capstone is Proficient or better on all six criteria. Ethics (criterion 5) must be at least Proficient regardless of other scores. A strategy that works but fails that check is incomplete.

The question to carry

What did I learn from what just happened?

My notes
